



Caching with Spring: Advanced Topics and Best Practices

Michael Plöd
@bitboss

I will talk about

Caching Types / Topologies
Best Practices for Caching in Enterprise Applications
Caching with Spring
JCache and Spring

I will NOT talk about

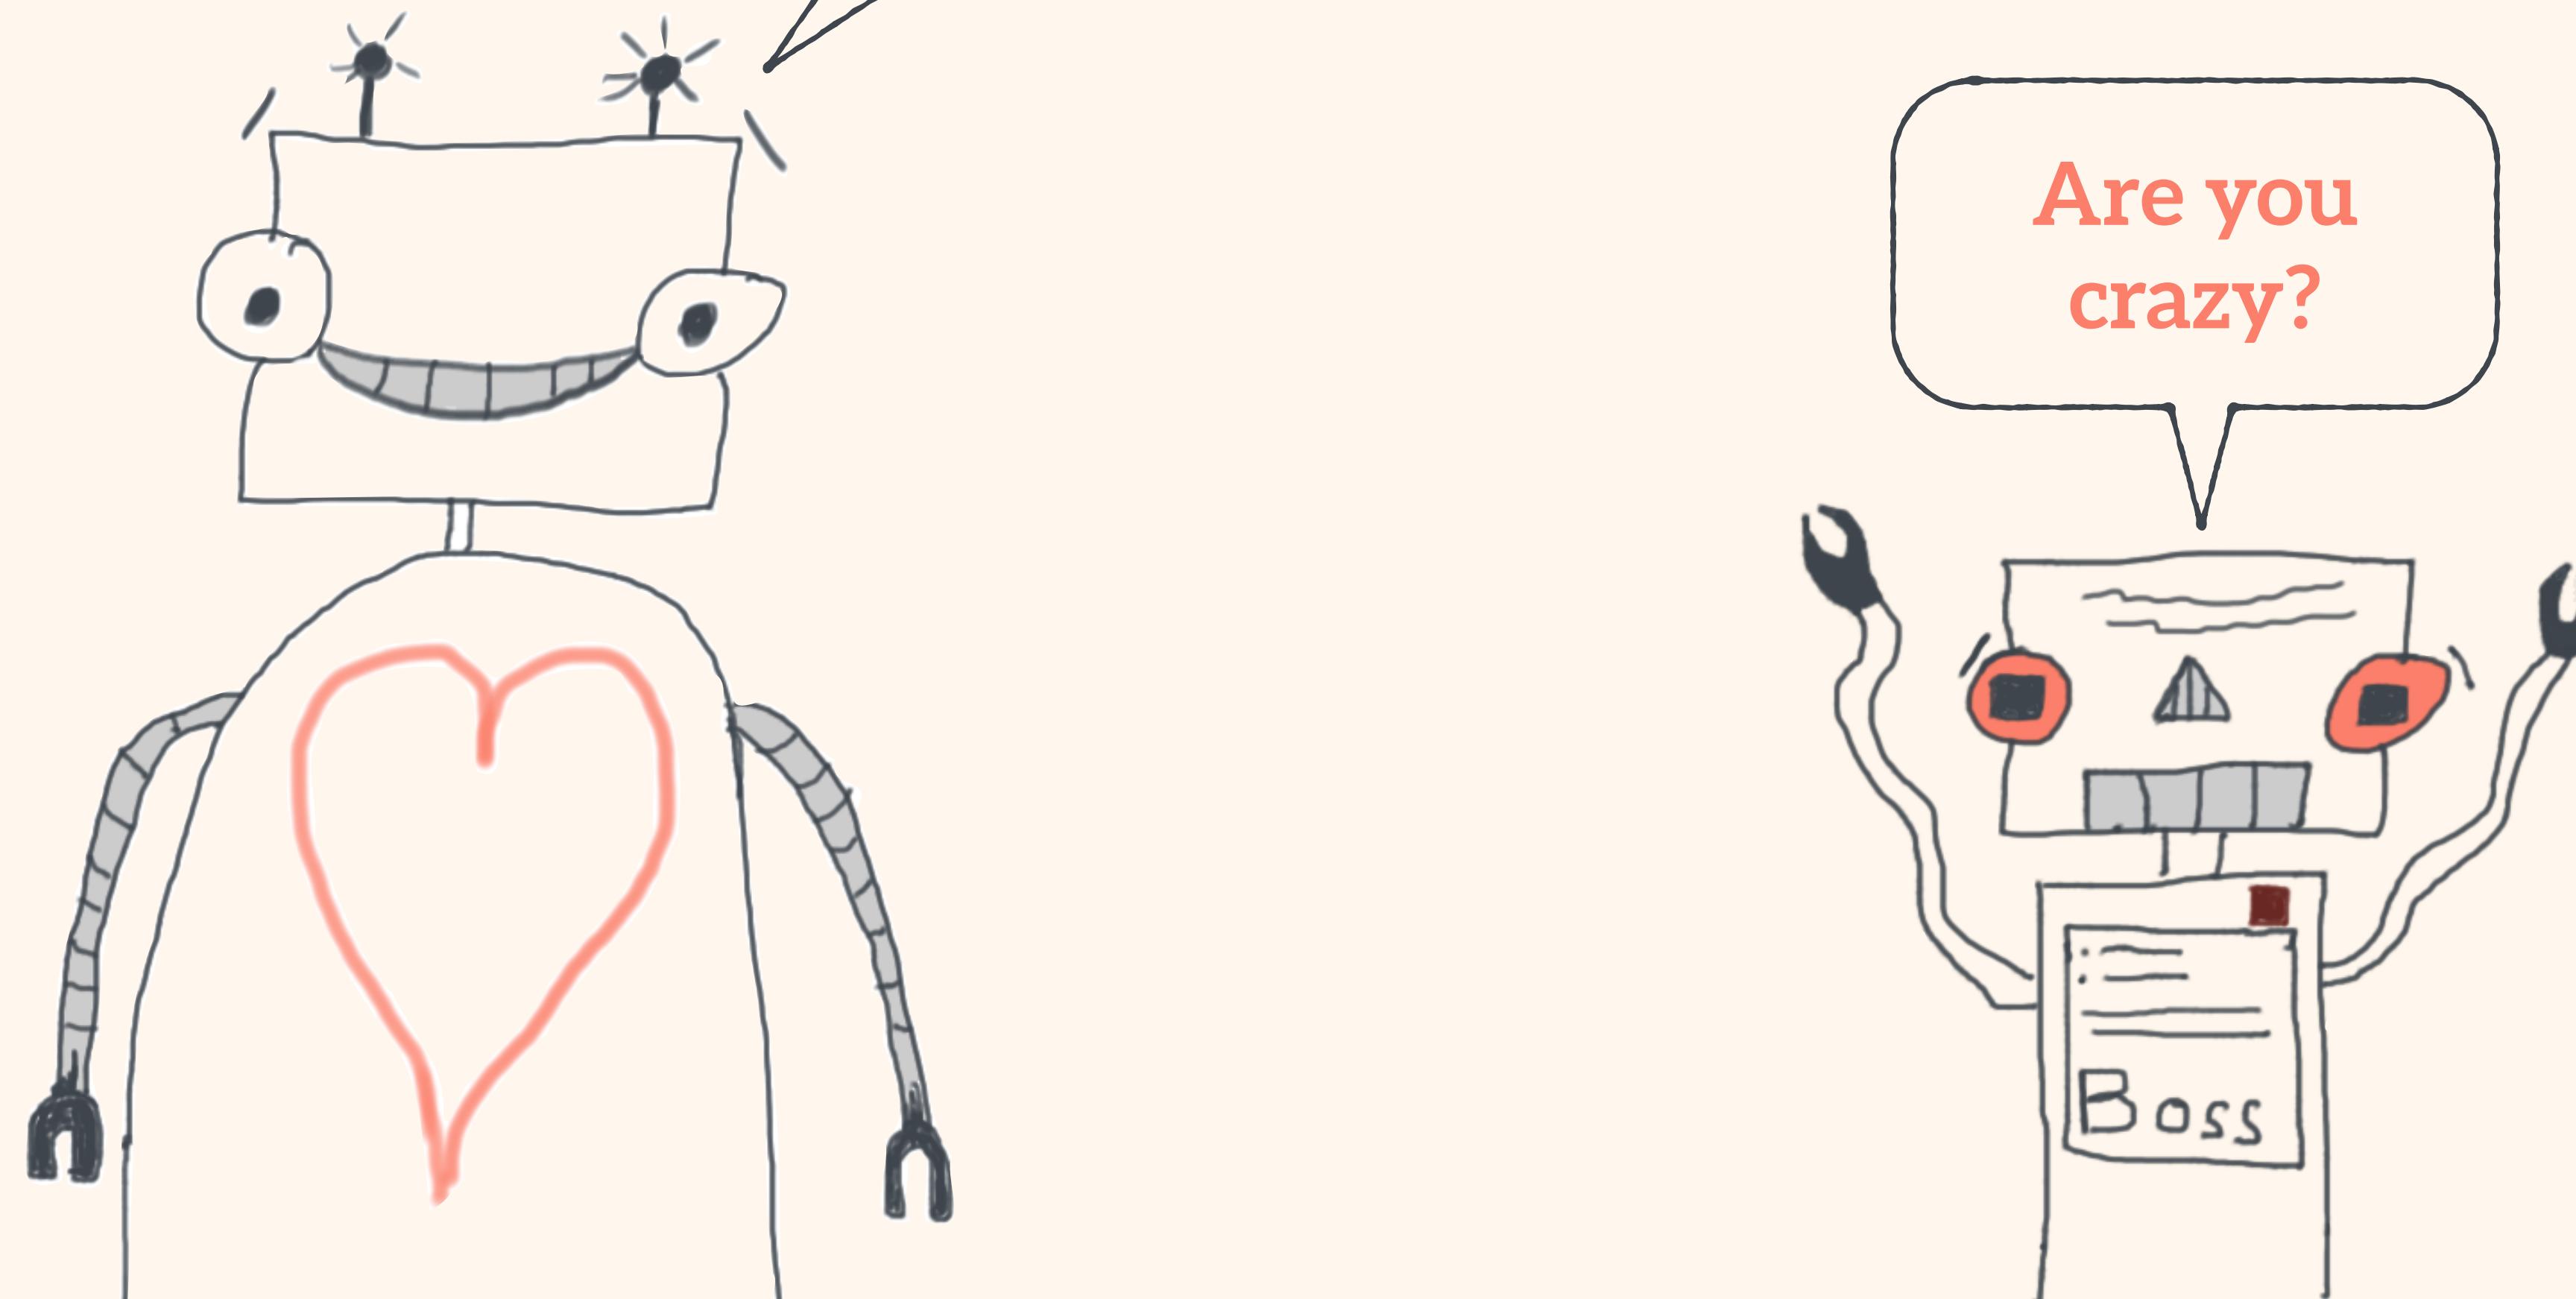
Latency / Synchronization discussion
What is the best caching product on the market
HTTP / Database Caching
Caching in JPA, Hibernate or other ORMs

Cache

/ kæʃ /

In computing, a cache is a component that **transparently stores data so that future requests for that data can be served faster**. The data that is stored within a cache might be values that have been computed earlier or duplicates of original values that are stored elsewhere. If requested data is contained in the cache (**cache hit**), this request can be served by simply reading the cache, which is comparatively faster. Otherwise (**cache miss**), the data has to be recomputed or fetched from its original storage location, which is comparatively slower. **Hence, the greater the number of requests that can be served from the cache, the faster the overall system performance becomes.**

That's awesome. Let's cache everything
and everywhere and distribute it all in
a Cluster in a transactional manner
ohhh by the way: Twitter has been
doing that for ages



Business-Applications



Twitter / Facebook & co.

*Many enterprise grade projects
are adapting caching too
defensive or too offensive and are
running into consistency or
performance issues because of
that*

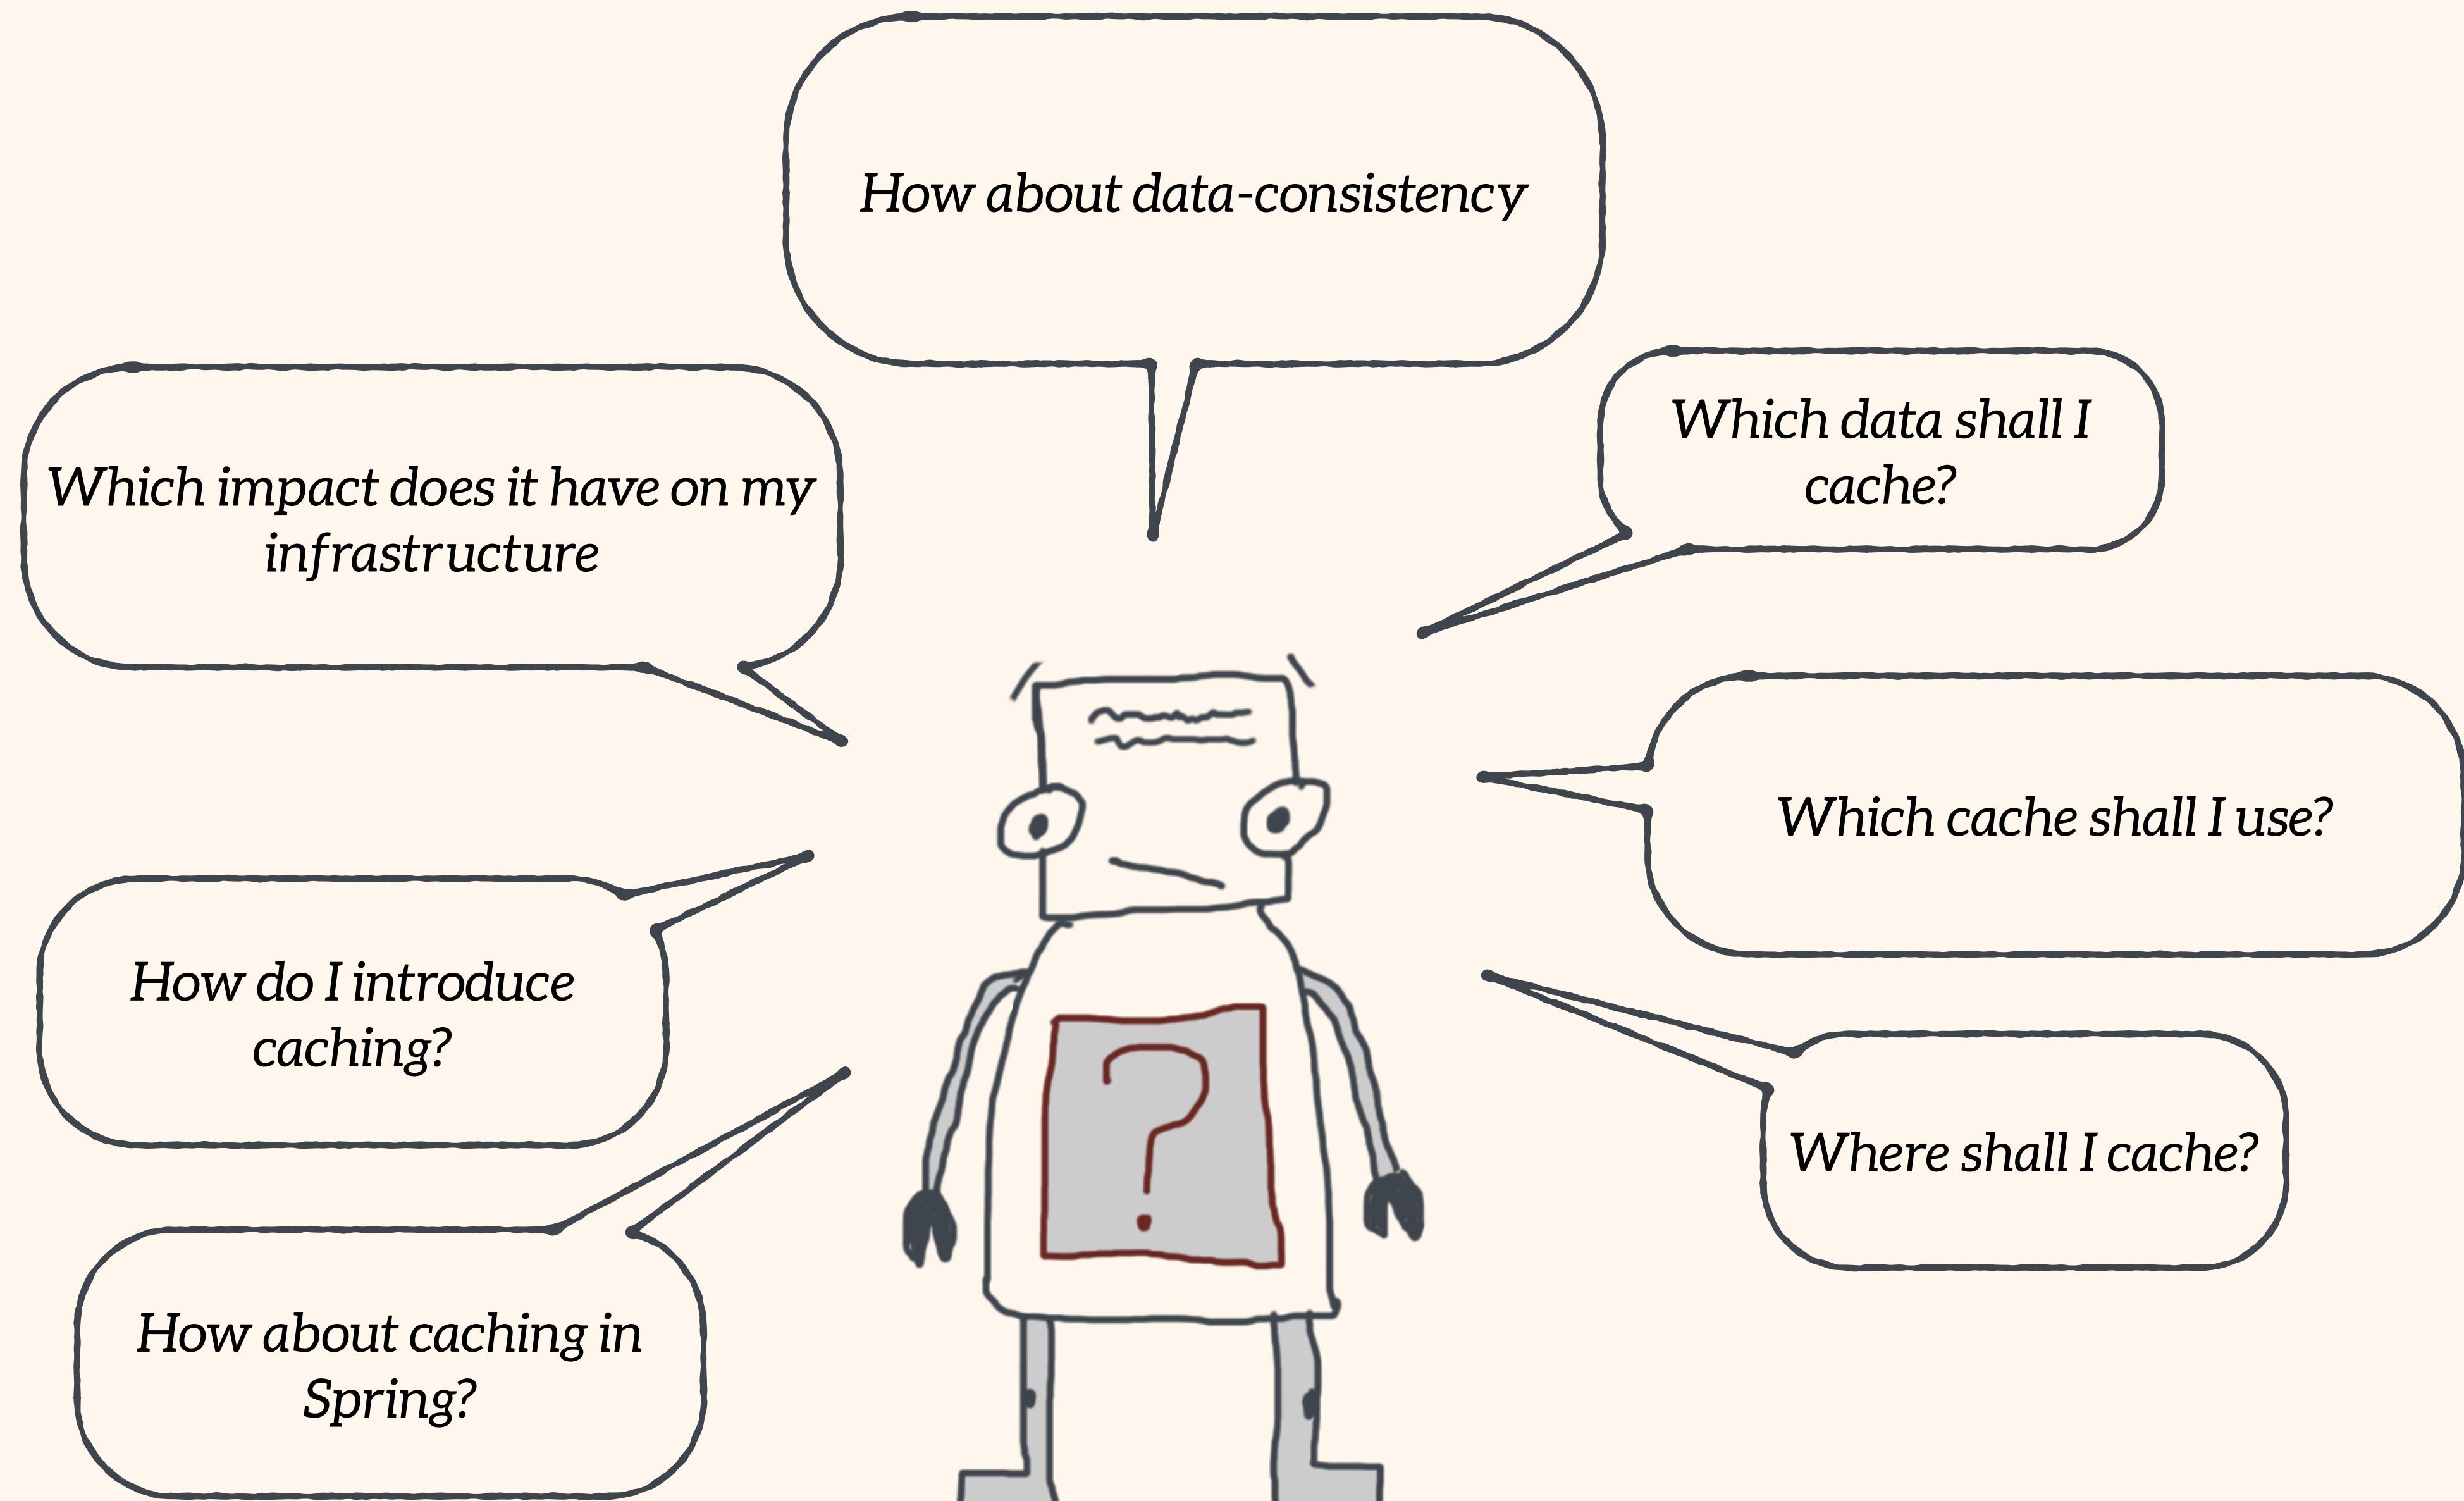
But with a well adjusted caching
strategy you will make your
application more scalable, faster
and cheaper to operate.

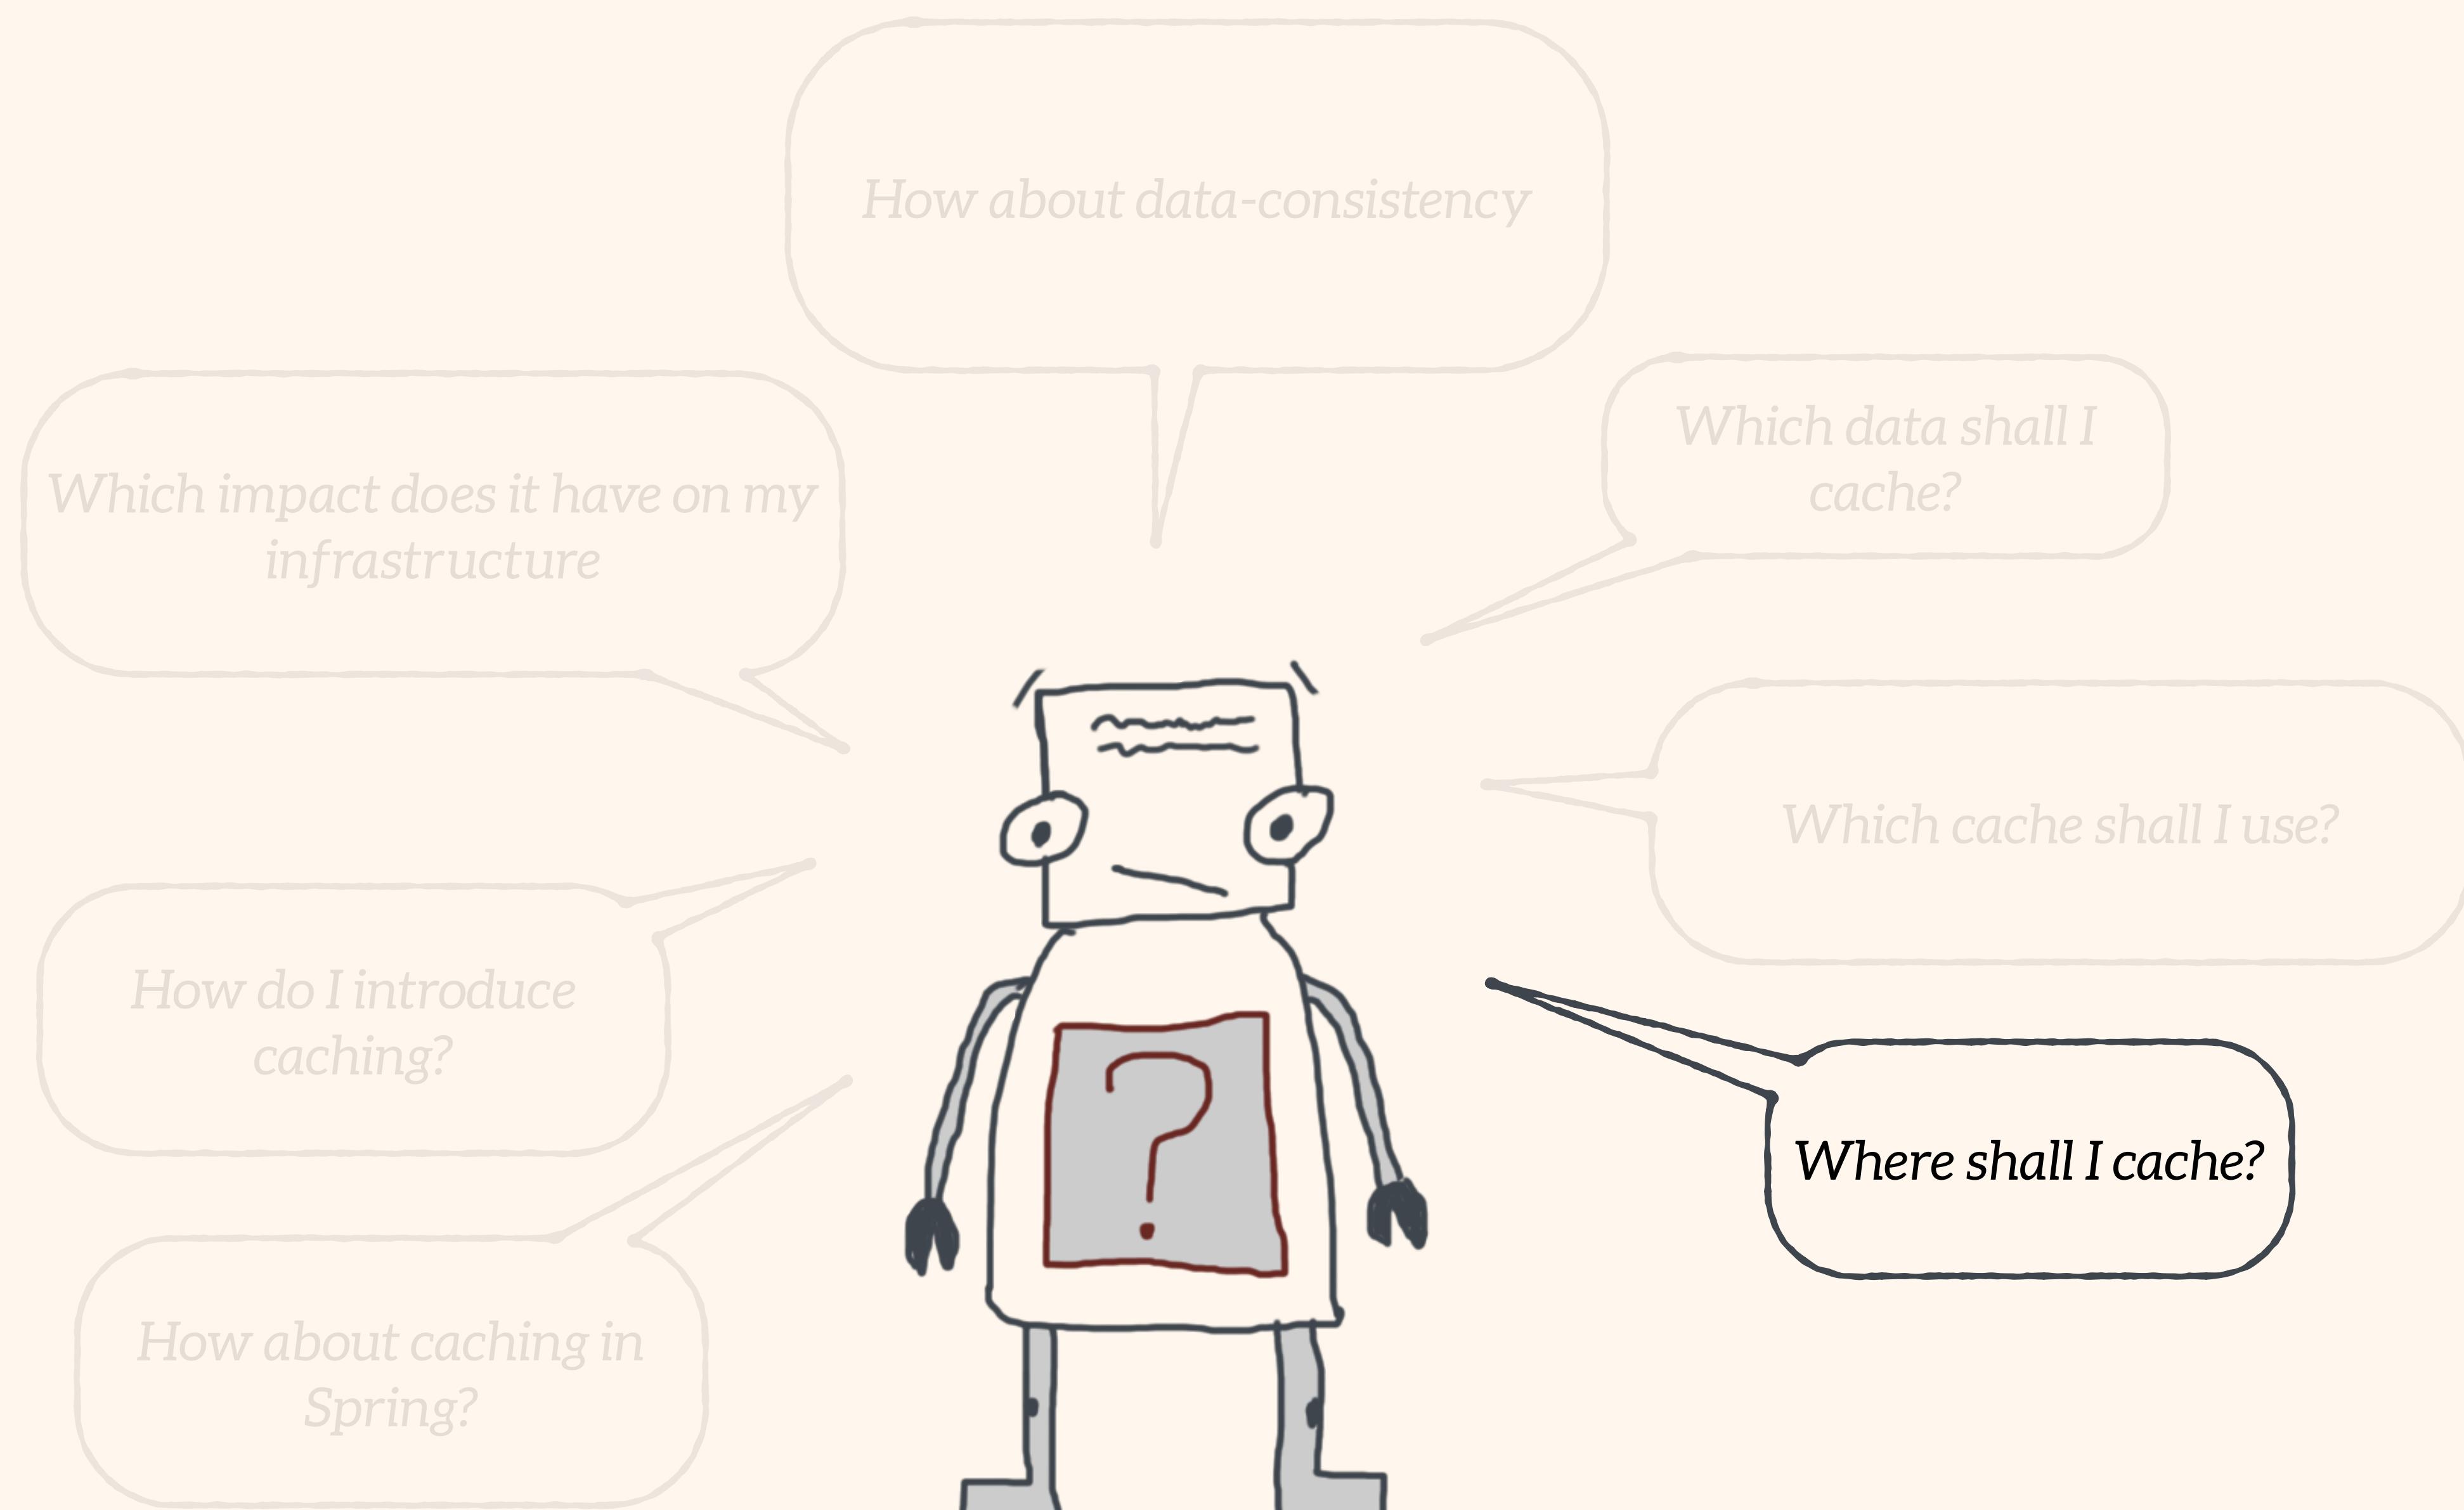
Local Cache, Data Grid, Document Store, JPA
First Level Cache, JPA Second Level Cache,
Hybrid Cache

Types Places of CACHES for

Database, Heap, HTTP Proxy, Browser,
Processor, Disk, Off Heap, Persistence-
Framework, Application

*We will focus on local and
distributed caching at the
application level with the Spring
Framework*



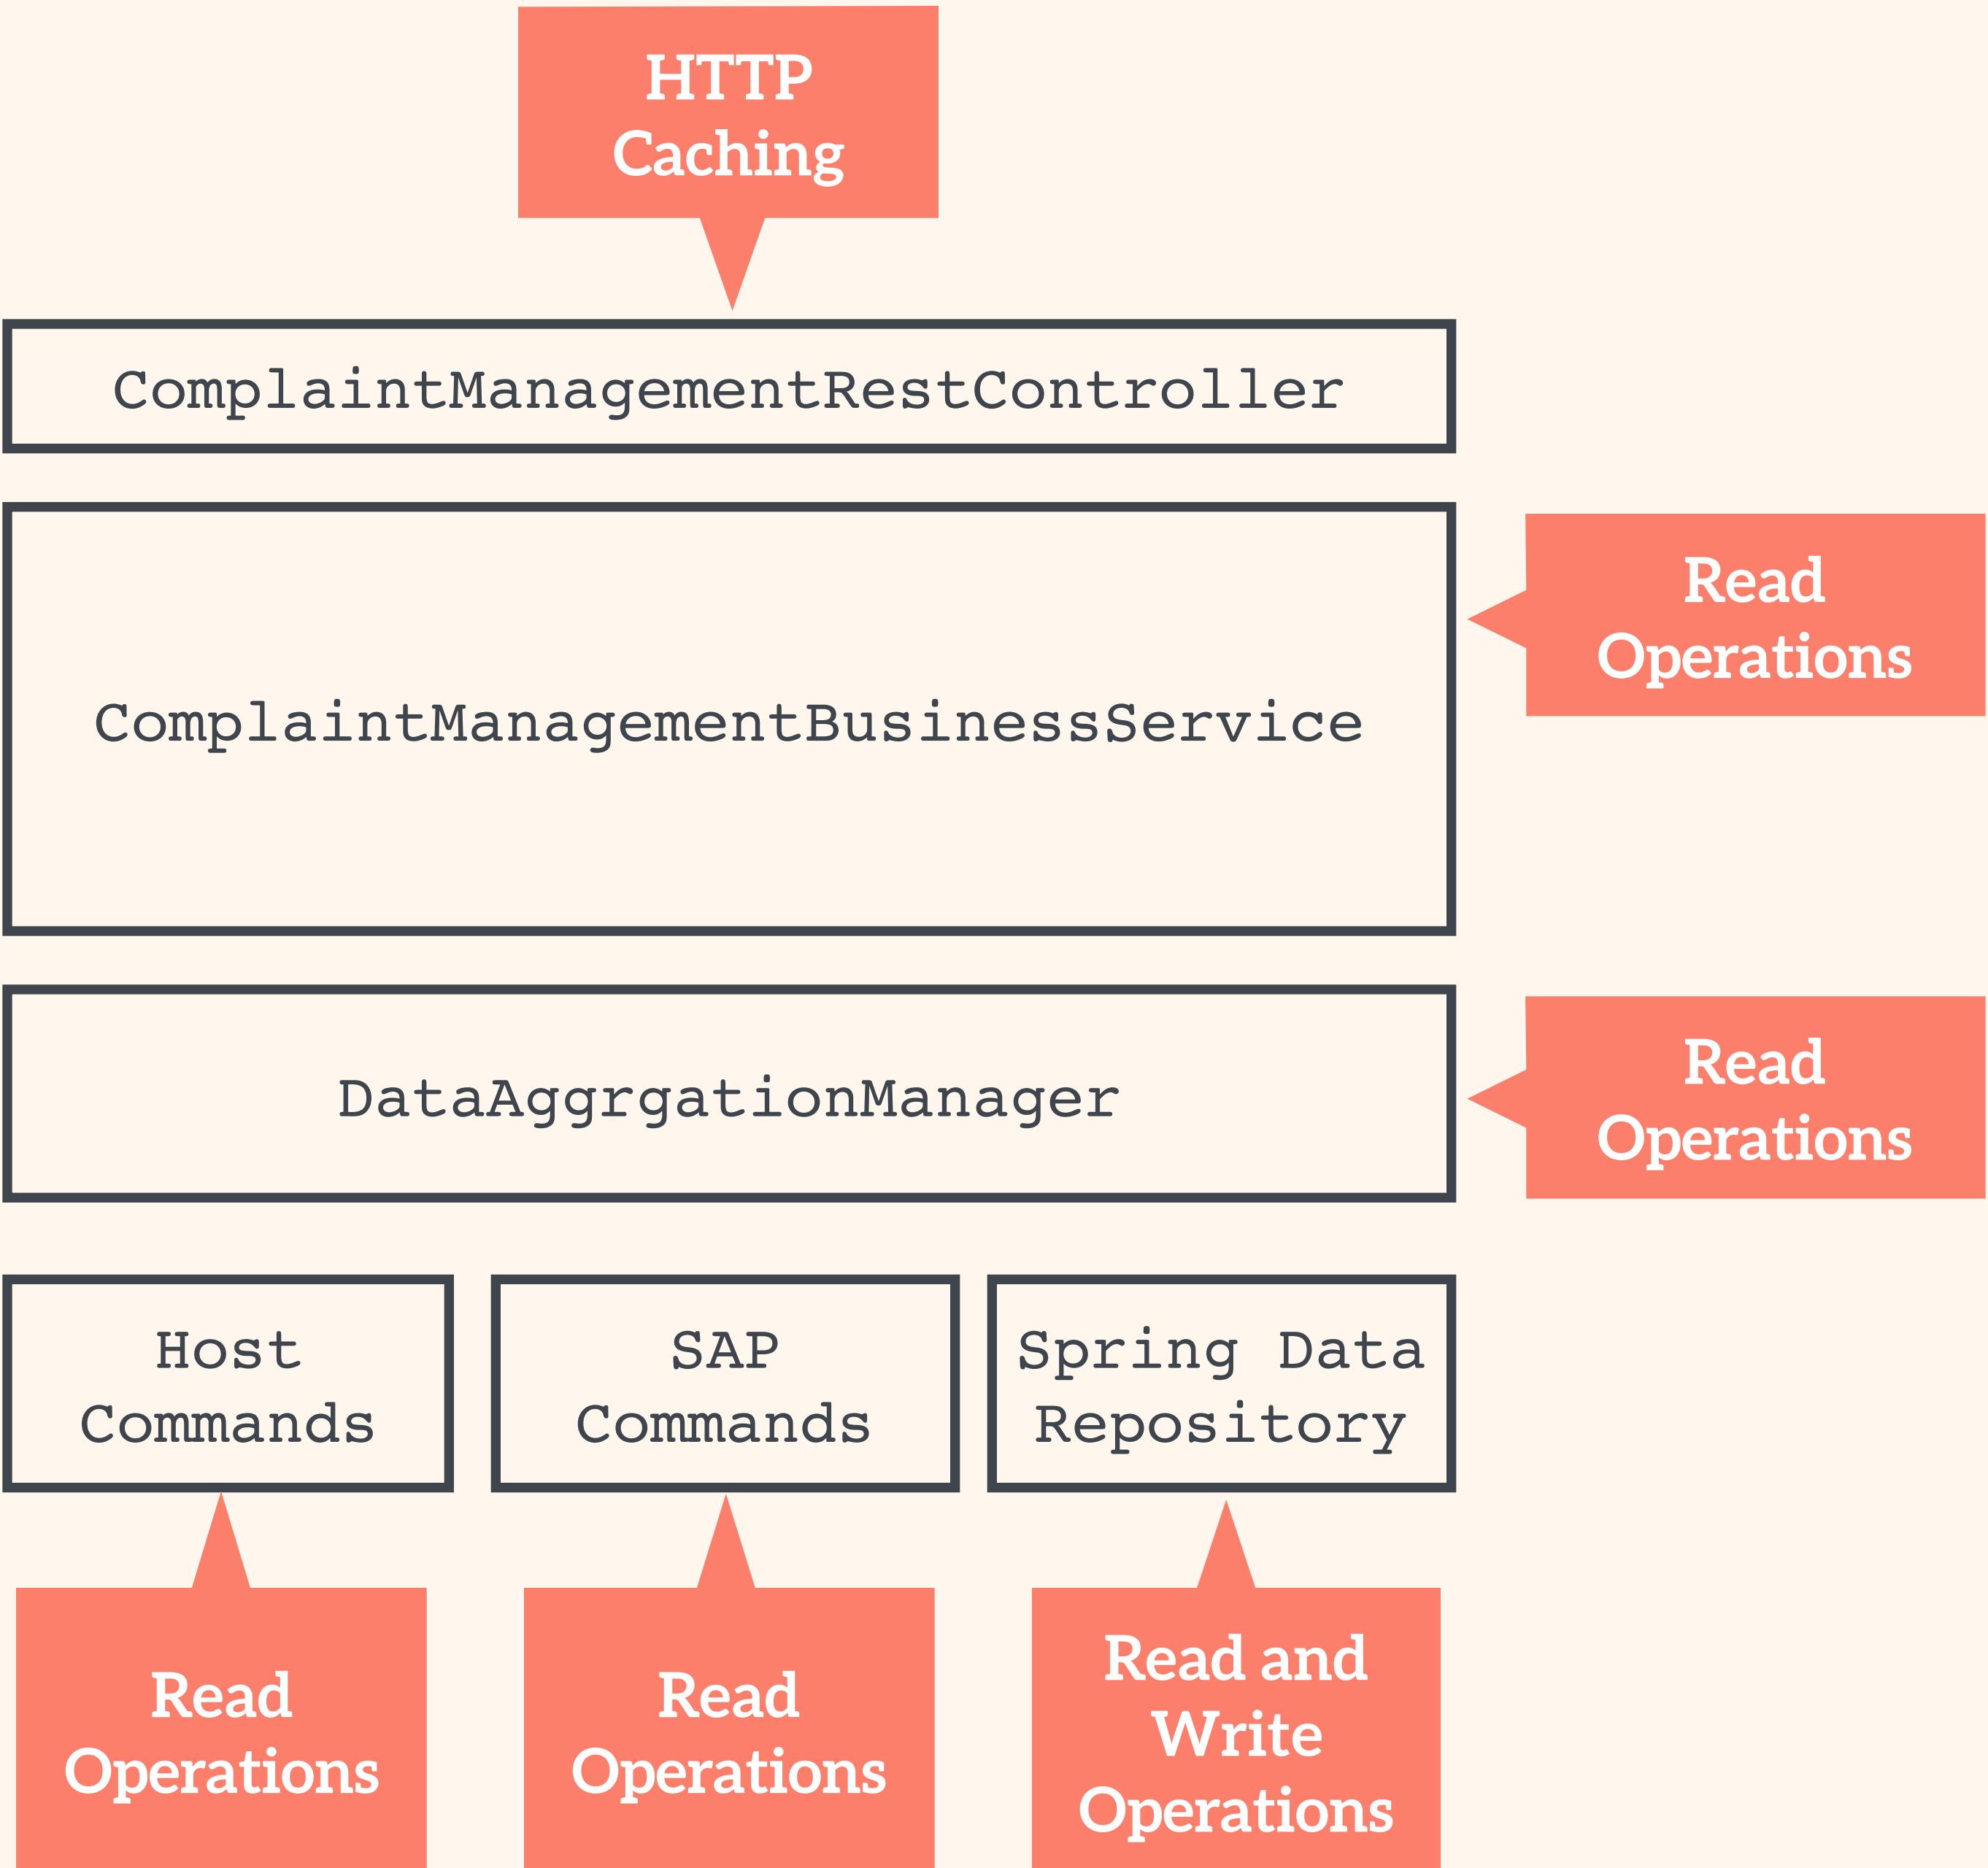




1

*Identify suitable layers for
caching*

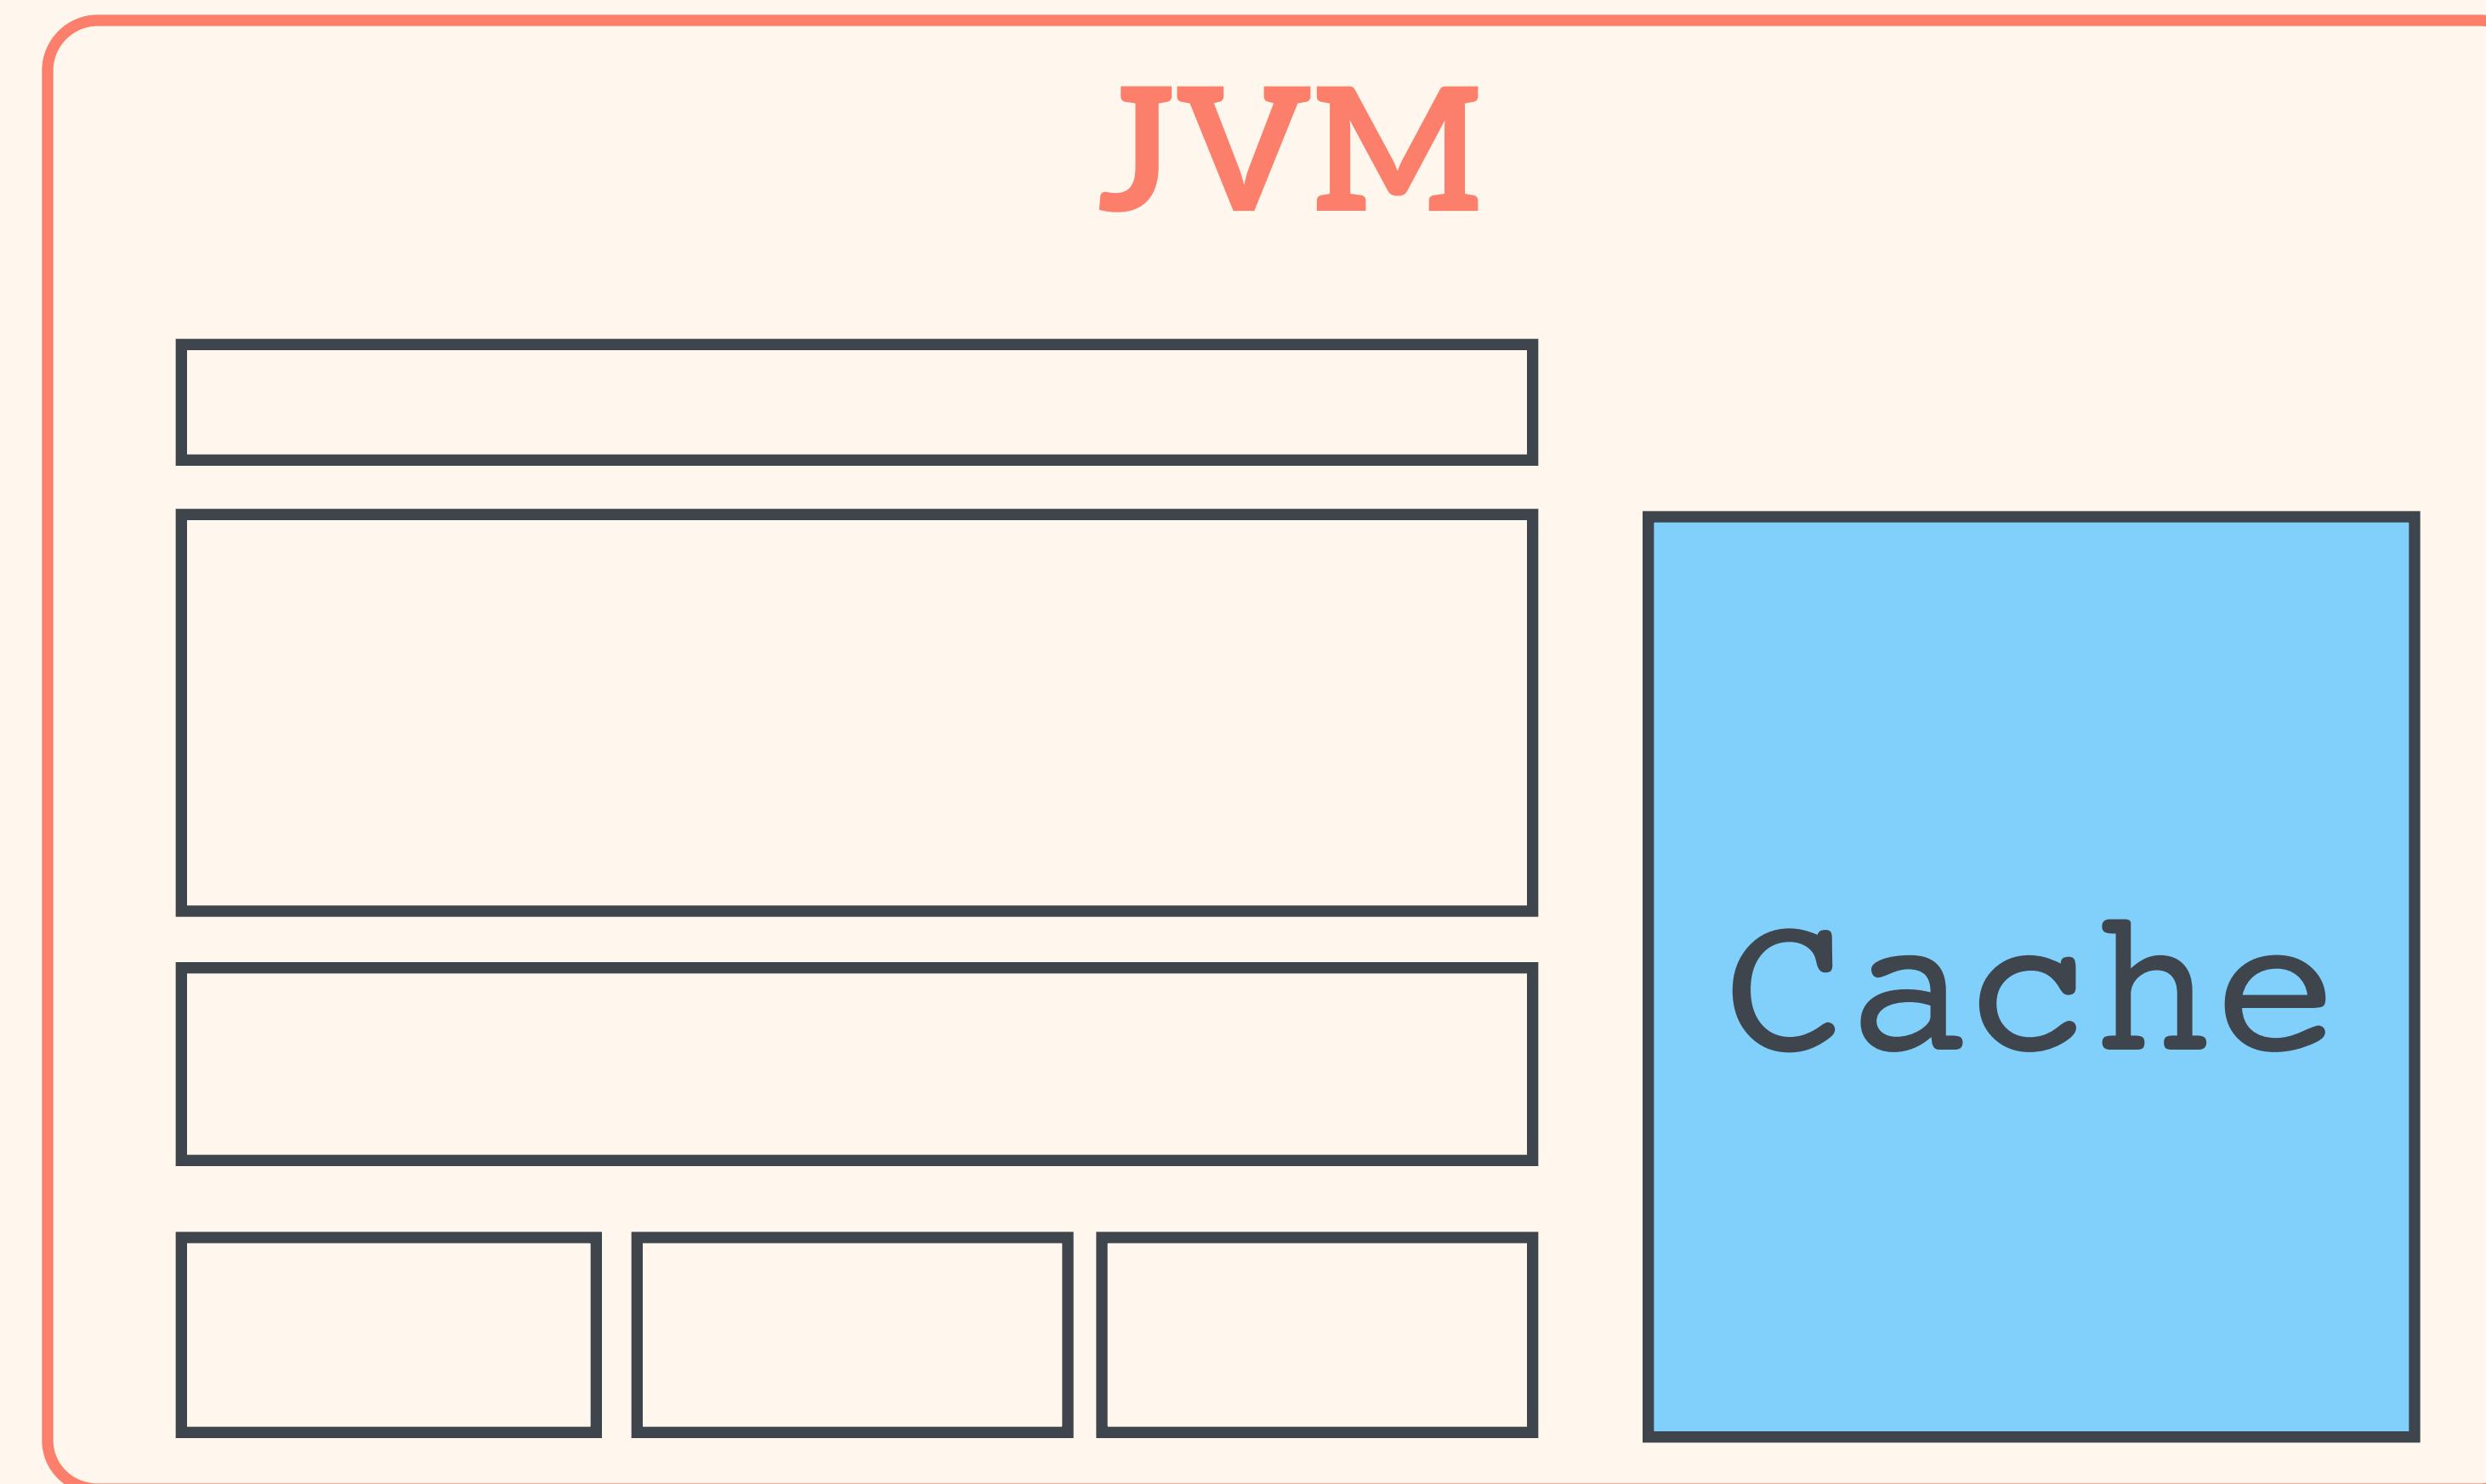
Suitable Layers for Caching



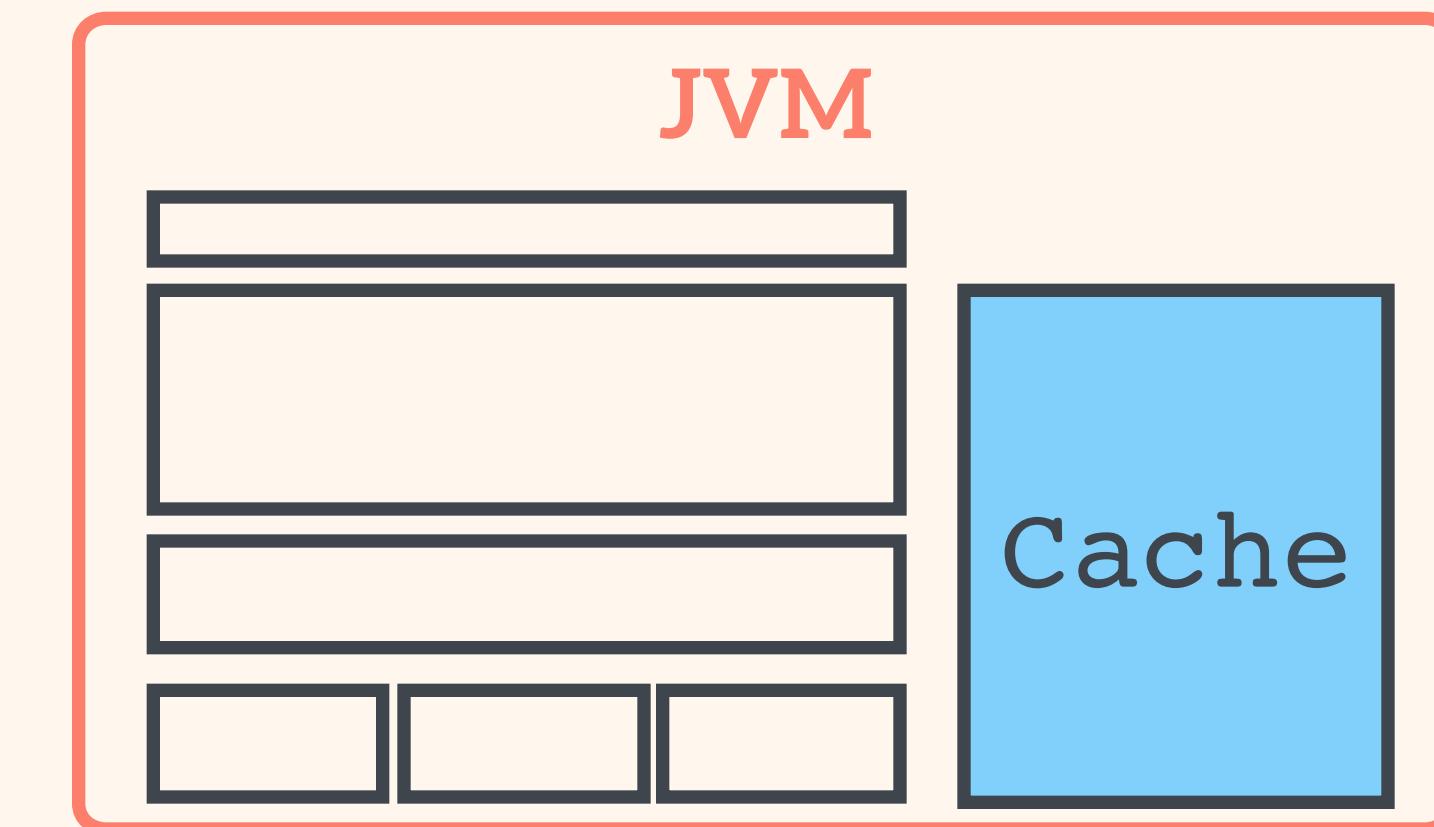
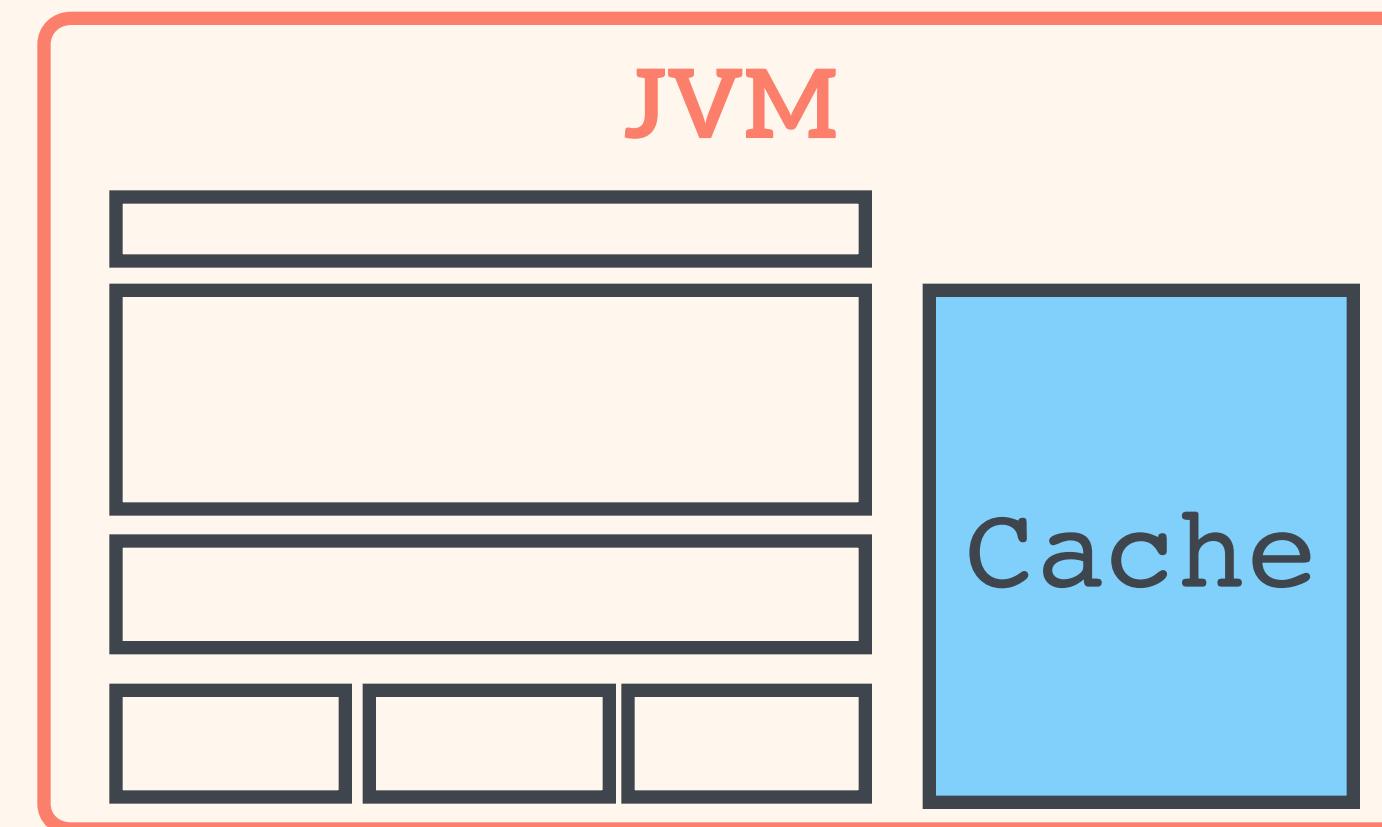
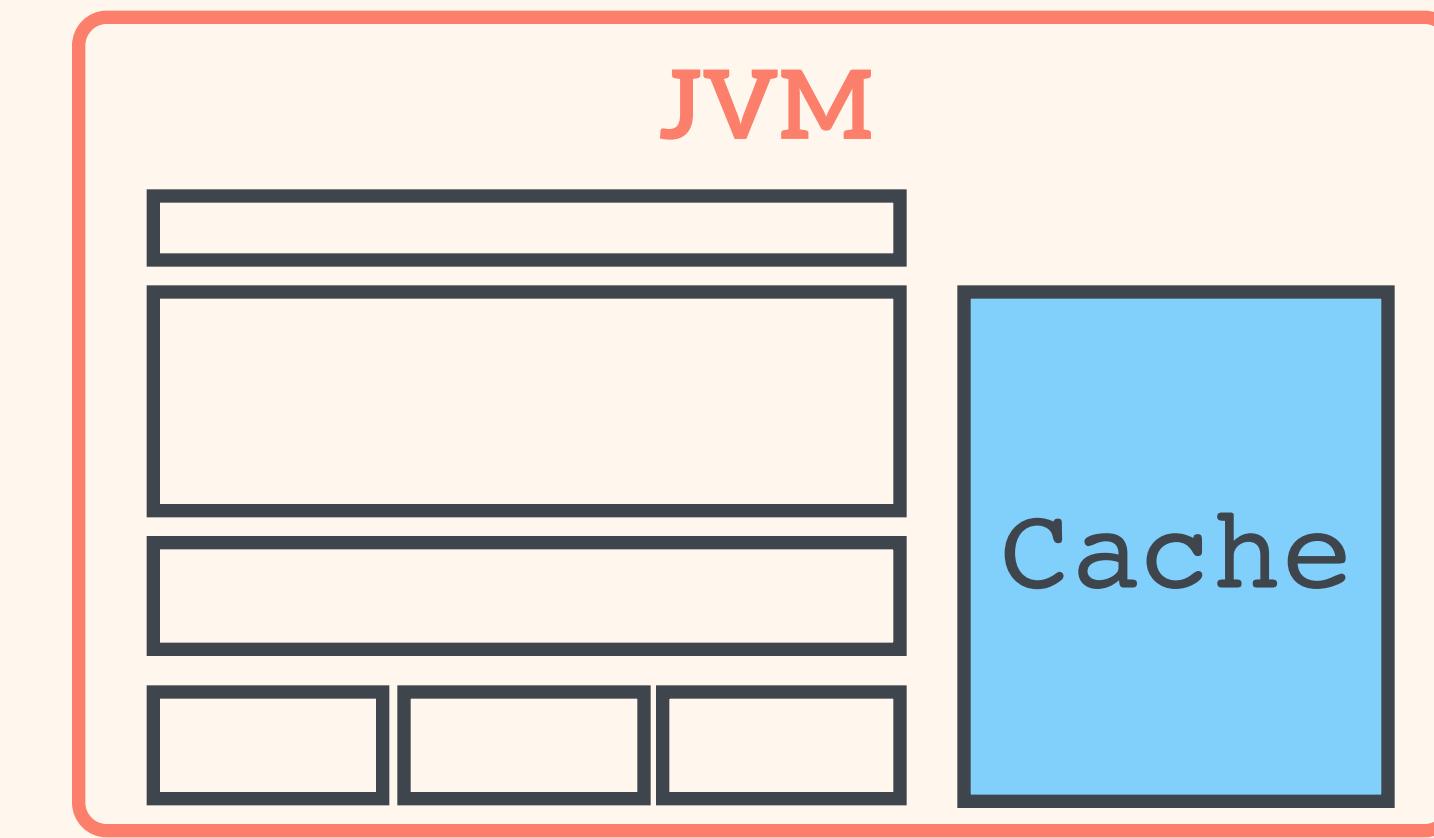
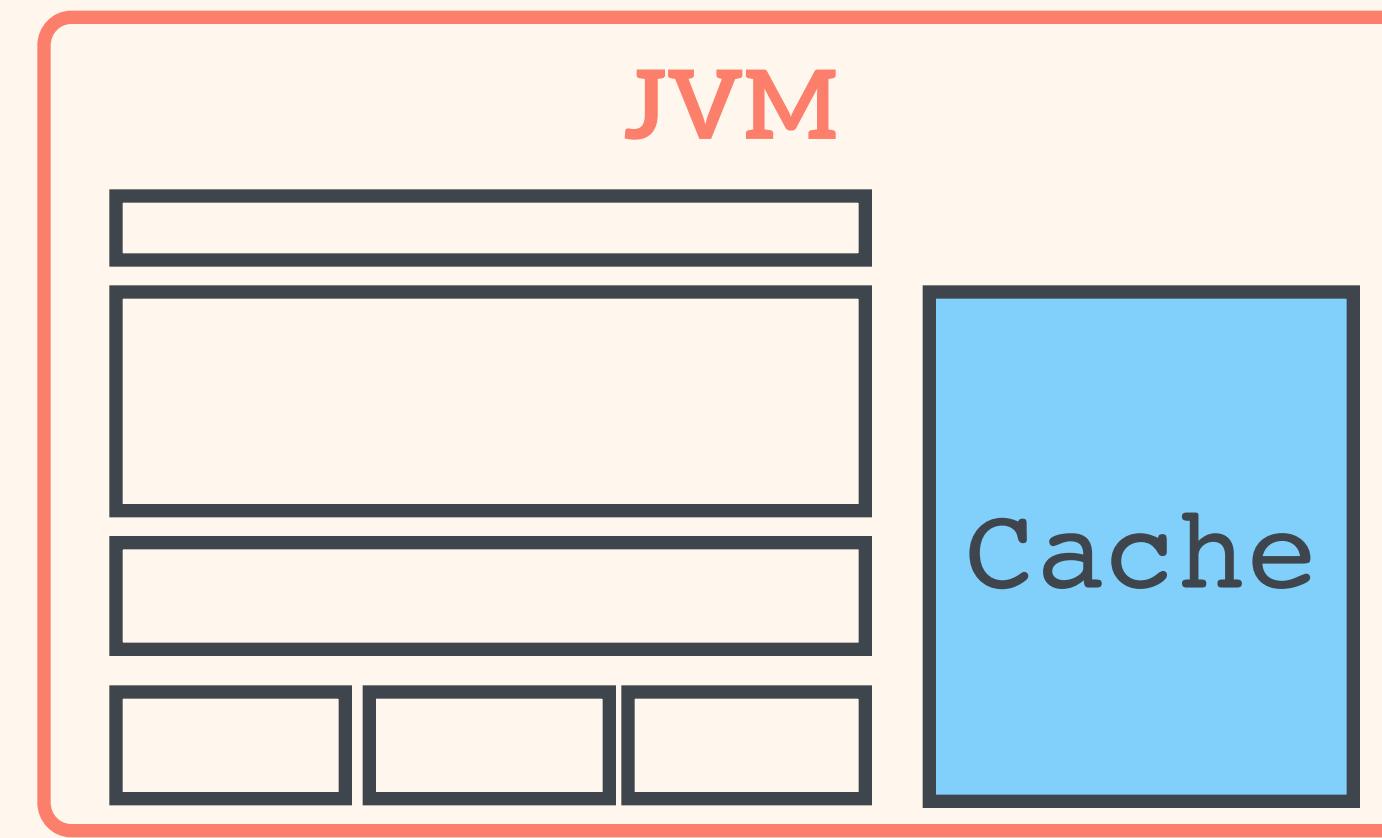
2

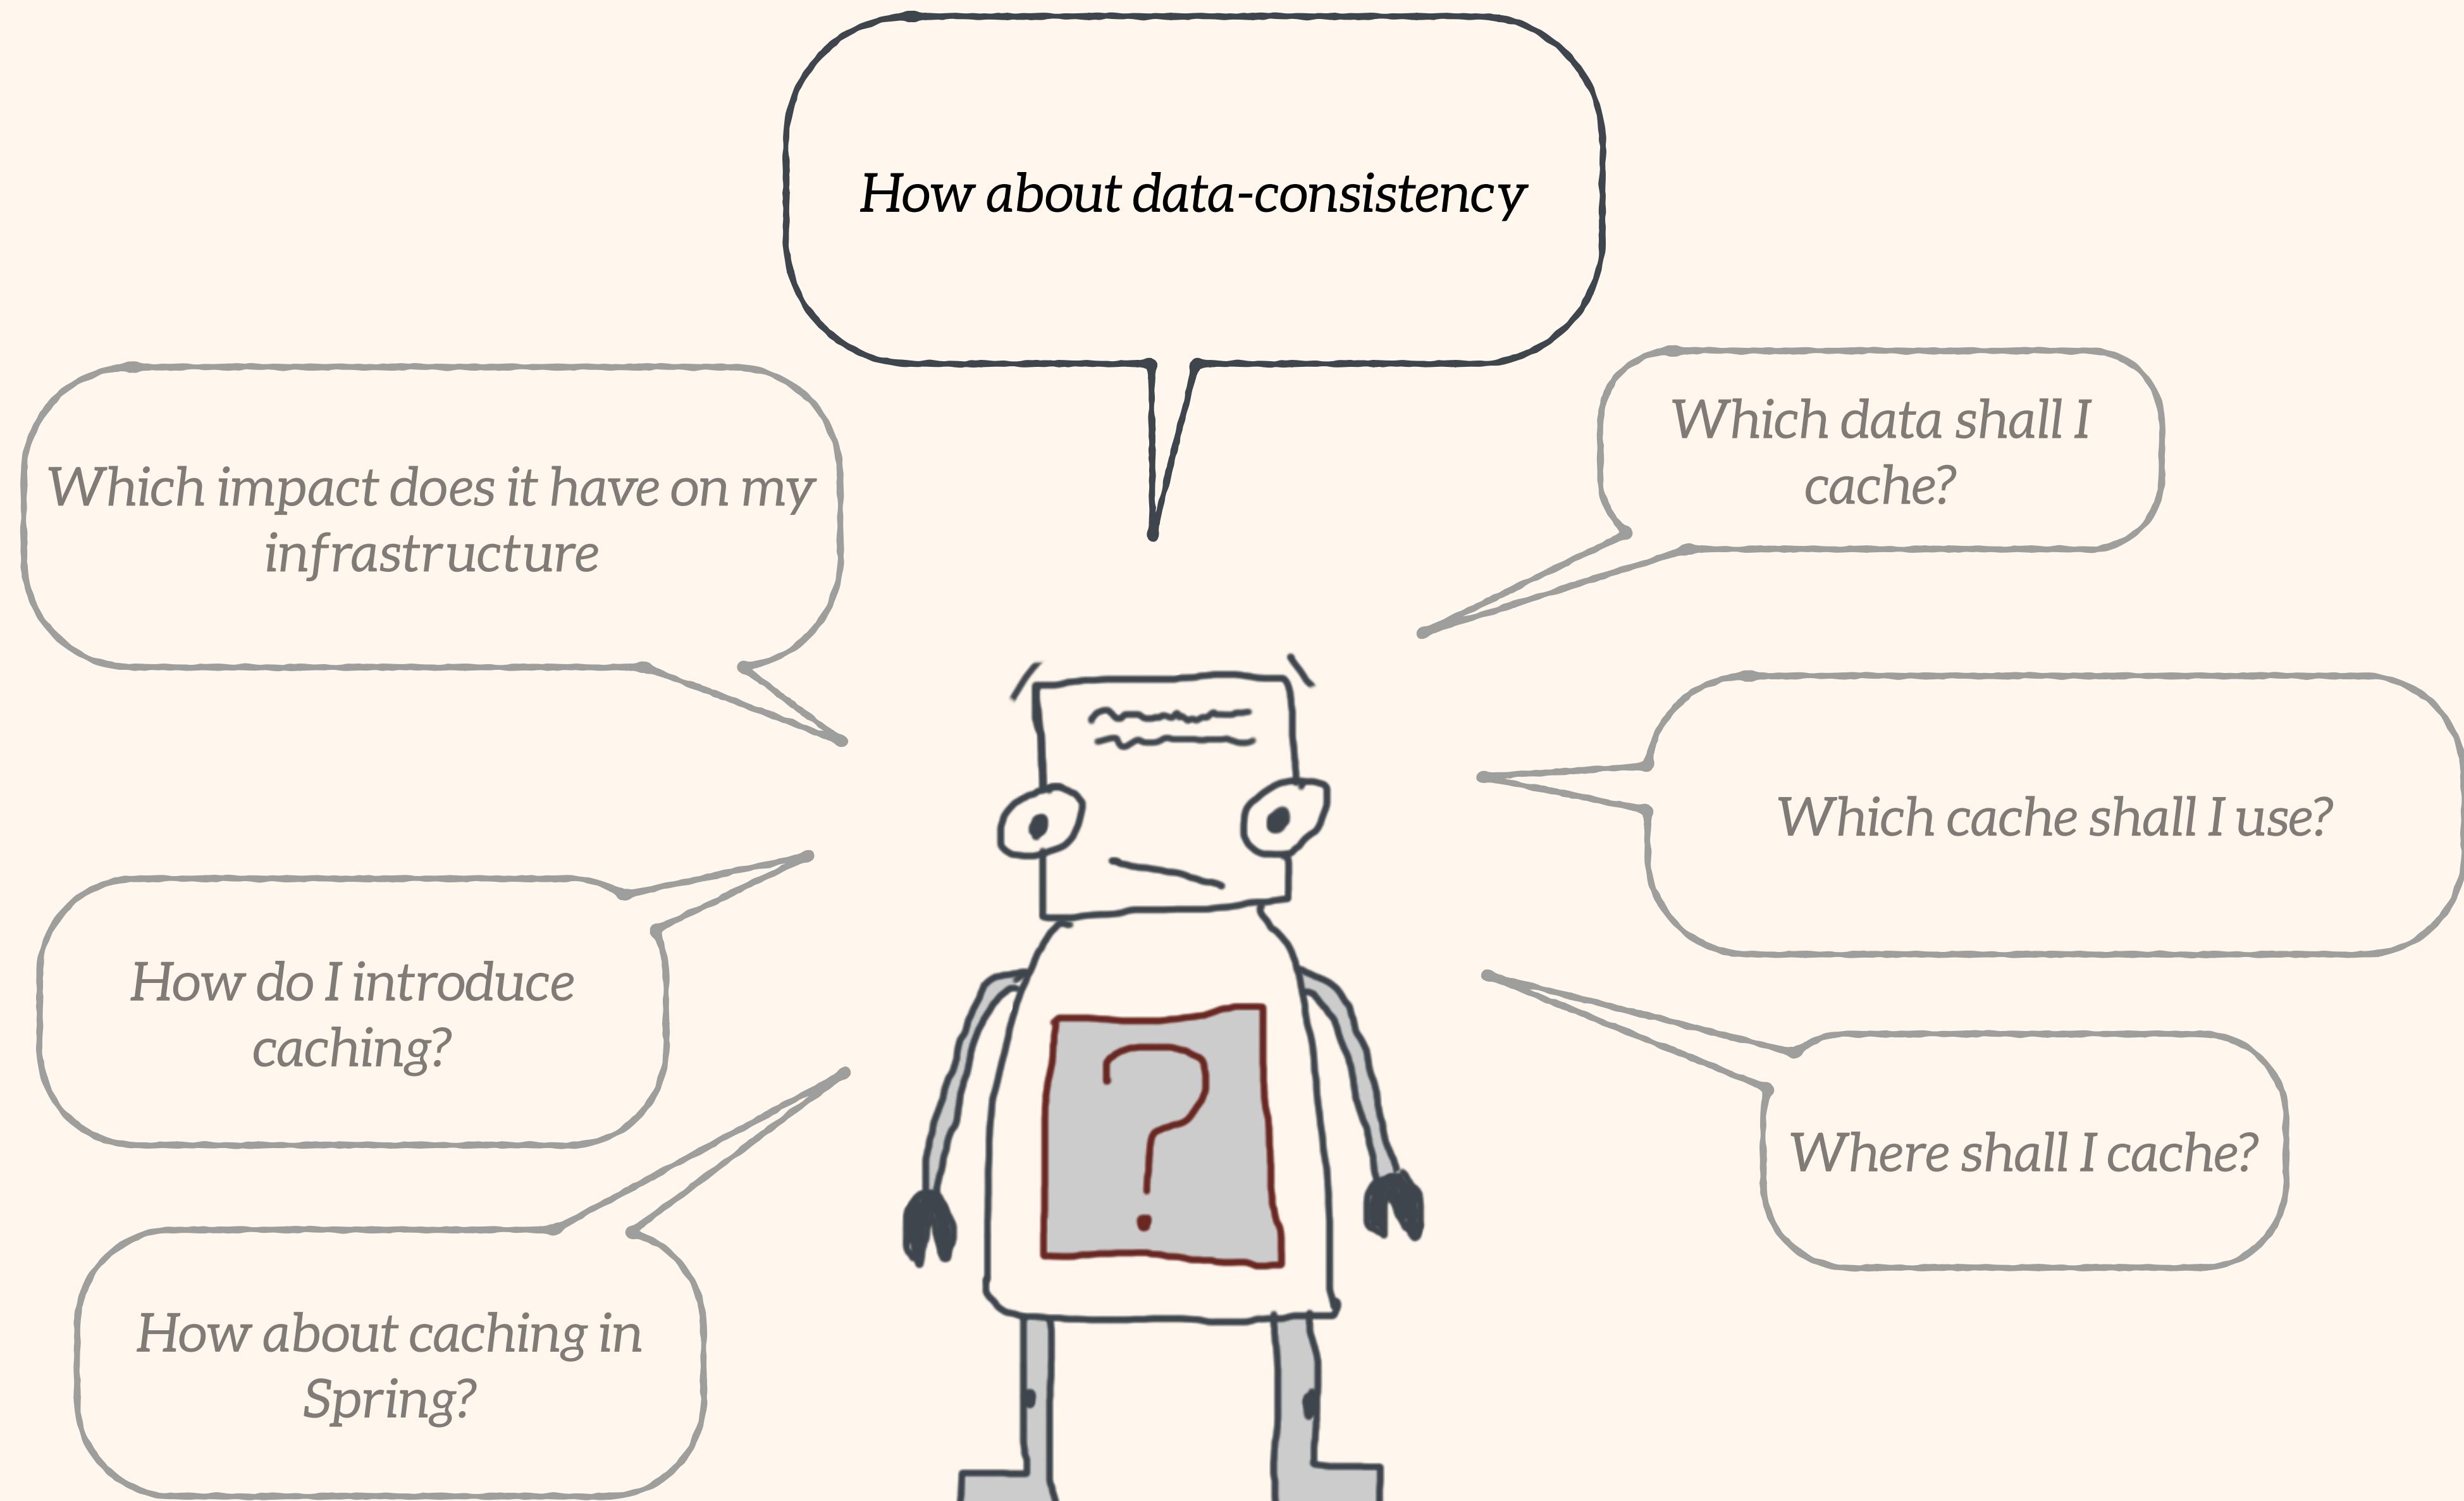
Stay local as long as possible

Lokal In-Memory

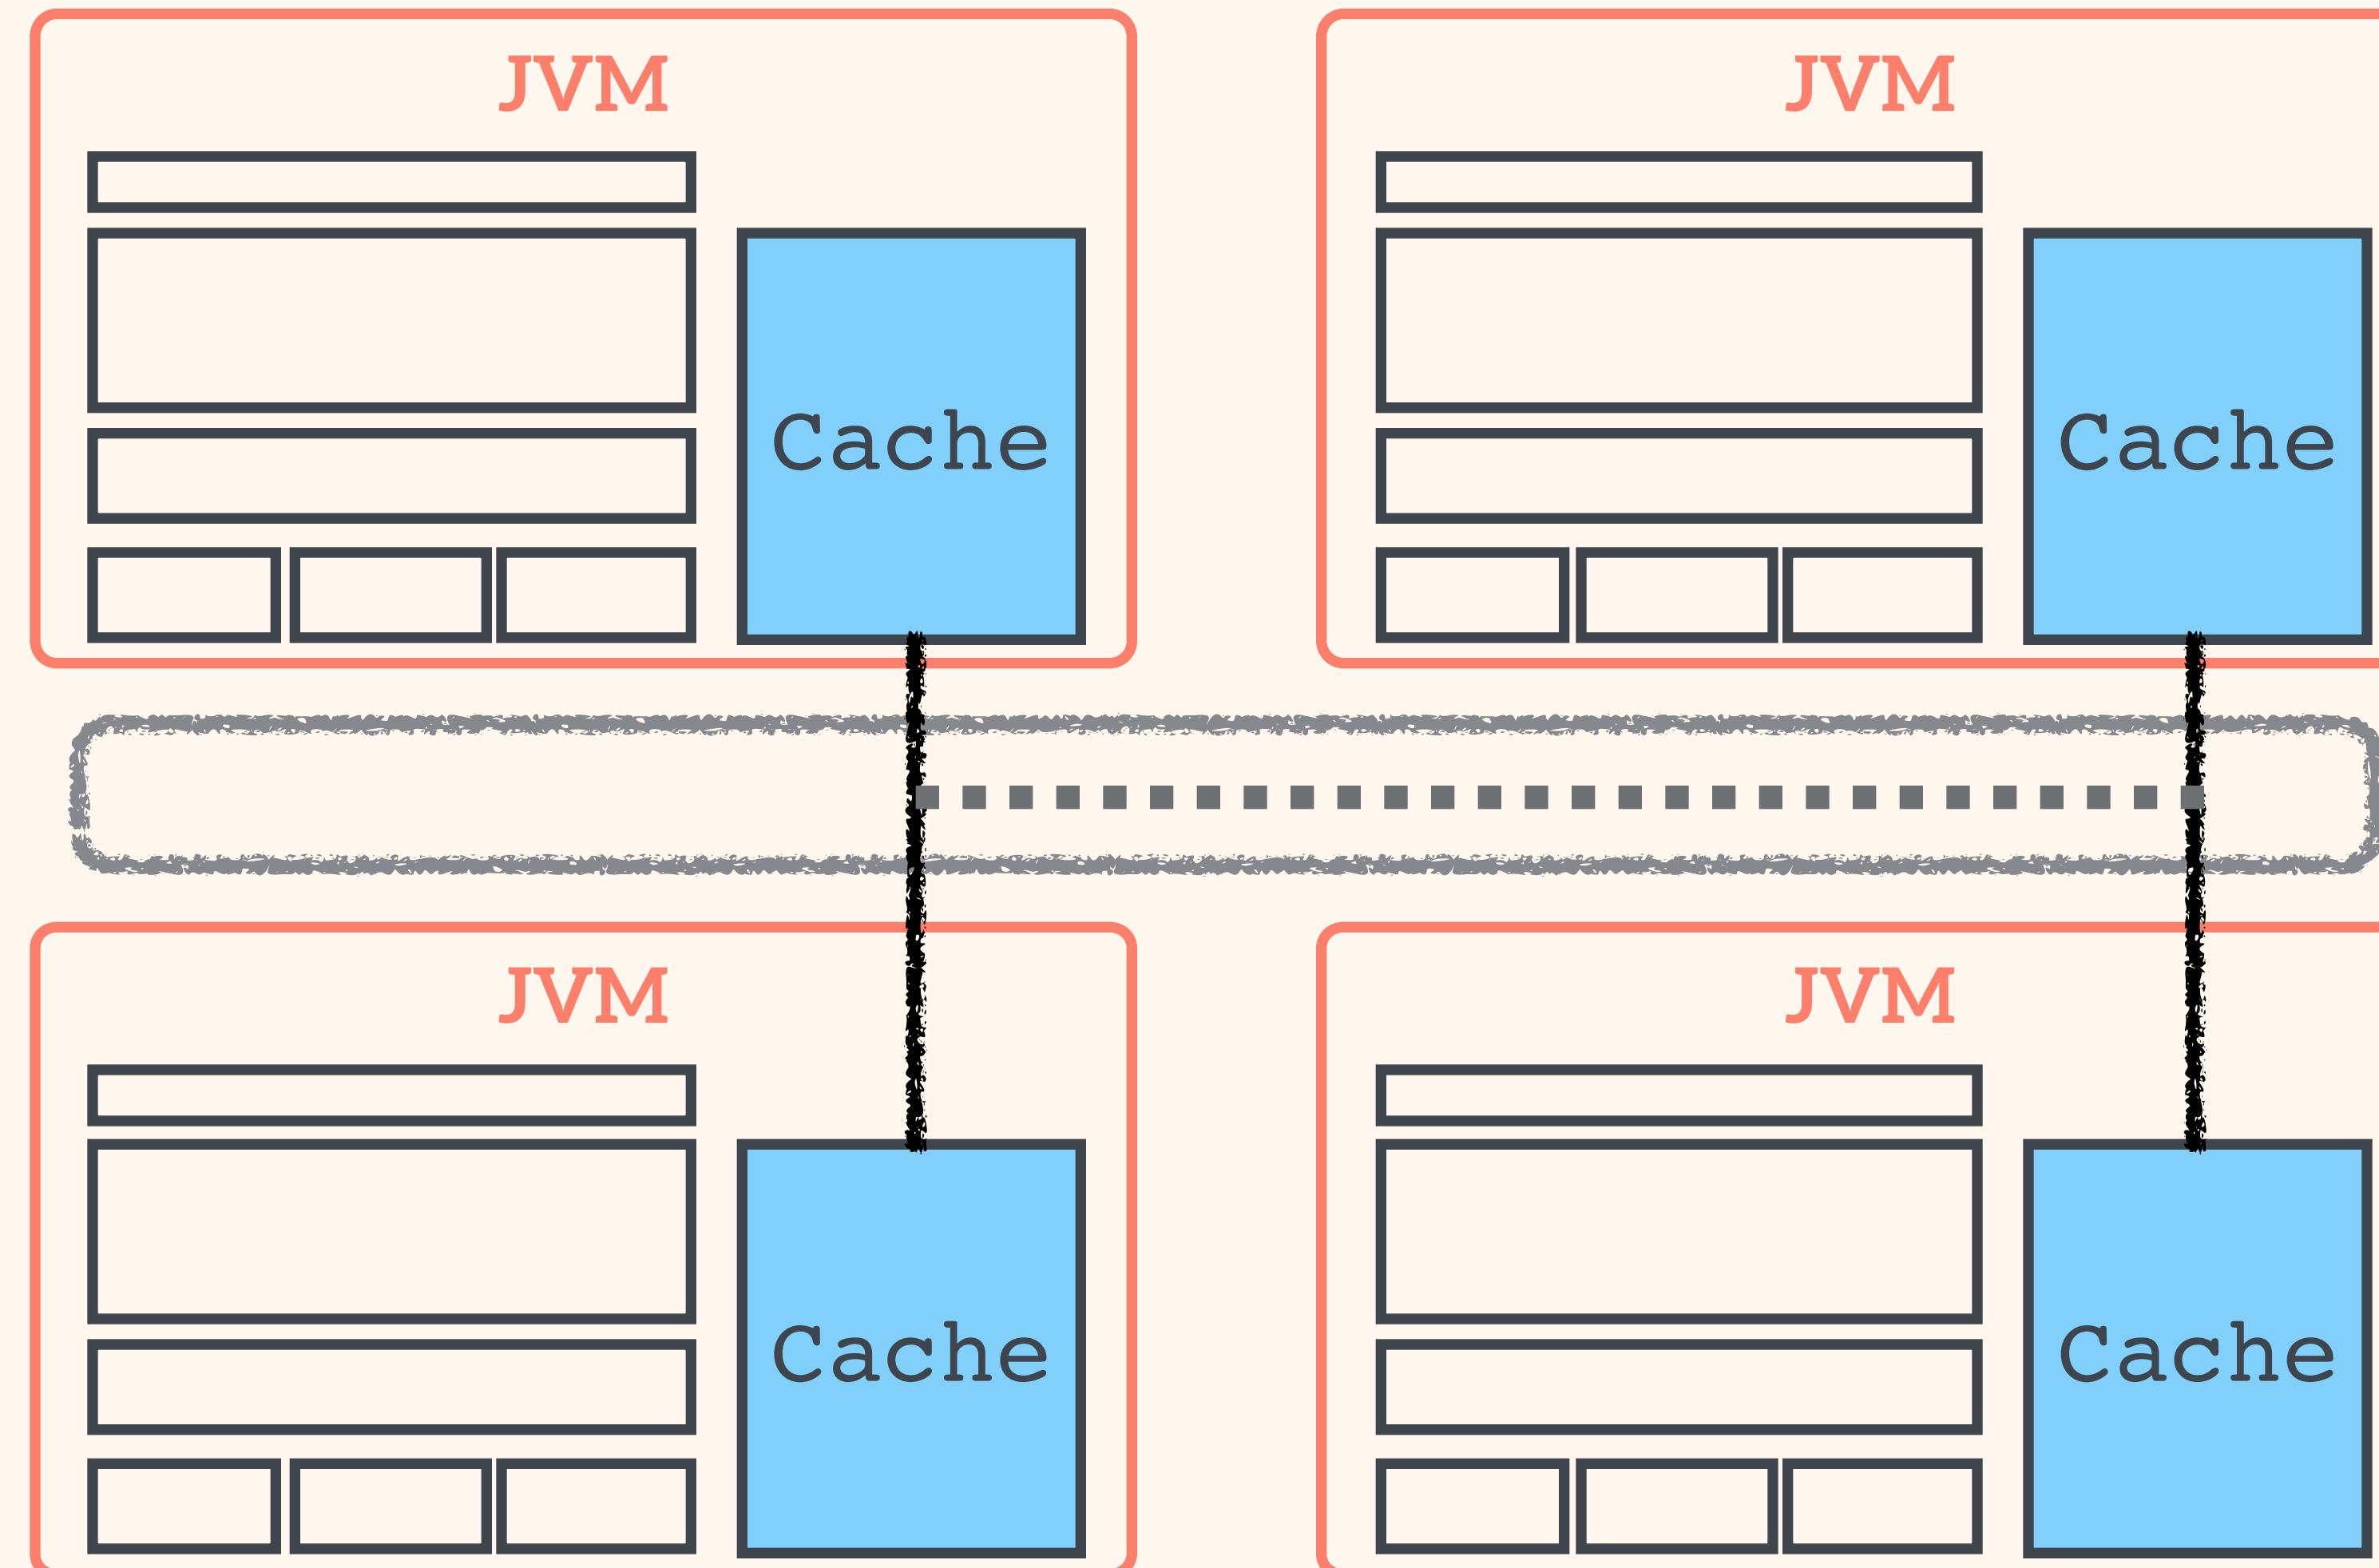


Clustered

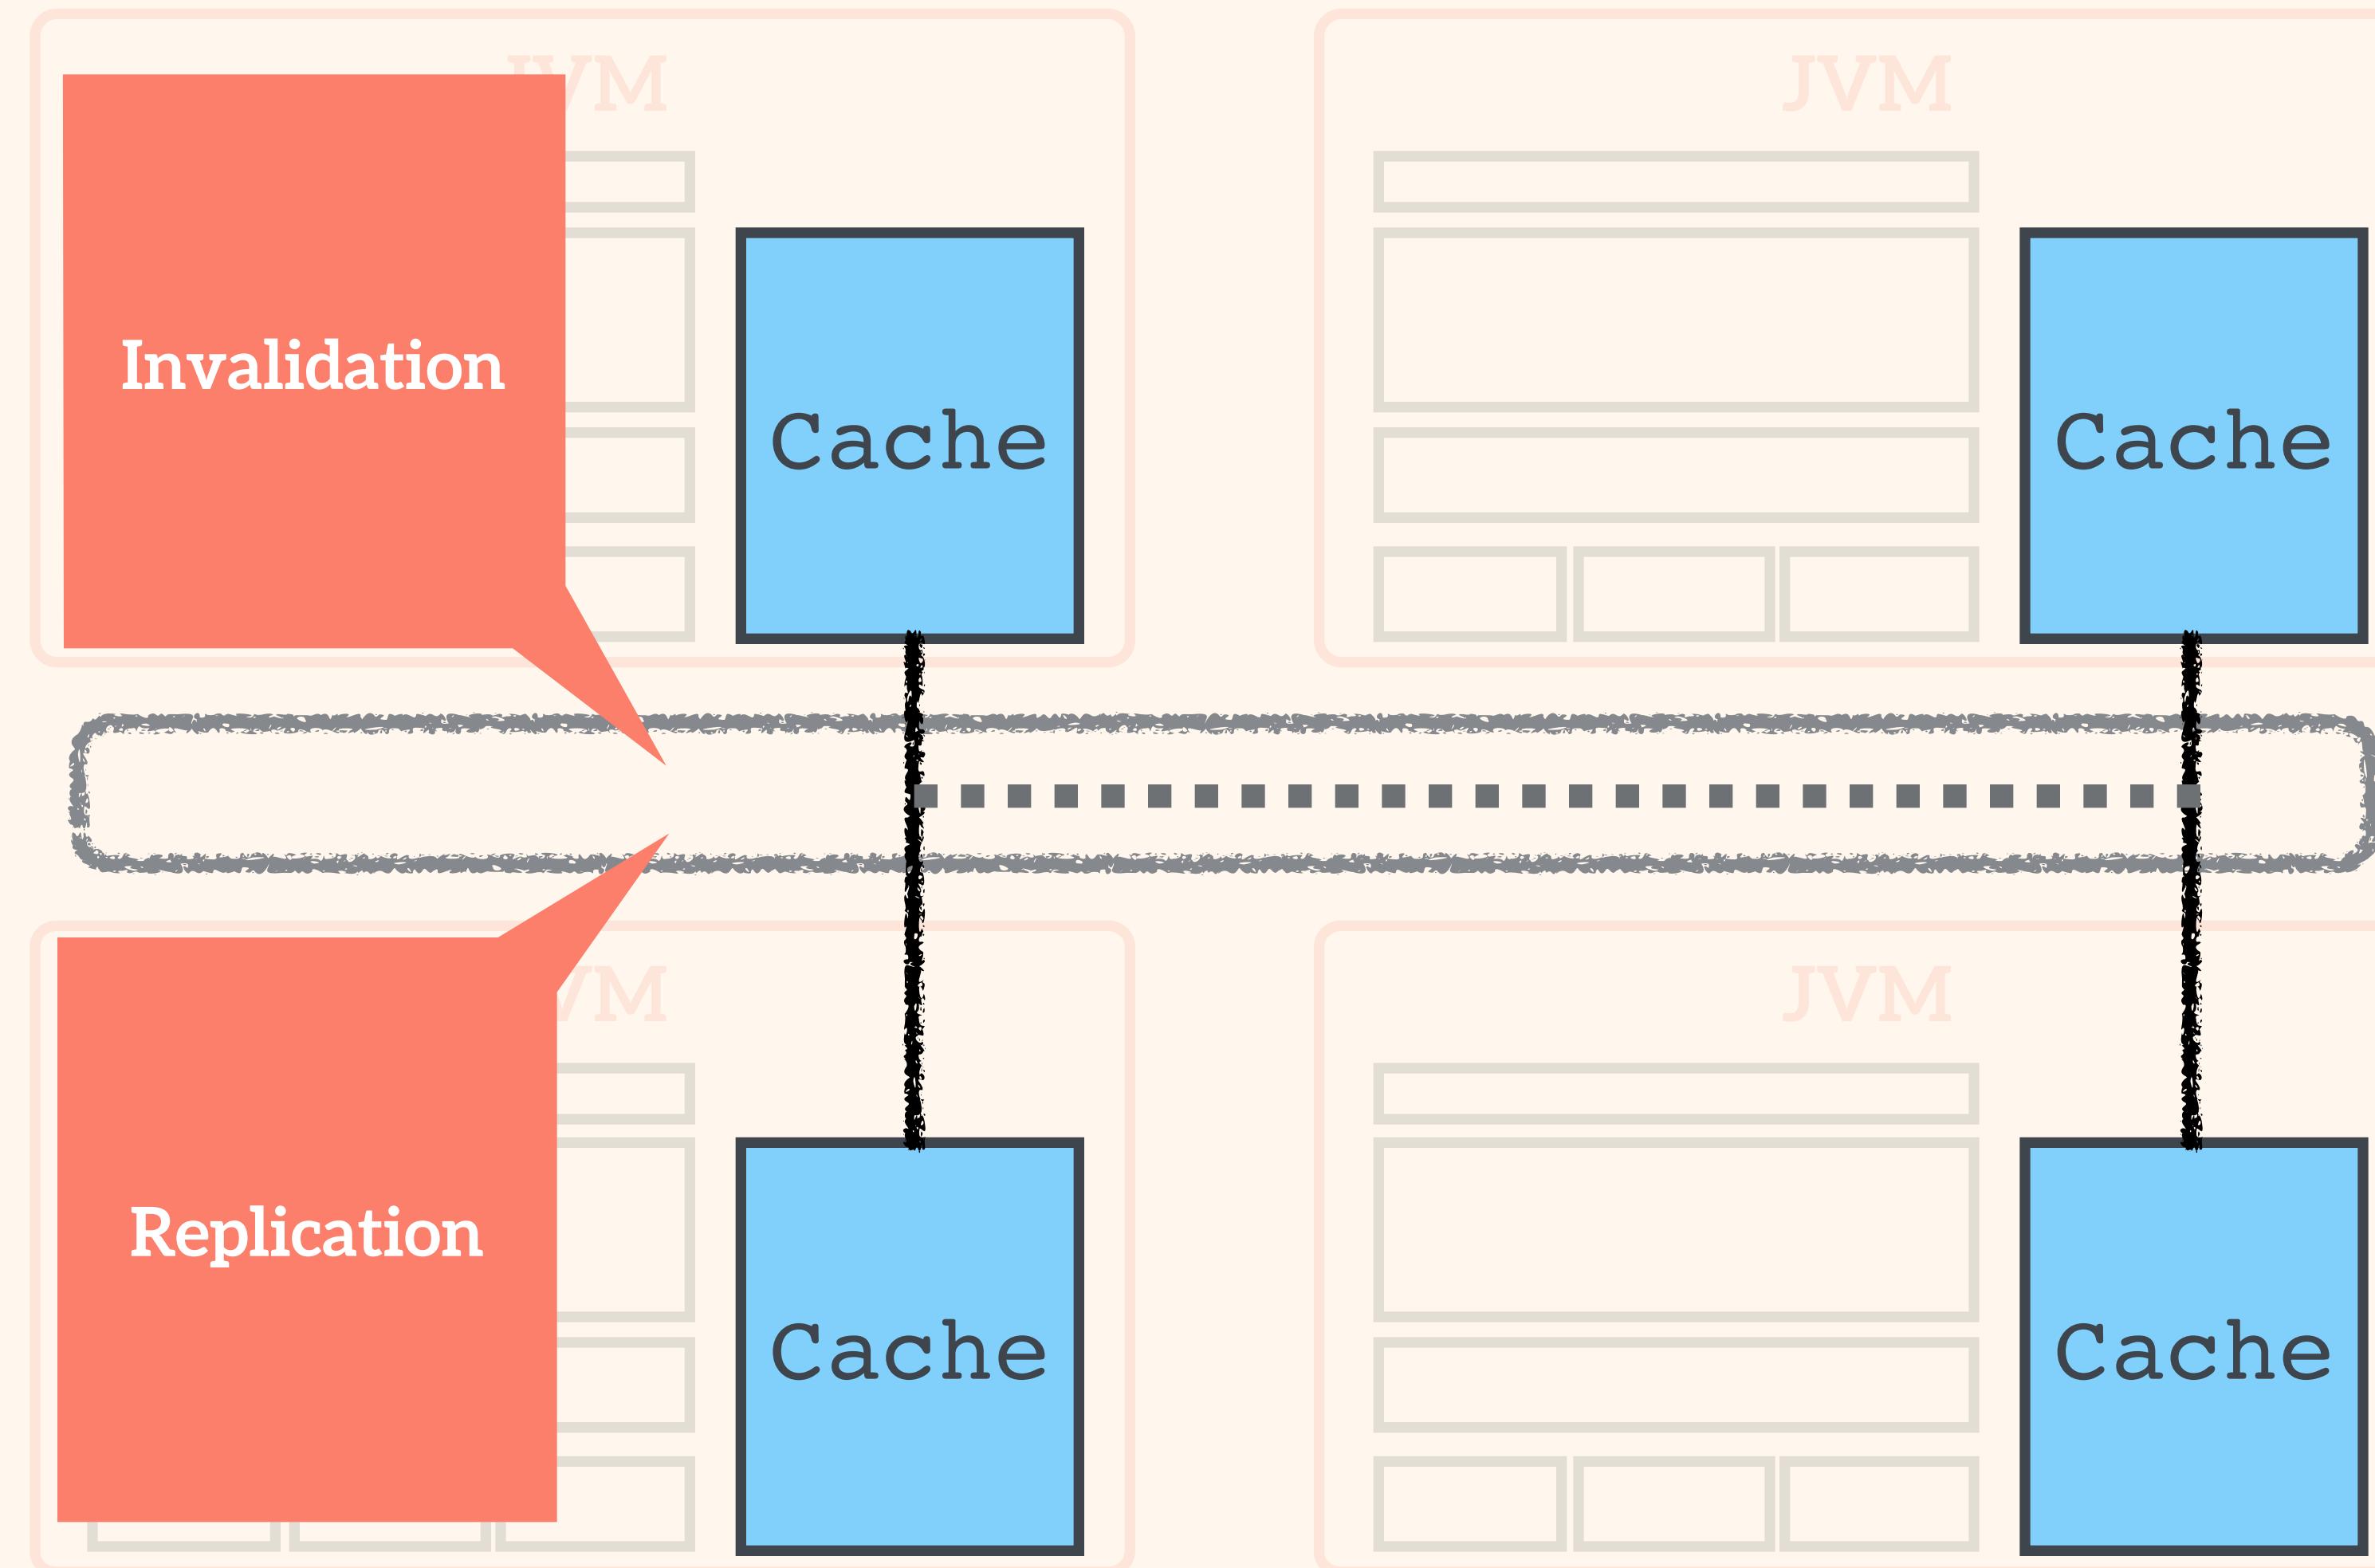




Clustered - with sync



Clustered - with sync

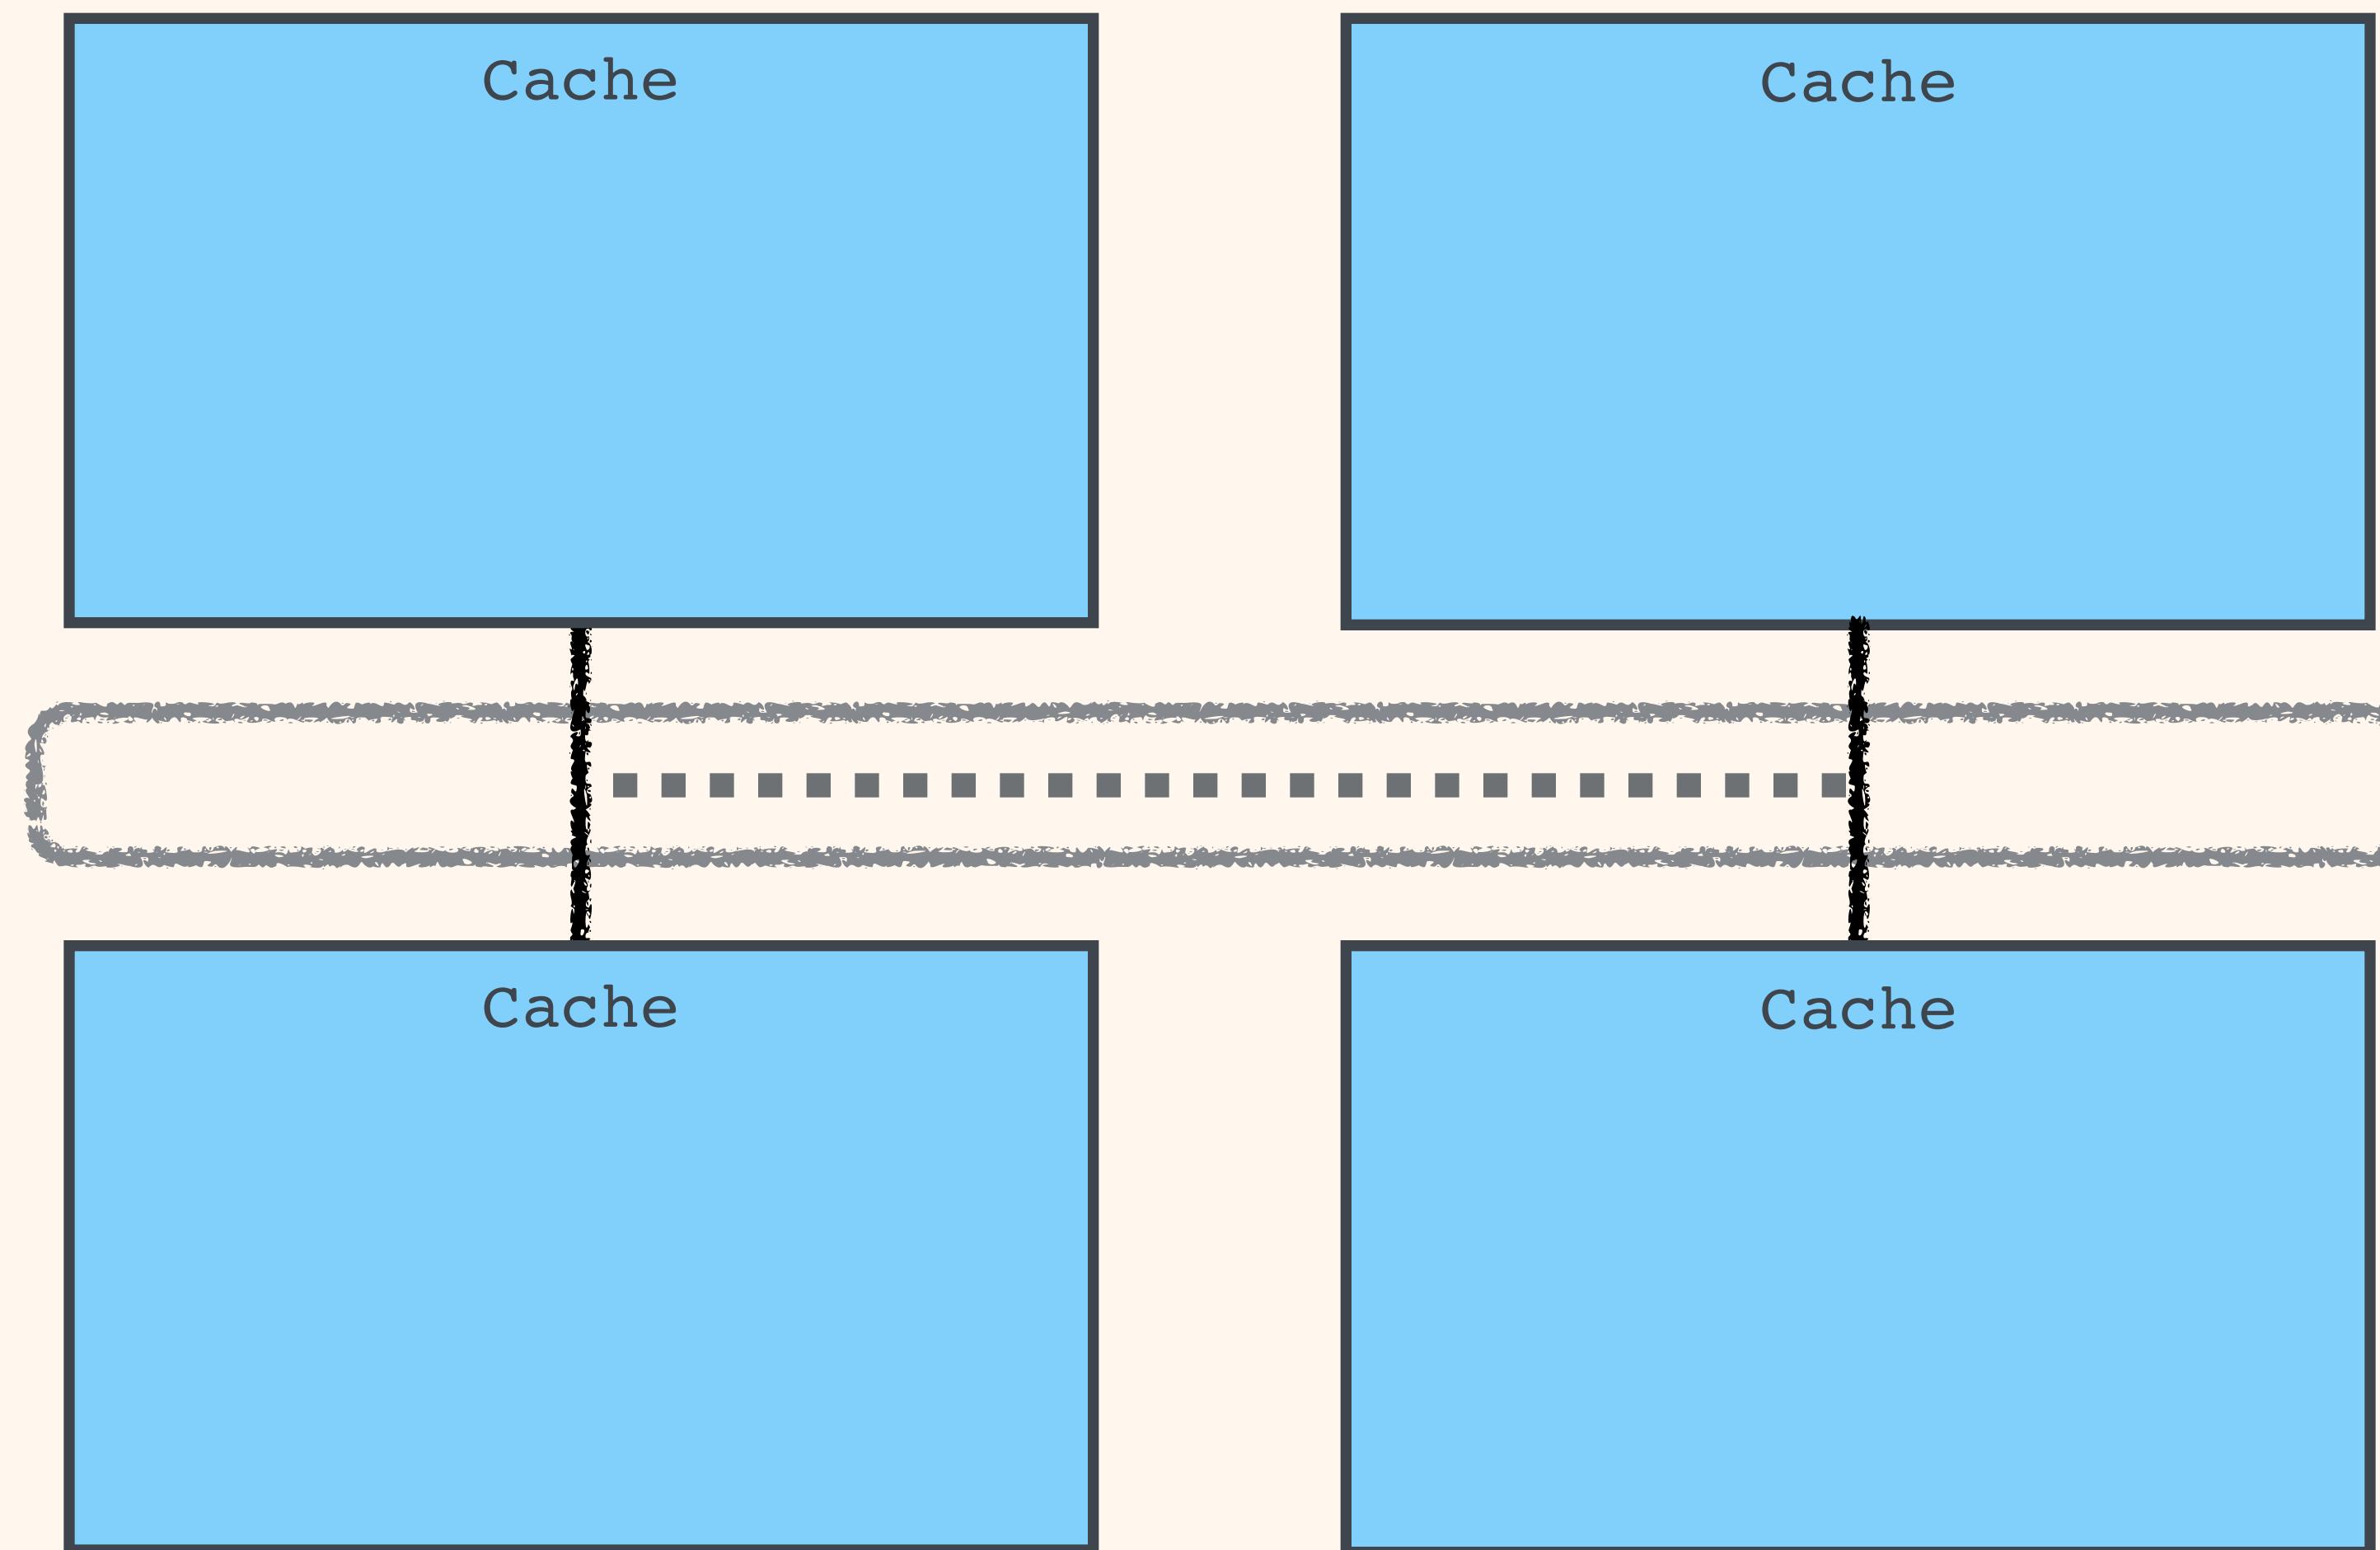




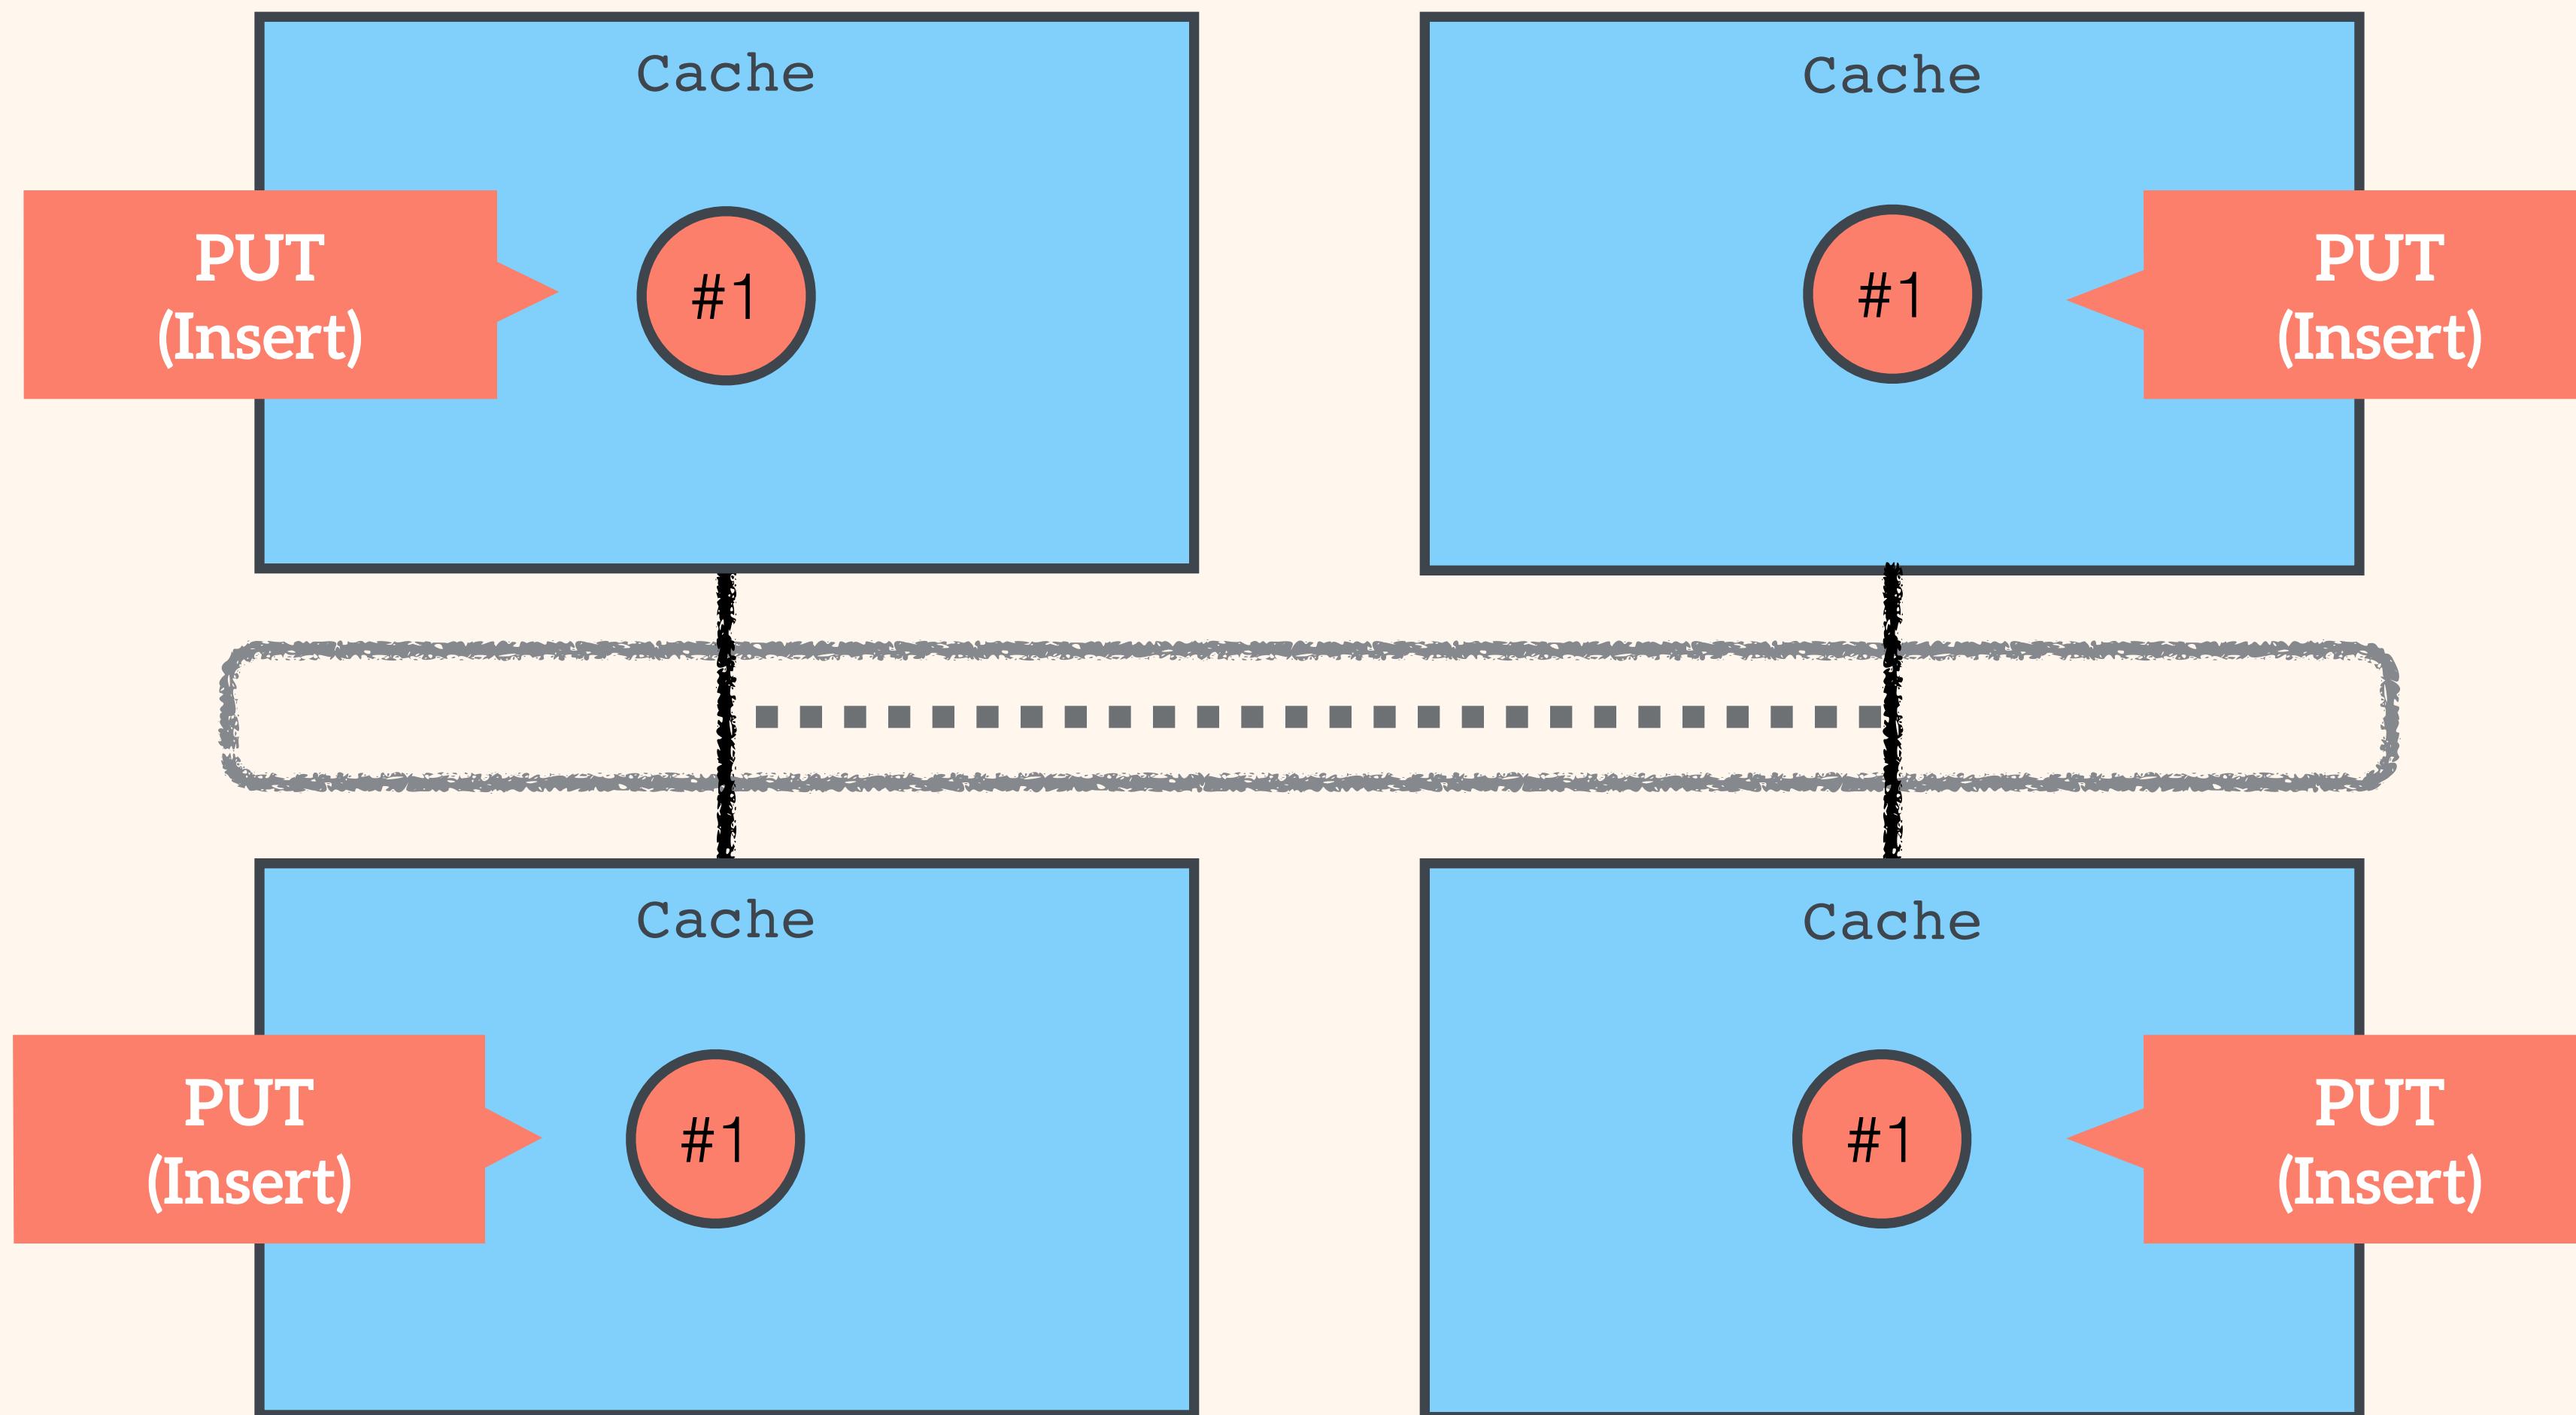
3

Avoid real replication where possible

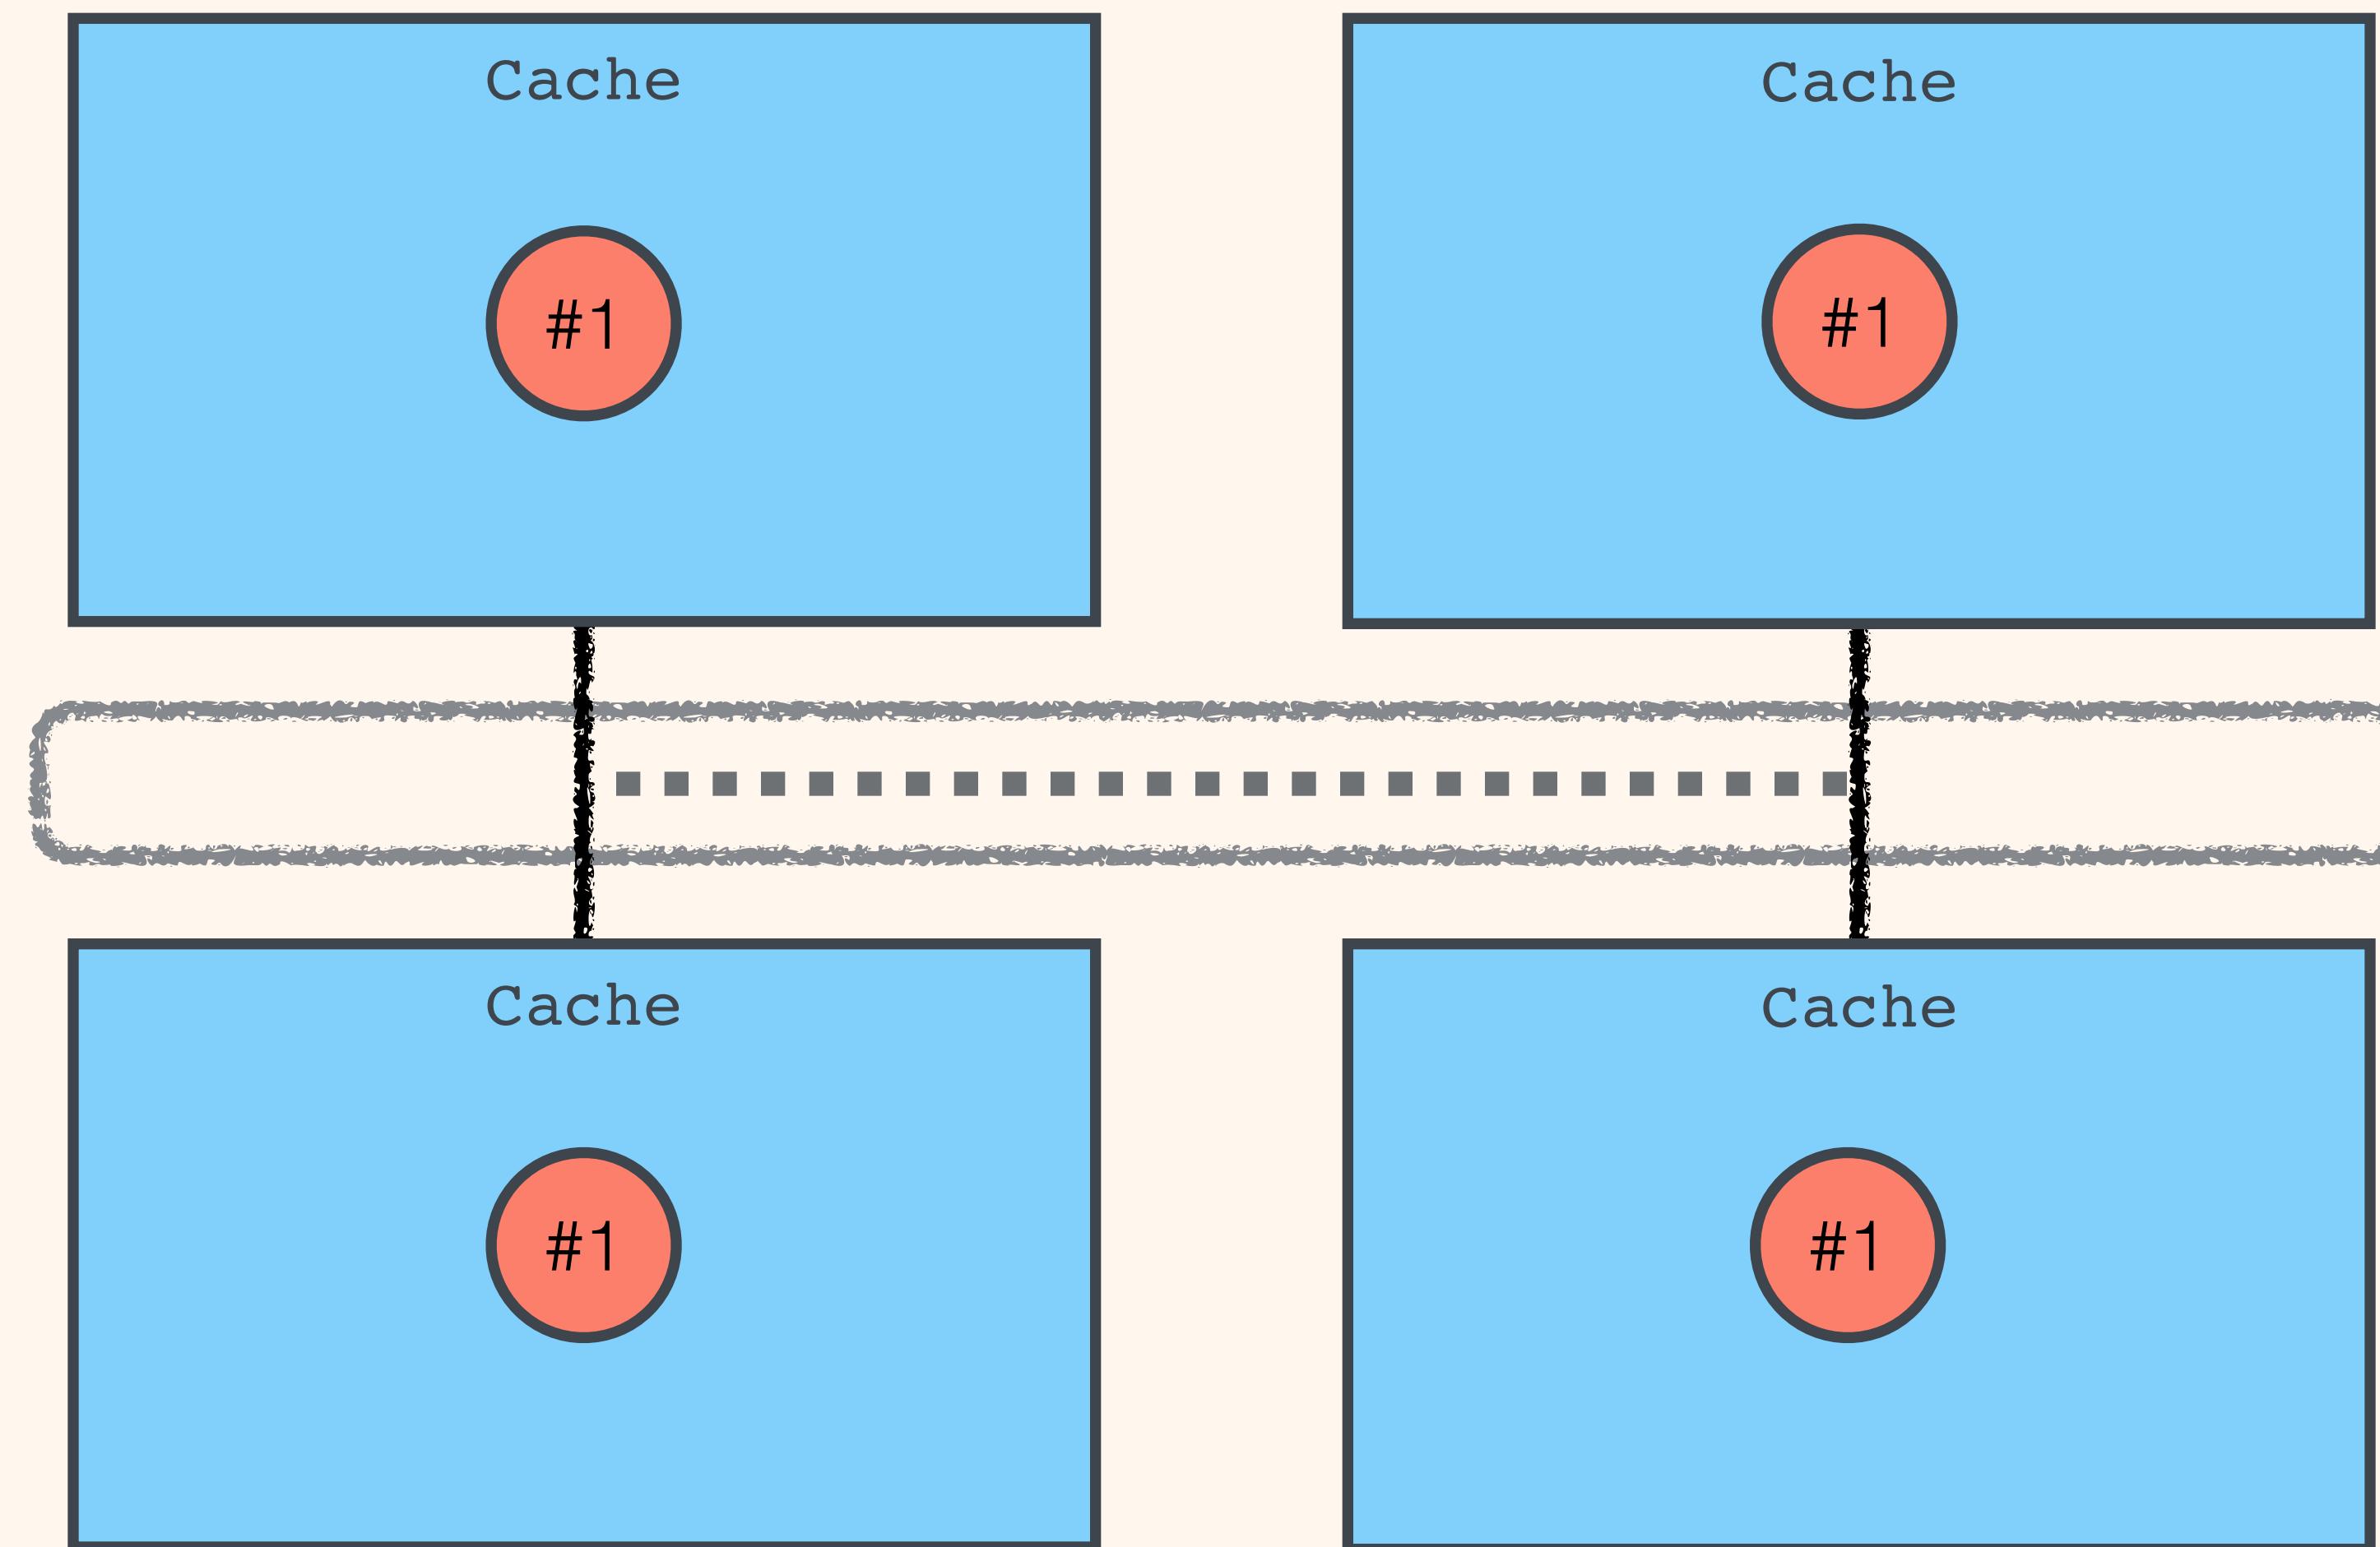
Invalidation - Option 1



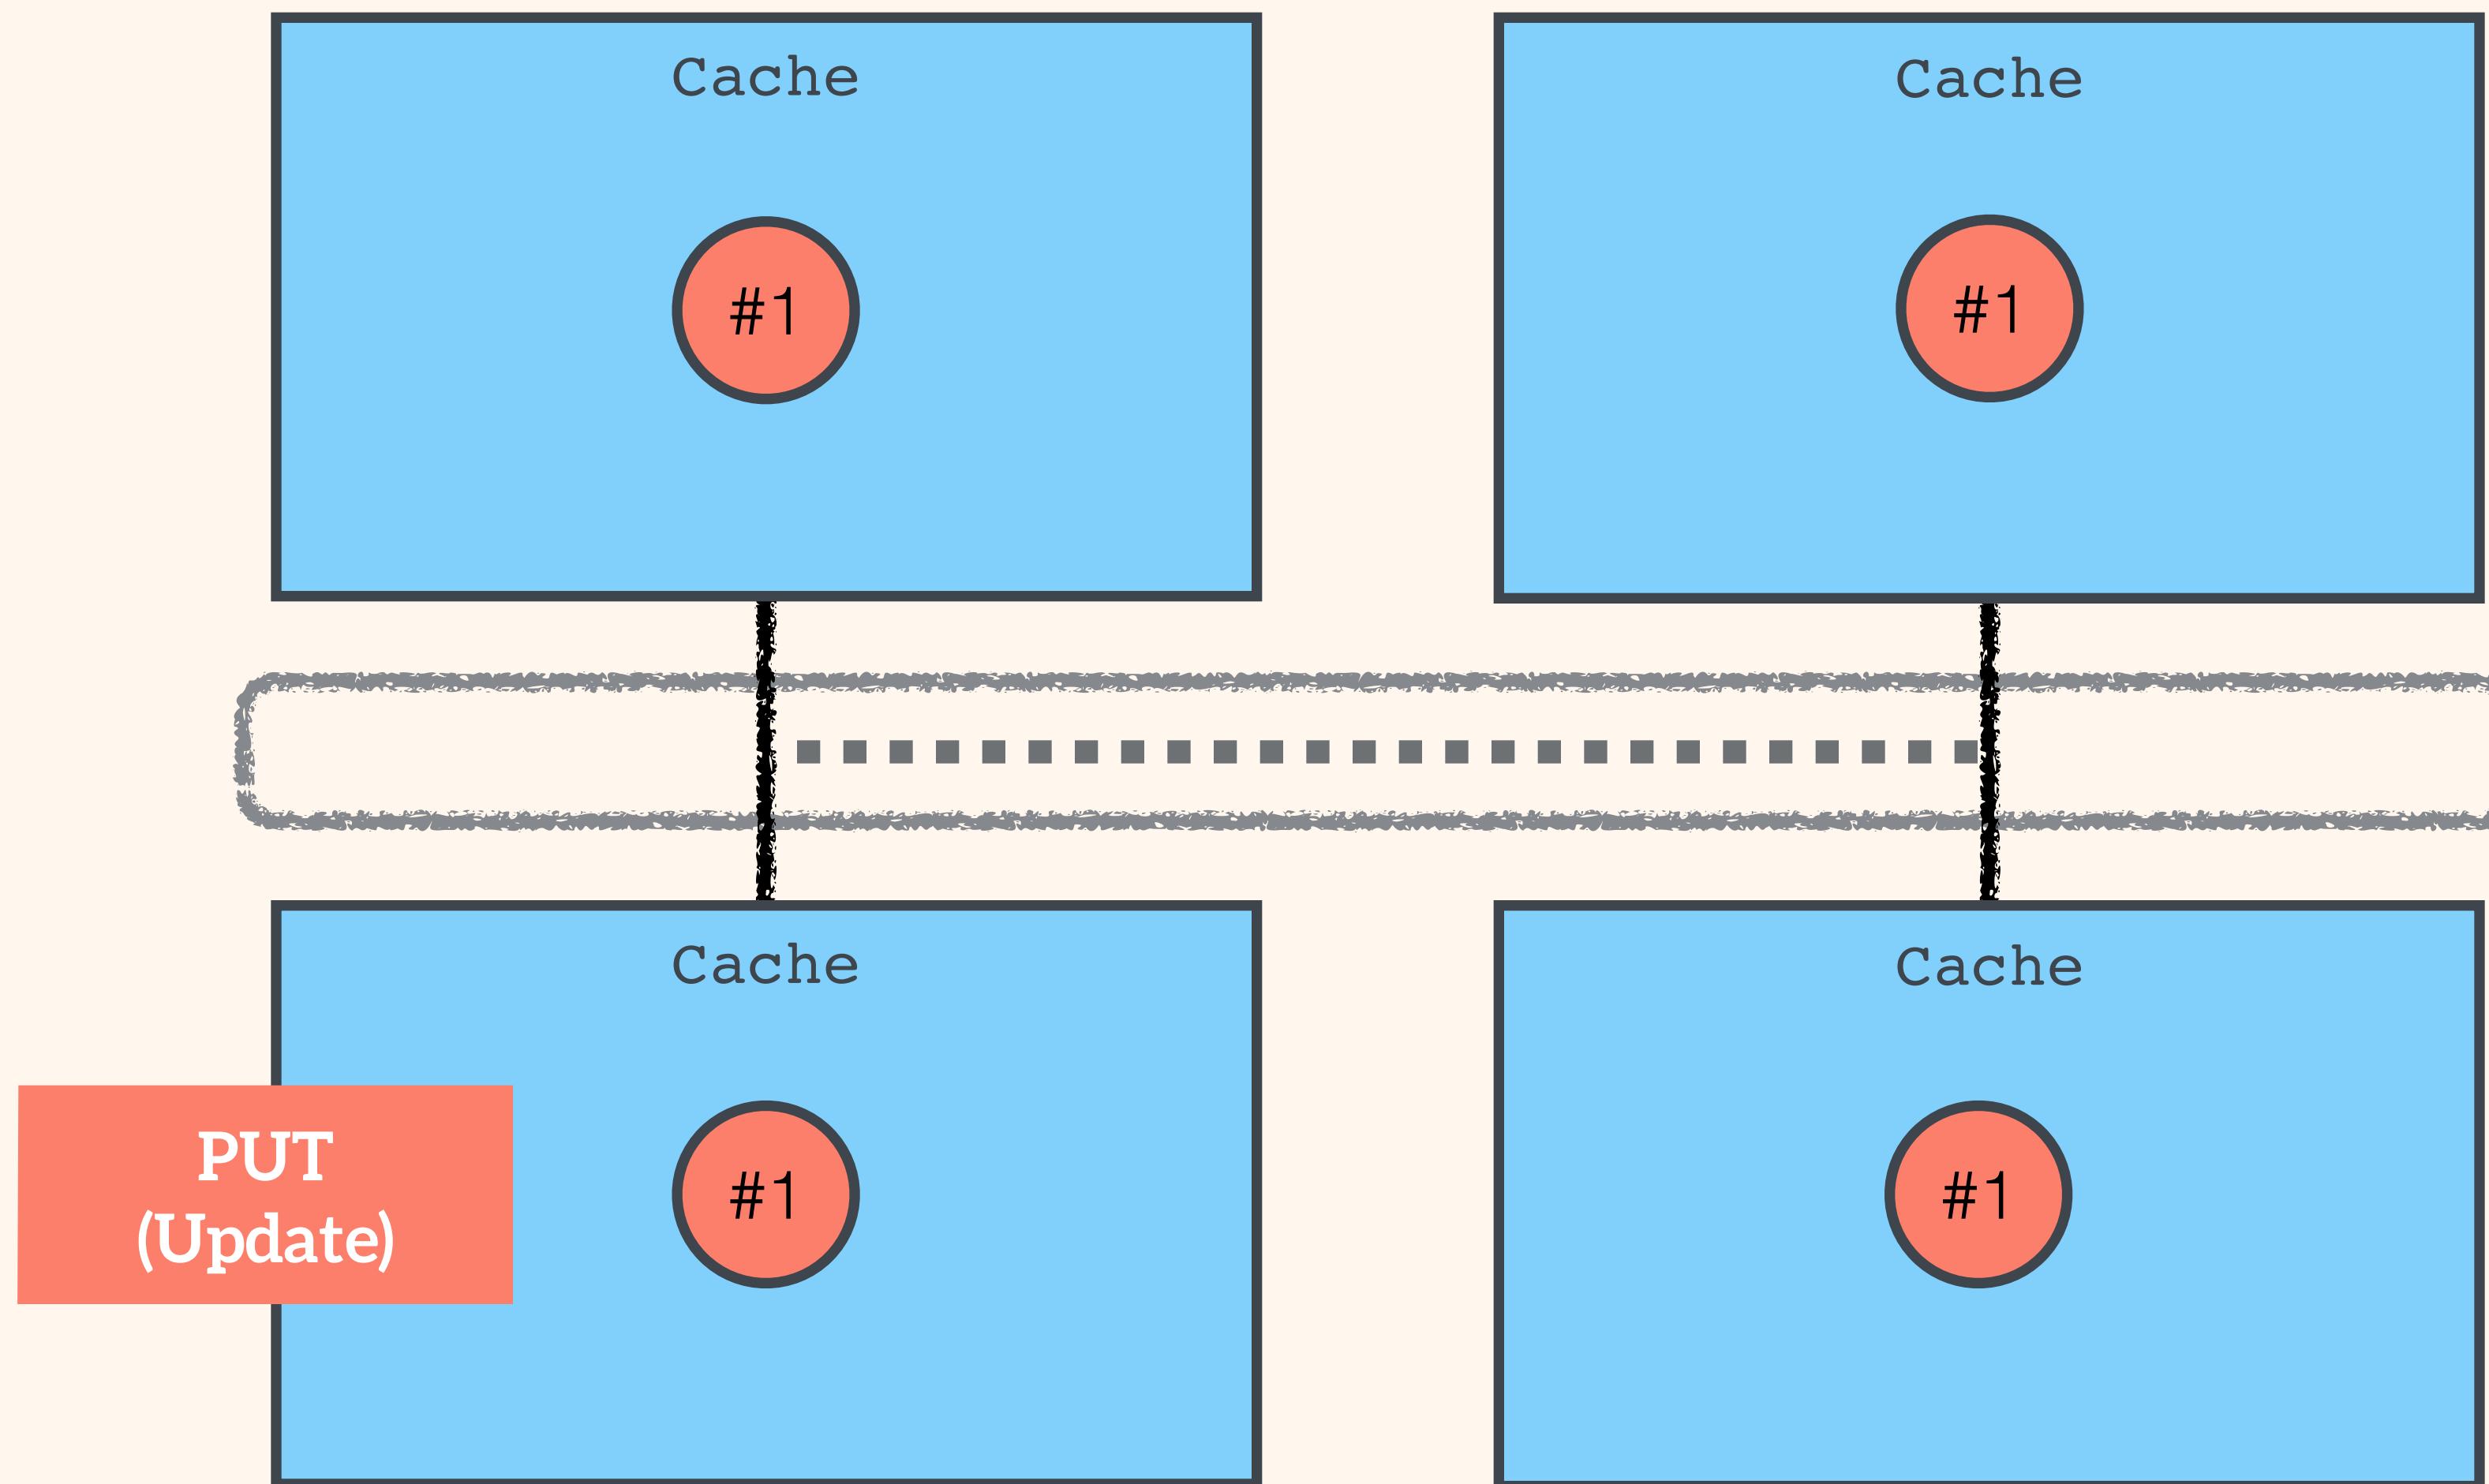
Invalidation - Option 1



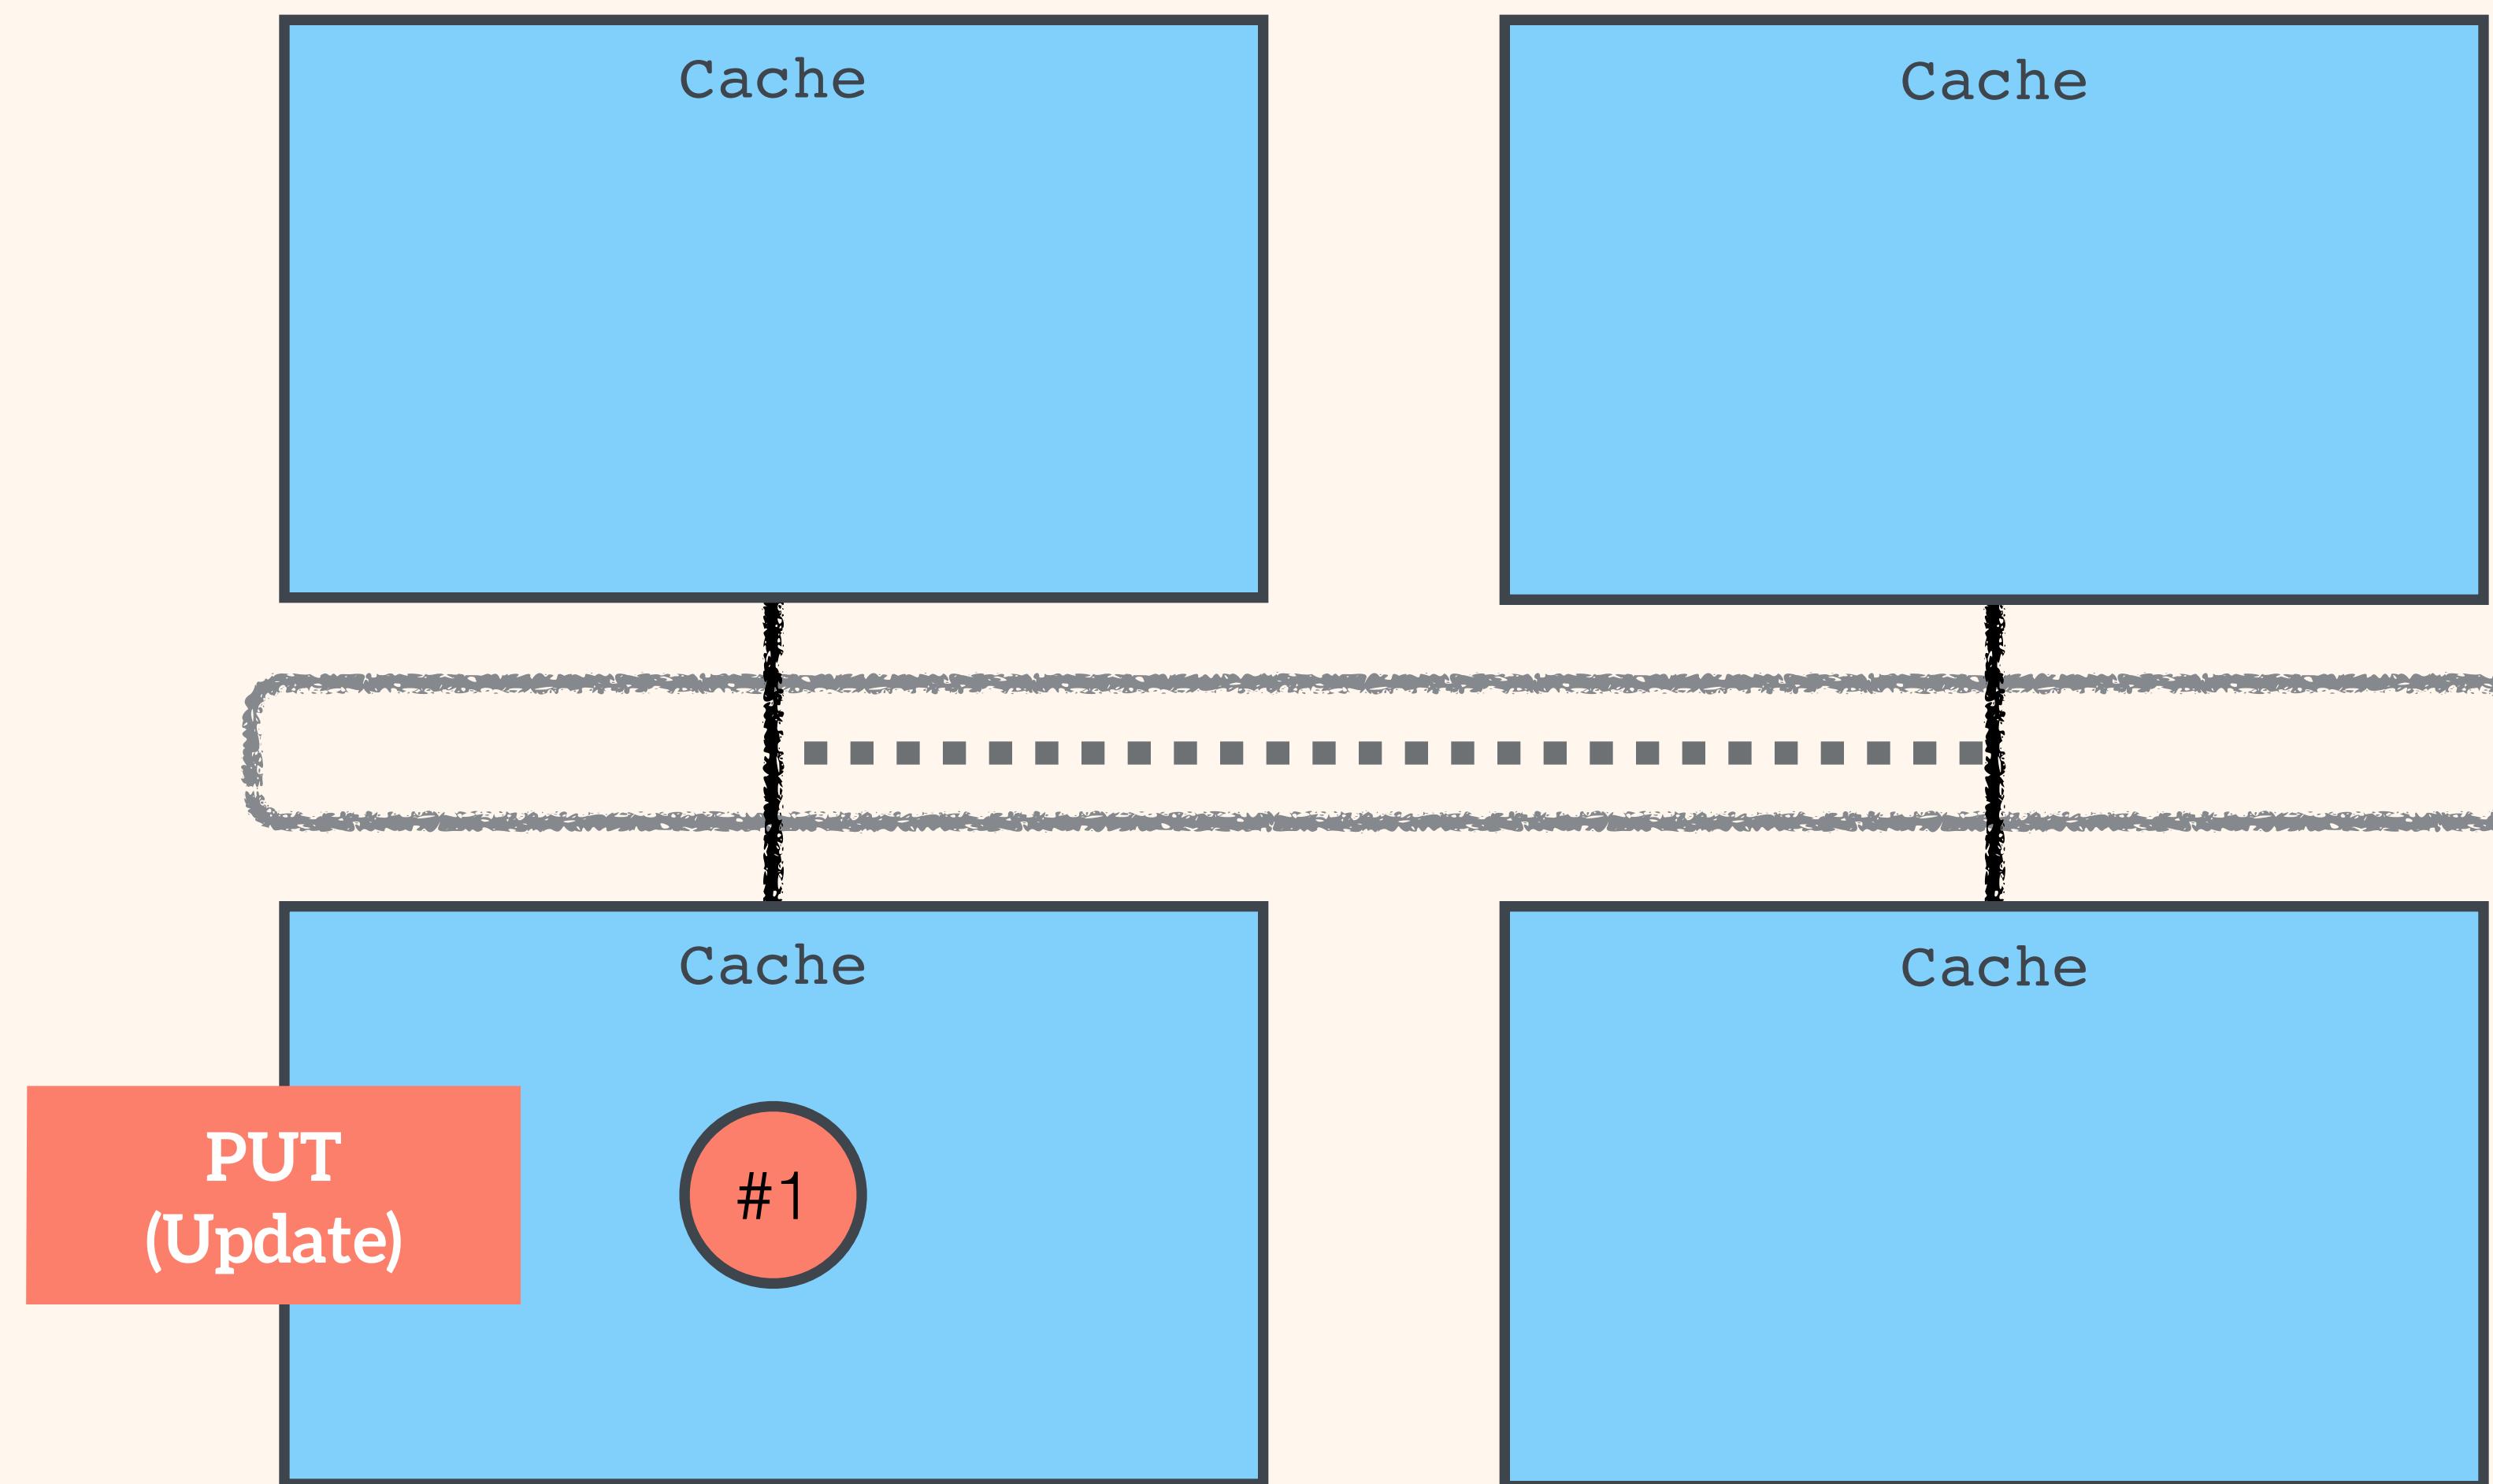
Invalidation - Option 1



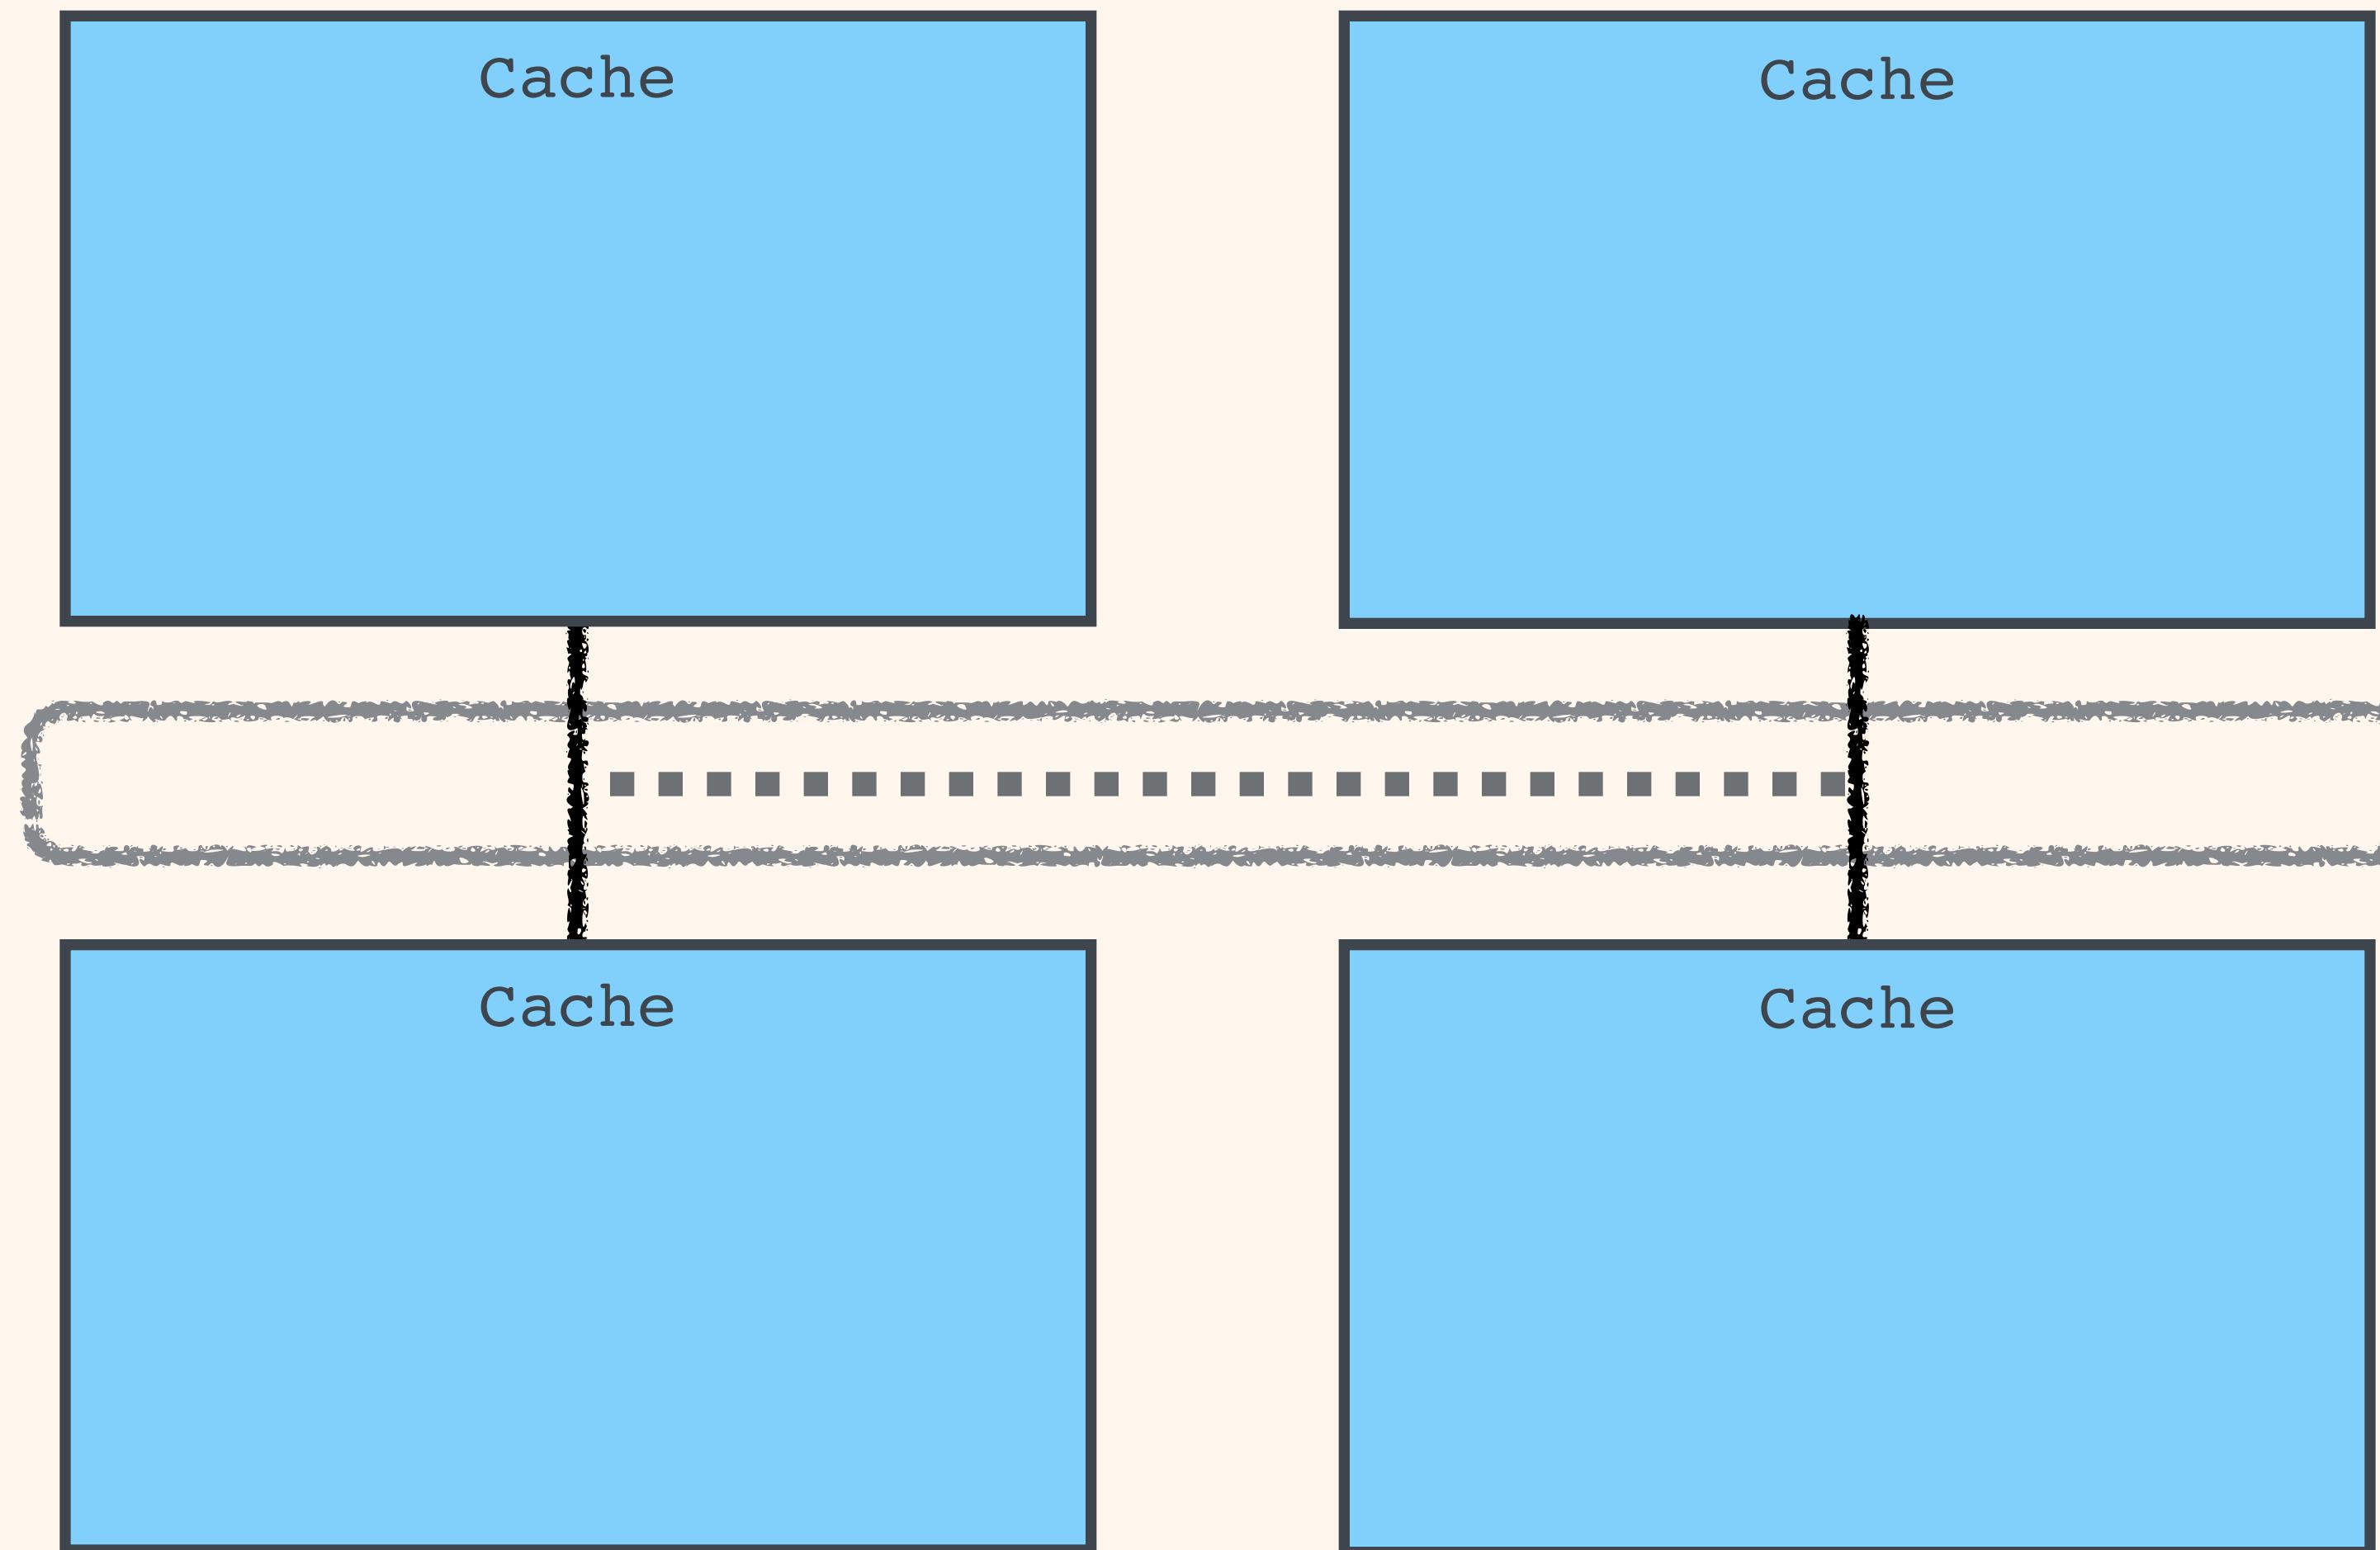
Invalidation - Option 1



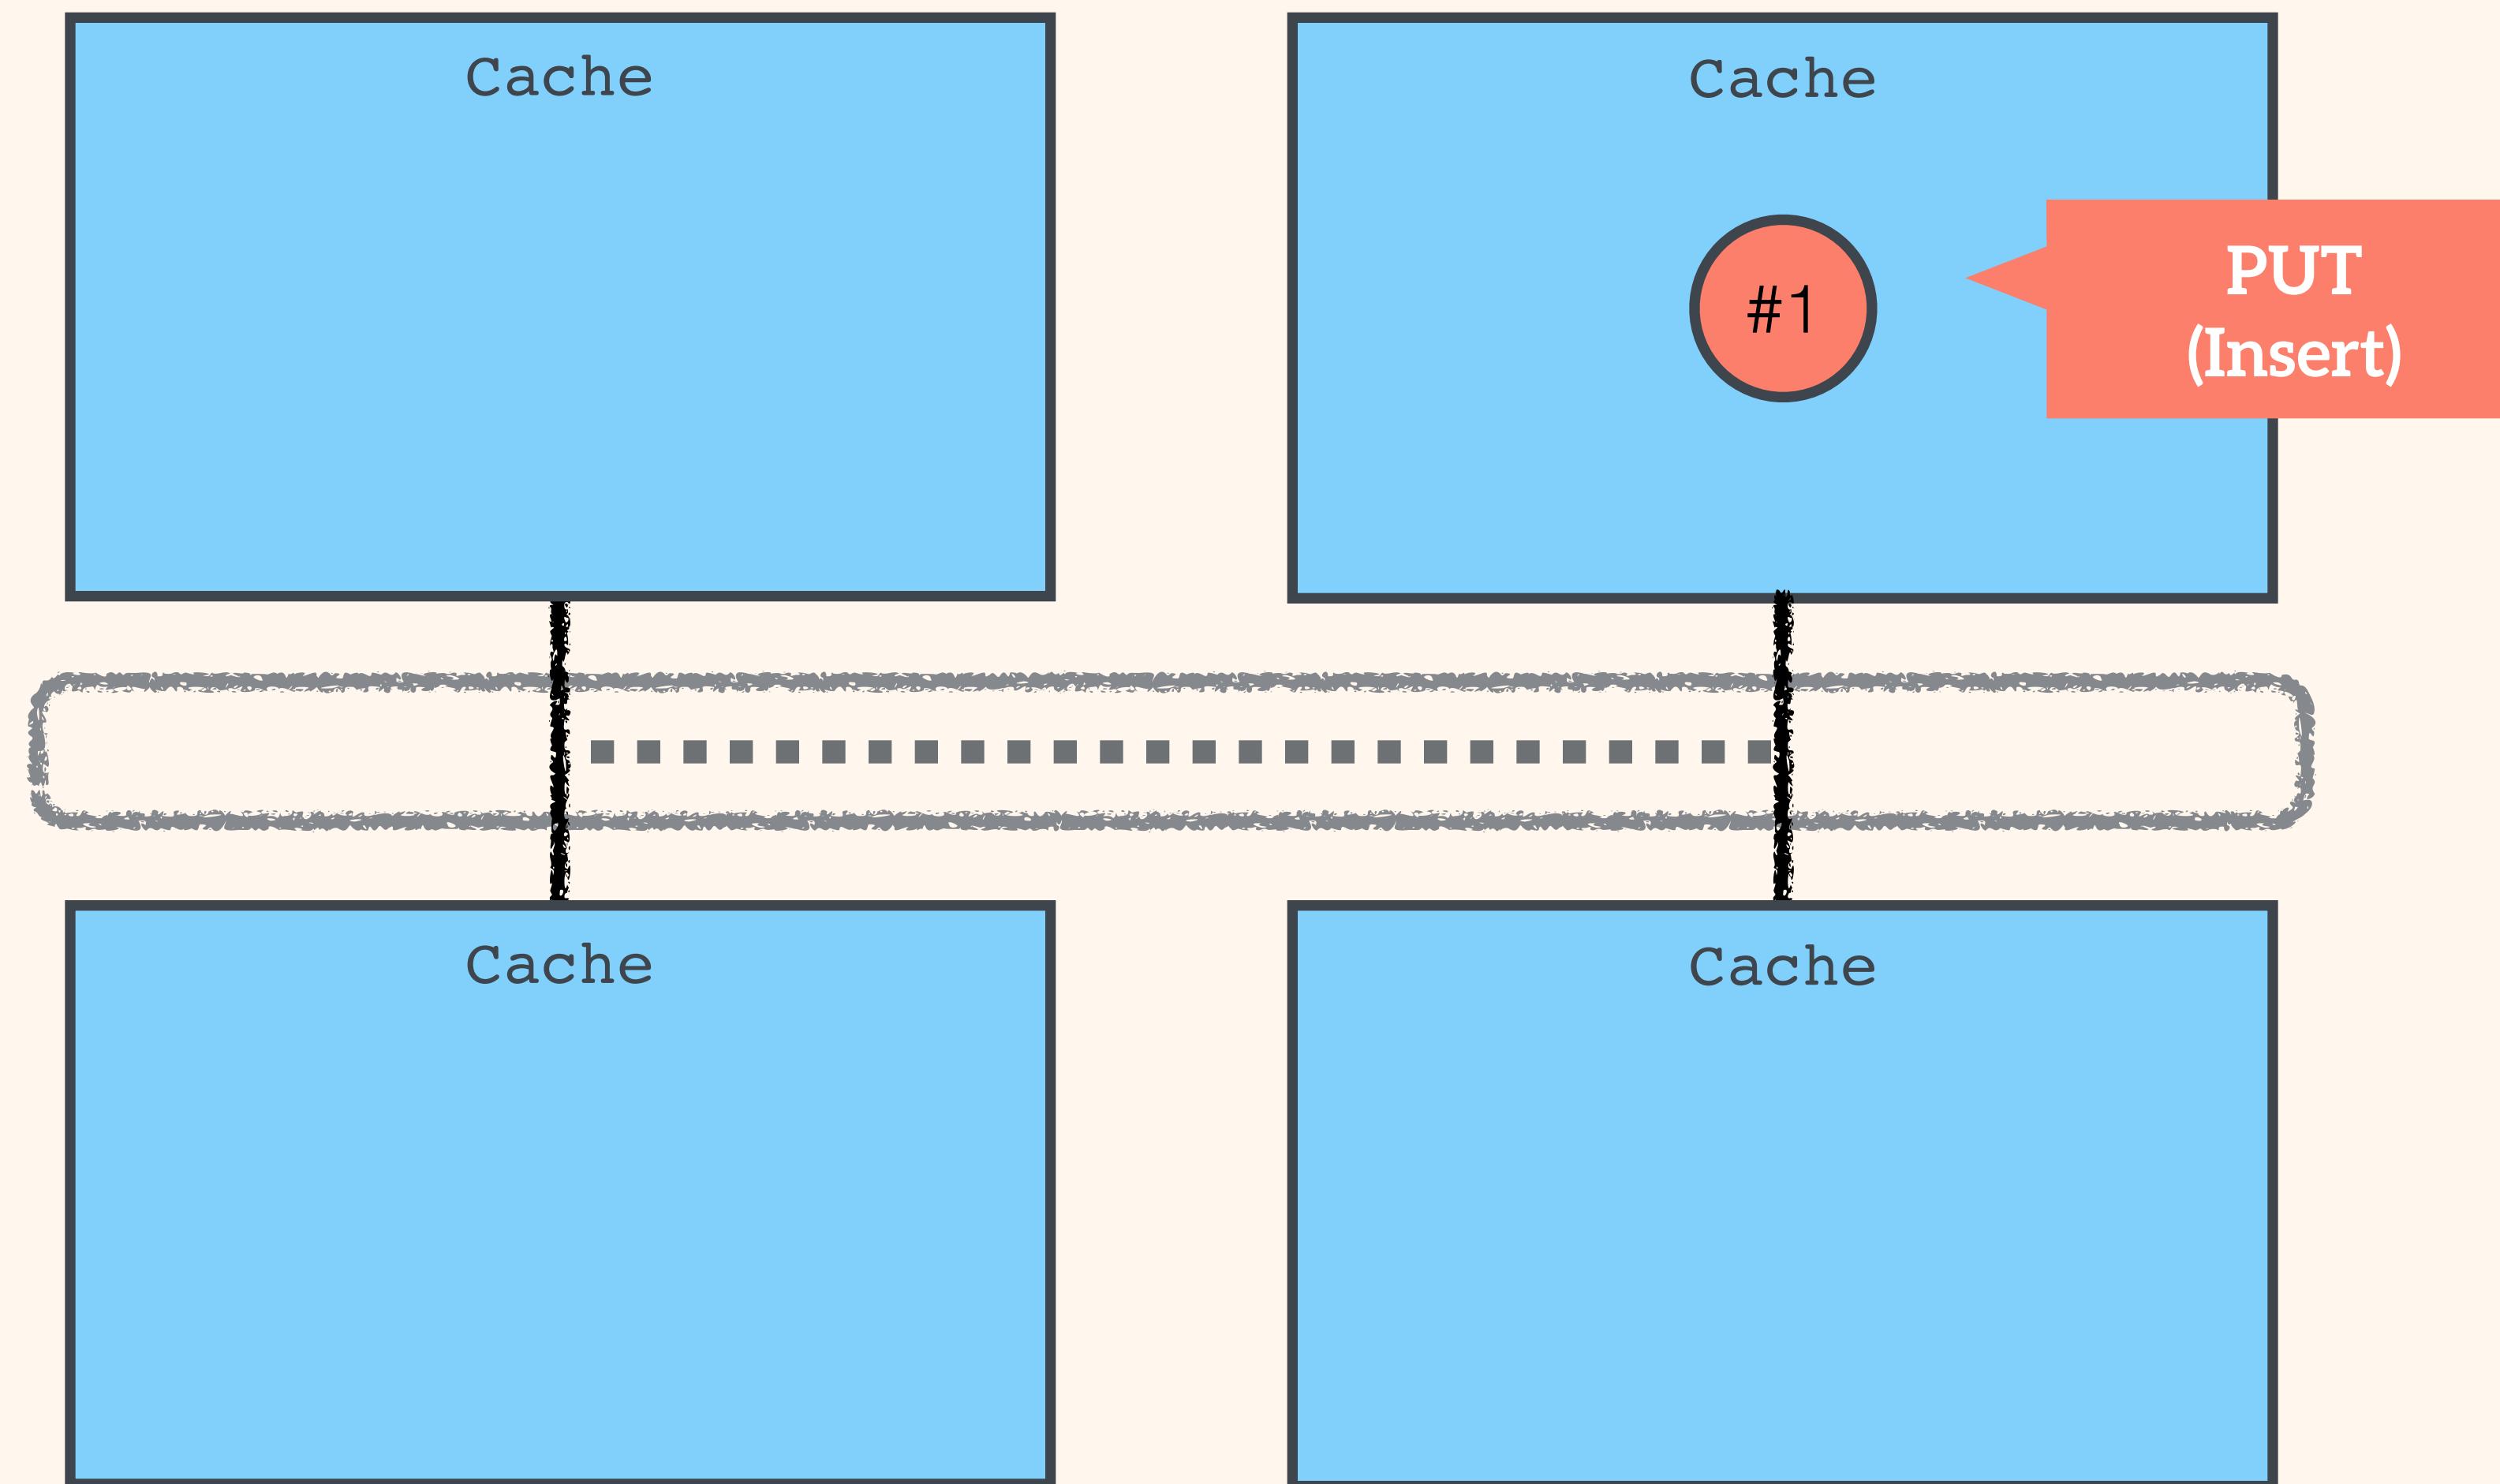
Invalidation - Option 1



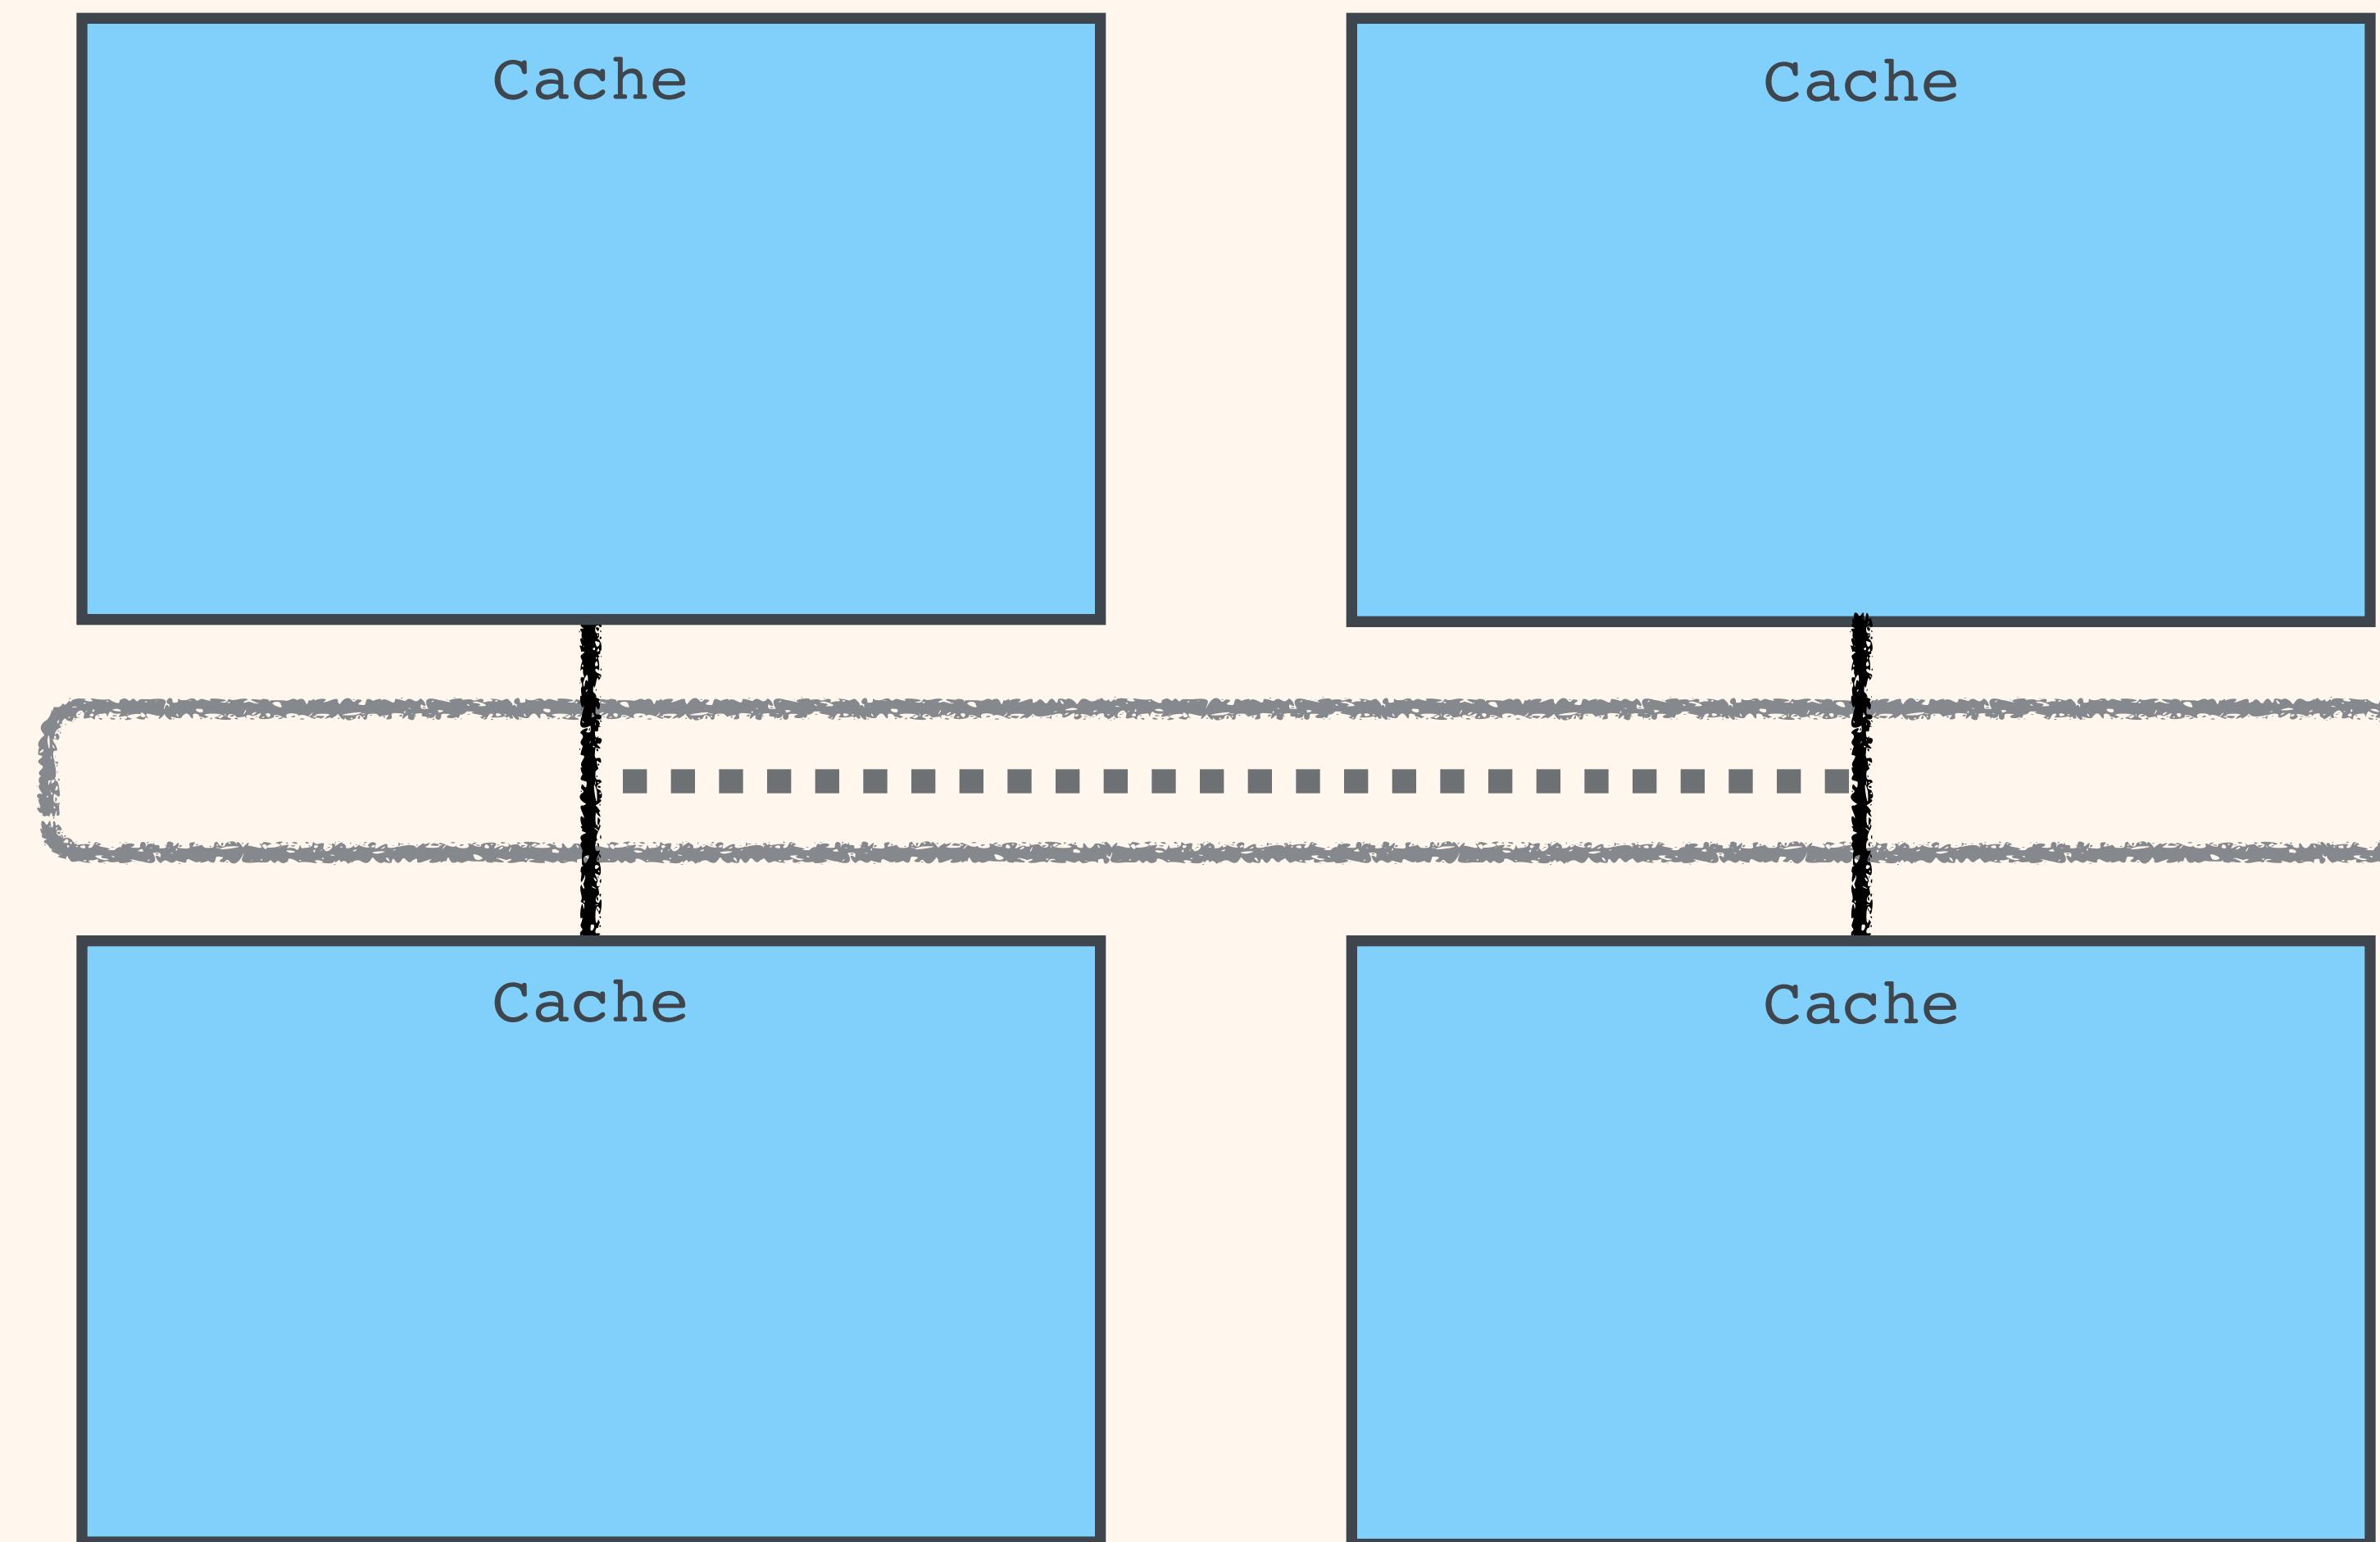
Invalidation - Option 2



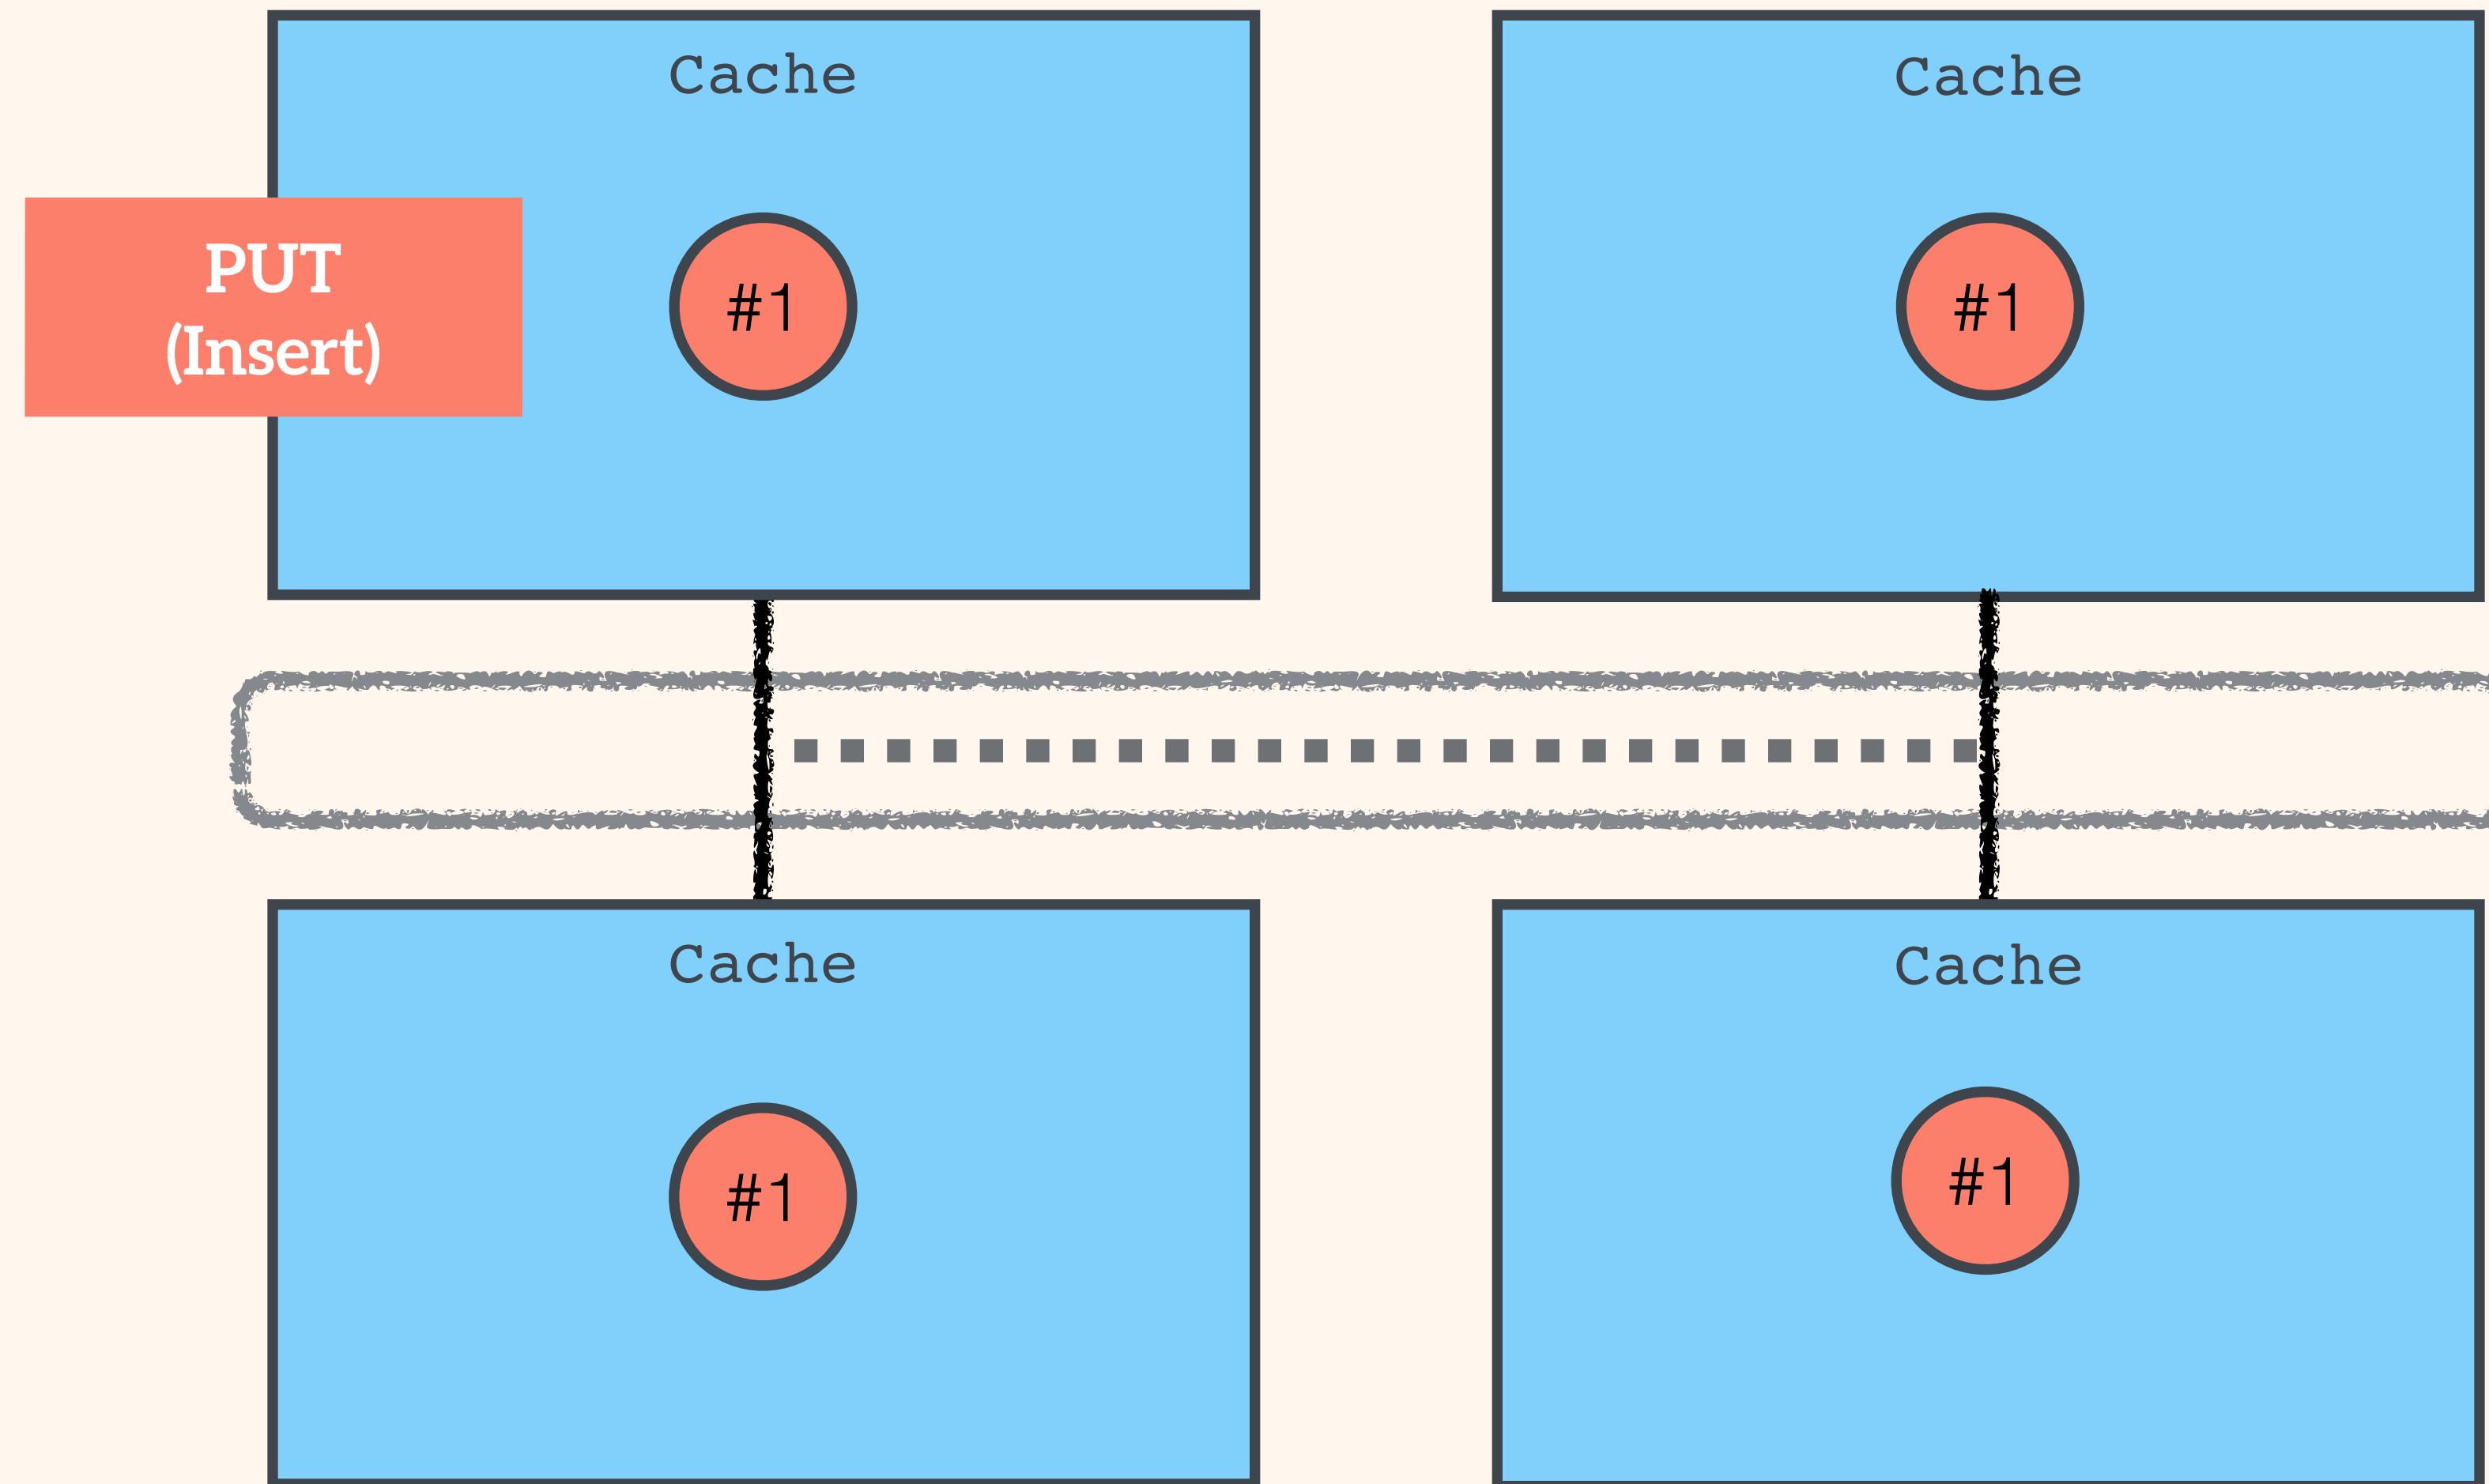
Invalidation - Option 2



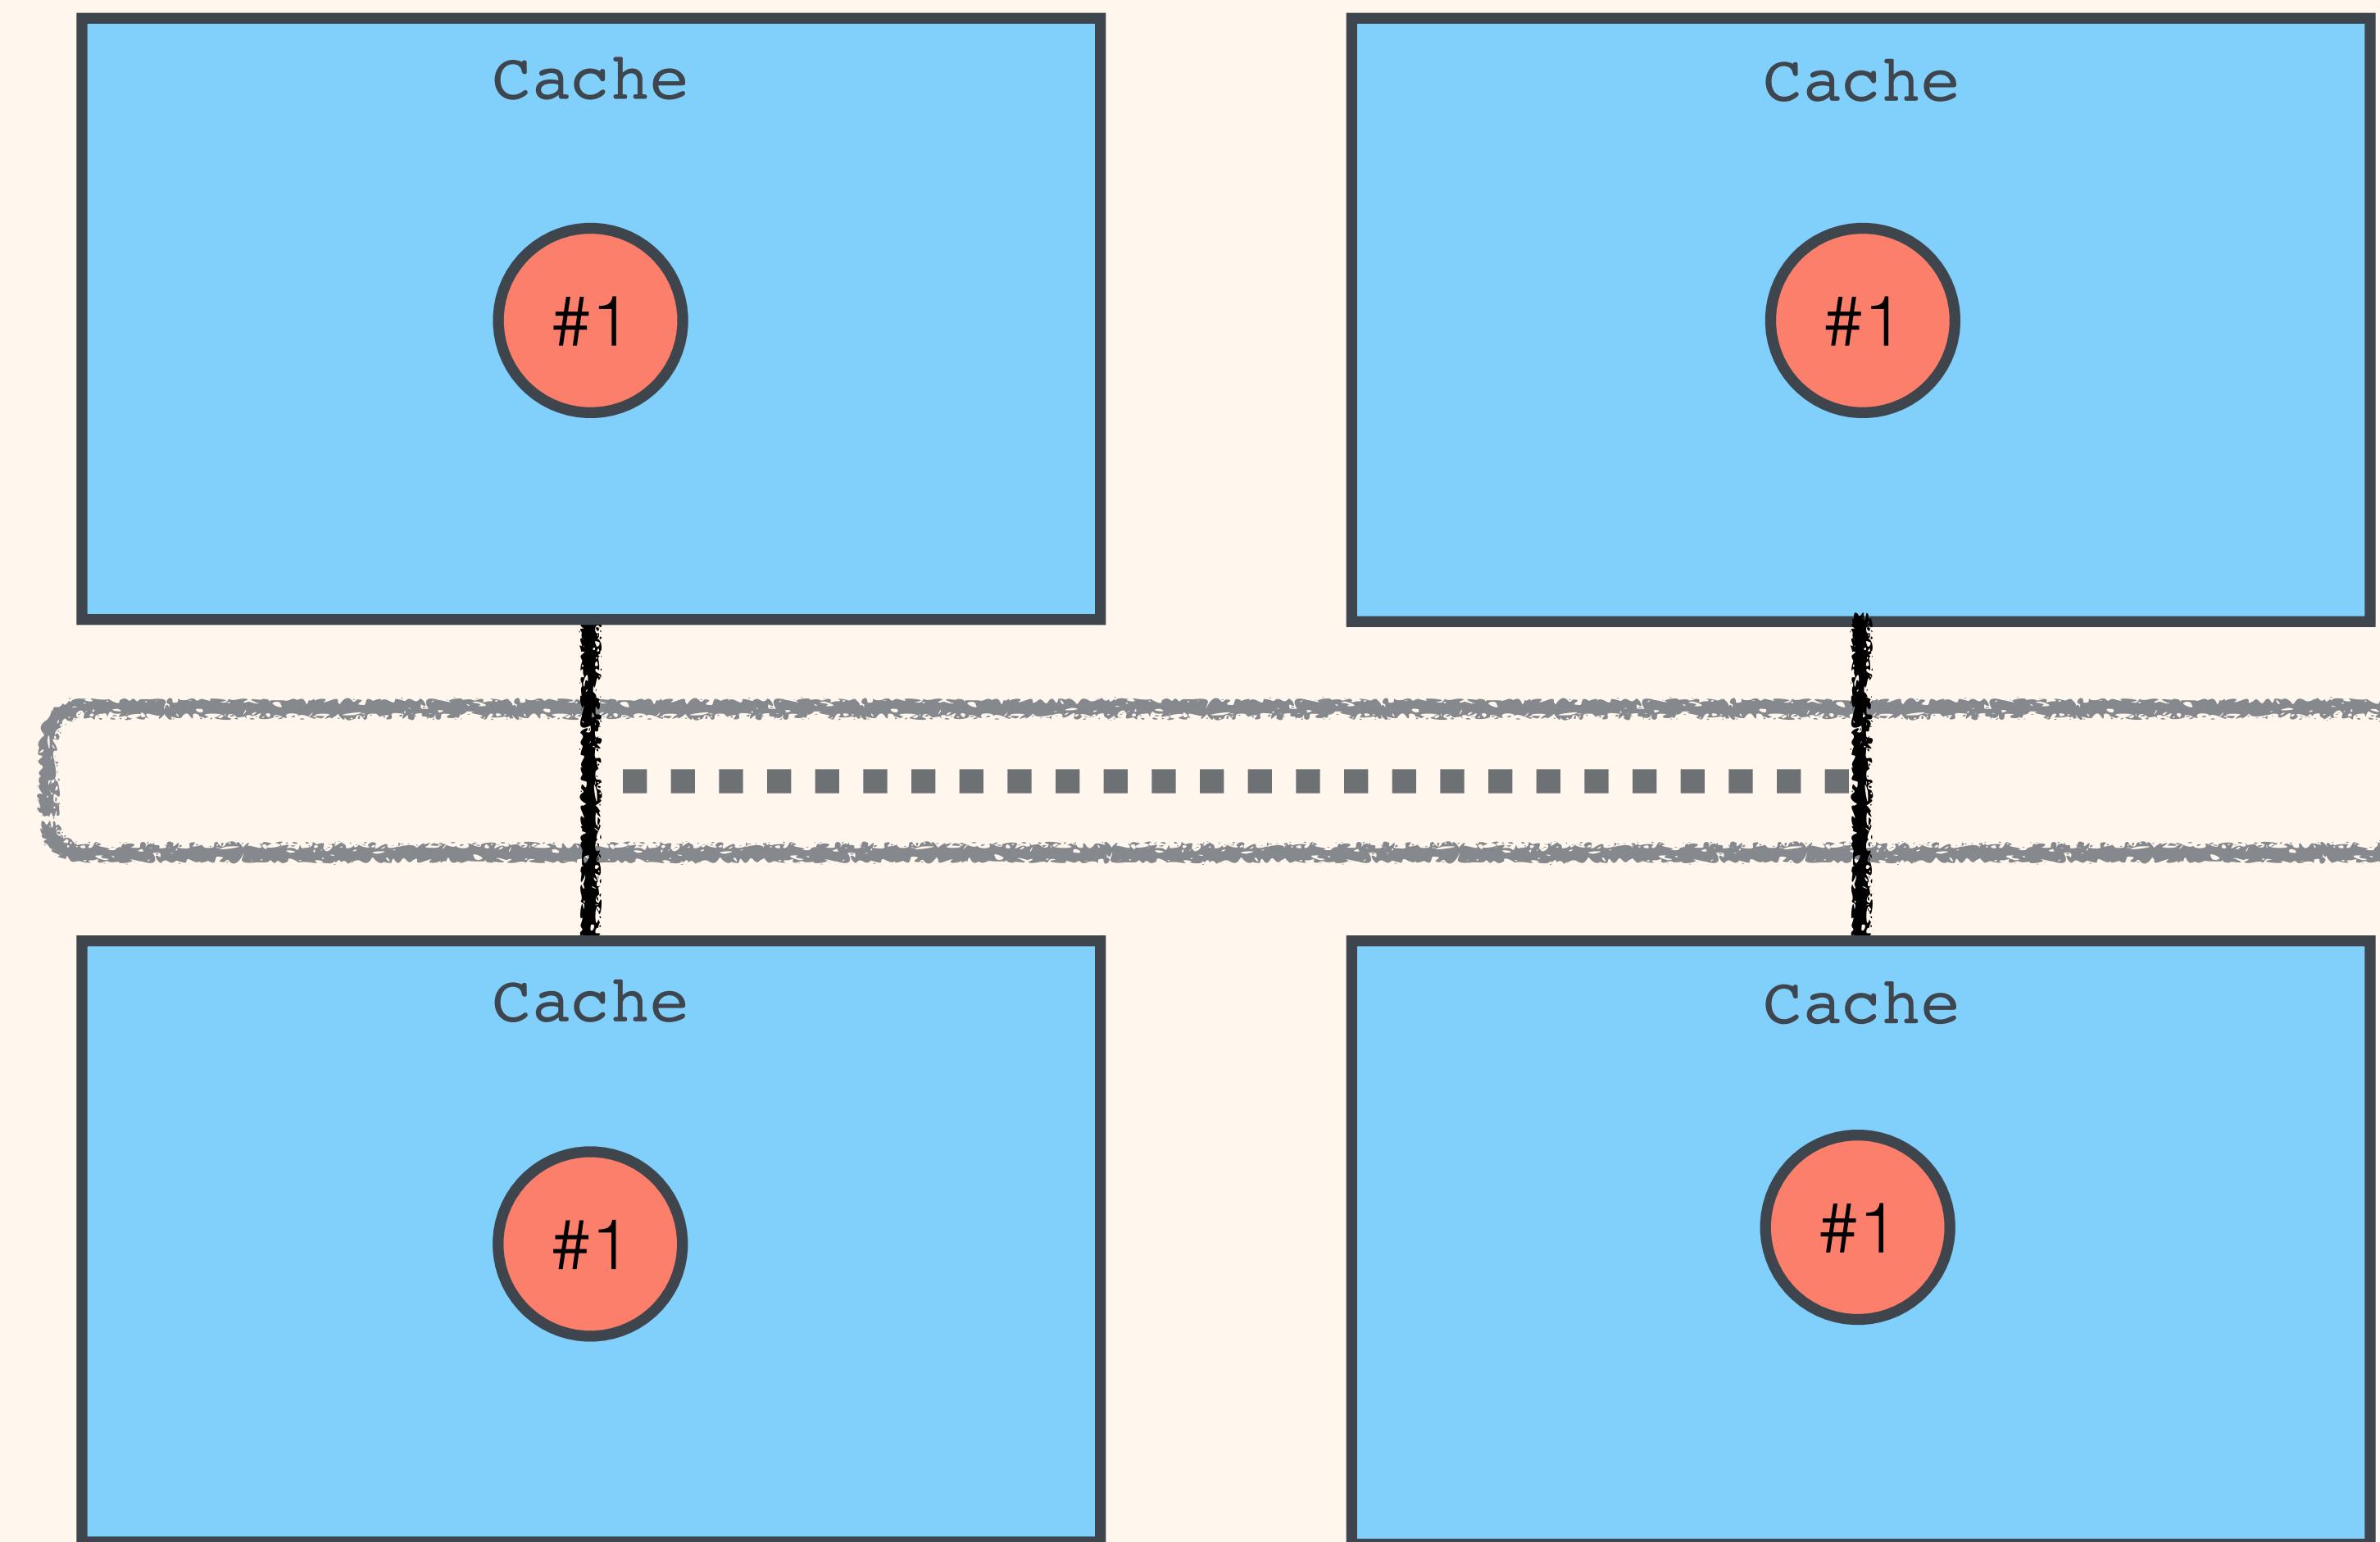
Replication



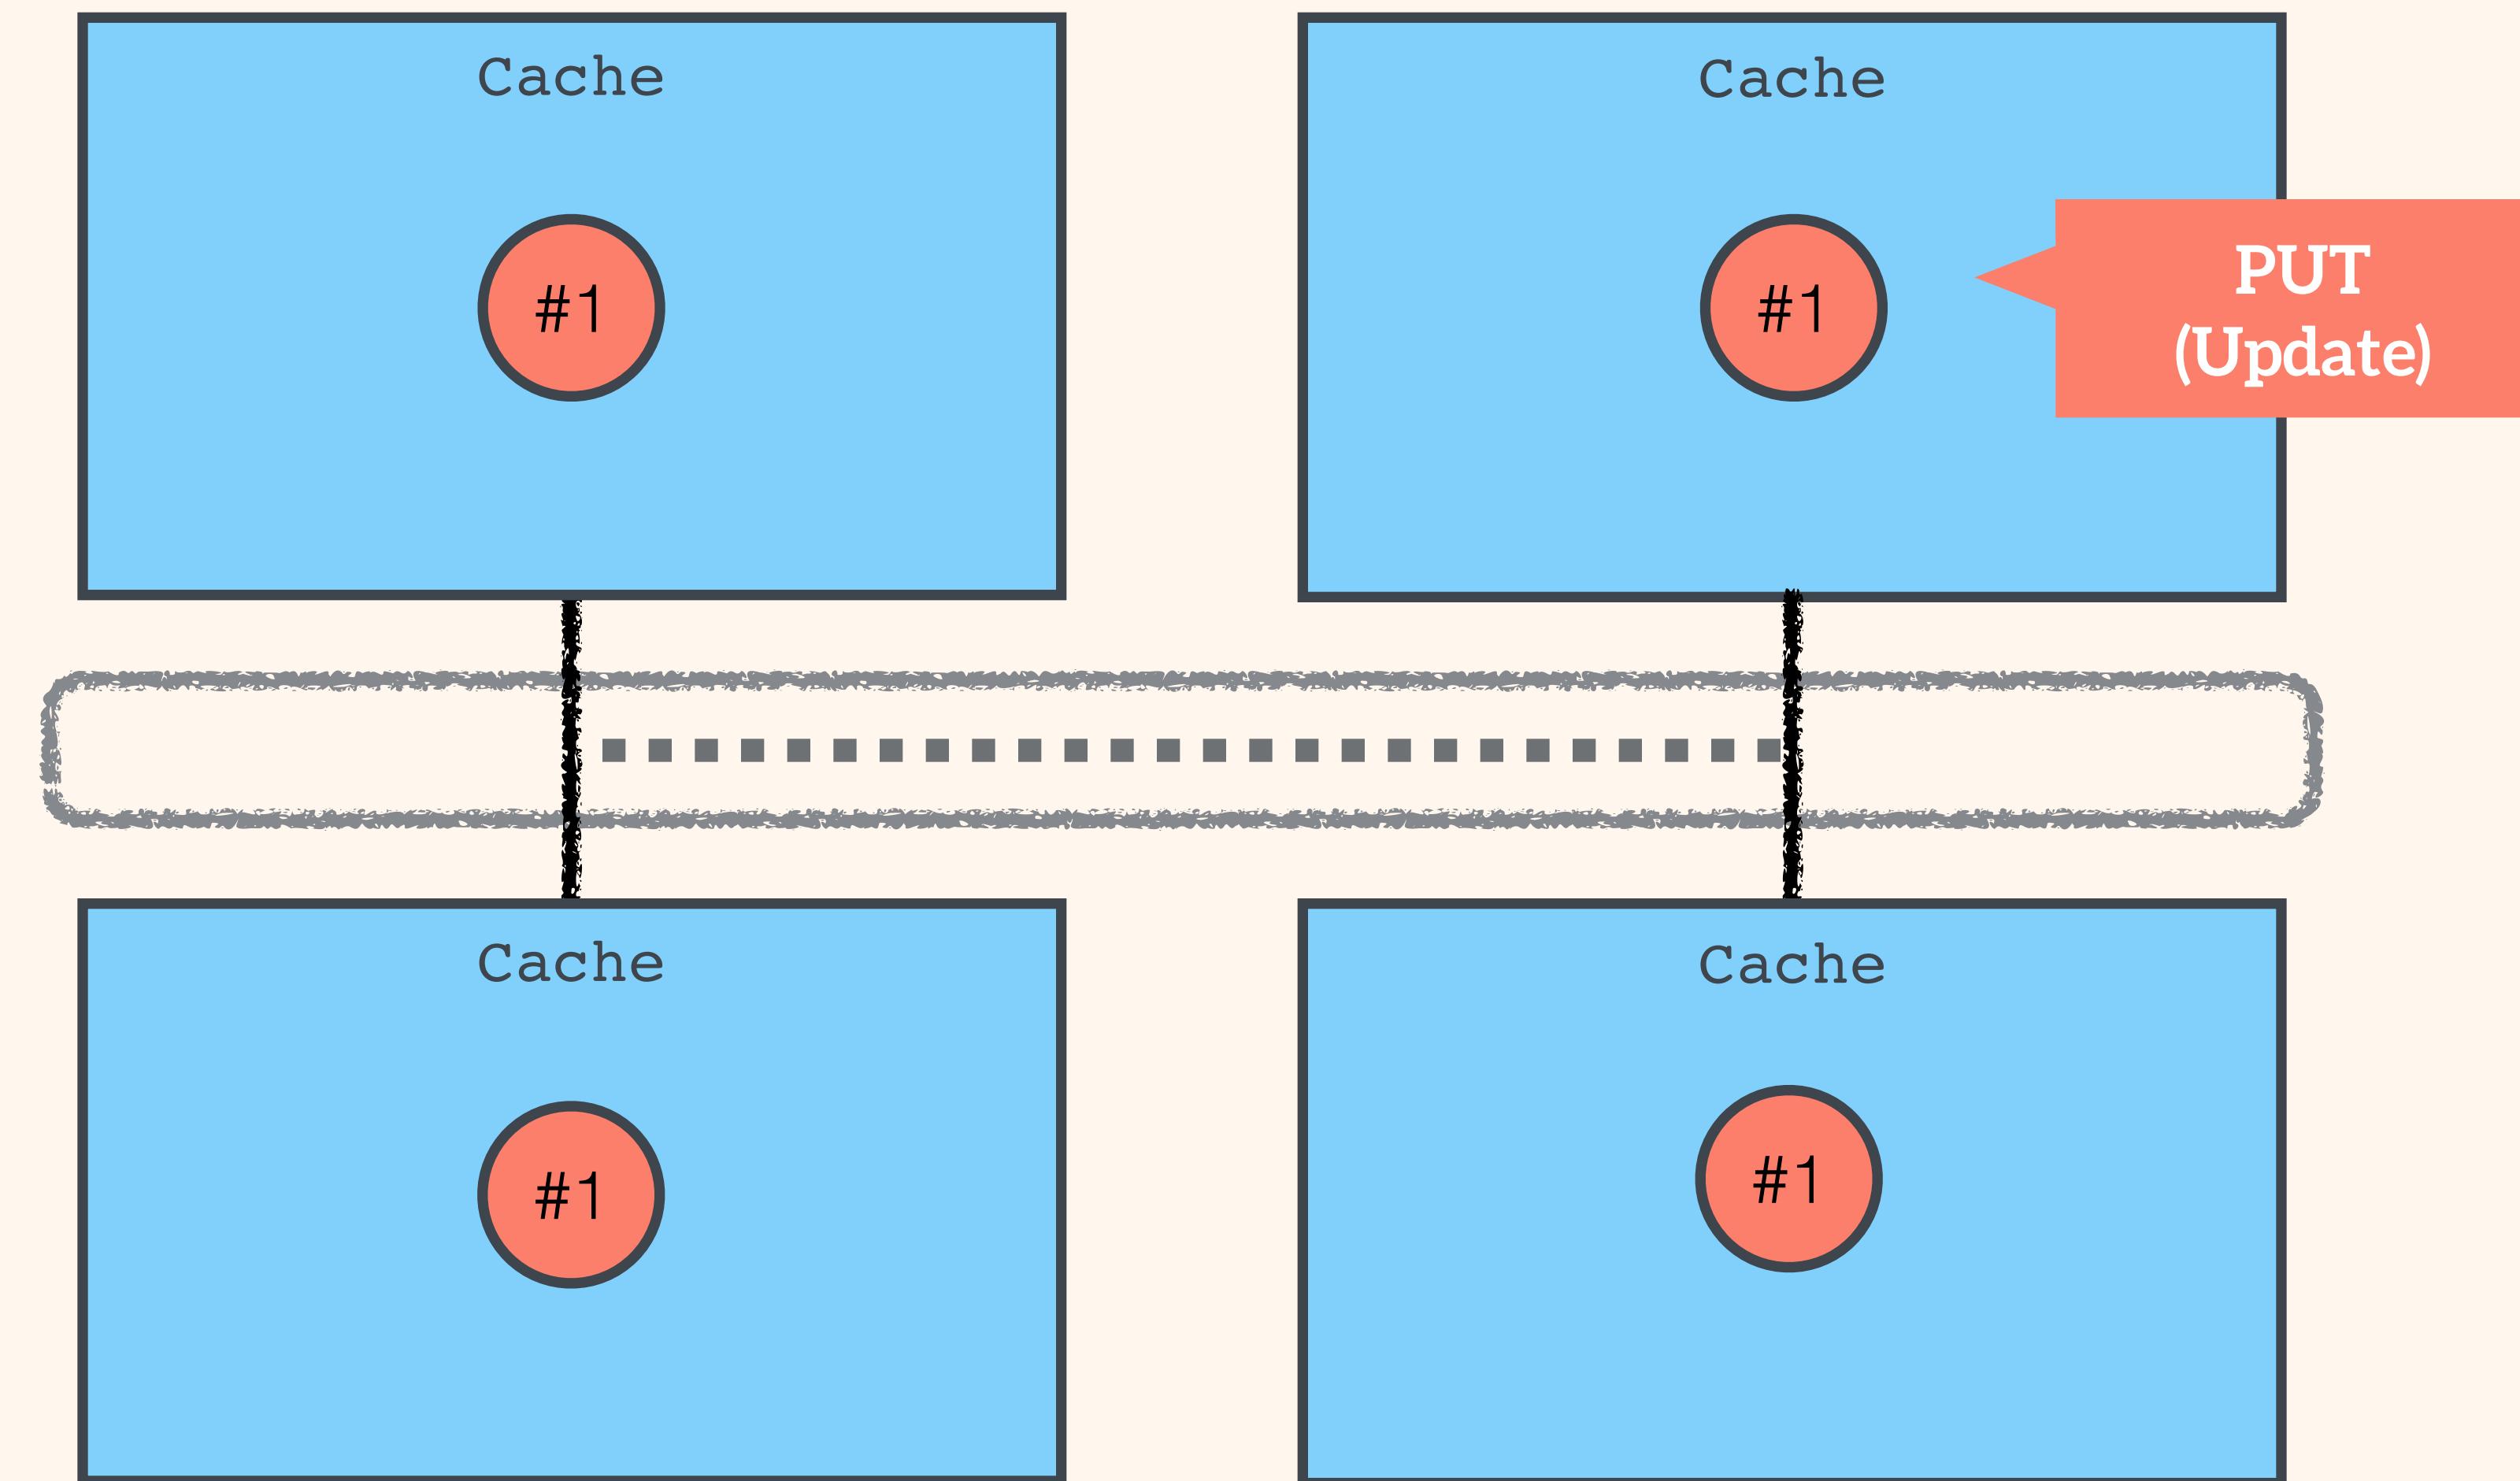
Replication



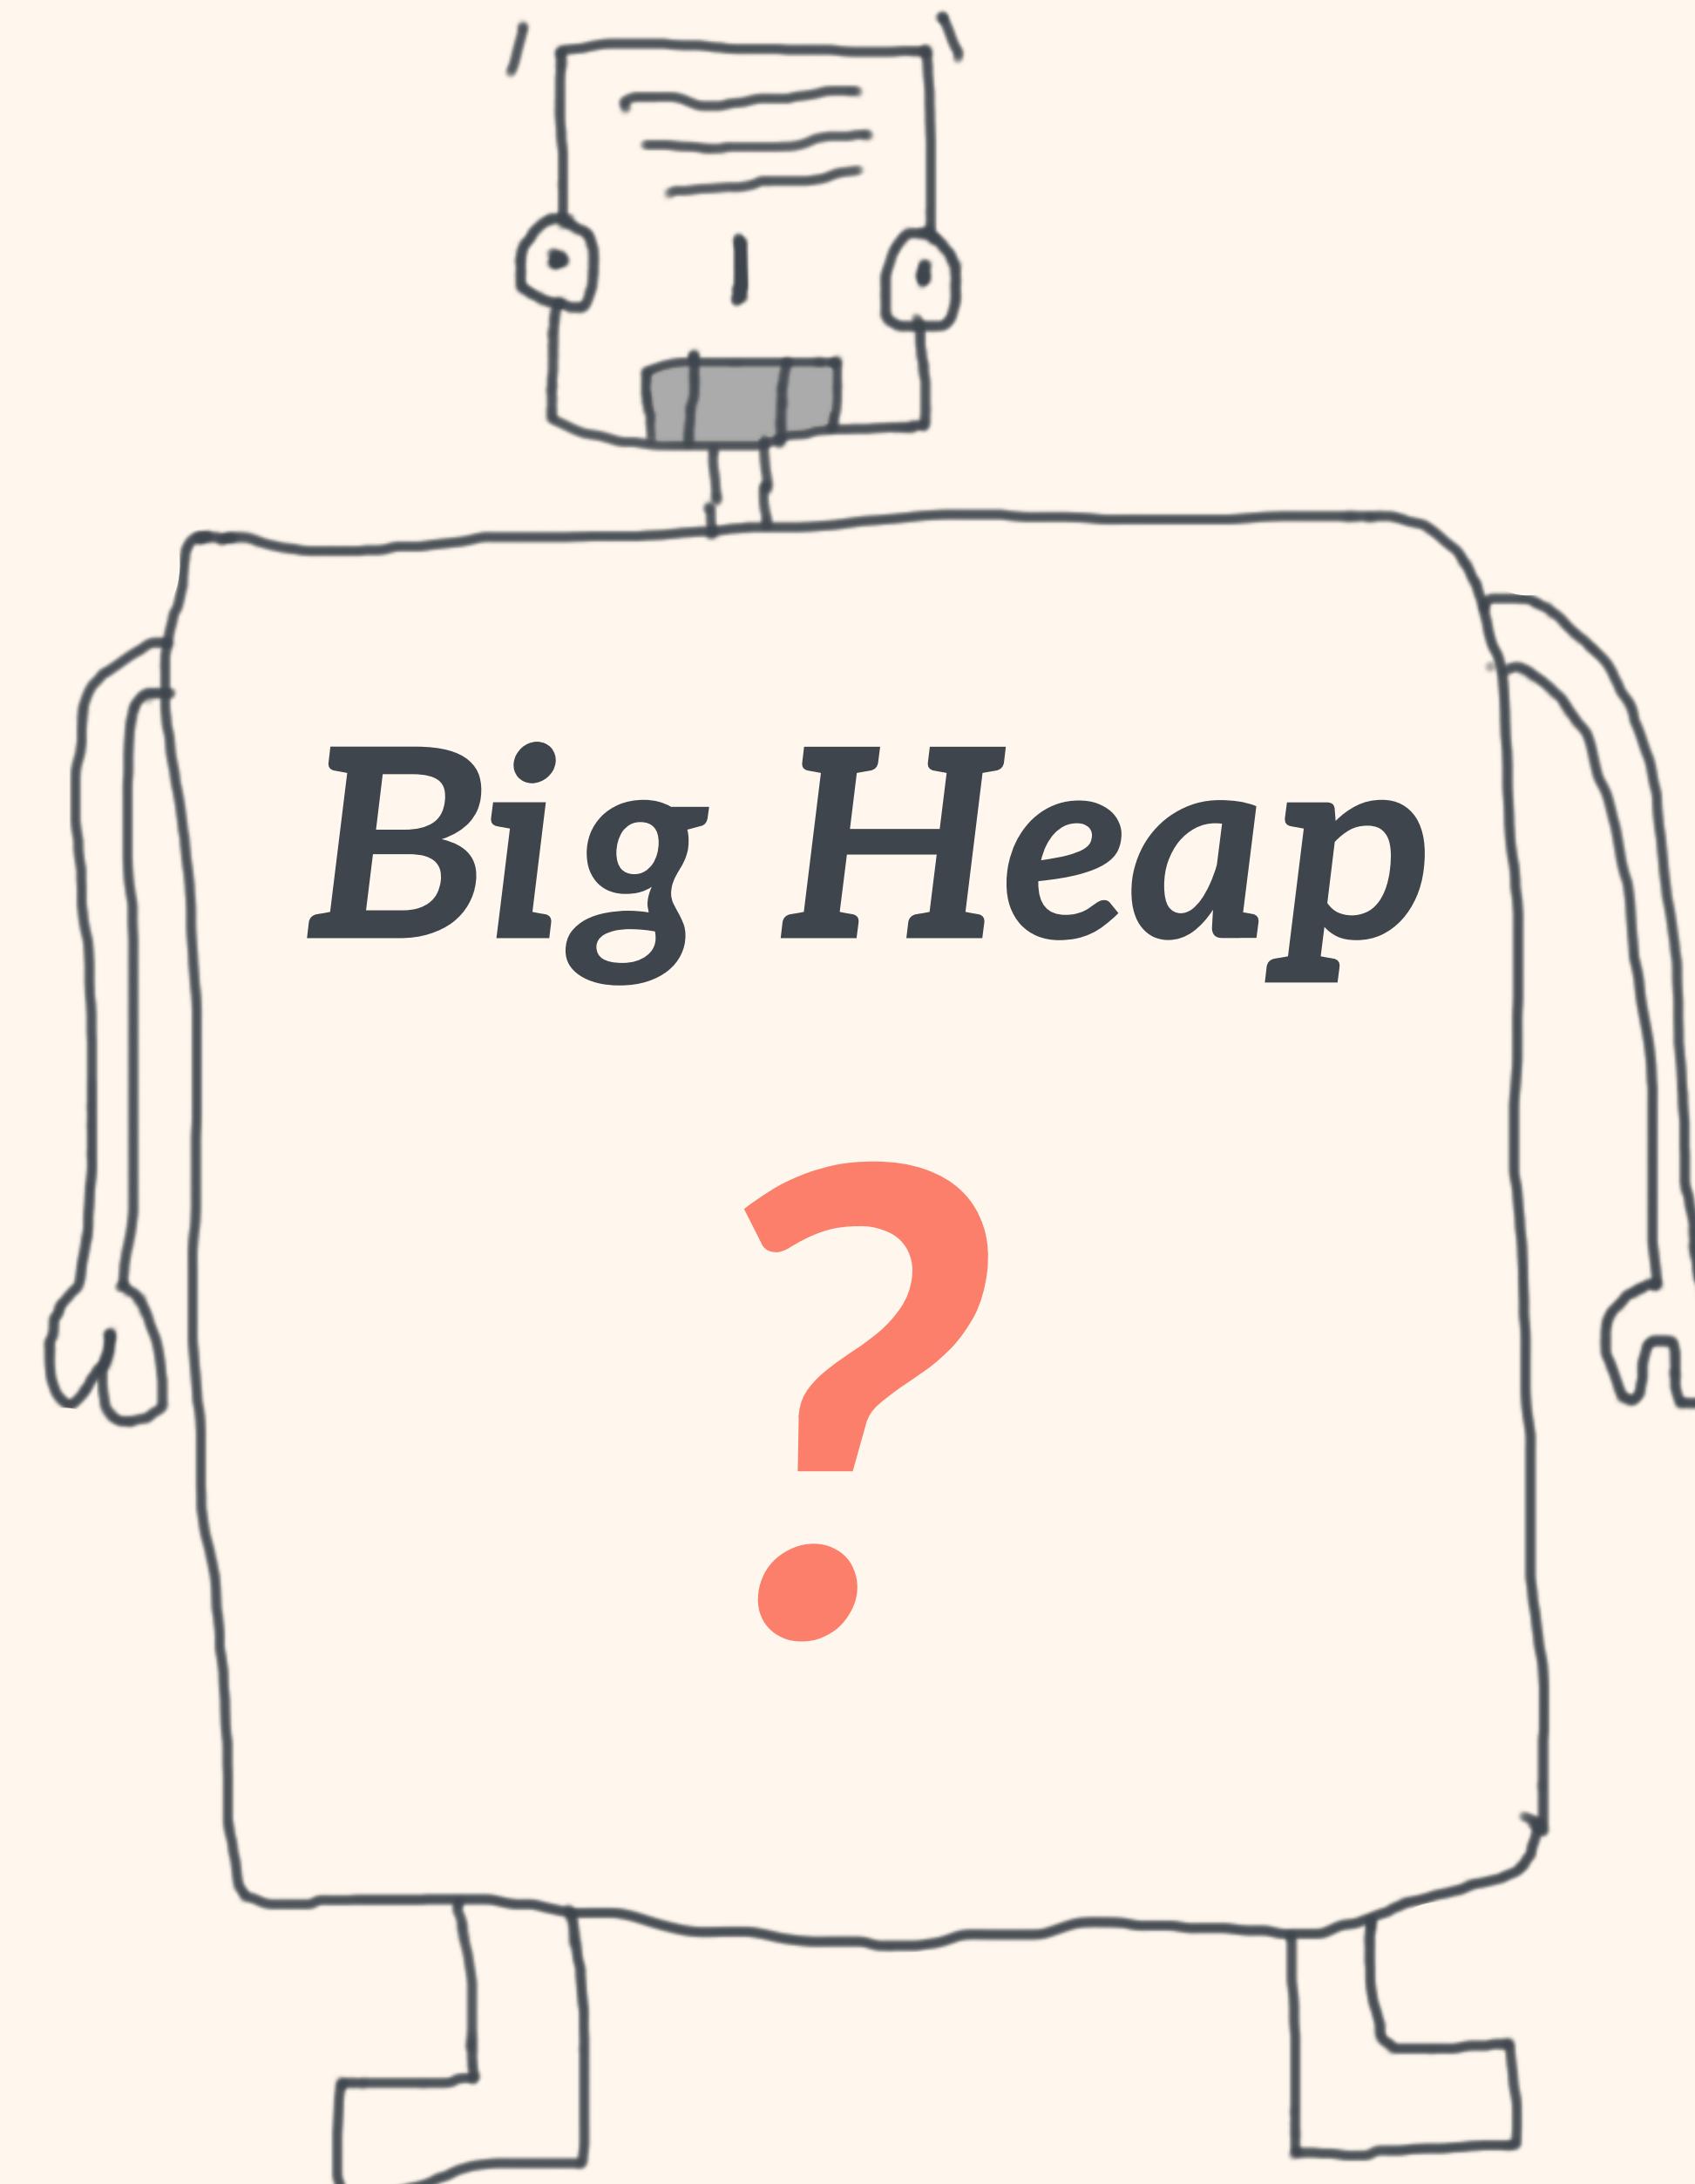
Replication



Replication

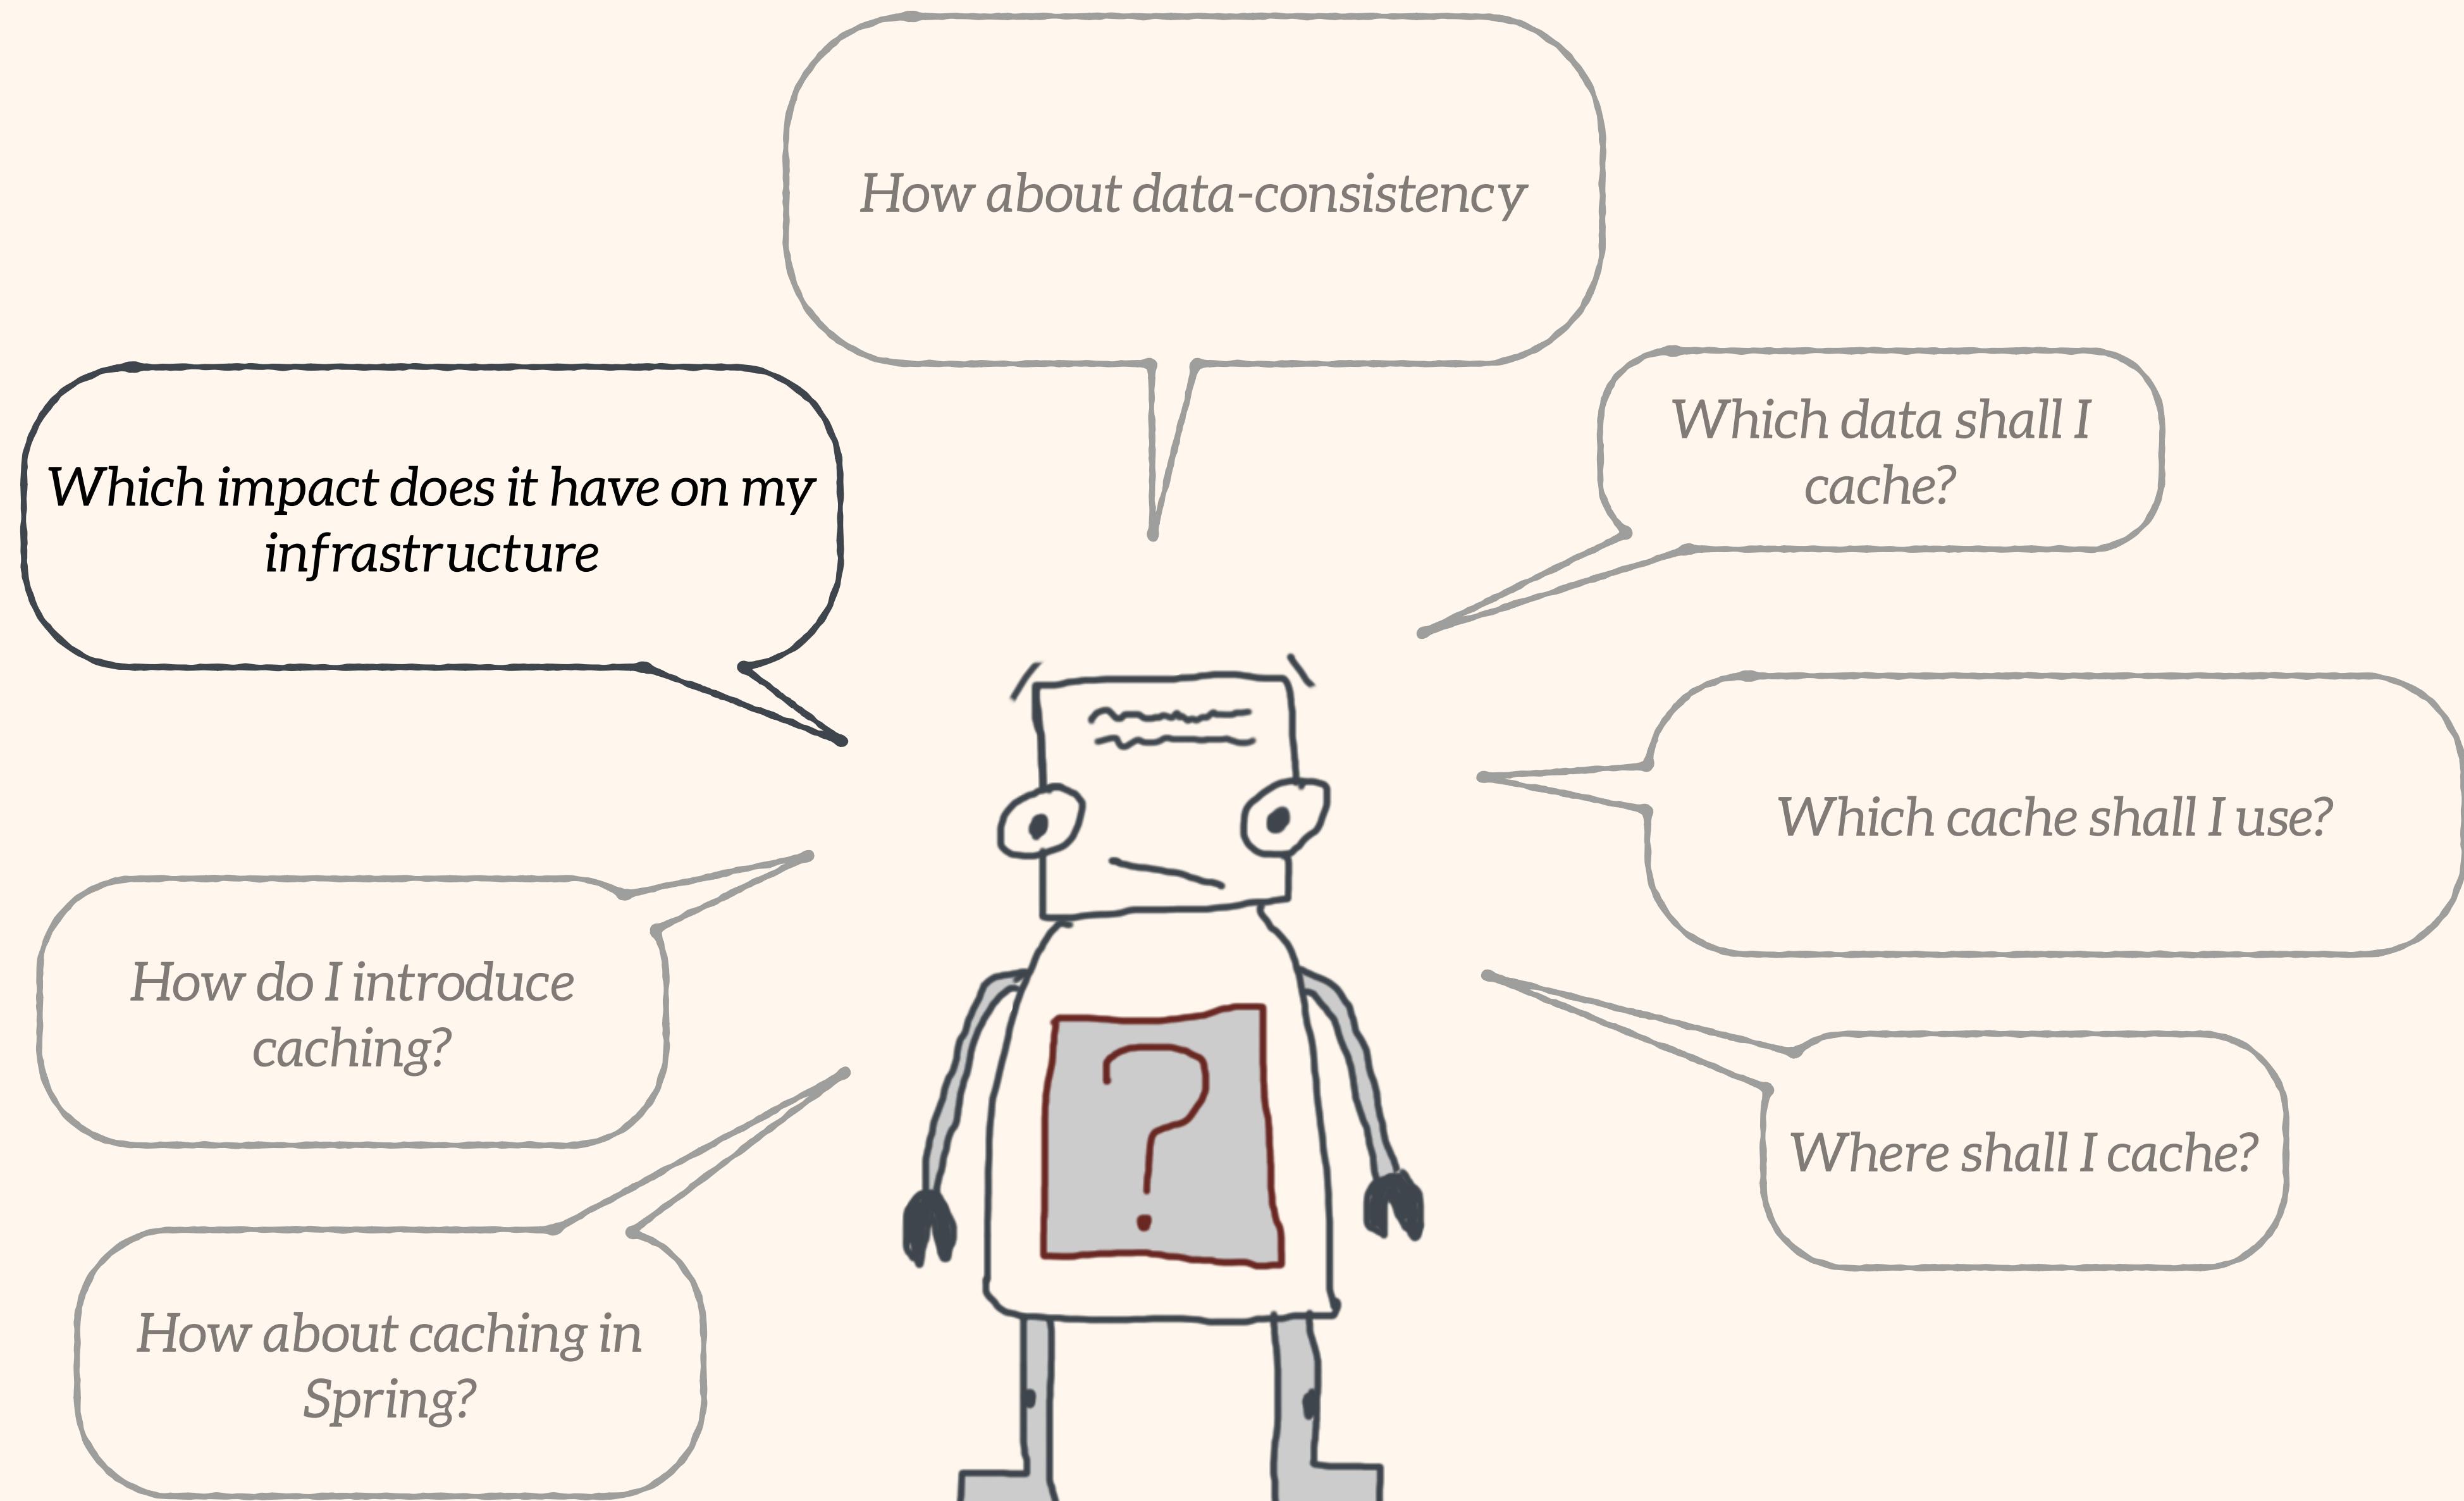


As of now every cache could potentially hold every data which consumes heap memory

A simple line-drawn cartoon character with a square head, two small ears, and a dark grey rectangular mouth. It has a single vertical line for a body and two simple legs. A speech bubble surrounds the character, containing the text "Big Heap".

Big Heap

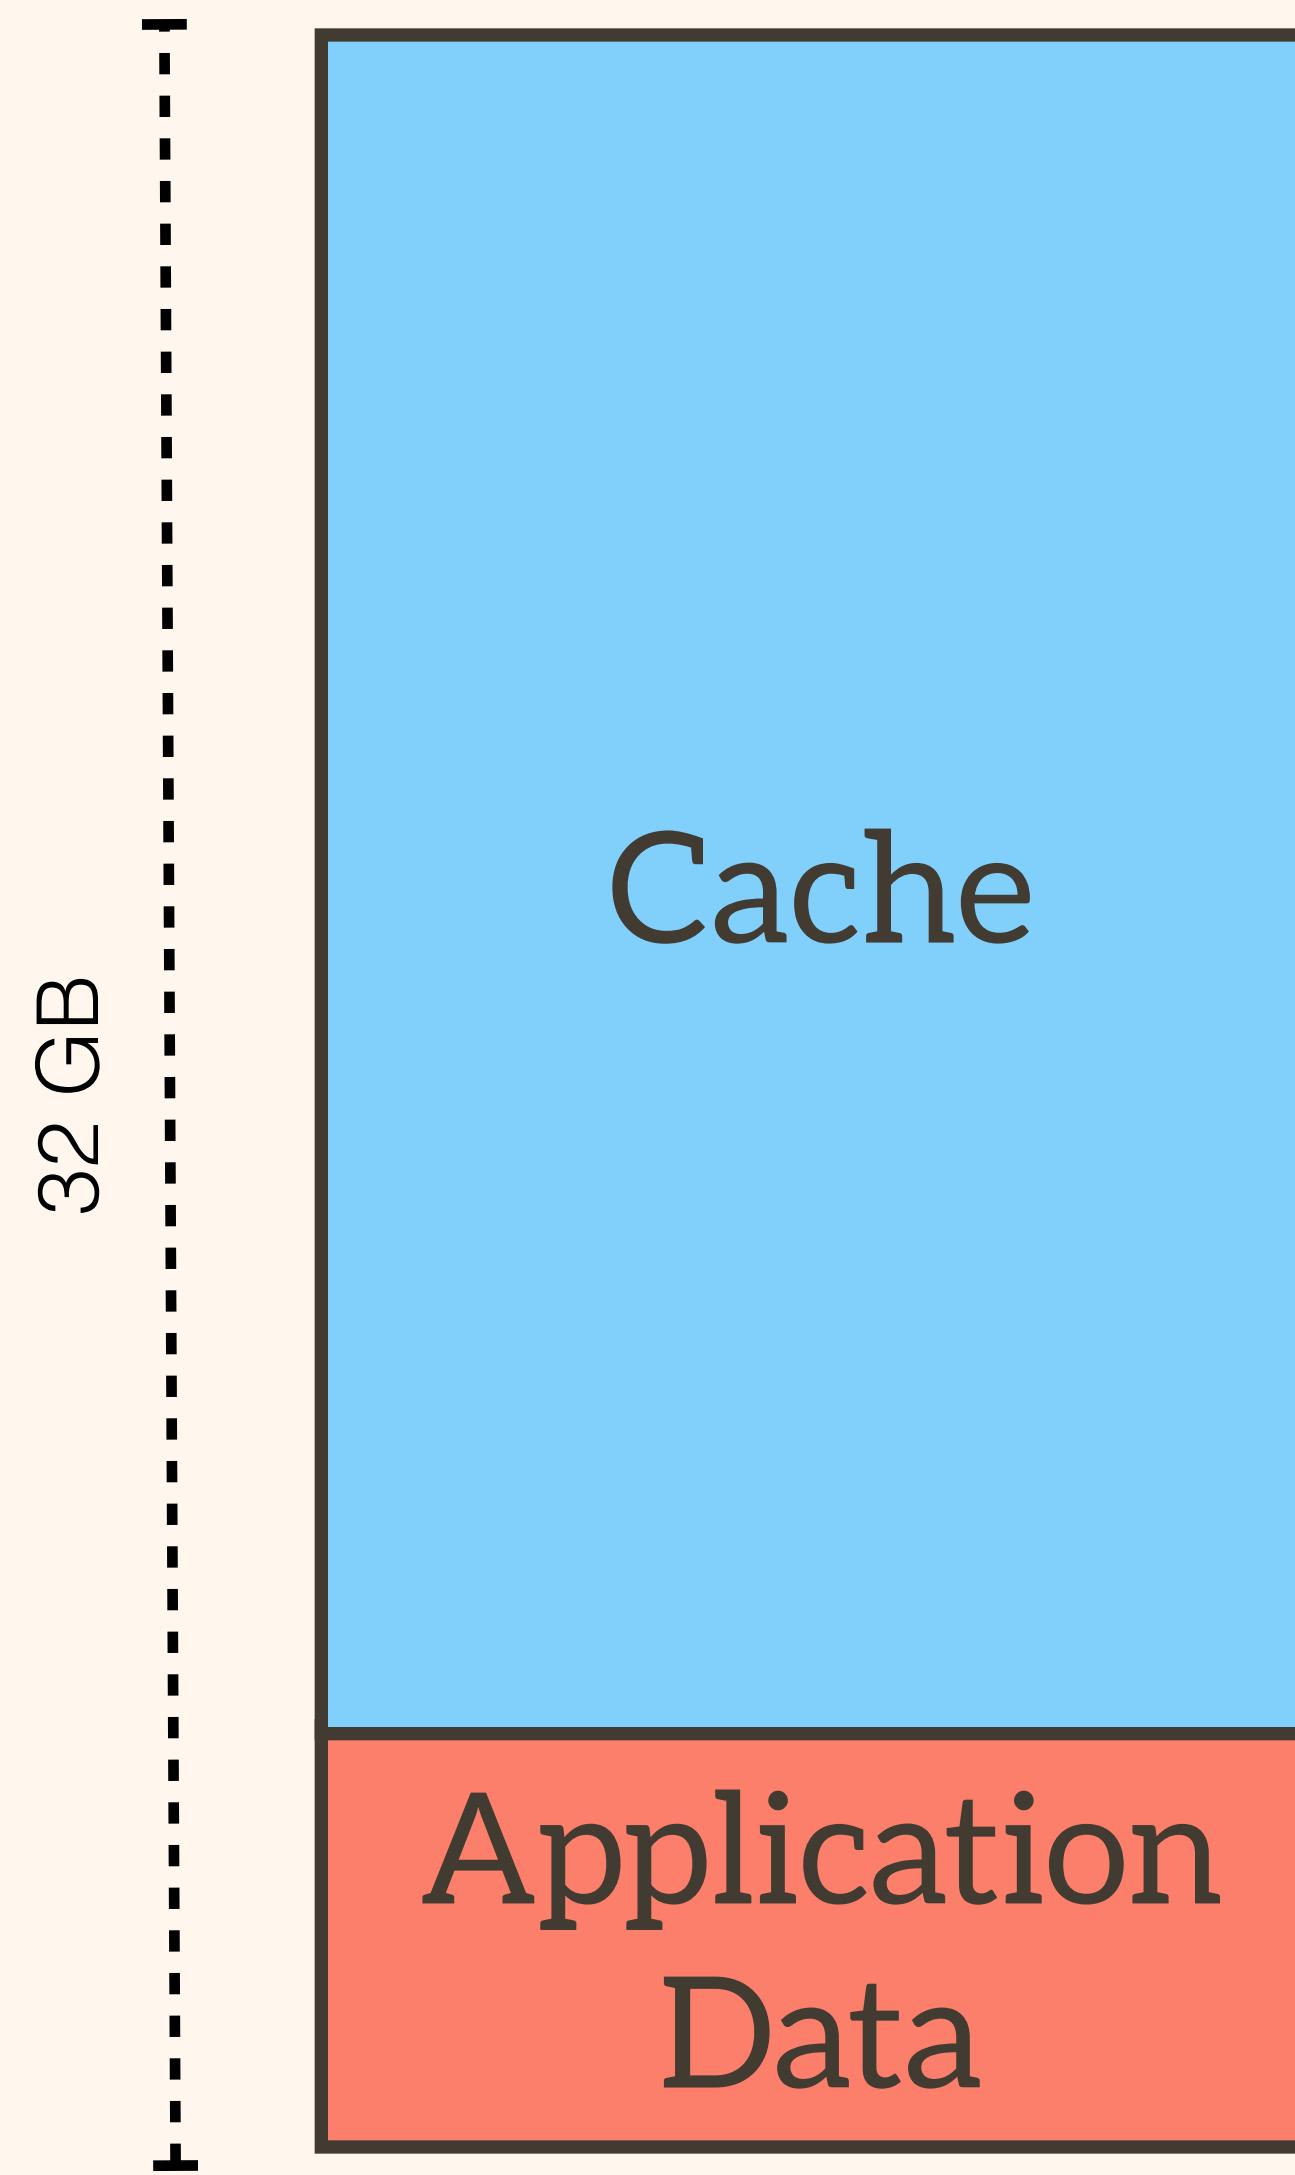
?





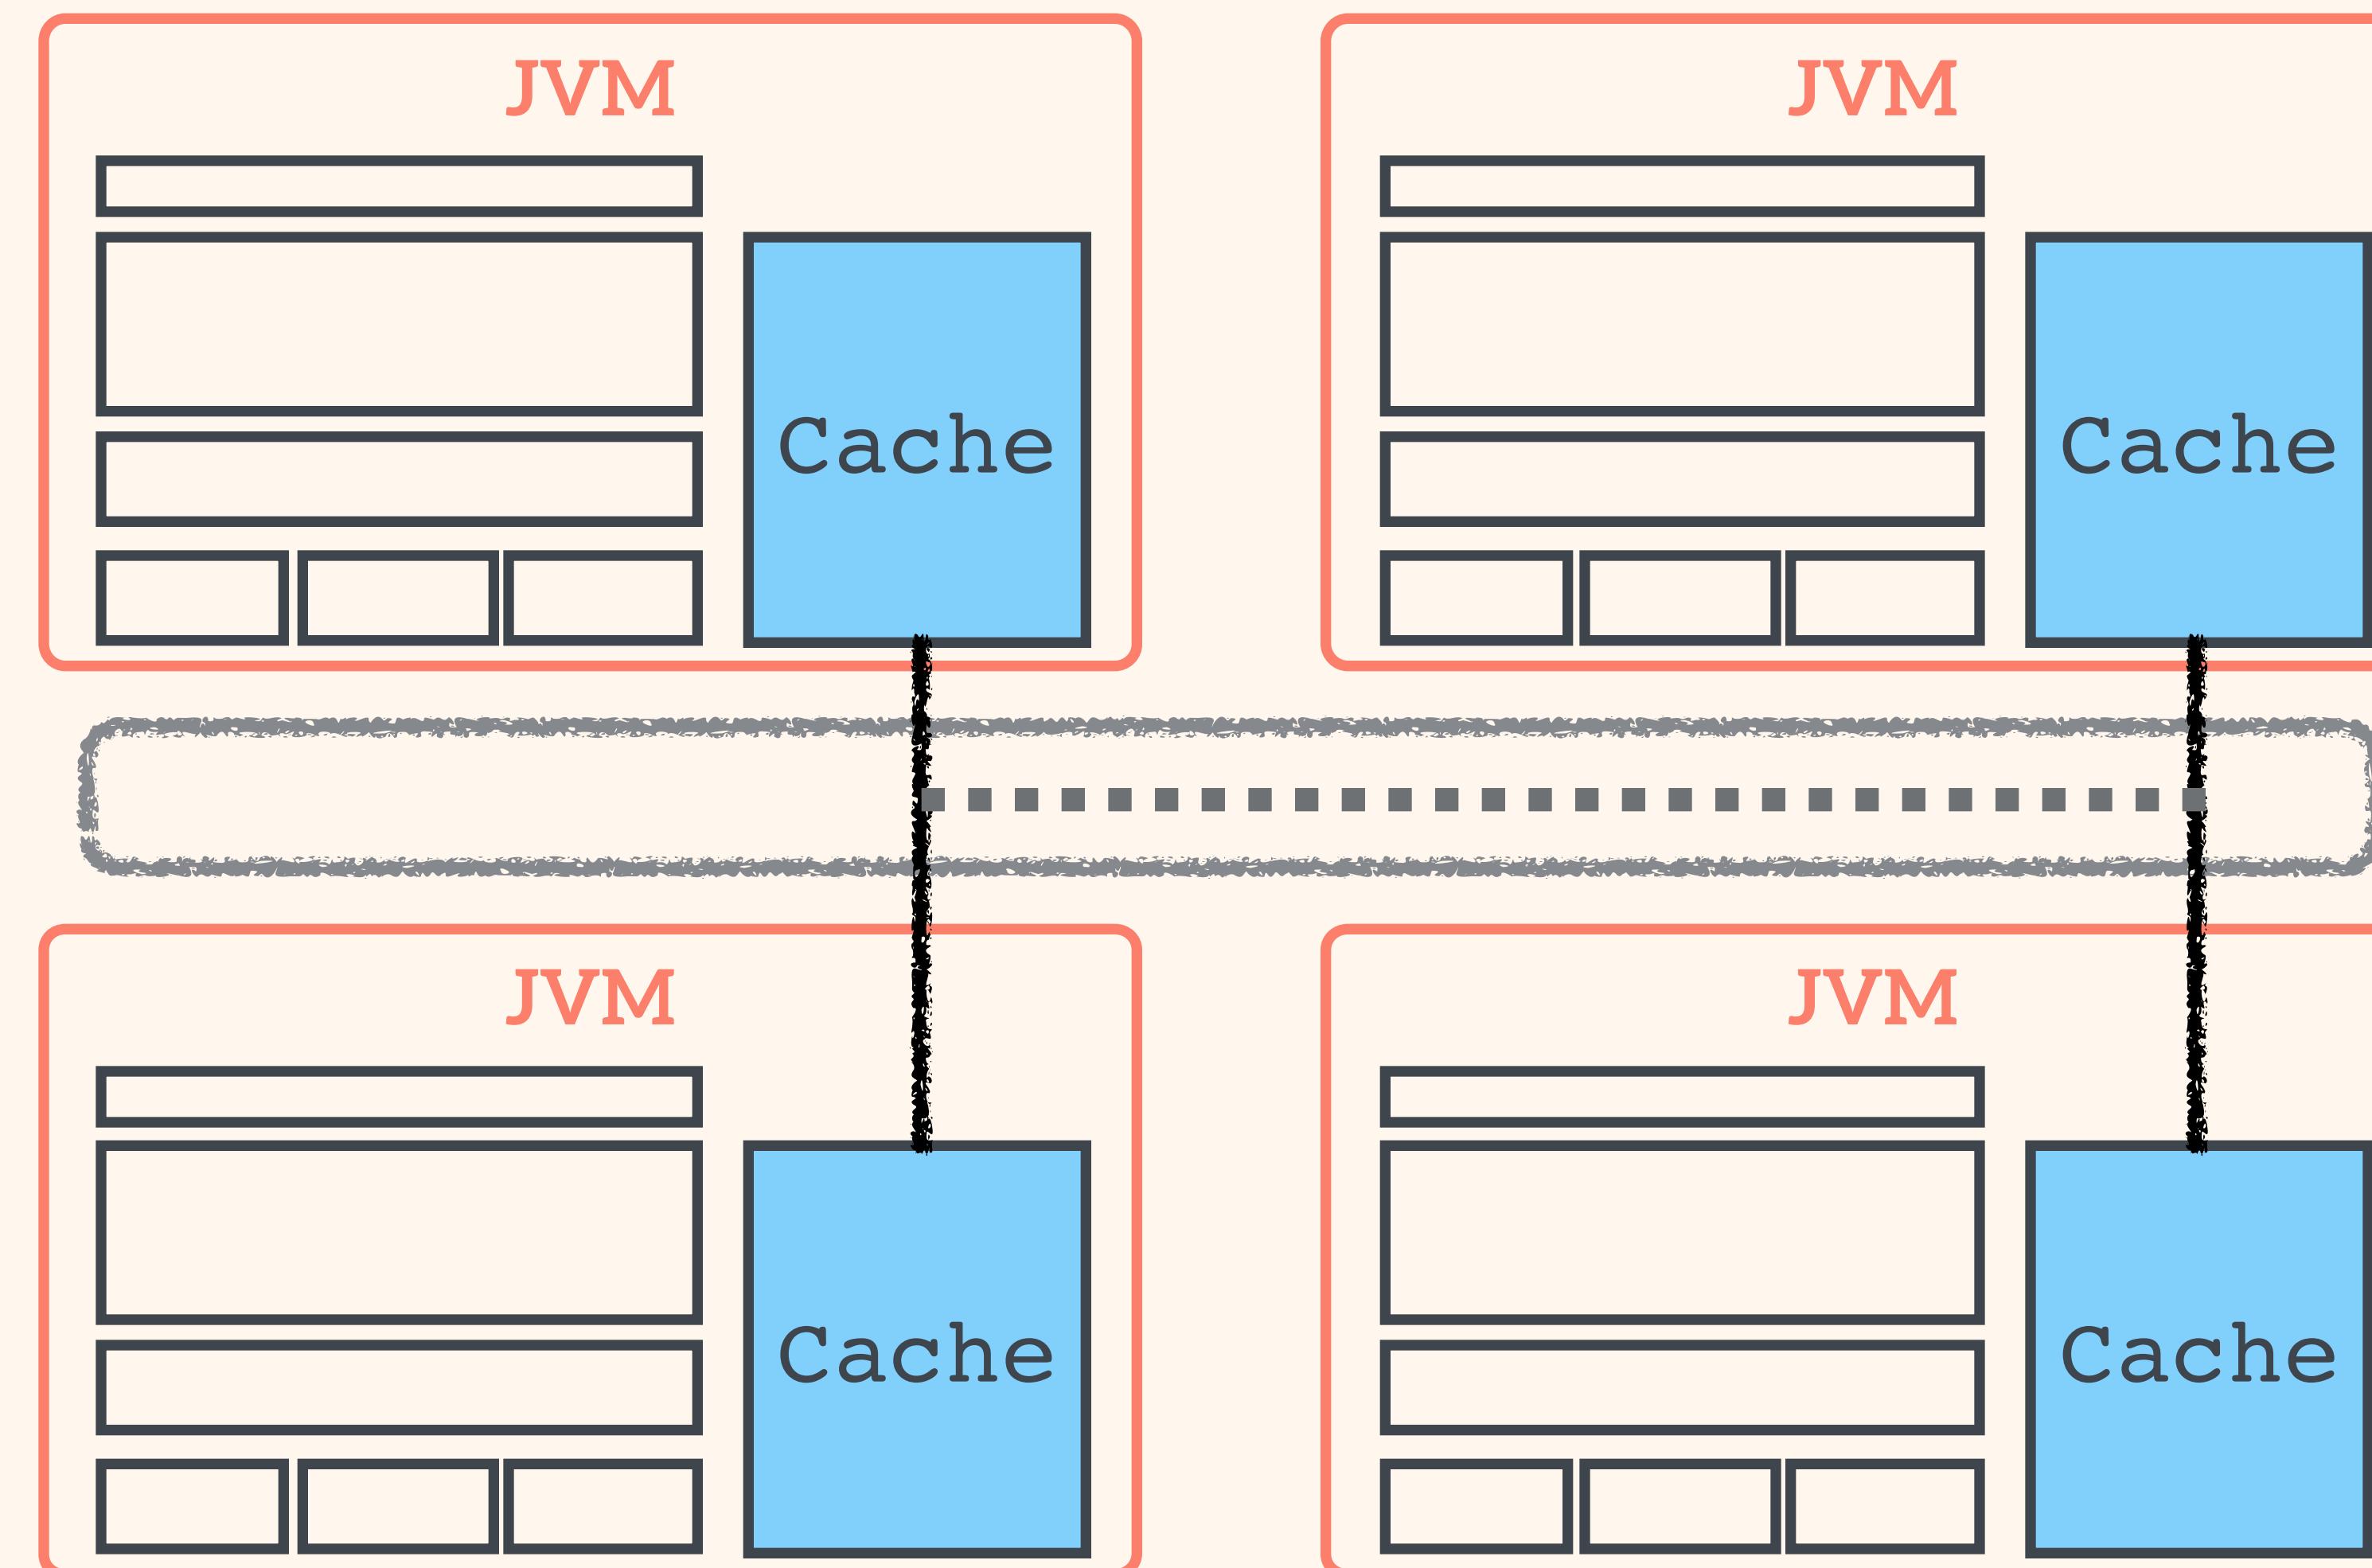
4

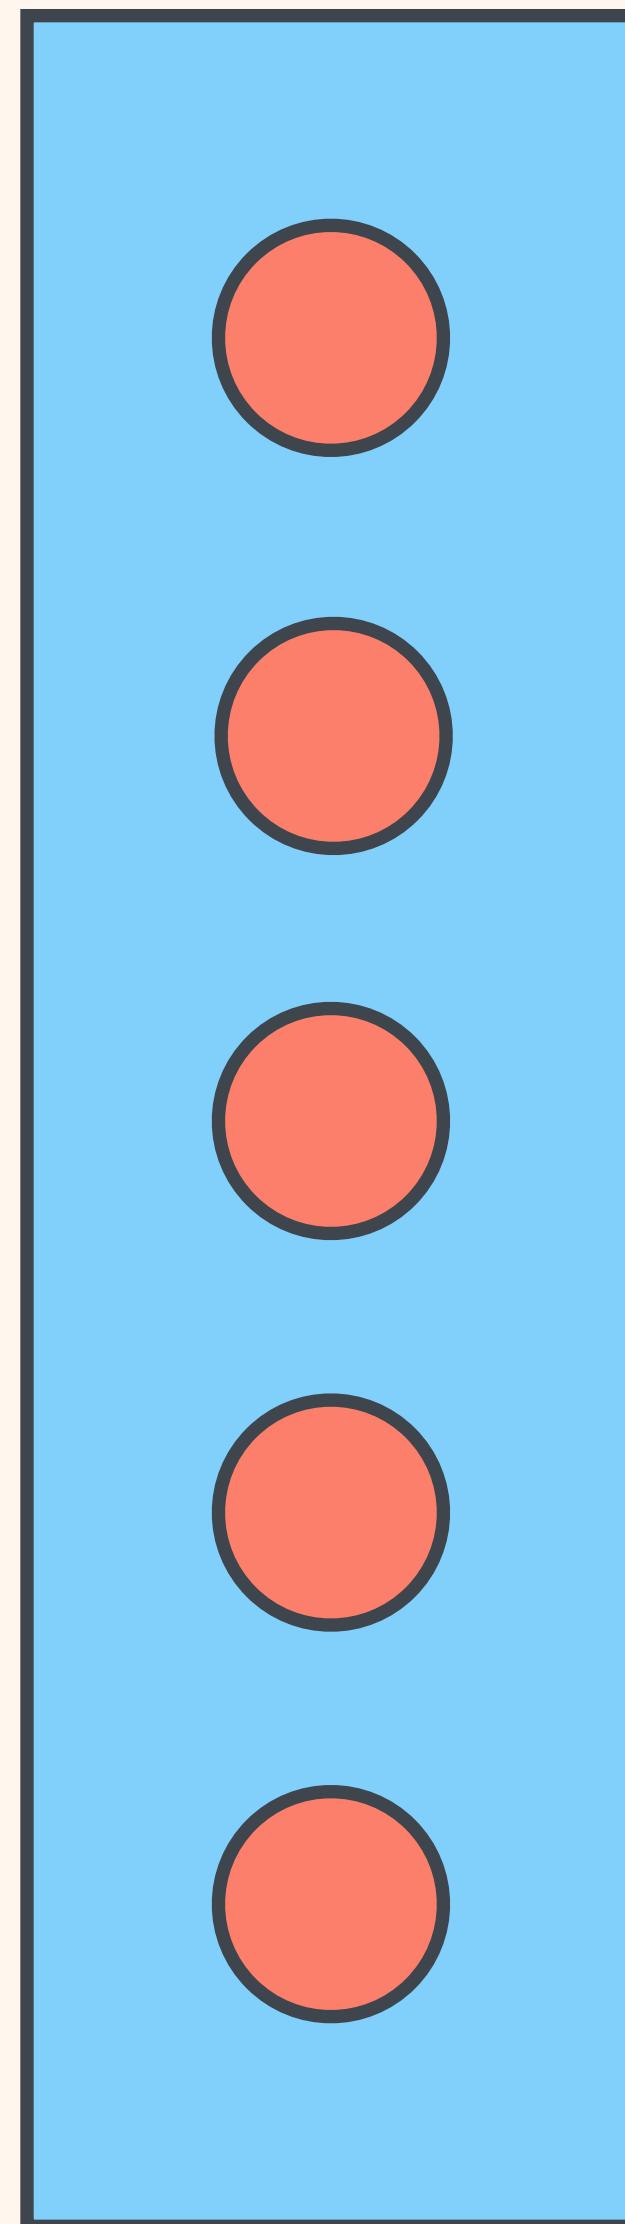
Avoid big heaps just for caching



*Big heap
leads to long
major GCs*

Long GCs can destabilize your cluster





***Small caches
are a bad idea!***

Many evictions, fewer hits,
no „hot data“.

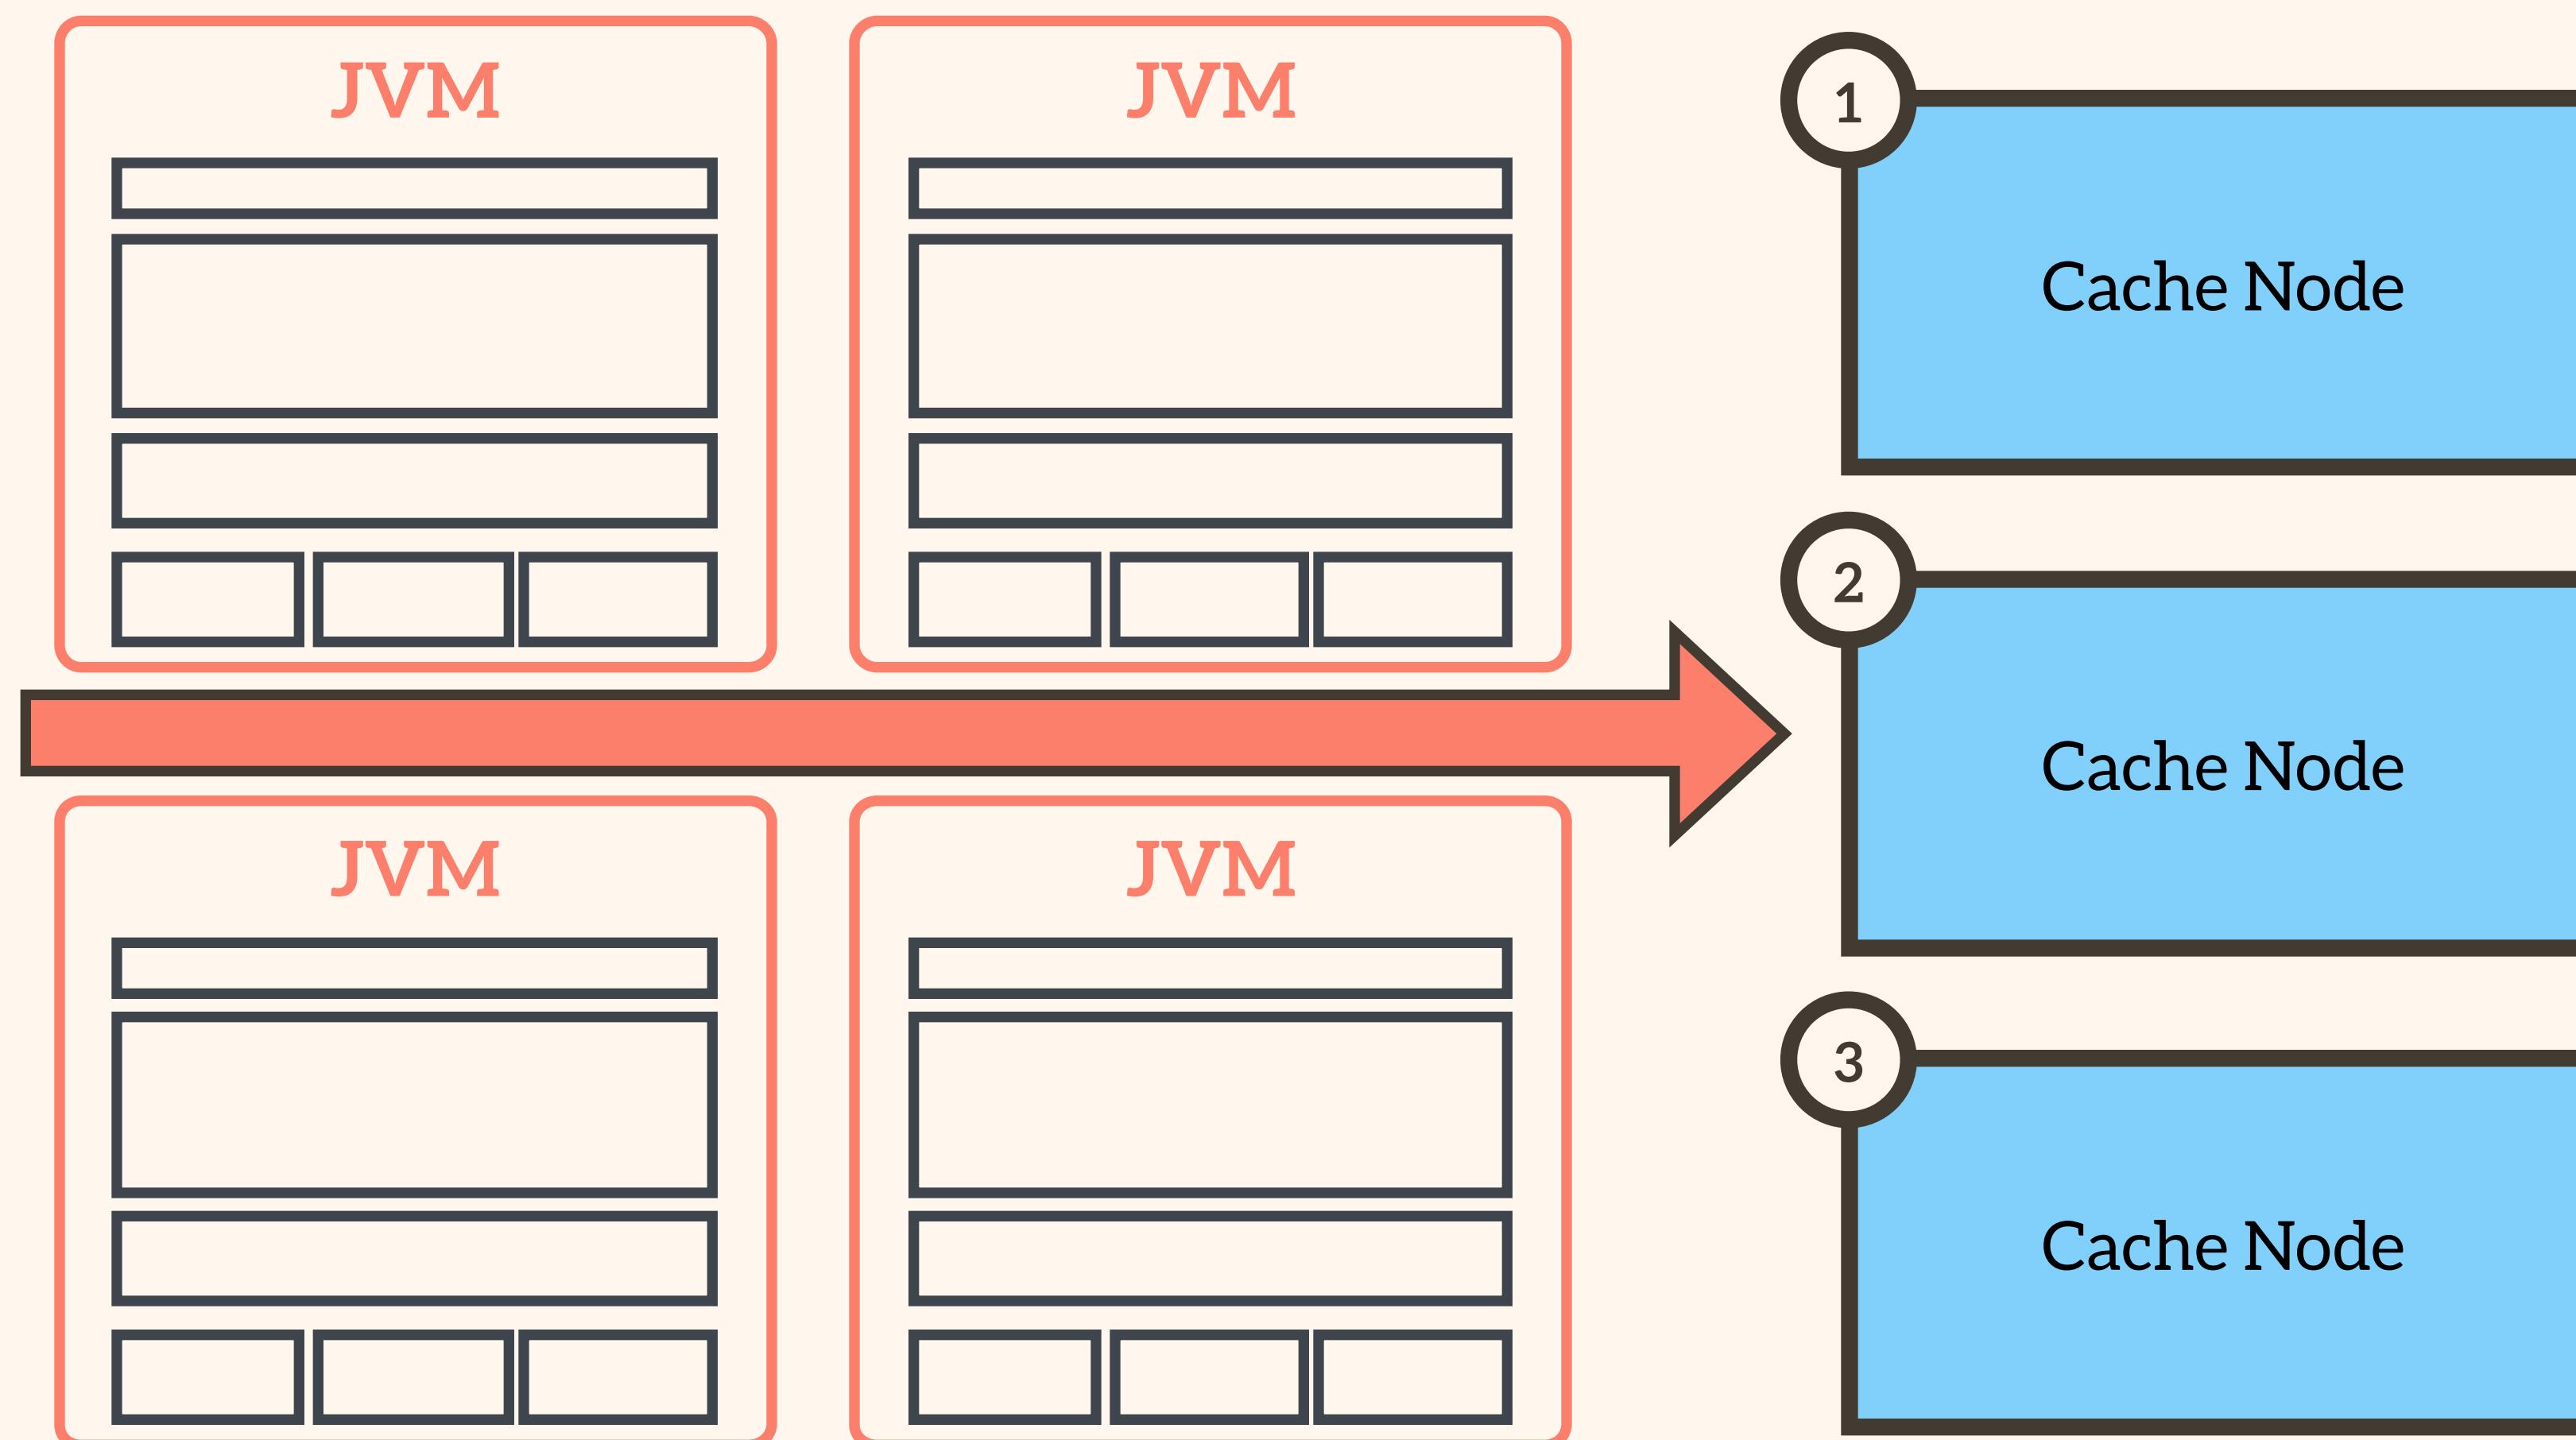
This is especially critical for
replicating caches.



5

*Use a distributed cache for
big amounts of data*

Distributed Caches



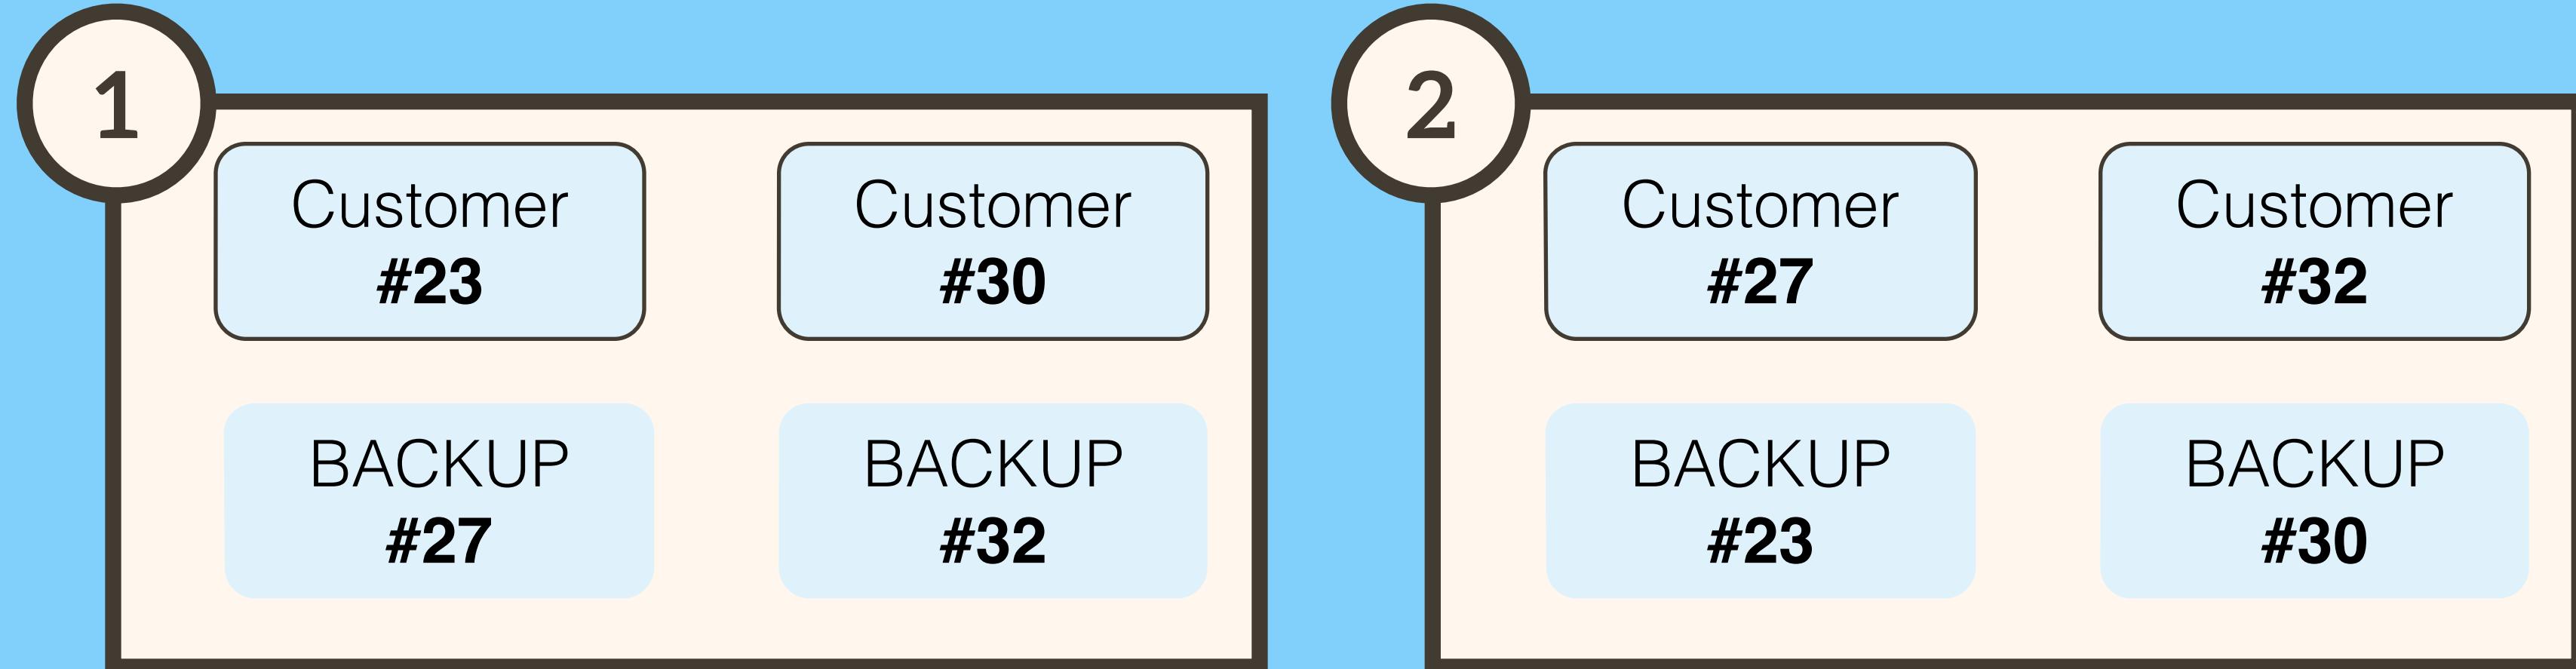
1

Customer
#23

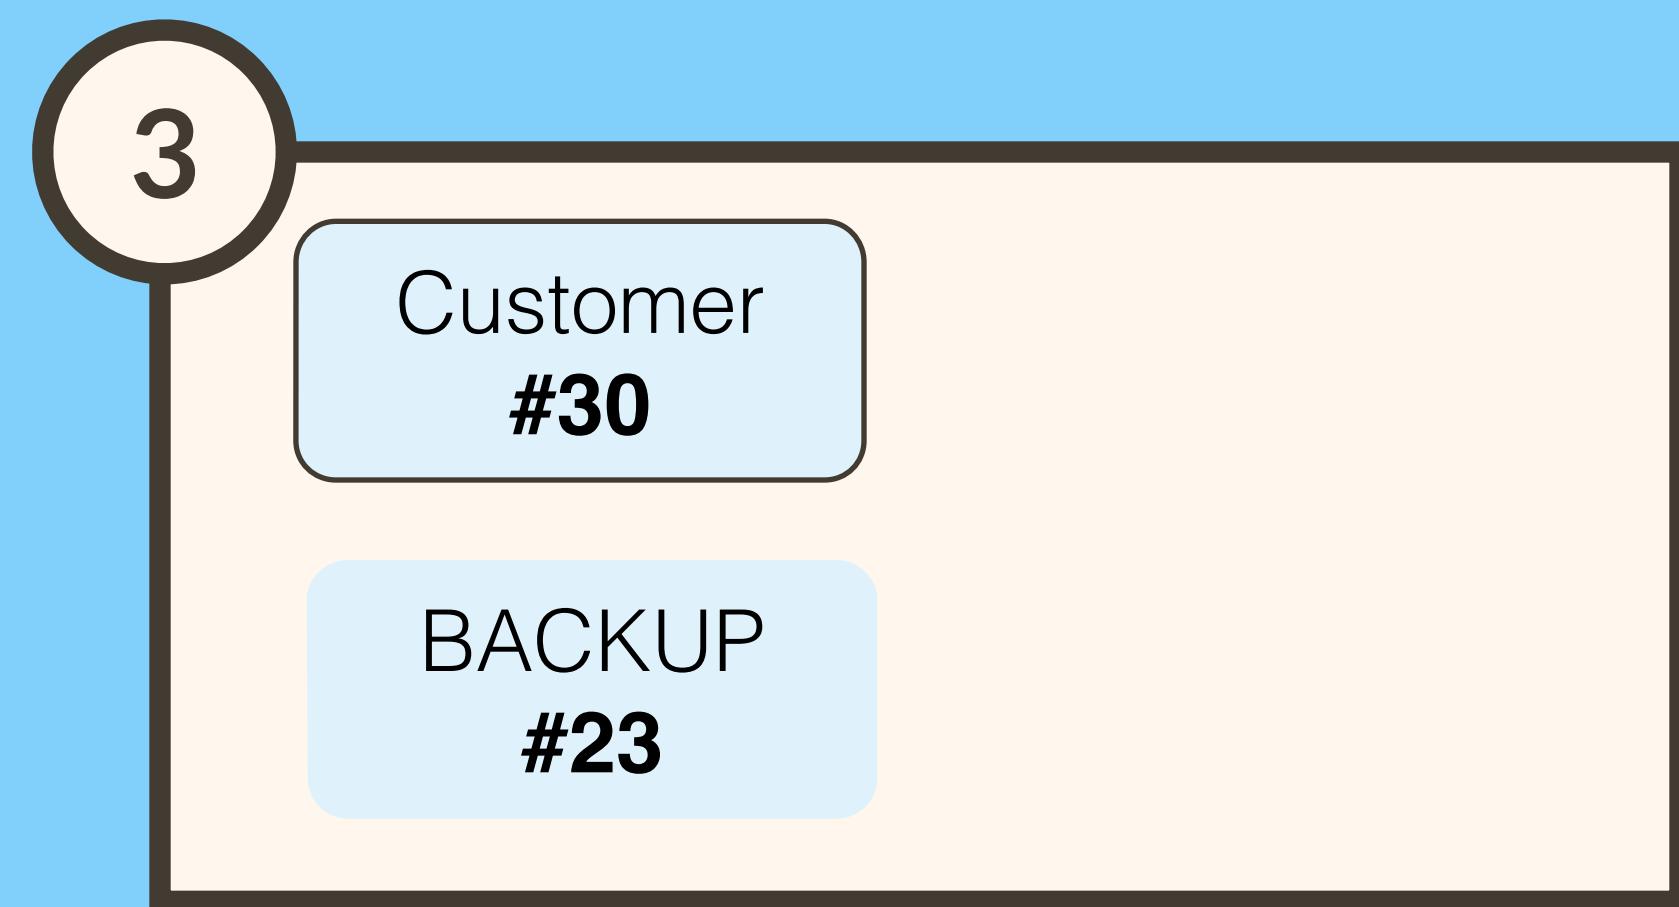
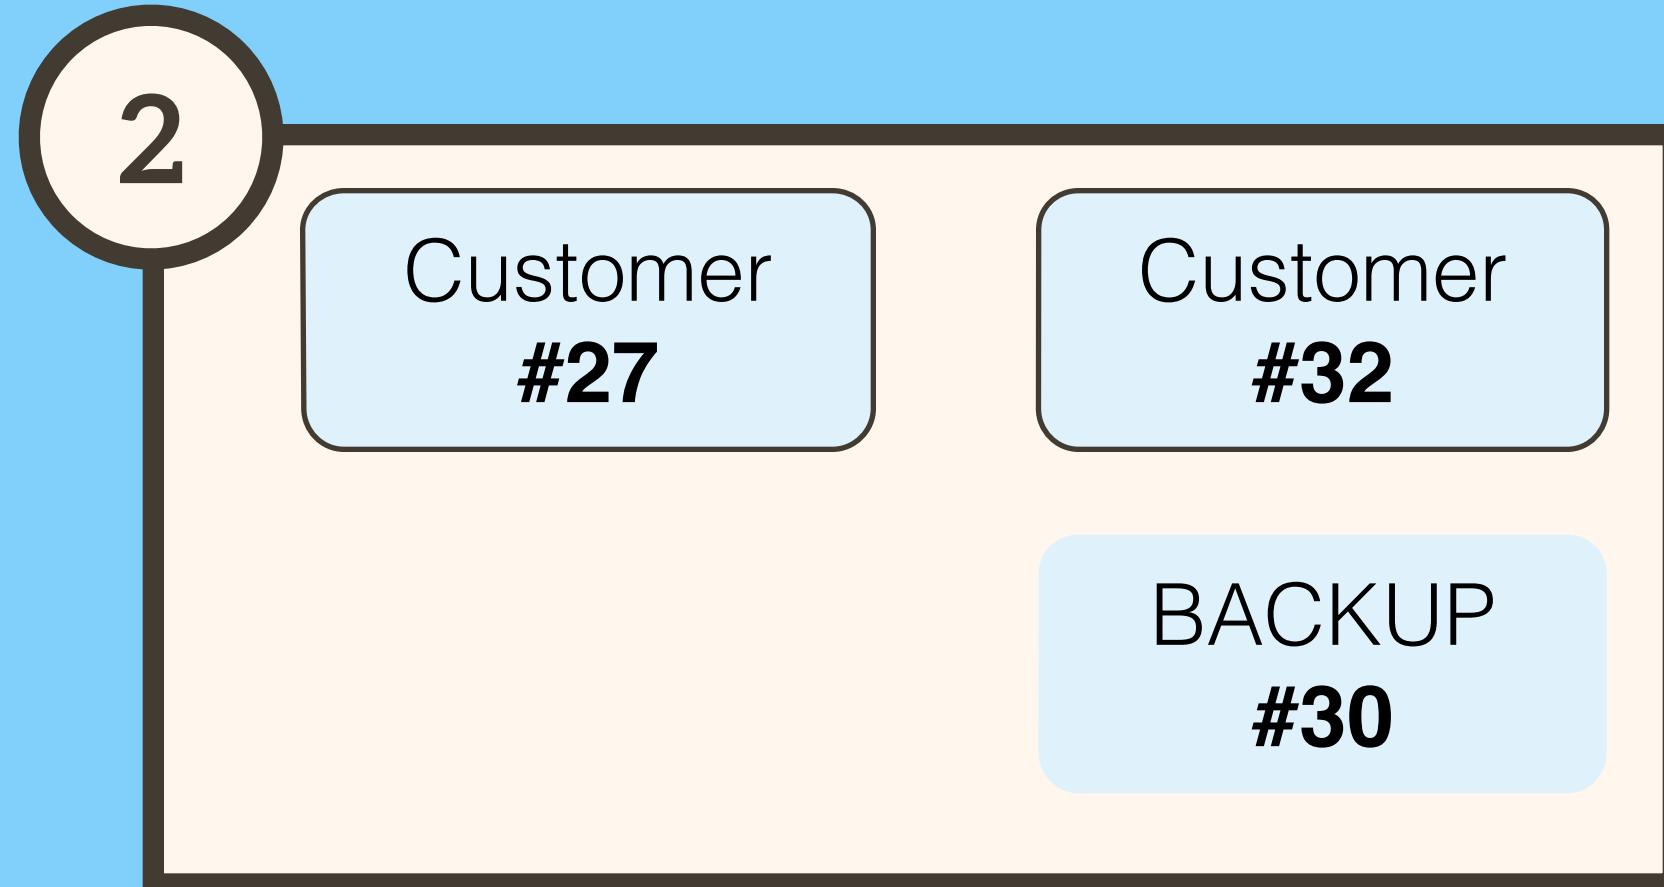
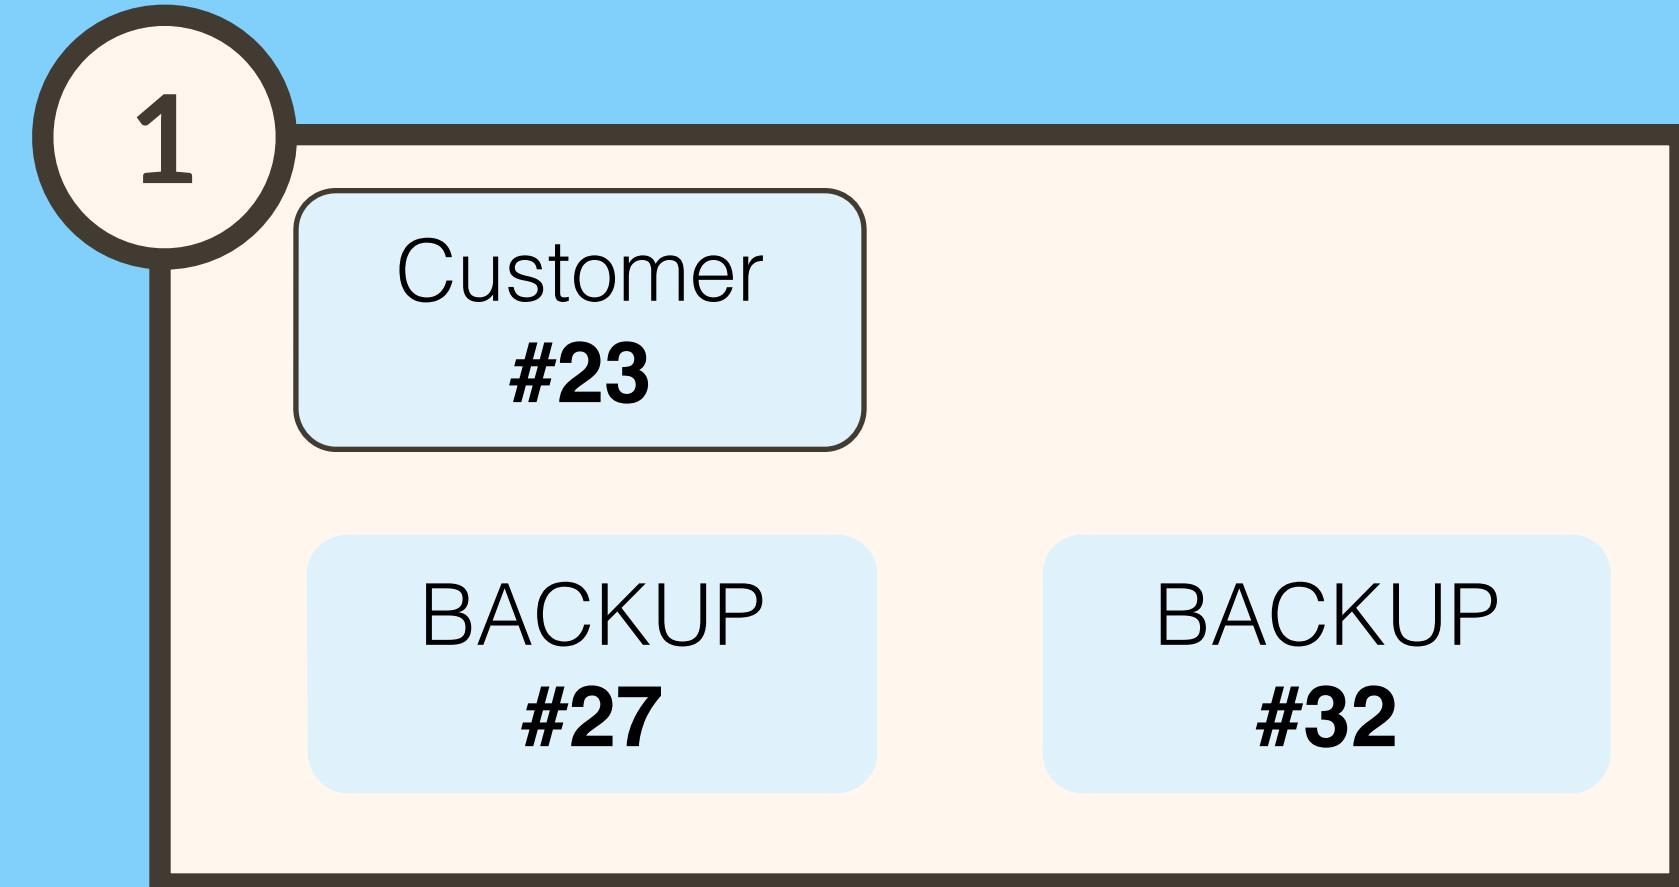
Customer
#30

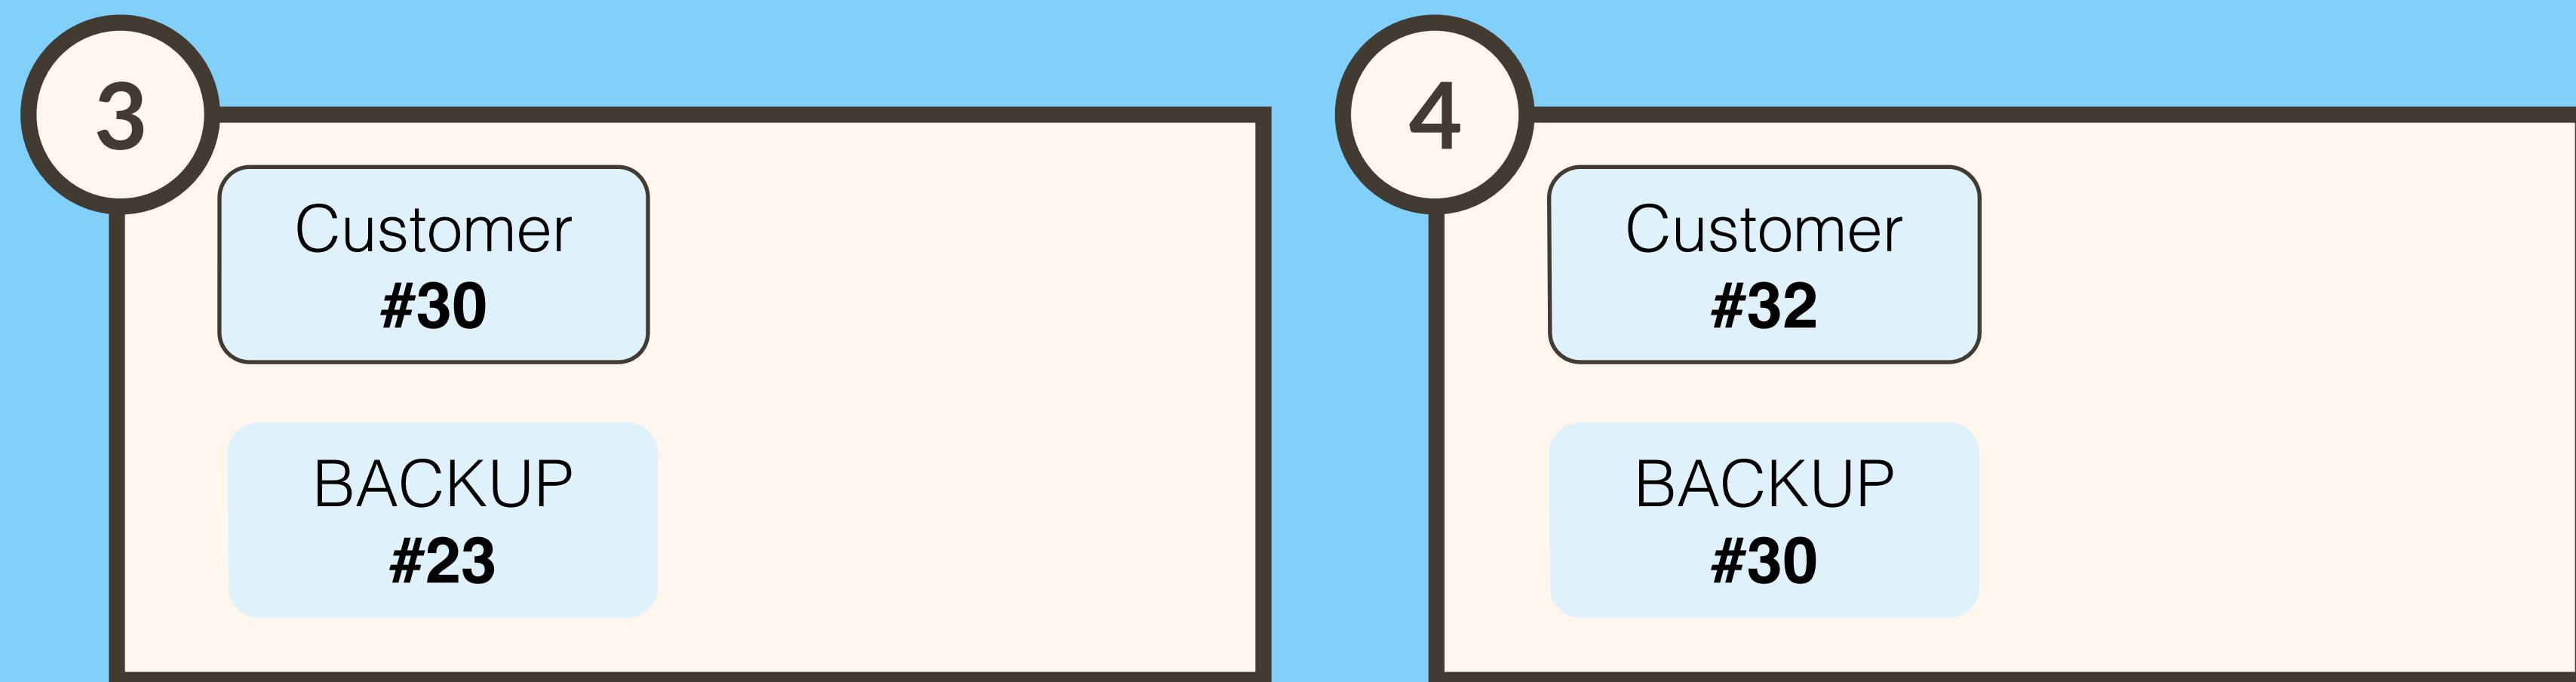
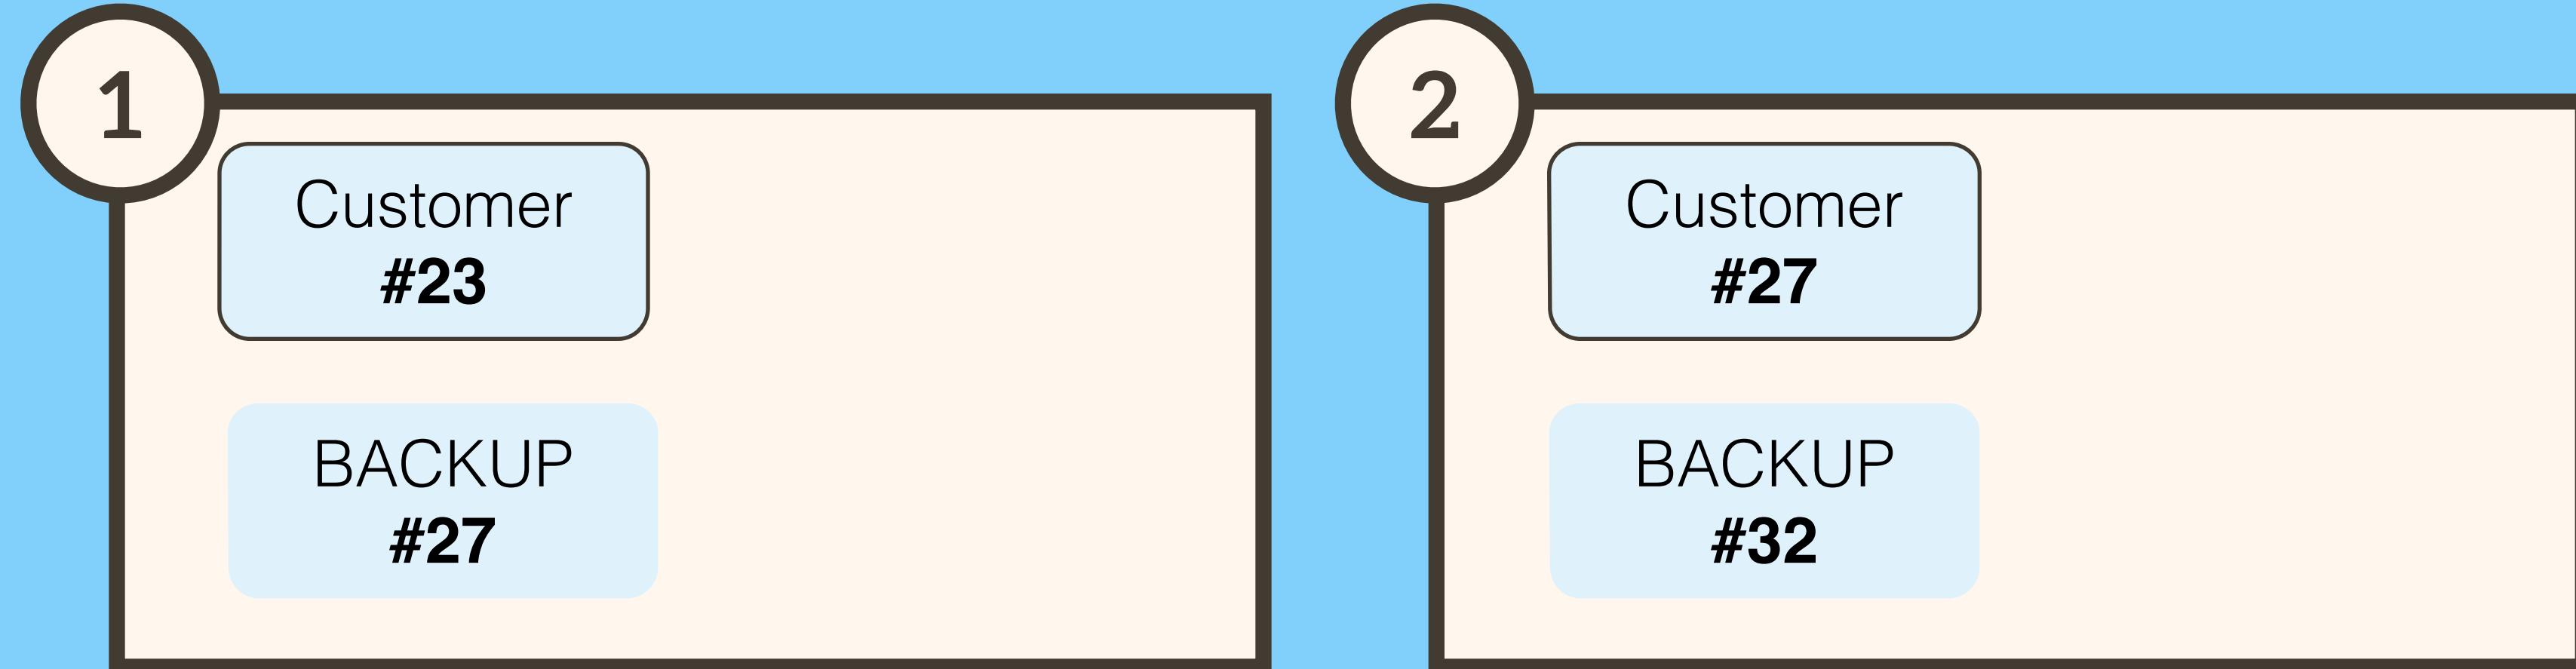
Customer
#27

Customer
#32



*Data is being
distributed and
backed up*

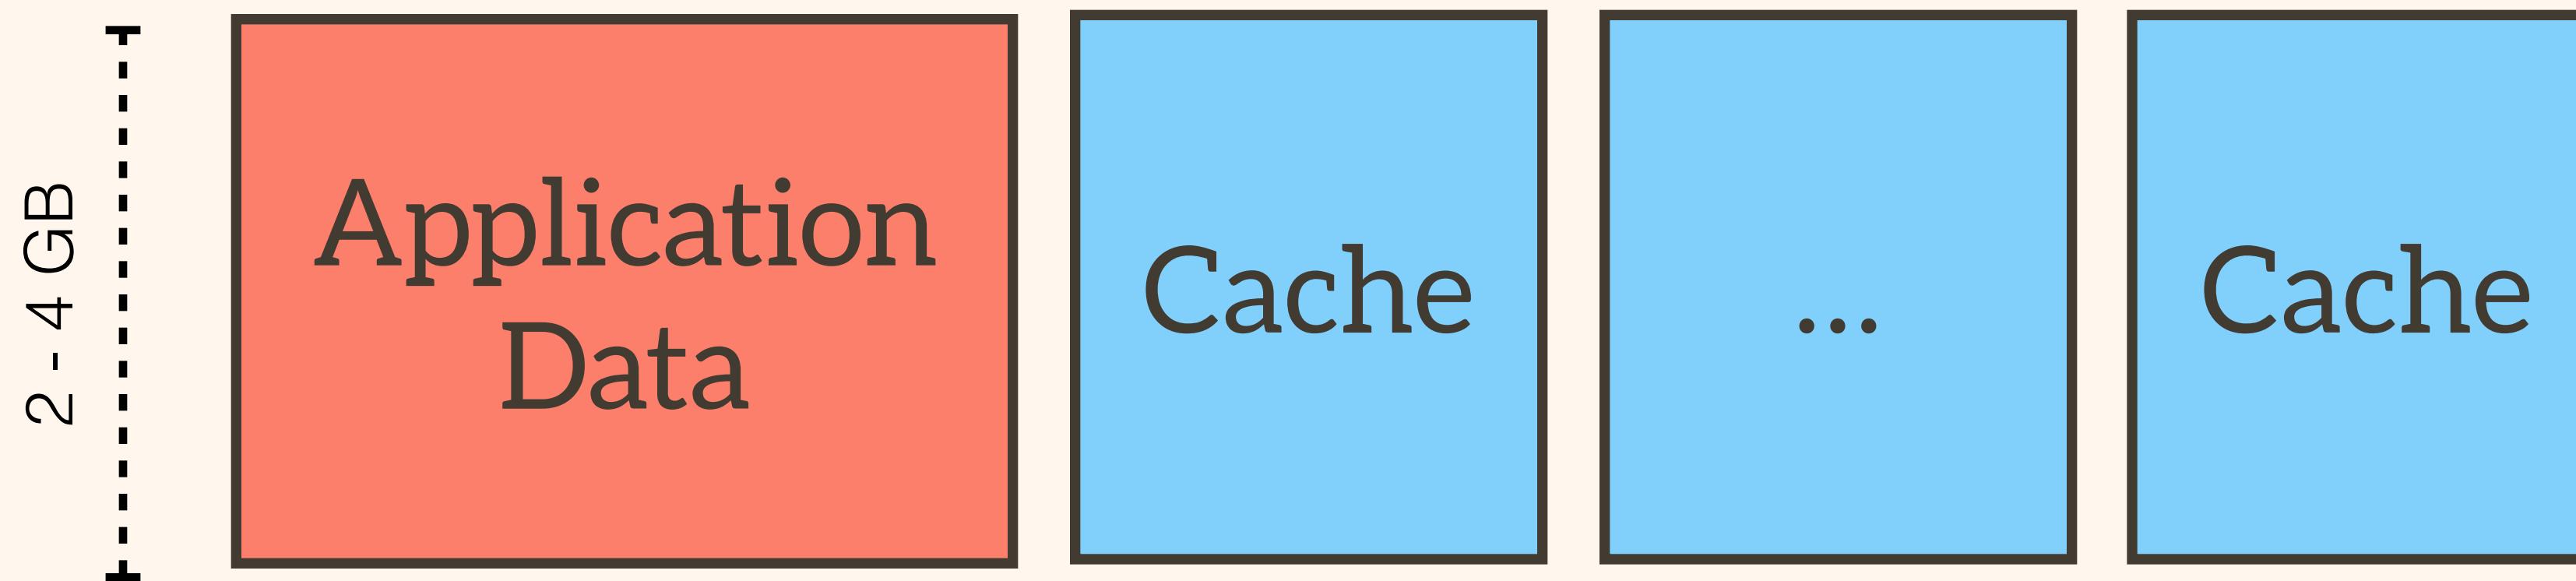




DEMO

Hazelcast

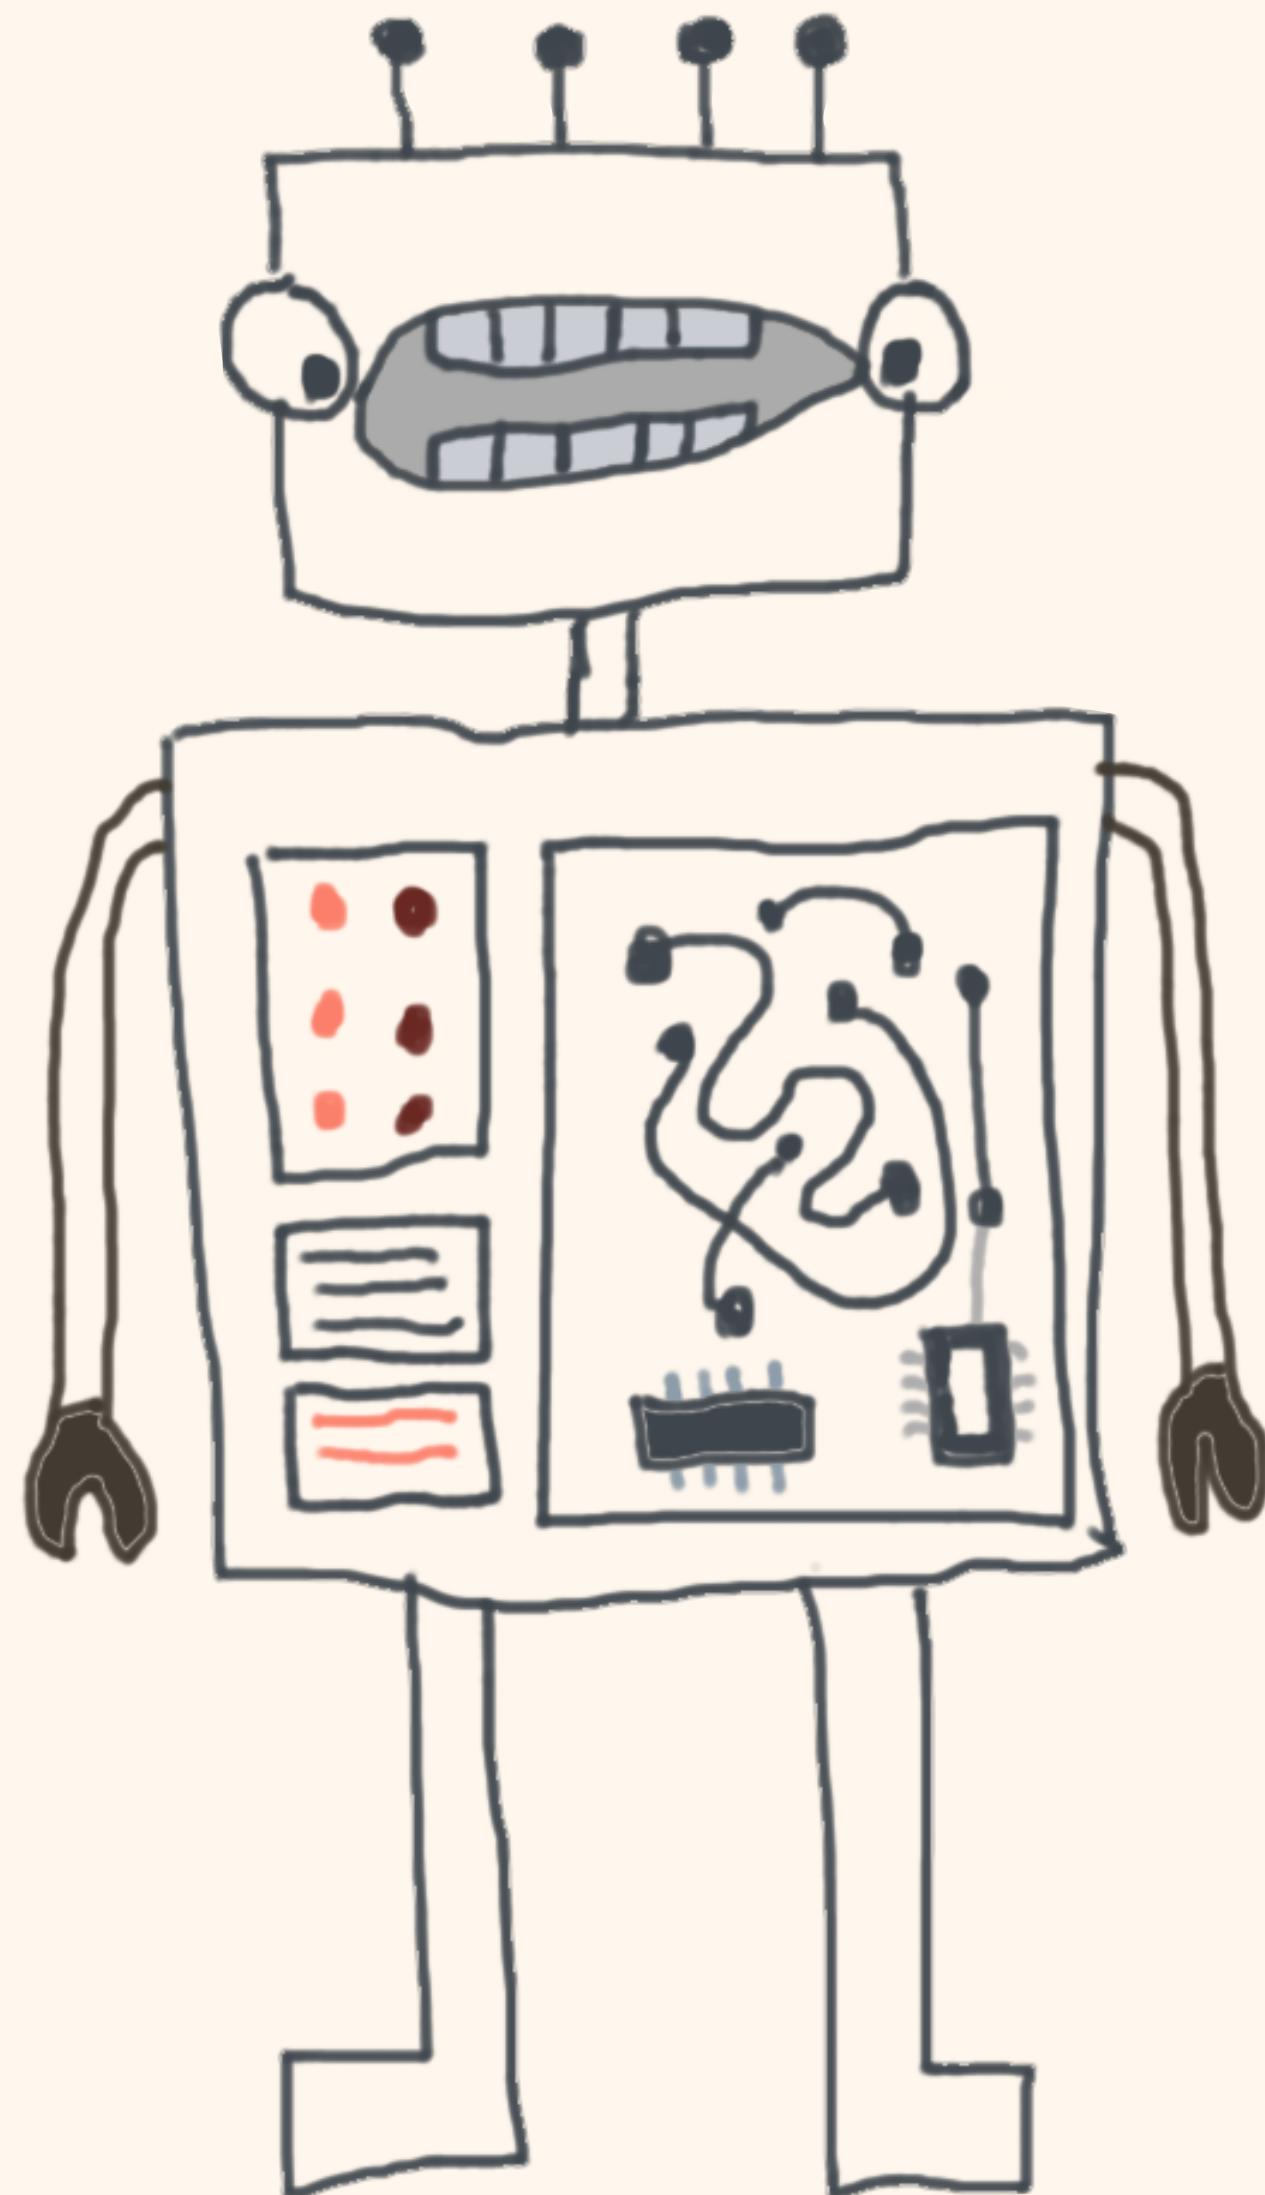
*A distributed cache leads to
smaller heaps, more capacity and
is easy to scale*



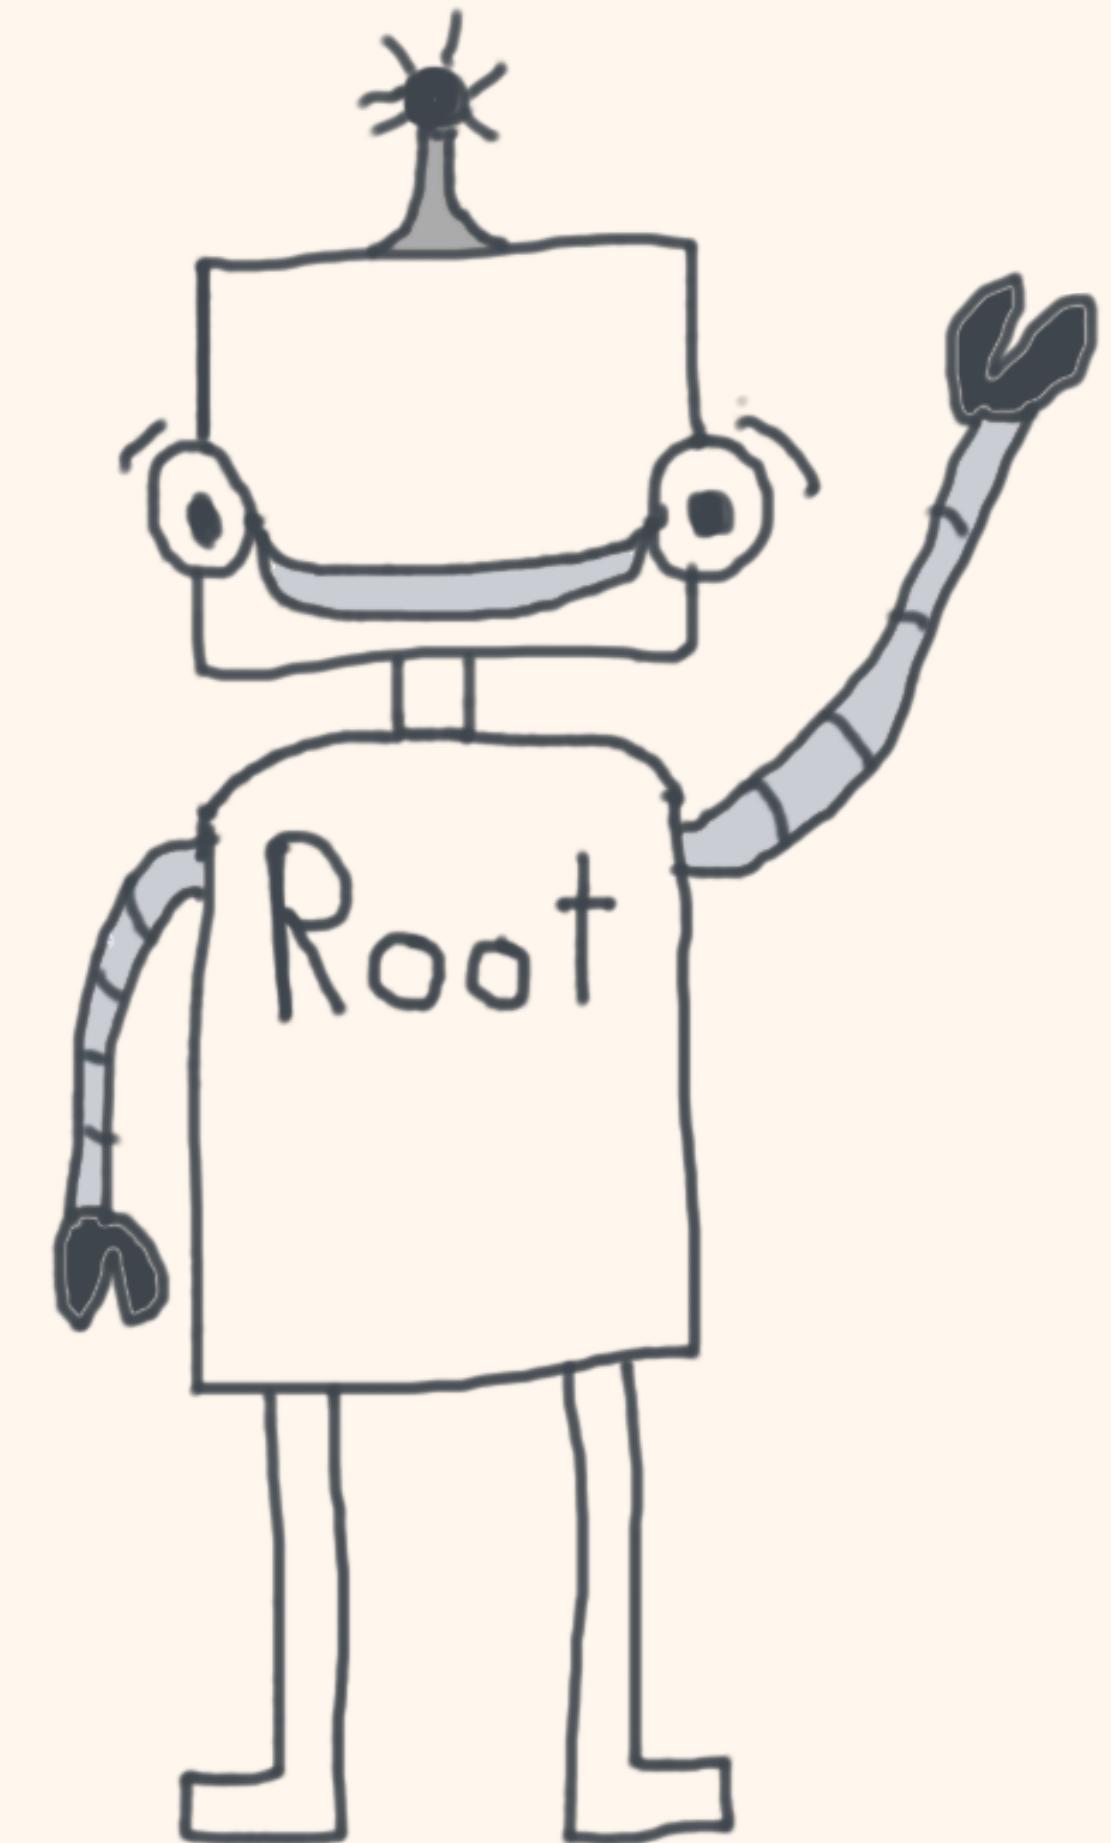


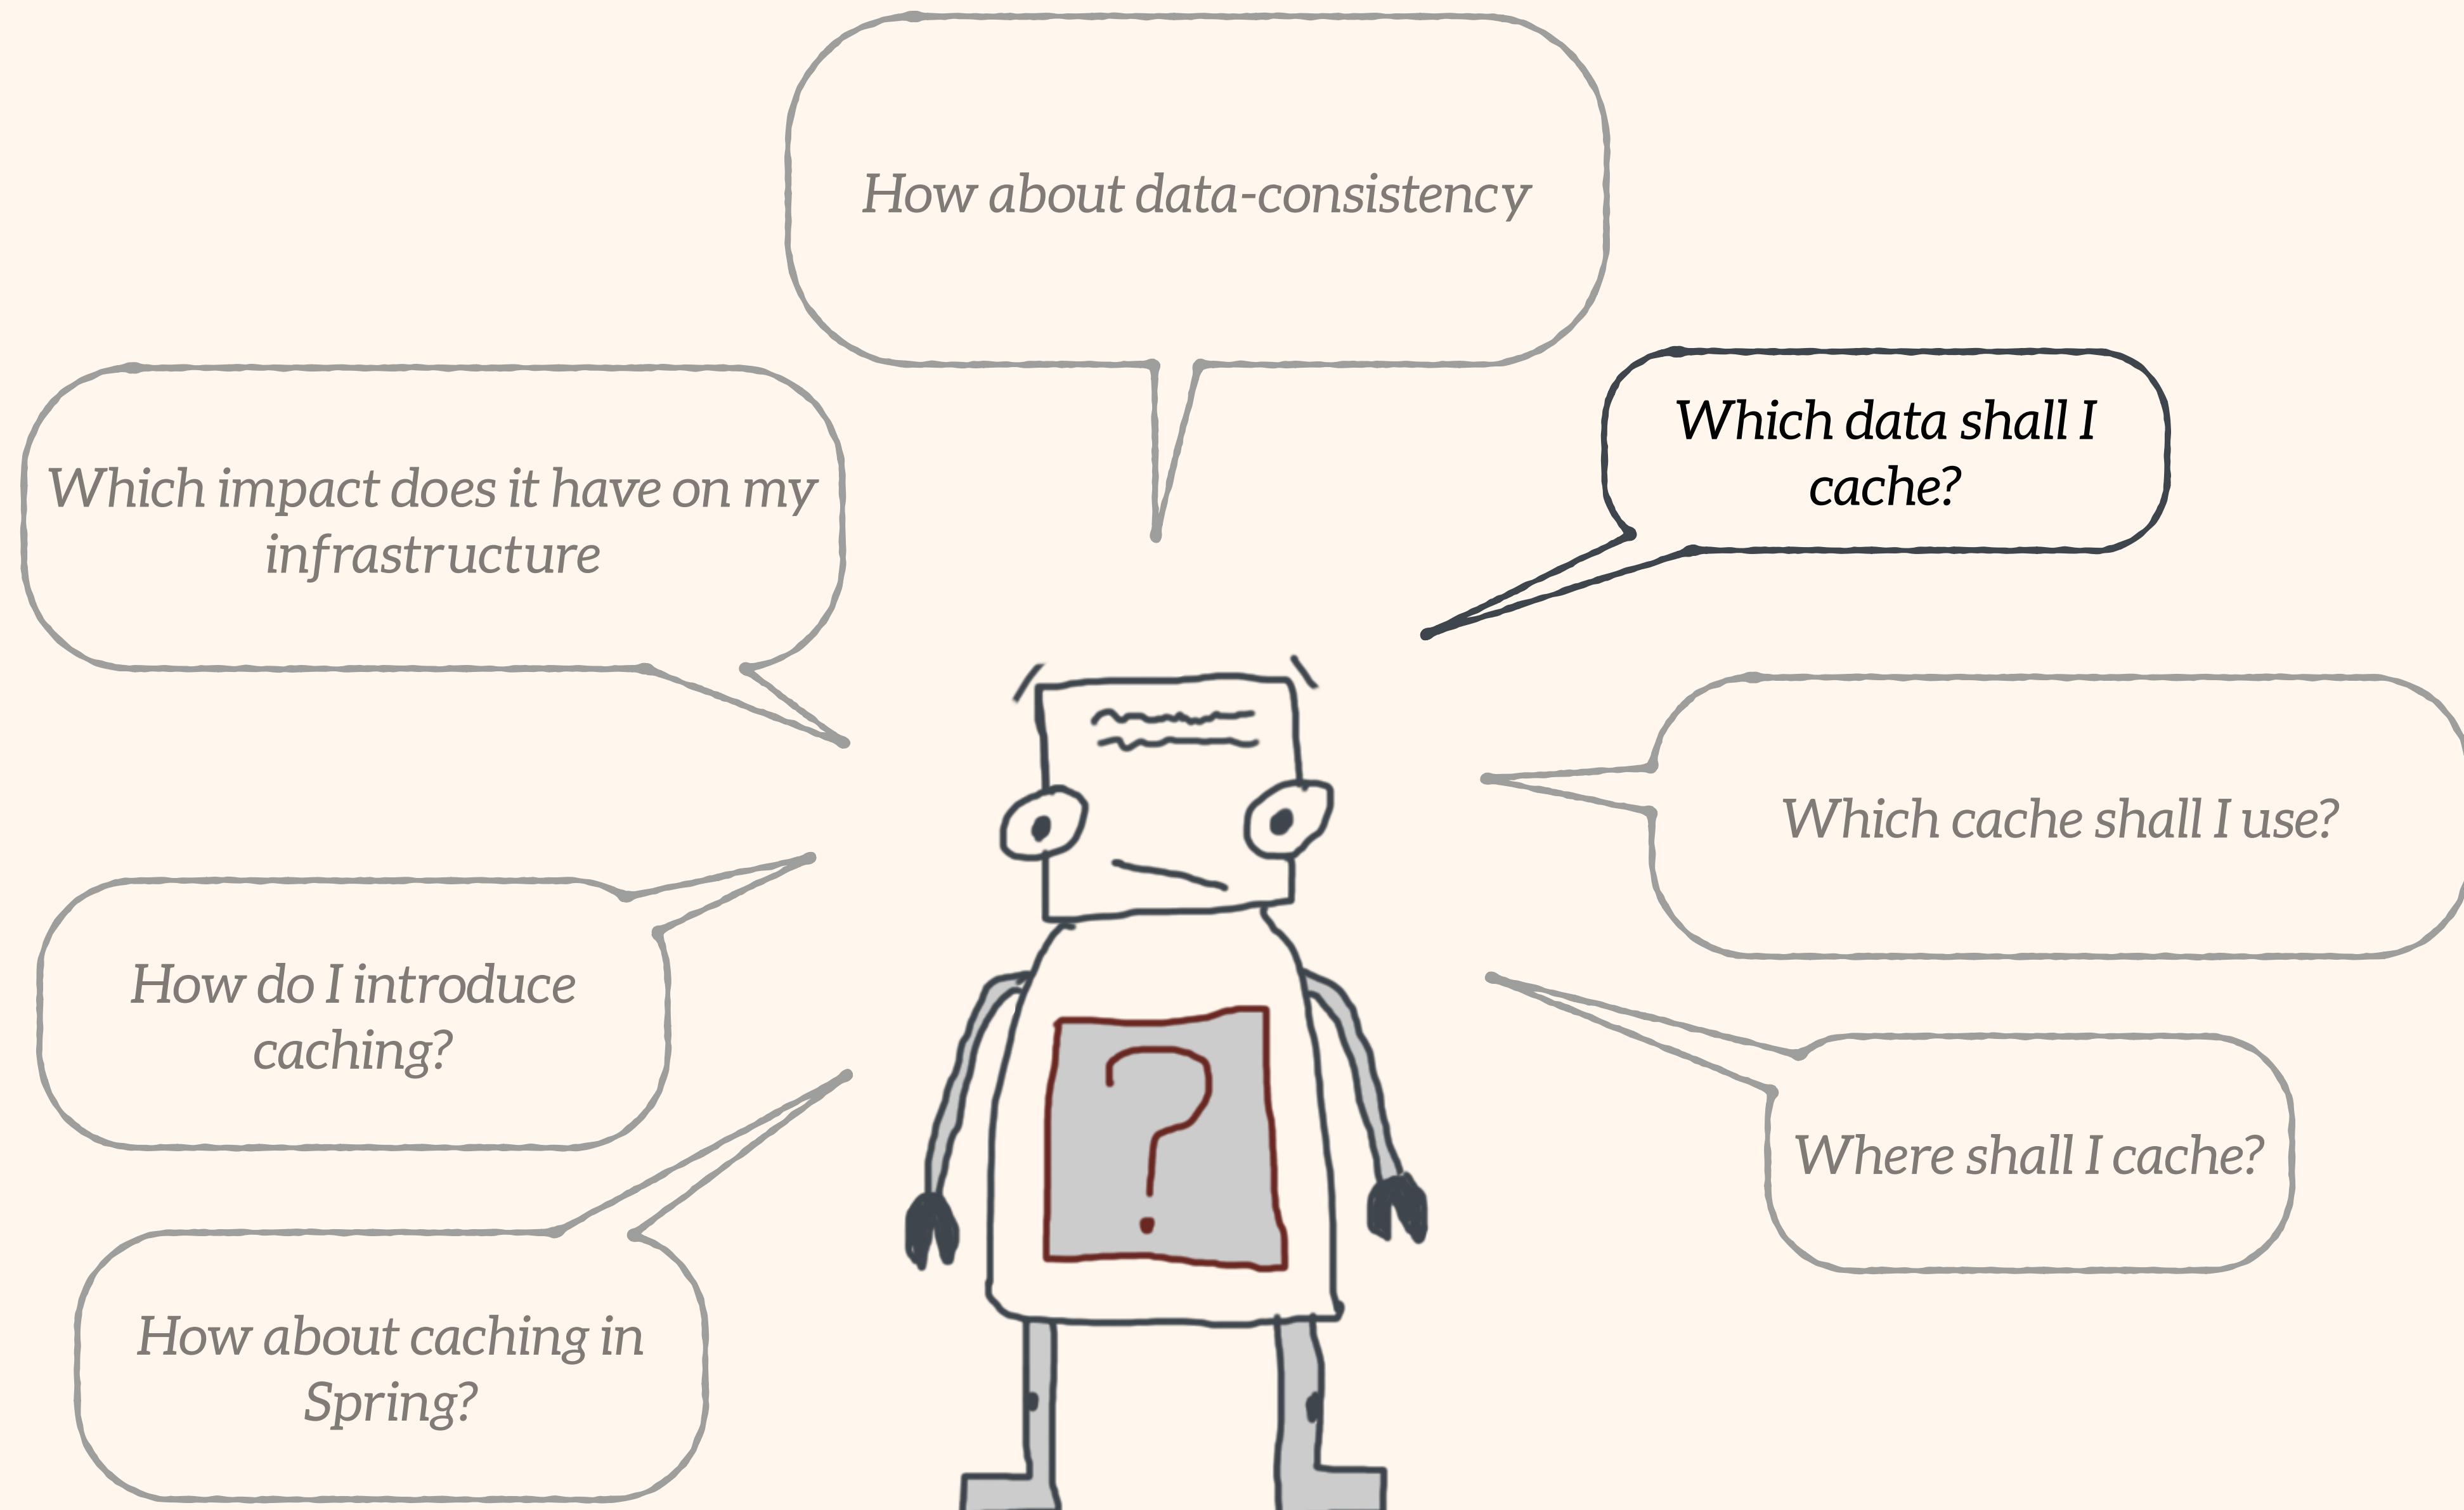
6

*The operations specialist is
your new best friend*



Clustered caches are complex. Please make sure that operations and networking are involved as early as possible.





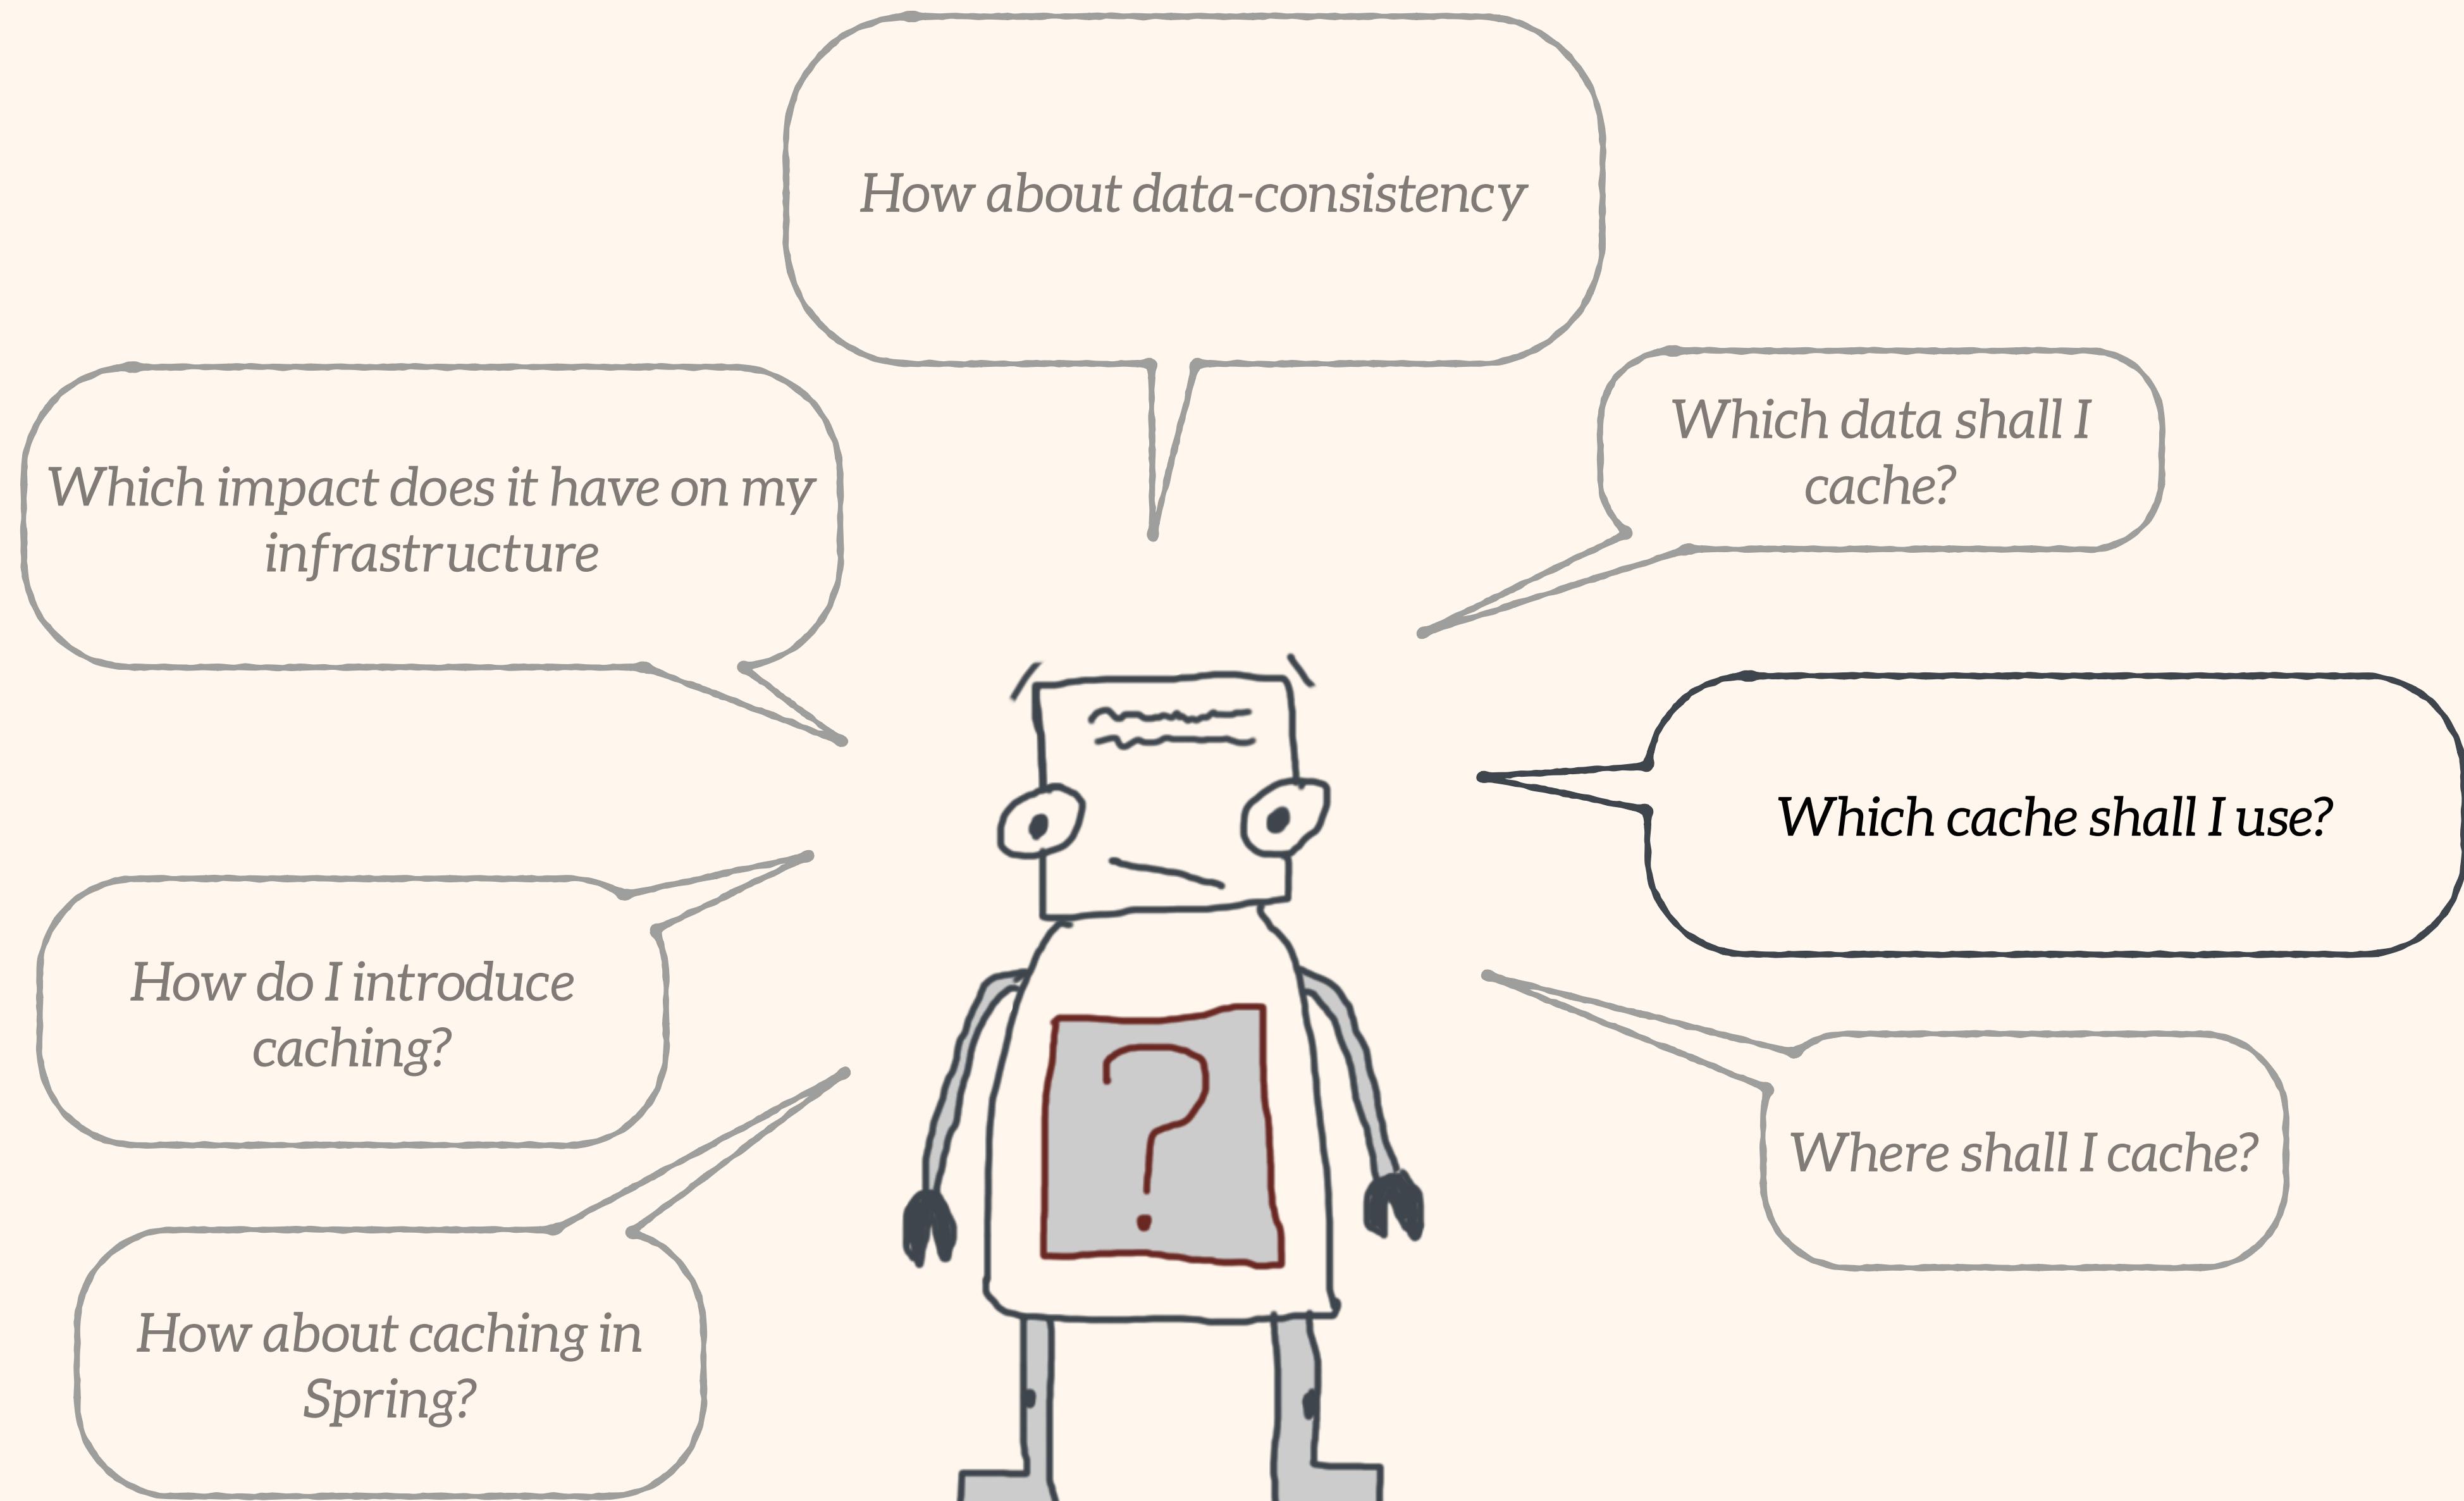


7

*Make sure that only suitable
data gets cached*

The best cache candidates are
read-mostly data, which are
expensive to obtain

If you urgently must cache write-intensive data make sure to use a distributed cache and not a replicated or invalidating one





8

*Only use existing cache
implementations*

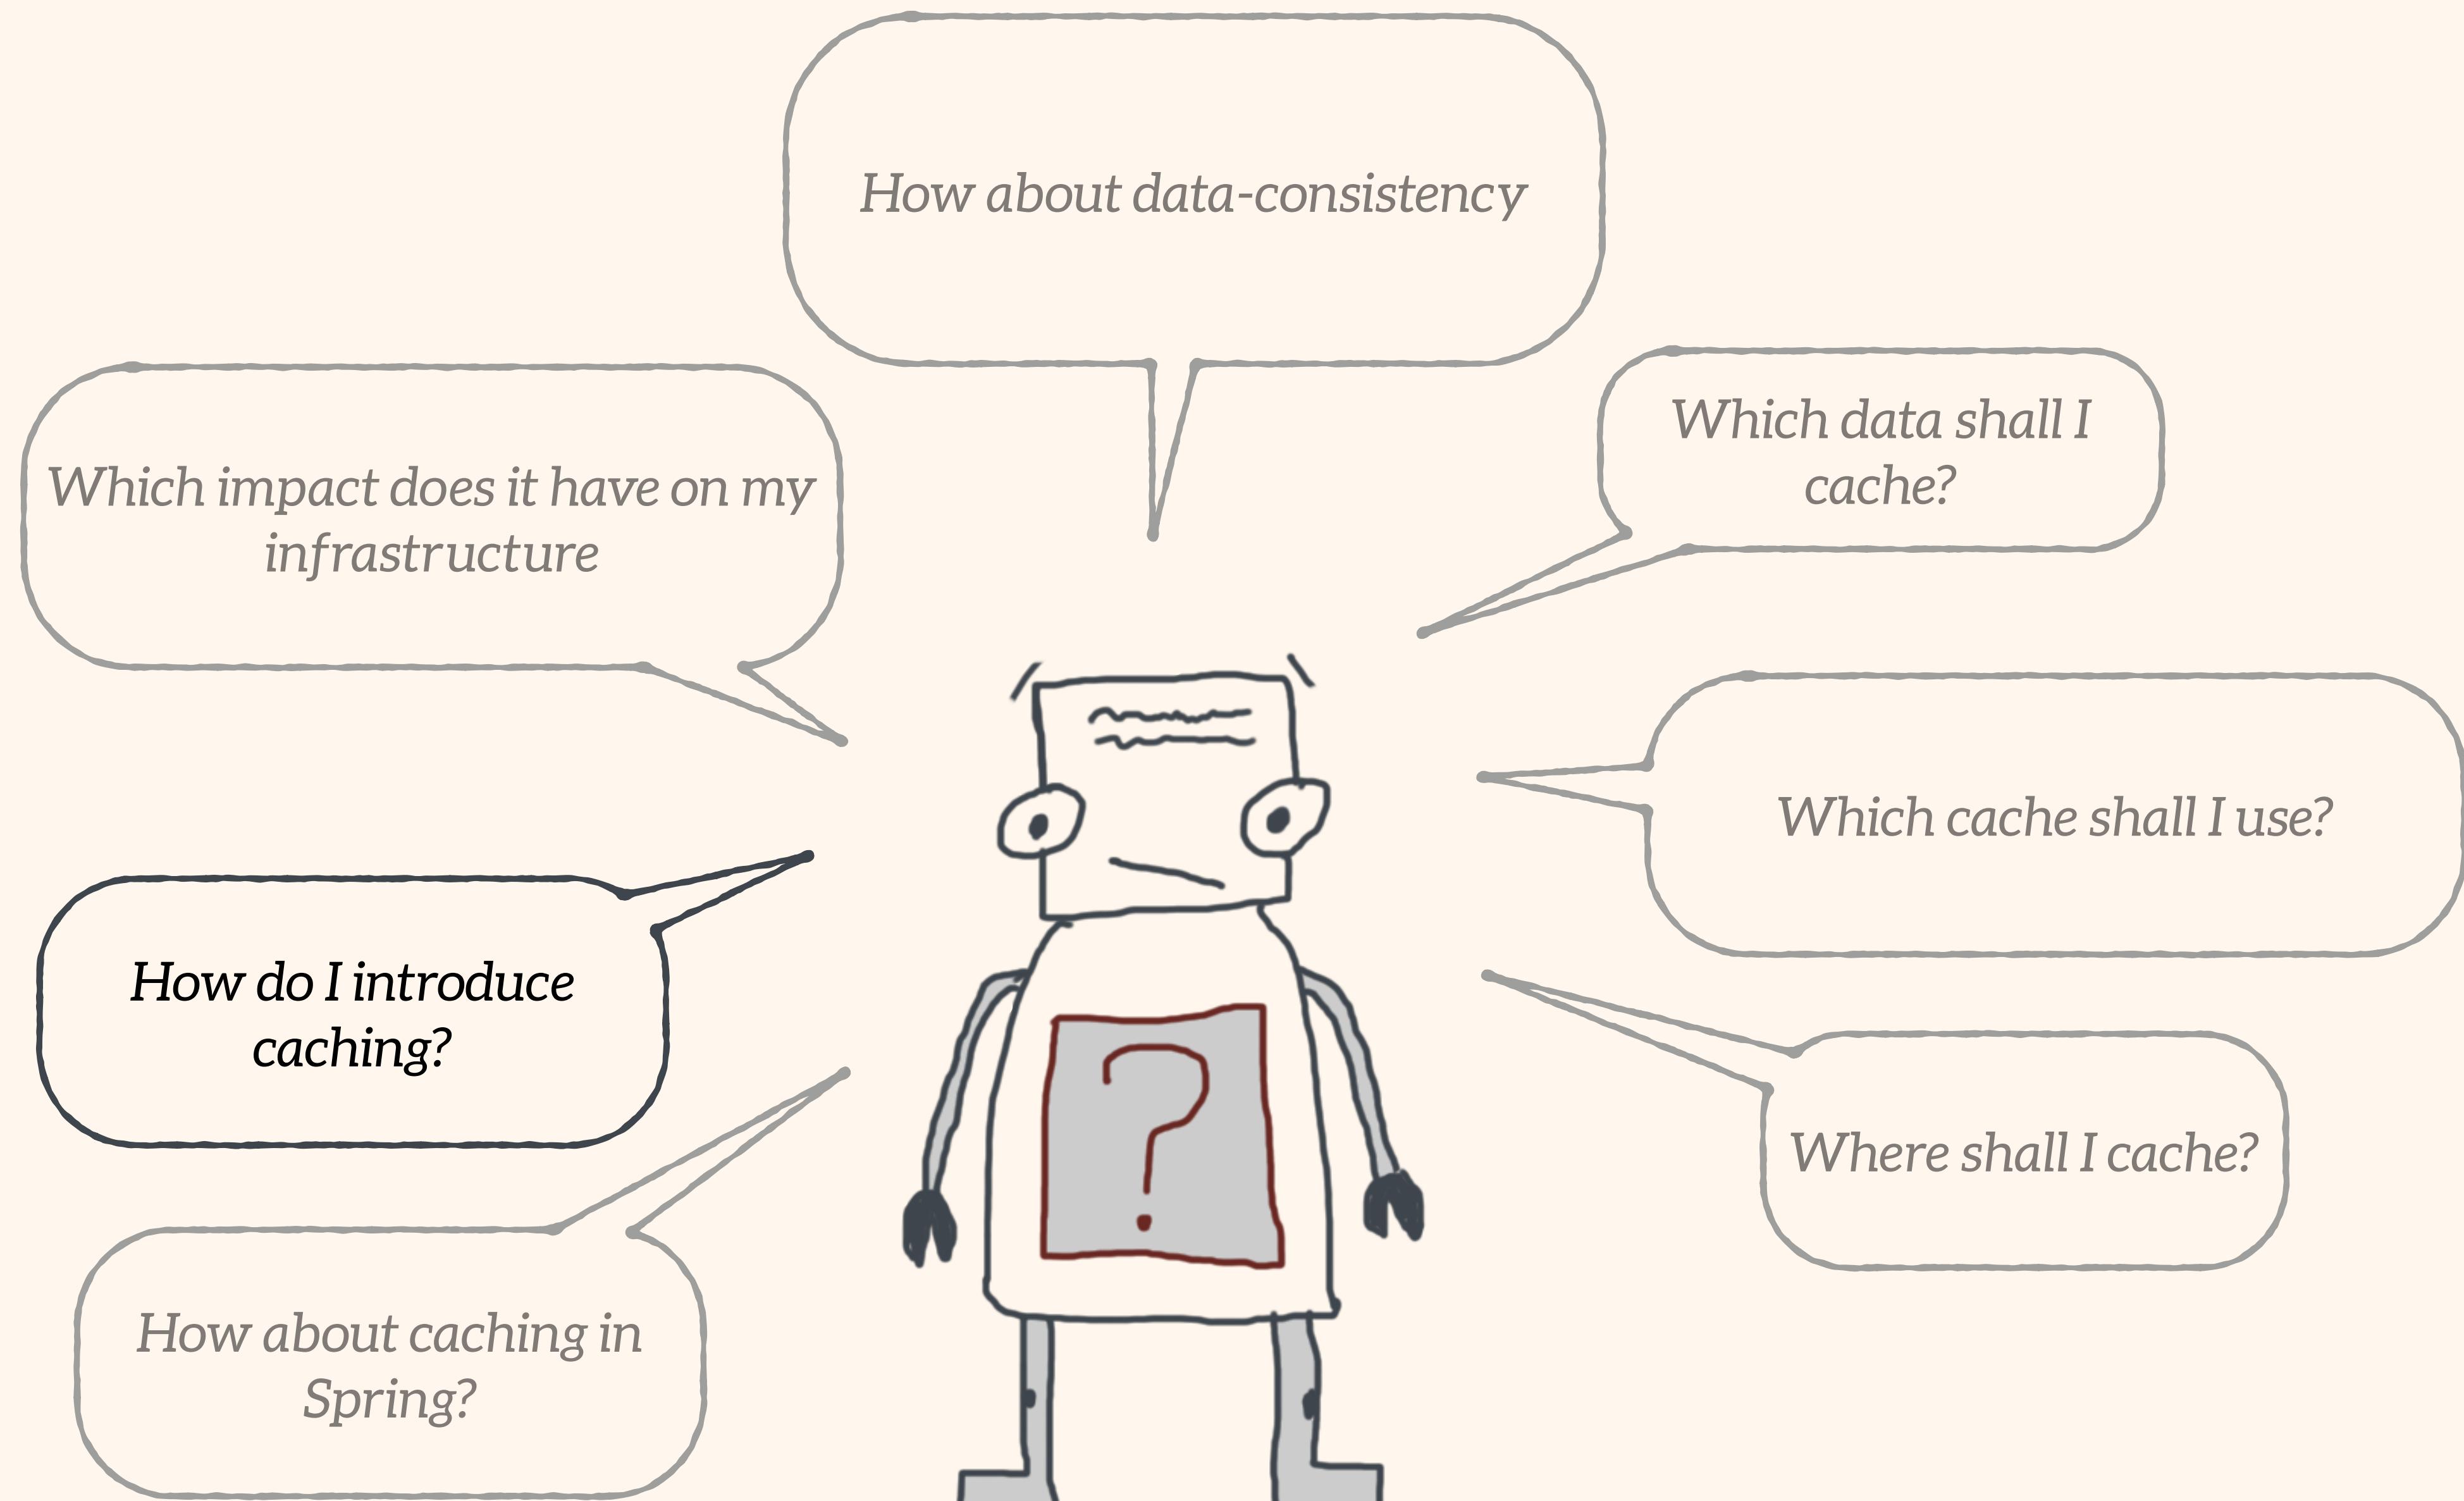
NEVER
write your own cache
implementation



Infinispan, EHCache, Hazelcast, Couchbase,
Memcache, OSCache, SwarmCache, Xtreme
Cache, Apache DirectMemory

CACHE Implementations

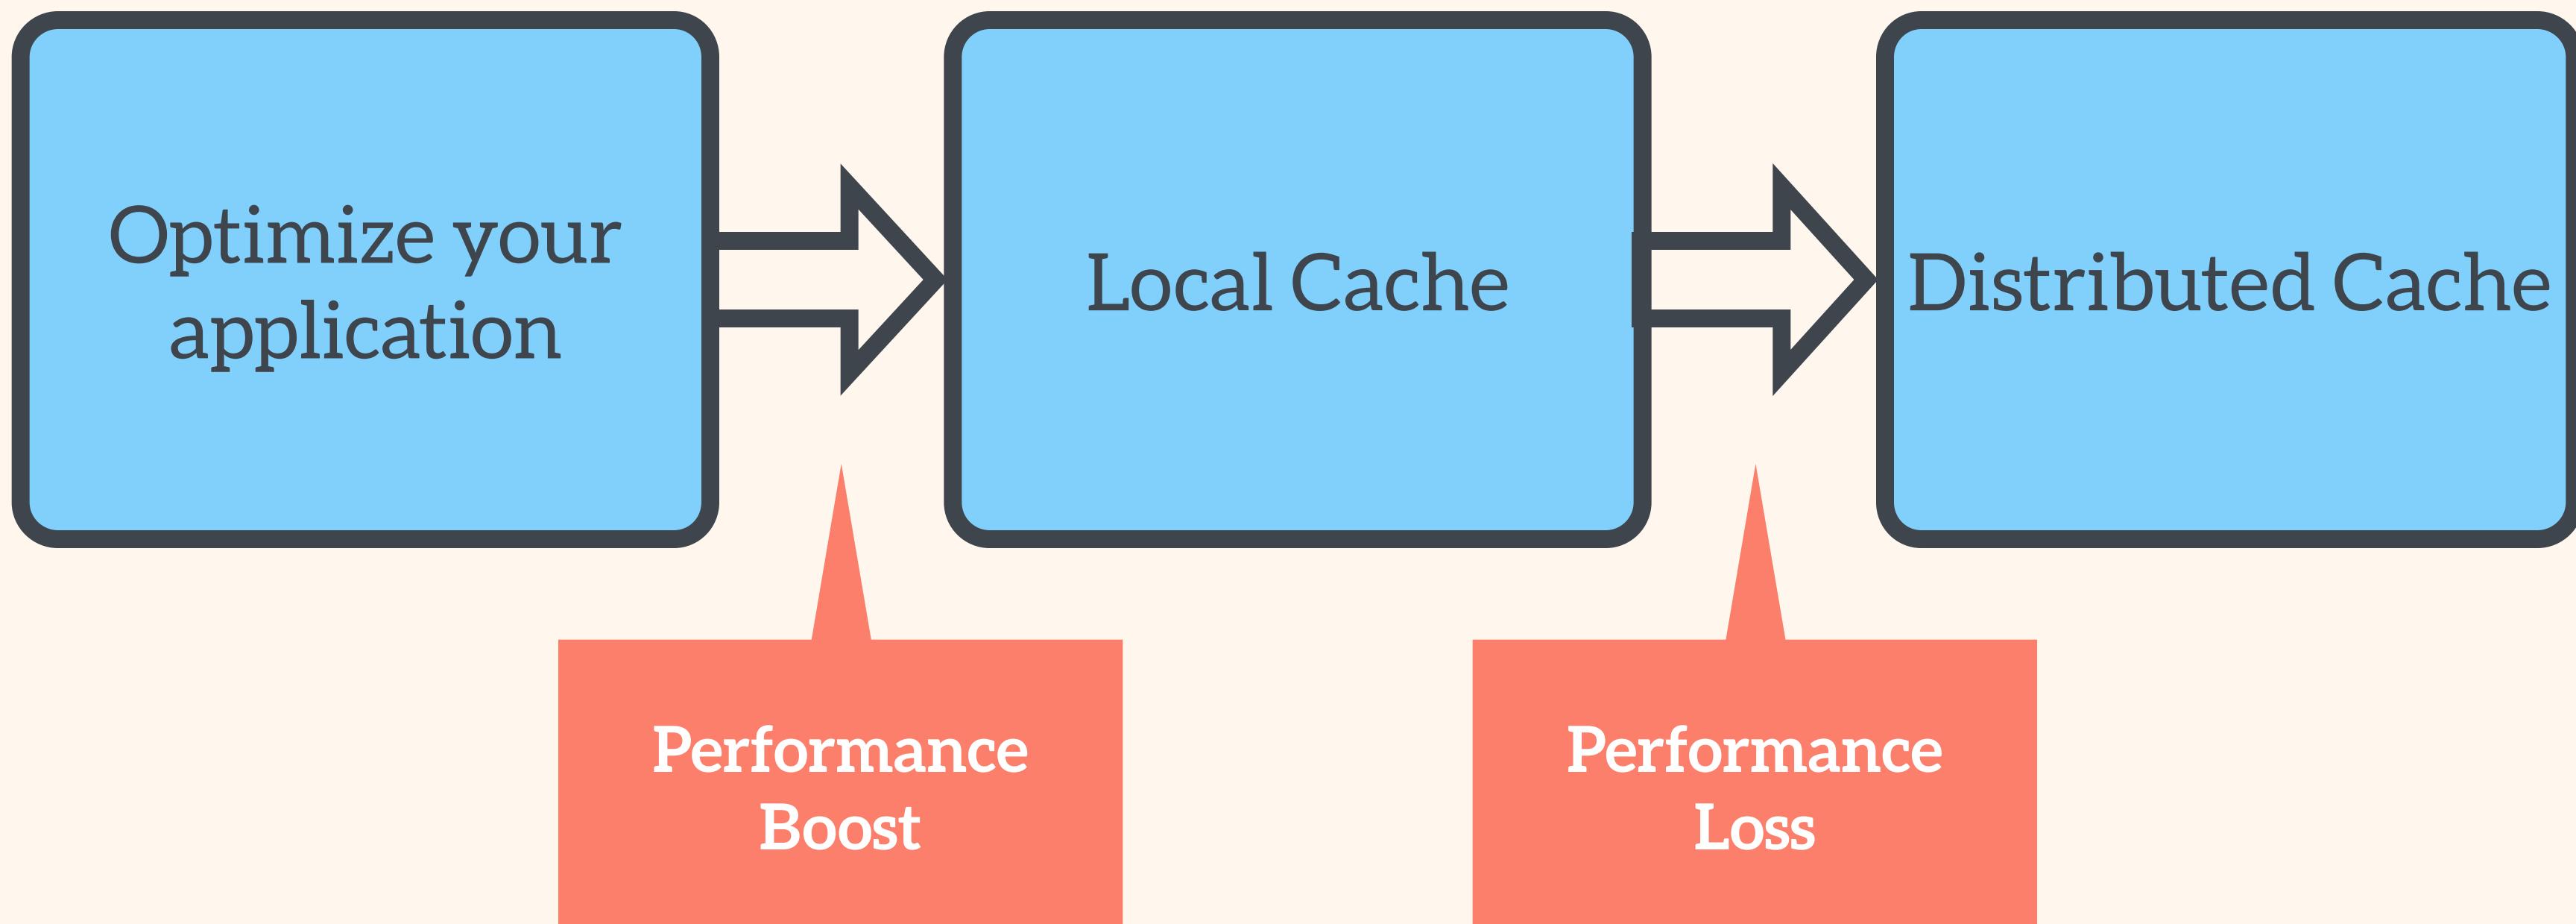
Terracotta, Coherence, Gemfire, Cacheonix,
WebSphere eXtreme Scale, Oracle 12c In
Memory Database





9

*Introduce Caching in three
steps*



10

Optimize Serialization

Example: Hazelcast

putting and getting 10.000 objects locally

	GET Time	PUT Time	Payload Size
Serializable	?	?	?
Data Serializable	?	?	?
Identifier Data Serializable	?	?	?

DEMO

Serialization

Example: Hazelcast

putting and getting 10.000 objects locally

	GET Time	PUT Time	Payload Size
Serializable	1287 ms	1220 ms	1164 byte
Data Serializable	443 ms	408 ms	916 byte
Identifier Data Serializable	264 ms	207 ms	882 byte

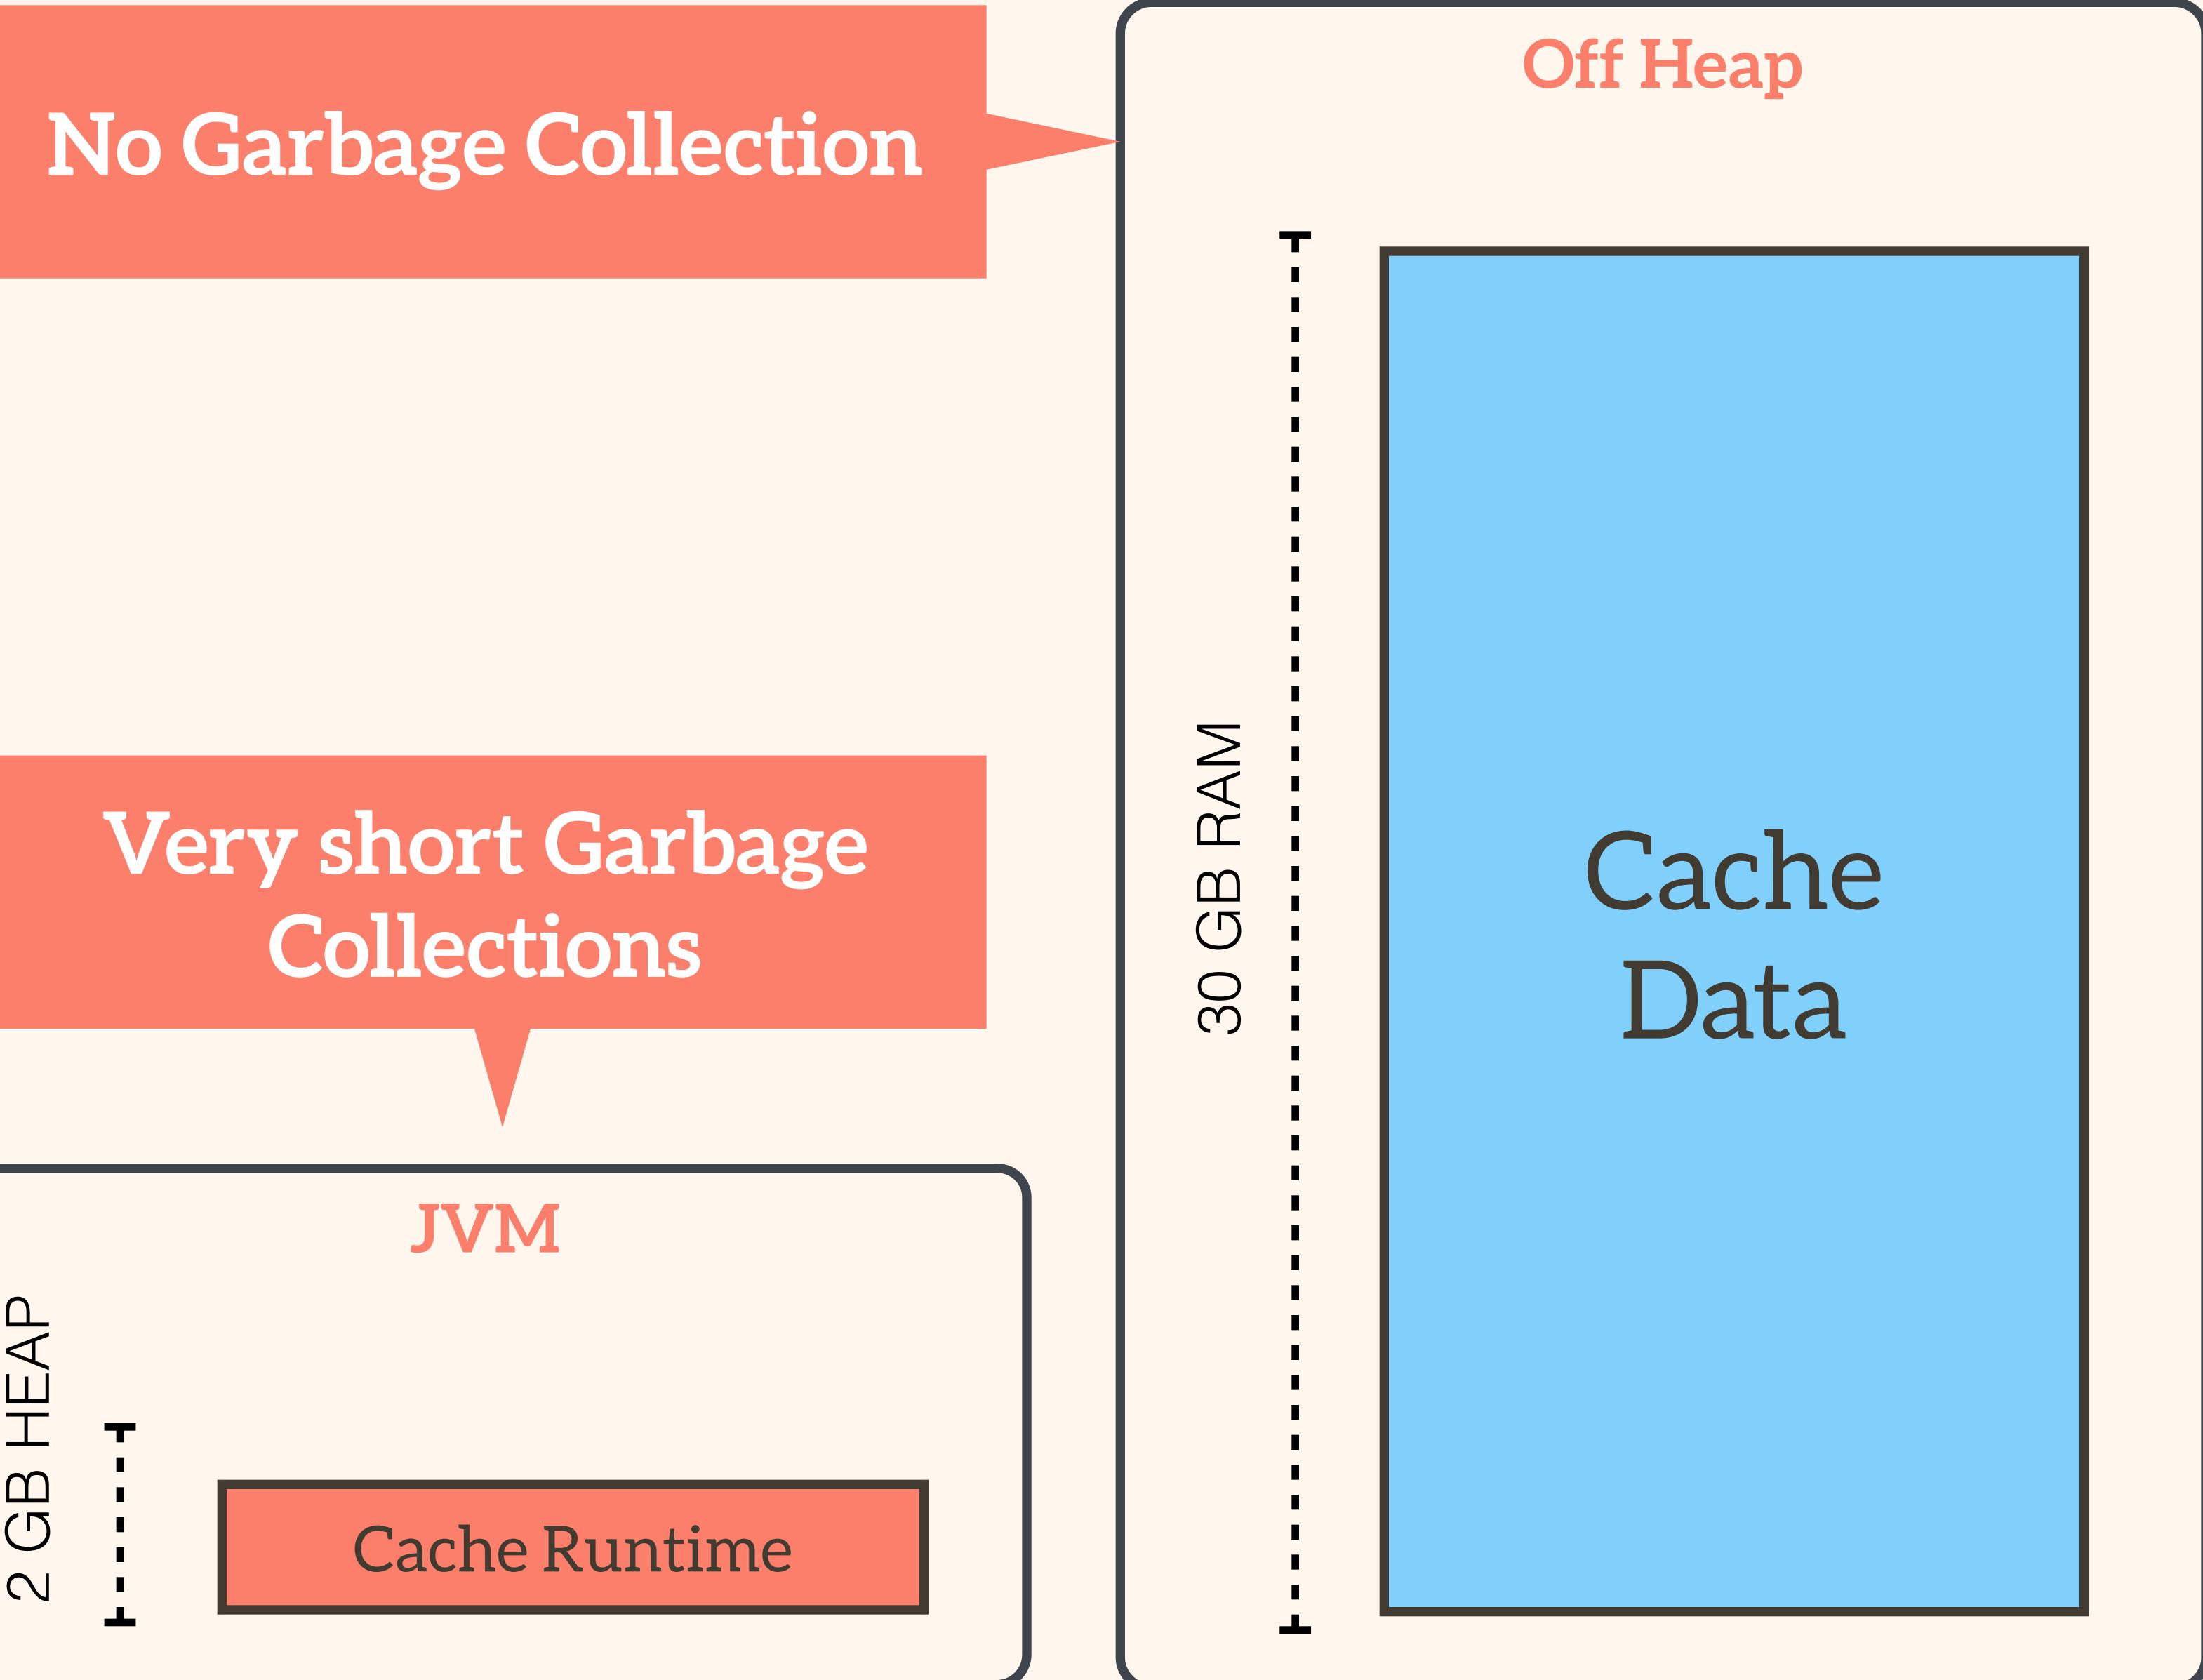
JAVA SERIALIZATION SUCKS

for Caching if alternatives are present



11

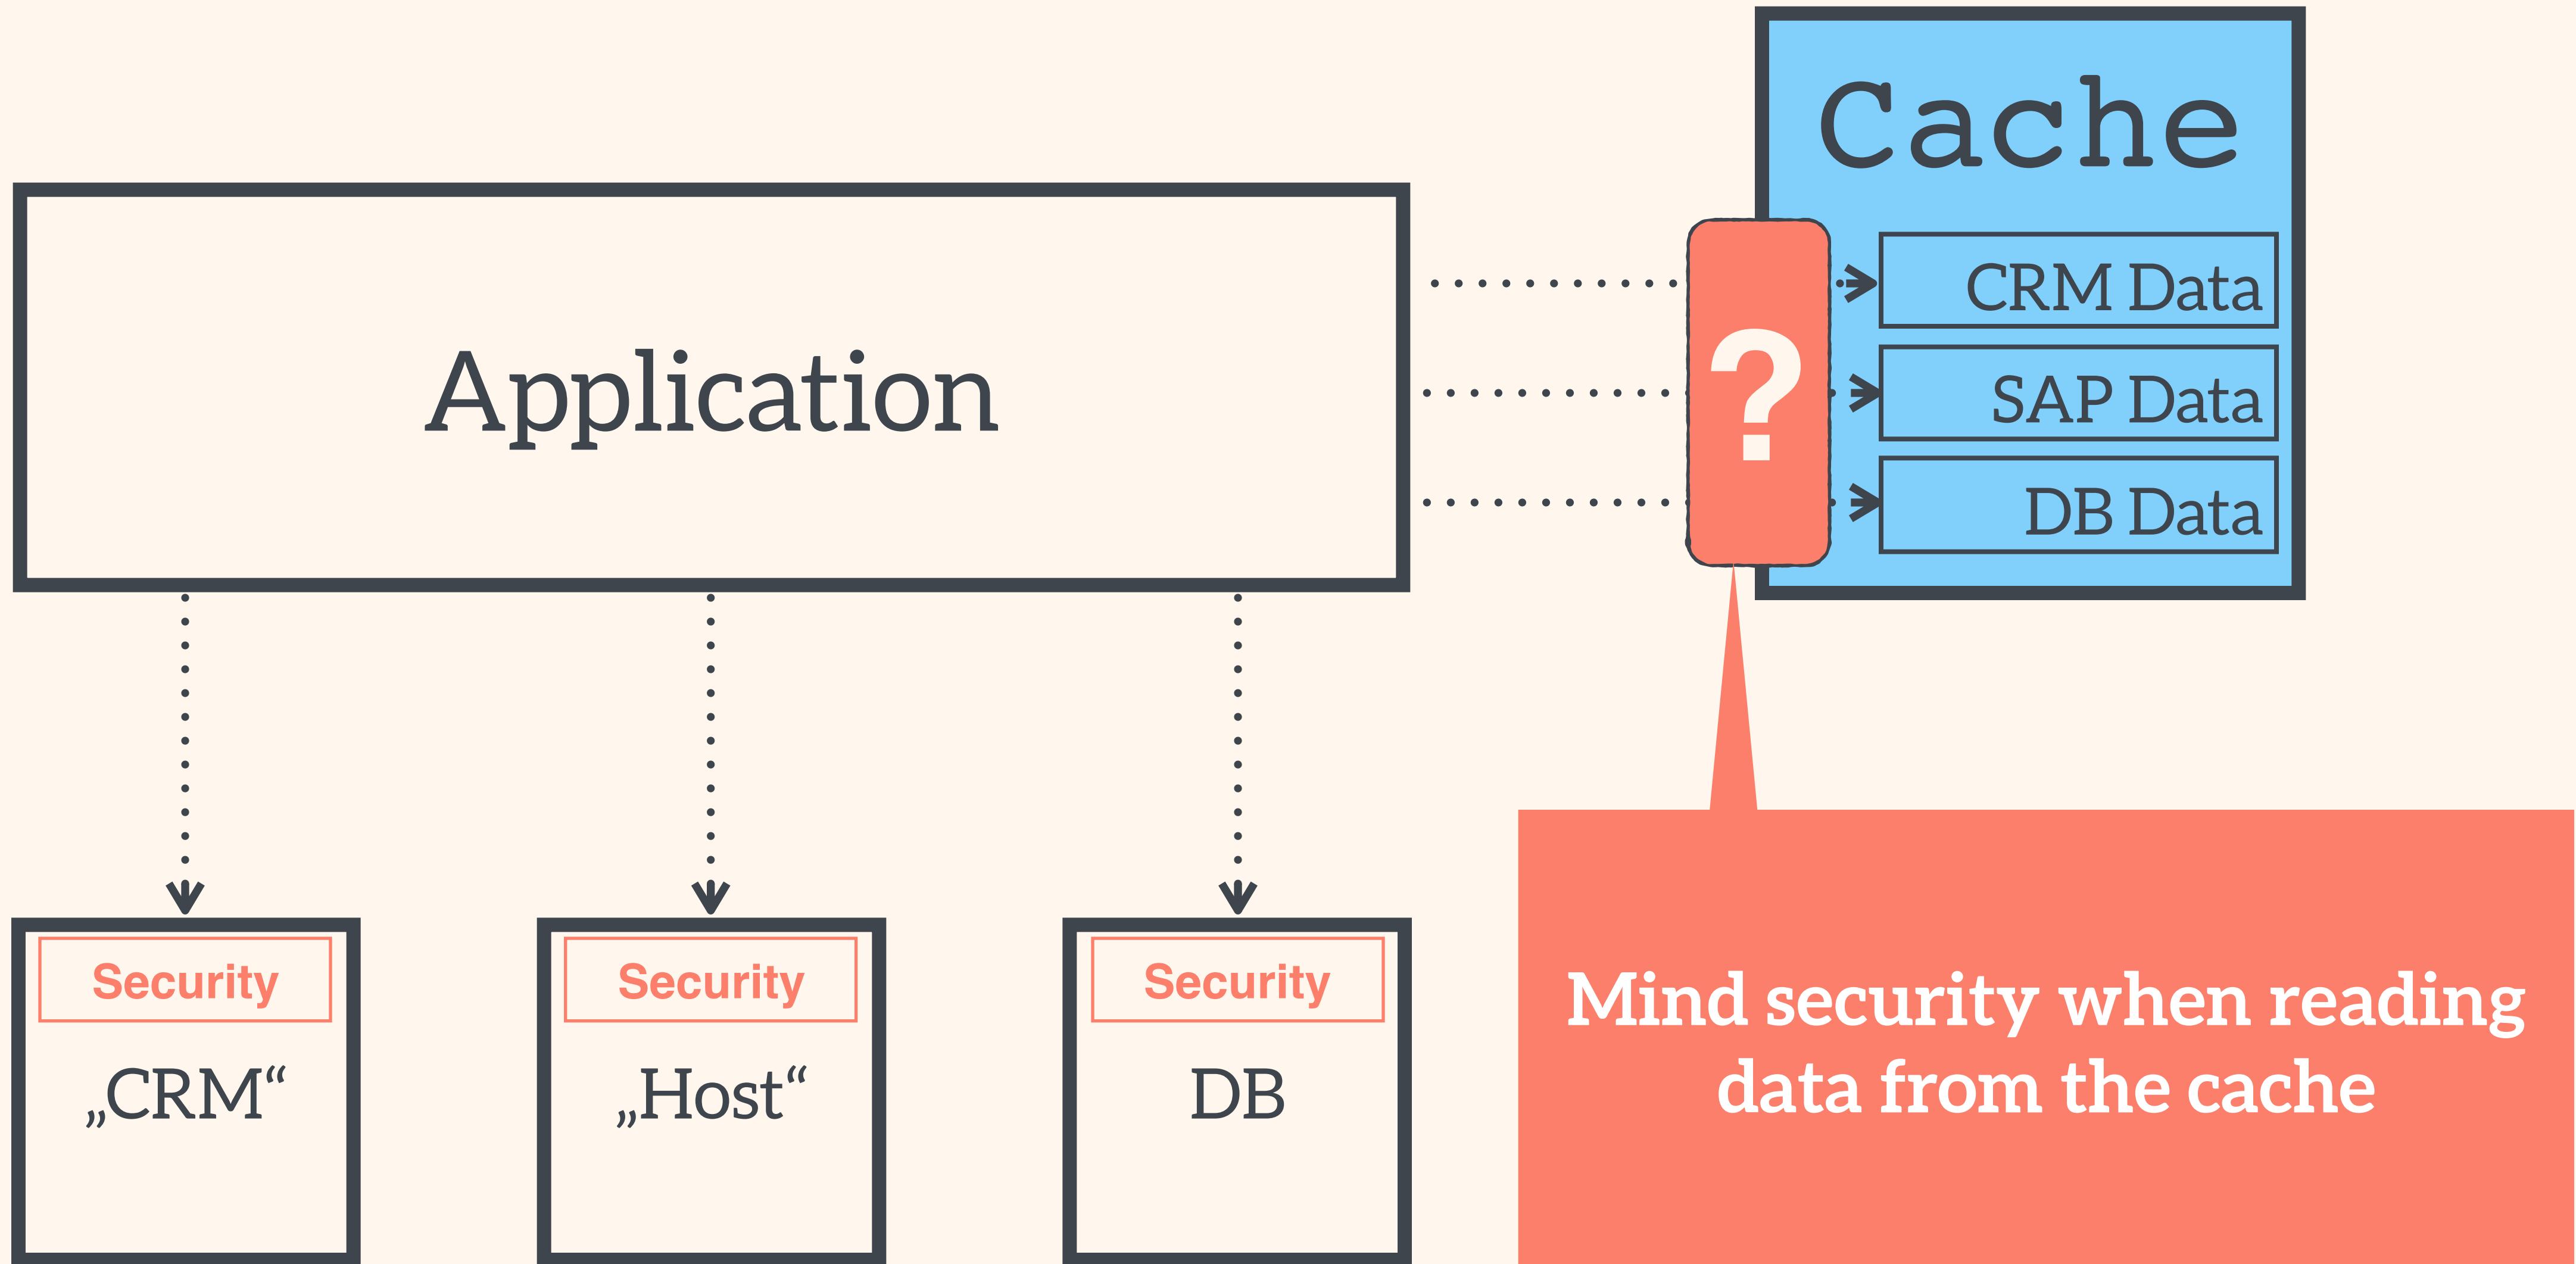
*Use Off-Heap Storage for
Cache instances with more
than 4 GB Heap Size*

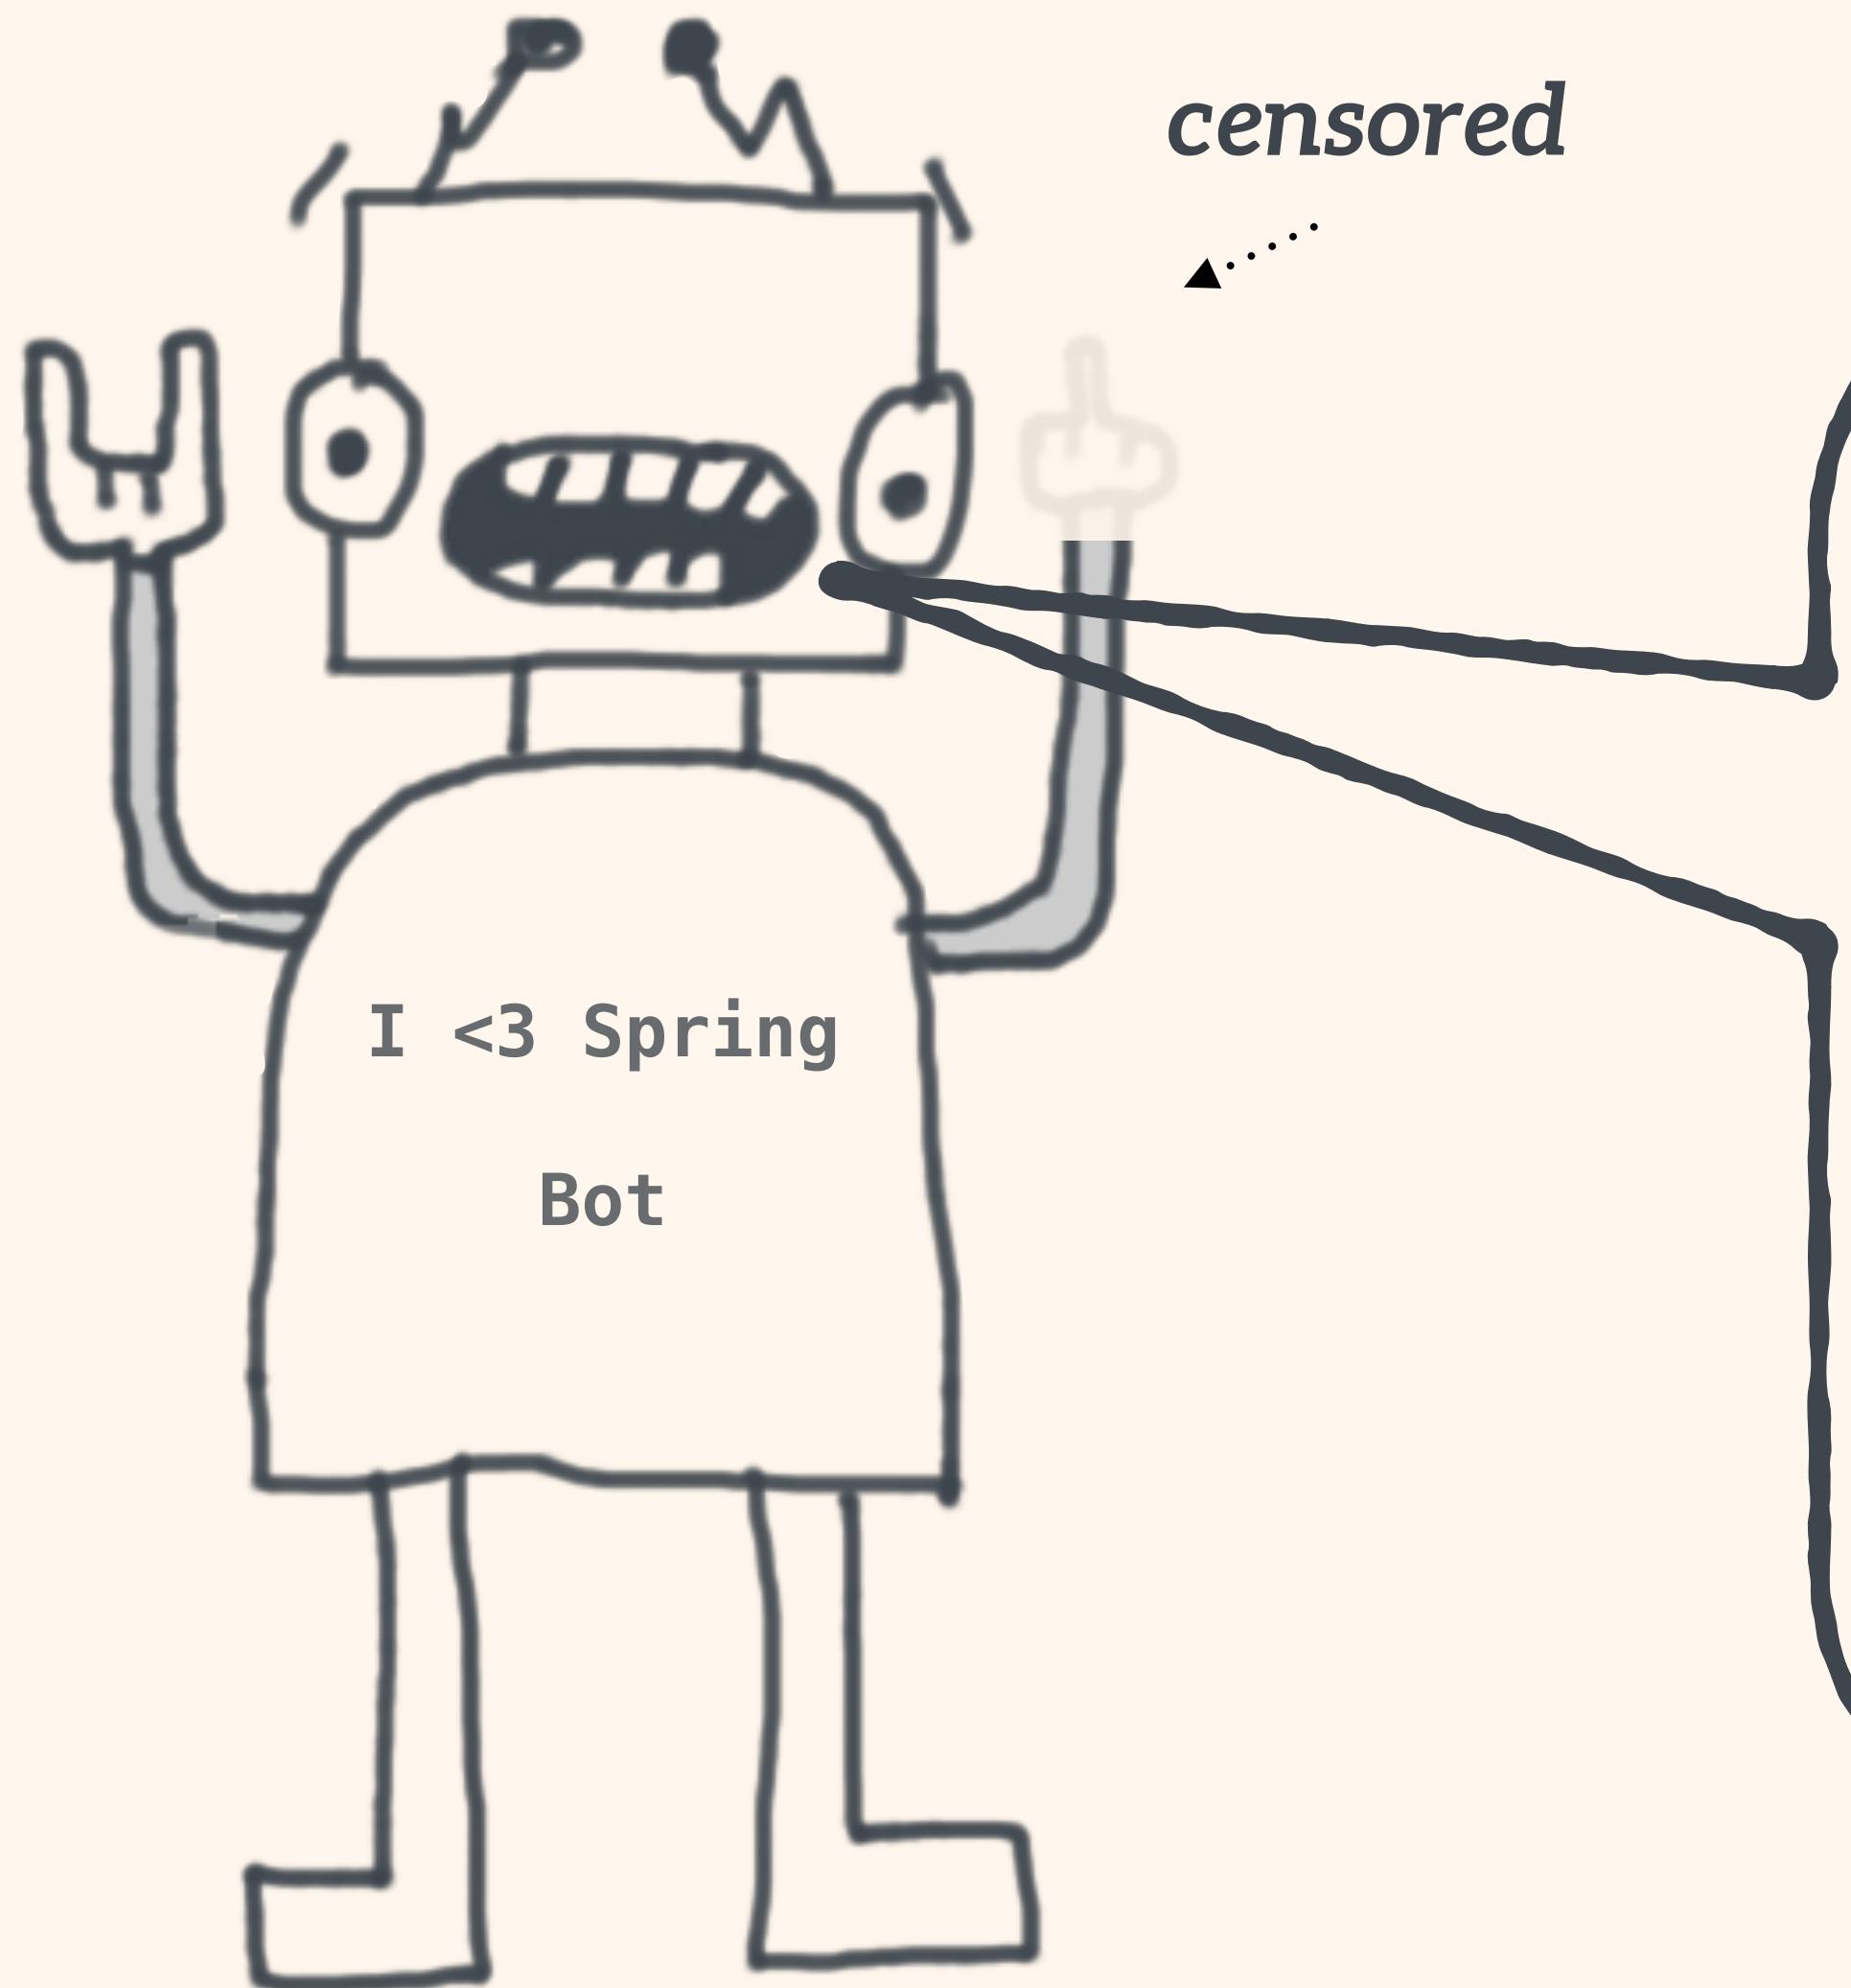




12

Mind the security gap





I'm at a Spring conference and this guy is 50 slides in and hasn't yet mentioned Spring even if he advertised



13

Abstract your cache provider

Tying your code to a cache provider is bad practice

```
public Account retrieveAccount(String accountNumber)
{
    Cache cache = ehCacheMgr.getCache("accounts");
    Account account = null;
    Element element = cache.get(accountNumber);
    if(element == null) {
        //execute some business logic for retrieval
        //account = result of logic above
        cache.put(new Element(accountNumber, account));
    } else {
        account = (Account)element.getObjectValue();
    }
    return account;
}
```

Try switching from EHCache to Hazelcast

```
public Account retrieveAccount(String accountNumber)
{
    Cache cache = ehCacheMgr.getCache("accounts");
    Account account = null;
    Element element = cache.get(accountNumber);
    if(element == null) {
        //execute some business logic for retrieval
        //account = result of logic above
        cache.put(new Element(accountNumber, account));
    } else {
        account = (Account)element.getObjectValue();
    }
    return account;
}
```

You will have to adjust these lines of code to the Hazelcast API

You can't switch cache providers between environments

```
public Account retrieveAccount(String accountNumber)
{
    Cache cache = ehCacheMgr.getCache("accounts");
    Account account = null;
    Element element = cache.get(accountNumber);
    if(element == null) {
        //execute some business logic for retrieval
        //account = result of logic above
        cache.put(new Element(accountNumber, account));
    } else {
        account = (Account)element.getObjectValue();
    }
    return account;
}
```

EHCache is tightly coupled to your code

You mess up your business logic with infrastructure

```
public Account retrieveAccount(String accountNumber)
{
    Cache cache = ehCacheMgr.getCache("accounts");
    Account account = null;
    Element element = cache.get(accountNumber);
    if(element == null) {
        //execute some business logic for retrieval
        //account = result of logic above
        cache.put(new Element(accountNumber, account));
    } else {
        account = (Account)element.getObjectValue();
    }
    return account;
}
```

This is all caching related code without any business relevance

Introducing Spring's cache abstraction

```
<cache:annotation-driven cache-manager="ehCacheManager"/>

<!-- EH Cache local -->
<bean id="ehCacheManager"
  class="org.springframework.cache.ehcache.EhCacheCacheManager"
  p:cacheManager-ref="ehcache"/>

<bean id="ehcache"
  class="org.springframework.cache.ehcache.EhCacheManagerFactoryBean"
  p:configLocation="/ehcache.xml"/>
```

```
@Cacheable("Customers")
public Customer getCustomer(String customerNumber) {
    ...
}
```

Spring's Caching Annotations

Annotation	Description
@Cacheable	Demarcates cachable methods, can read and write to the cache(s)
@CacheEvict	Demarcates methods that perform cache eviction, that is methods that act as triggers for removing data from the cache.
@CachePut	Updates the cache with the annotated method's return value. Will always execute the method.
@Caching	Allows multiple nested @Cacheable, @CacheEvict and @CachePut annotations to be used on the same method
@CacheConfig	Class-level annotation that allows to share the cache names, the custom KeyGenerator, the custom CacheManager and finally the custom CacheResolver. Does not enable caching.

Default Key Generation Strategy

Annotation

```
@Cacheable("Customers")
public Customer getCustomer(String customerNumber) {
    ...
}
```

Key

customerNumber

```
@Cacheable("CustomerList")
public List<Customer> listCustomers(int start, int count) {
    ...
}
```

SimpleKey containing
start and count

```
@Cacheable("MonthlyReport")
public Report getMonthlyReport() {
    ...
}
```

SimpleKey.EMPTY

You need a custom default KeyGenerator?

```
public class MyOwnKeyGenerator implements KeyGenerator {  
    @Override  
    public Object generate(Object target, Method method, Object... params) {  
        if (params.length == 0) {  
            return new SimpleKey("EMPTY");  
        }  
        if (params.length == 1) {  
            Object param = params[0];  
            if (param != null && !param.getClass().isArray()) {  
                return param;  
            }  
        }  
        return new SimpleKey(params);  
    }  
}
```

```
<cache:annotation-driven cache-manager="hazelcastCacheManager"  
    keyGenerator="myOwnKeyGenerator" />
```

SpEL in Caching Annotations

Annotation

```
@Cacheable("concerts", key="#location.id")
public List<Concert> findConcerts(Location location)
```

Effect

Key: id of location

```
@Cacheable("concerts",
key="T(someType).hash(#location)")
public List<Concert> findConcerts(Location location)
```

Key: hashCode of location

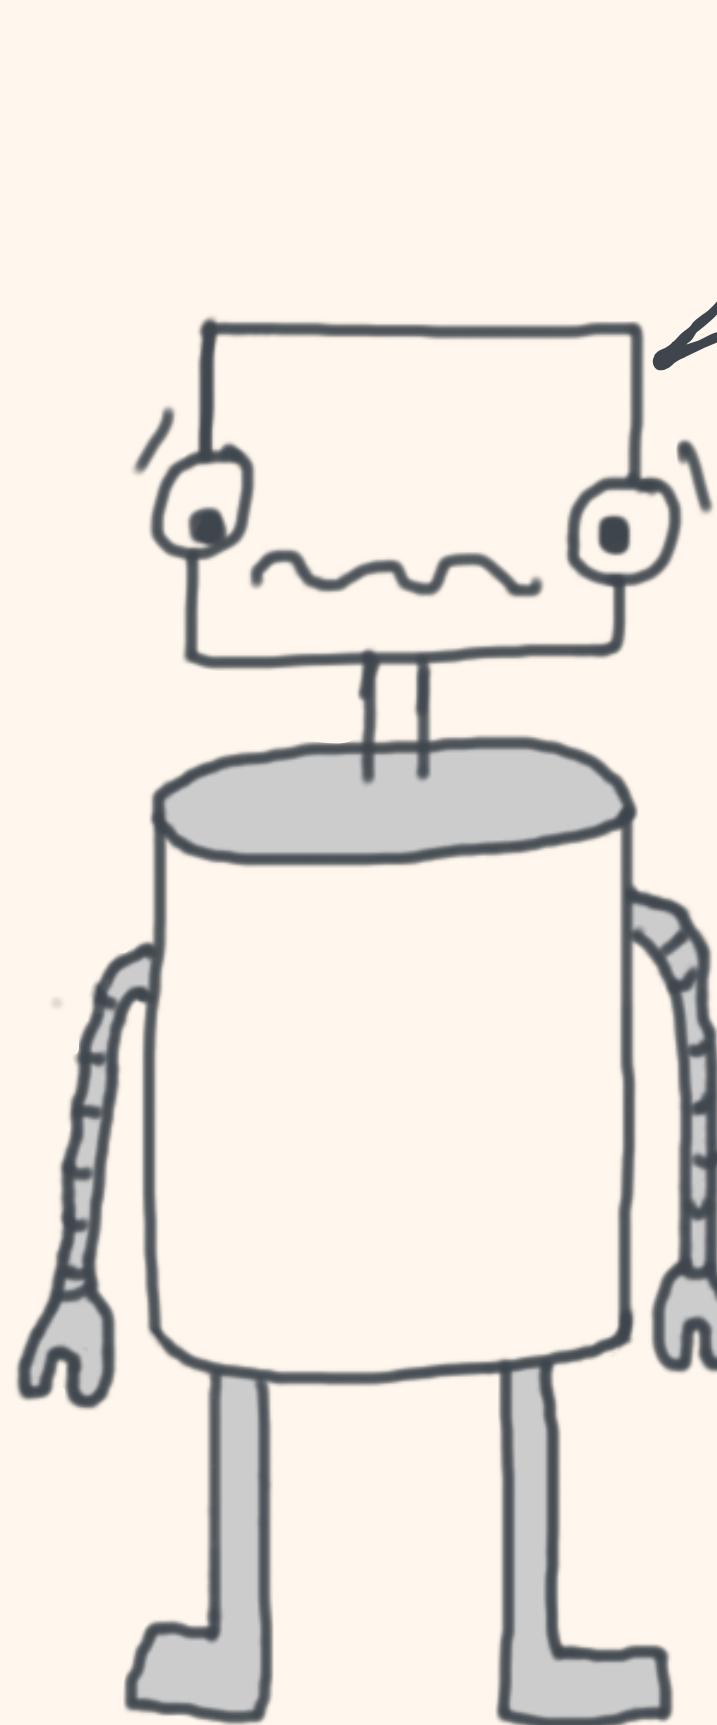
```
@Cacheable("concerts",
    condition="#location.city == 'Dallas')",
    unless="#location.outOfBusiness")
public List<Concert> findConcerts(Location location)
```

Conditional Caching if Location
is in Dallas in operating

```
@CachePut("locations", key="#result.id")
public Location saveLocation(Location location)
```

Key: generated id of result

I have multiple Caches and Cache Managers!



```
@Cacheable("concerts", cacheManager="hazelCastCacheManager")  
public List<Concert> findConcerts(Location location)
```

Manual Assignment

```
@Cacheable("bands", cacheManager="gemfireCacheManager"))  
public List<Band> listBand(int start, int count)
```

Manual Assignment

```
@Cacheable("bands", cacheResolver="myOwnCacheResolver"))  
public List<Band> listBand(int start, int count)
```

Programmatic resolution through an
implementation of the CacheResolver Interface

Working with CacheResolvers

```
@Cacheable("bands", cacheResolver="myOwnCacheResolver"))
public List<Band> listBand(int start, int count)
```

```
public class MyOwnCacheResolver extends AbstractCacheResolver {
    @Autowired
    public MyOwnCacheResolver(CacheManager cacheManager) {
        super(cacheManager);
    }

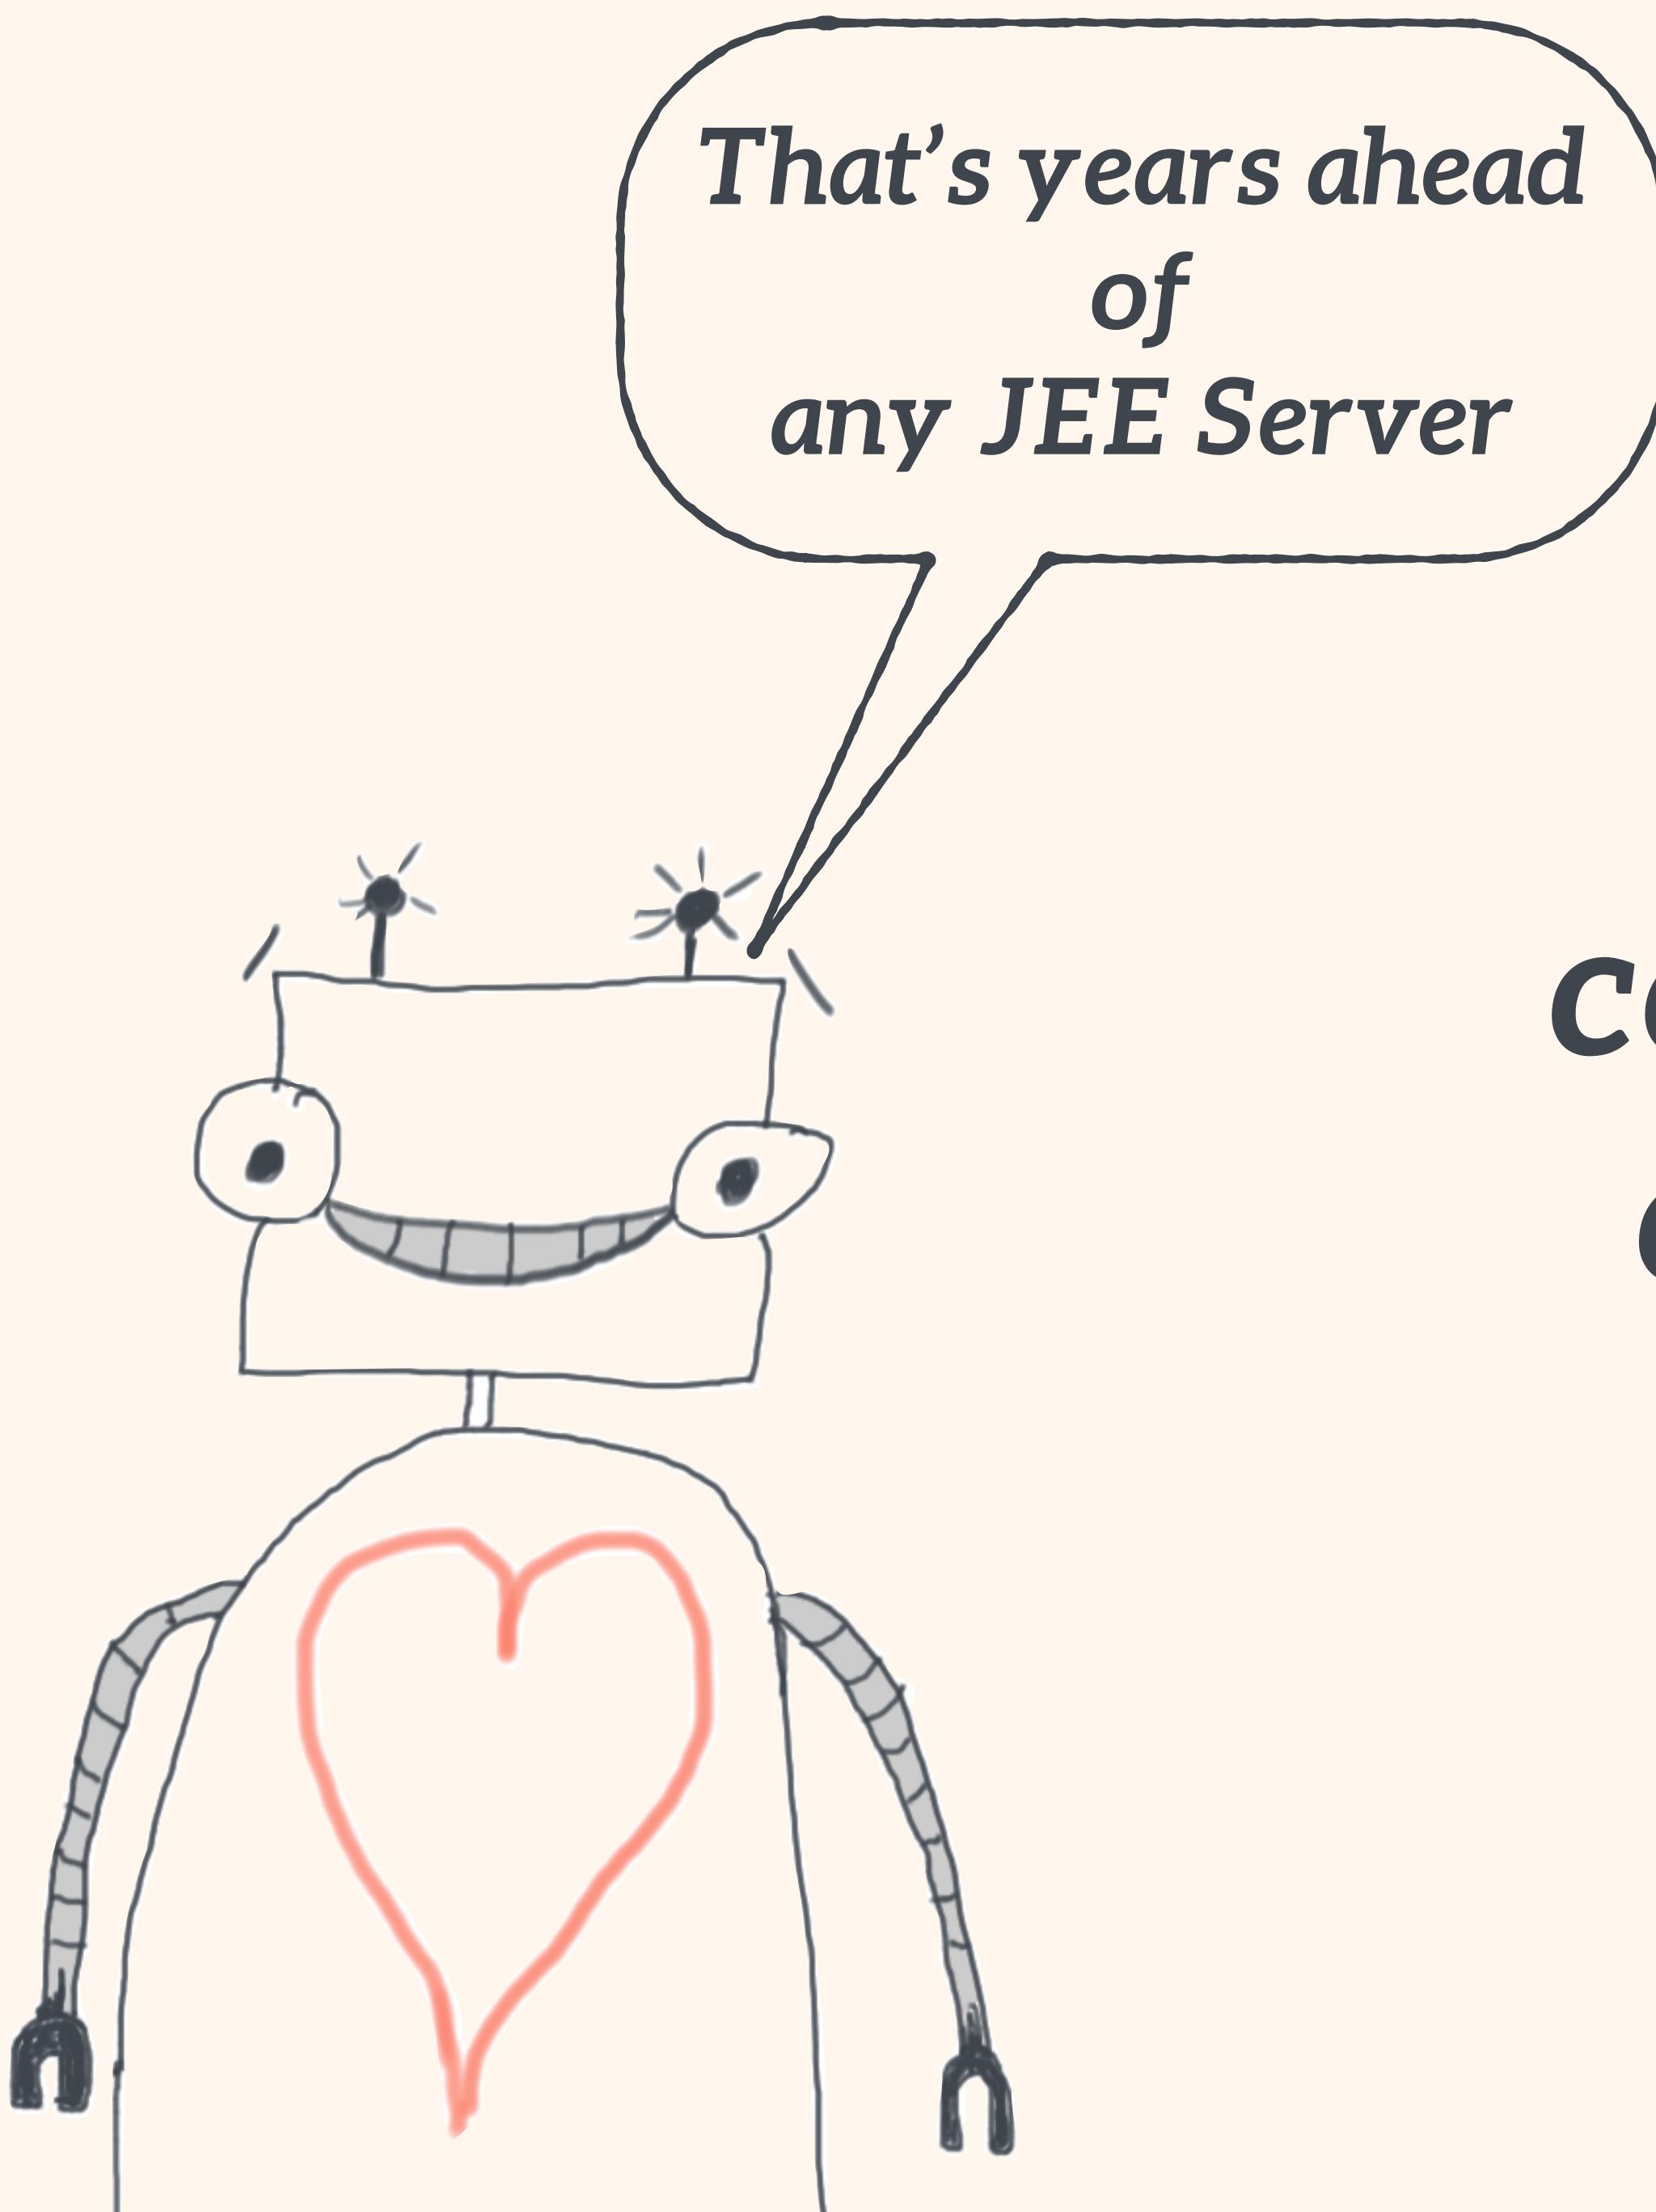
    protected Collection<String> getCacheNames(CacheOperationInvocationContext<?> context) {
        return getCacheNames(context.getTarget().getClass());
    }

    private getCacheNames(Class<?> businessServiceClass) {
        ...
    }
}
```

You can use your own custom Annotations

```
@Retention(RetentionPolicy.RUNTIME)
@Target({ElementType.METHOD})
@Cacheable("concerts", key="id")
public @interface DefaultConcertCacheable { }
```

```
@DefaultConcertCacheable
public Concert getConcert(Long id)
```

A cartoon illustration of a character with a large head, a wide smile showing many teeth, and antennae-like features on its head. A speech bubble originates from its mouth, containing the text.

*That's years ahead
of
any JEE Server*

***Spring 4.x is the first
commercially supported
container with JCache
(JSR-107) Support!***

Spring vs JCache Annotations

Spring	JCache	Description
<code>@Cacheable</code>	<code>@CacheResult</code>	Similar, but <code>@CacheResult</code> can cache Exceptions and force method execution
<code>@CacheEvict</code>	<code>@CacheRemove</code>	Similar, but <code>@CacheRemove</code> supports eviction in the case of Exceptions
<code>@CacheEvict (removeAll=true)</code>	<code>@CacheRemoveAll</code>	Same rules as for <code>@CacheEvict</code> vs <code>@CacheRemove</code>
<code>@CachePut</code>	<code>@CachePut</code>	Different semantic: cache content must be annotated with <code>@CacheValue</code> . JCache brings Exception caching and caching before or after method execution
<code>@CacheConfig</code>	<code>@CachePut</code>	Identical

*Except for the dependencies
JCache API and spring-context-
support no further steps need to
be taken to enable JCache*

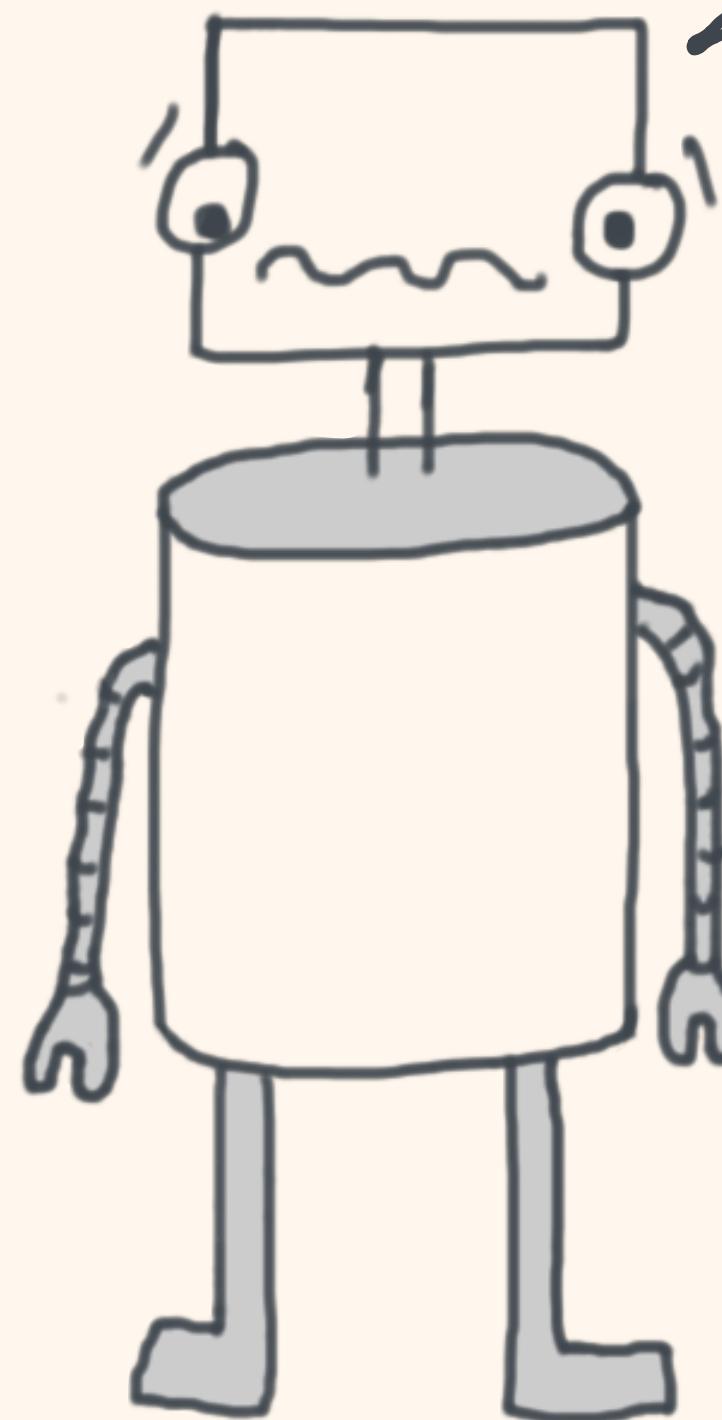
*Annotations in Spring
Applications*



Your Cache Provider has no
integration for Spring or JCache

Code walkthrough **CacheManager & Cache API**

How do I disable caching for Unit Tests?



```
<bean id="cacheManager"
  class="org.springframework.cache.support.CompositeCacheManager">
  <property name="cacheManagers">
    <list>
      <ref bean="guavaCache"/>
      <ref bean="ehCache"/>
    </list>
  </property>
  <property name="fallbackToNoOpCache" value="true"/>
</bean>
```

Thank you!

Michael Plöd - Freelance Consultant



<https://github.com/mploed>



@bitboss



<http://slideshare.net/mploed>