



# Implementing the Lambda Architecture with Spring XD

By Carlos Queiroz

# Me



Carlos Queiroz

Advanced Field Engineer

**Pivotal.**



# Agenda



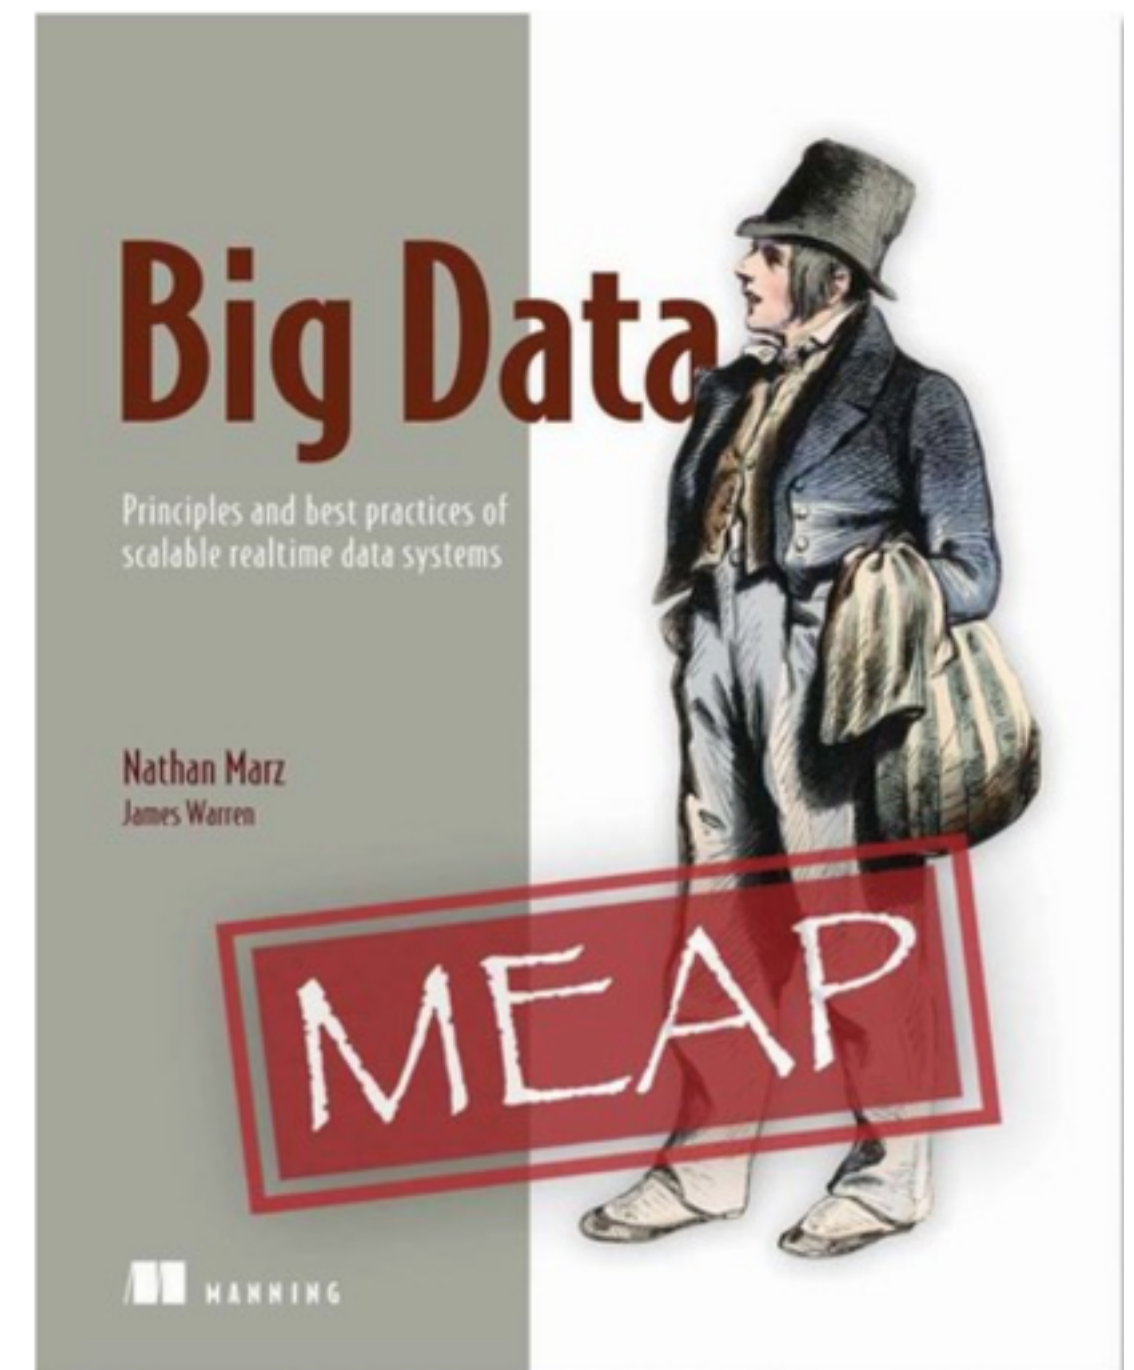
- Introduction to the Lambda Architecture
- Applying Lambda architecture to a business case
- Implementation details

# What is Lambda Architecture?

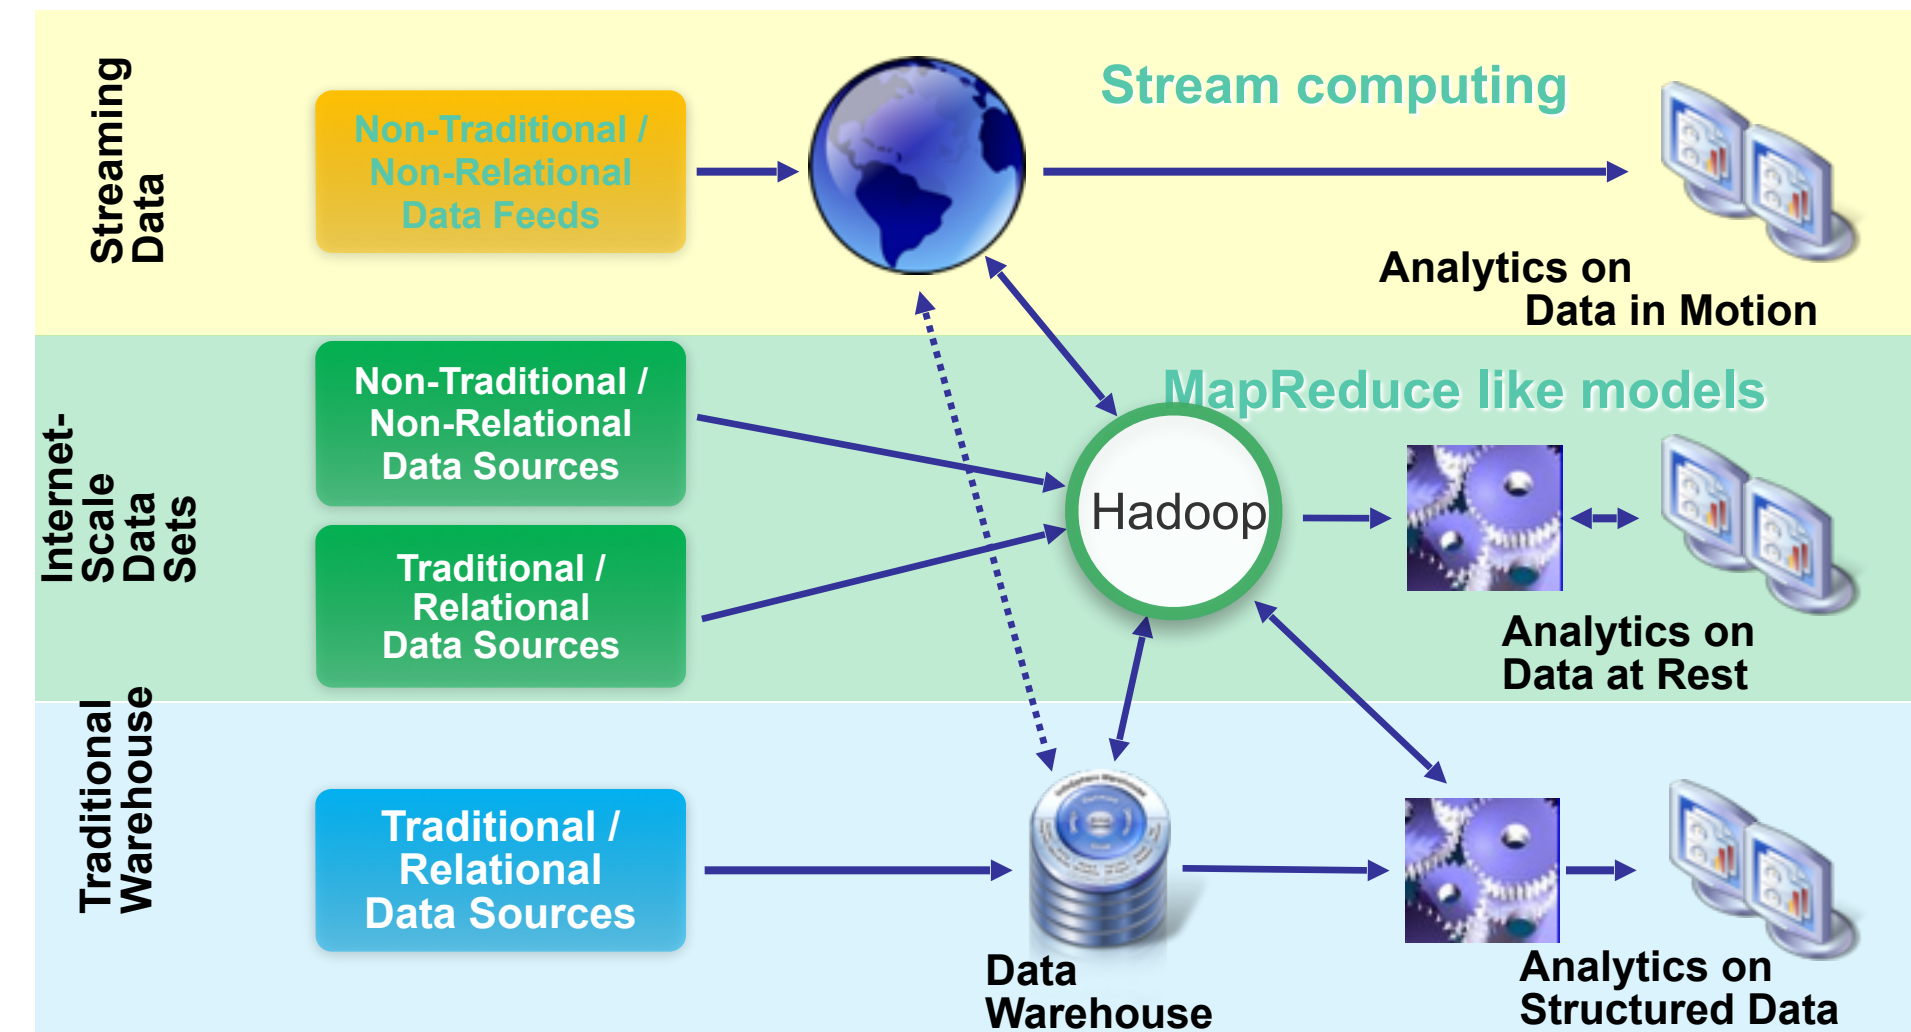
**Implementation** of a set of **desired properties** on general purpose big data systems.

**Generic architecture** addressing common requirements in *big data* applications.

Set of **design patterns** for dealing with historical and operational data.



# Approach is new, ideas not so...



## The Big Data Ecosystem

# Why Lambda Architecture?

To build a data system that can answer questions by running functions that take at the **entire dataset** as input. **A general purpose data system**

# Desired Properties of such systems

Fault-tolerance

Generic

Scale linearly, horizontally.

Extensible

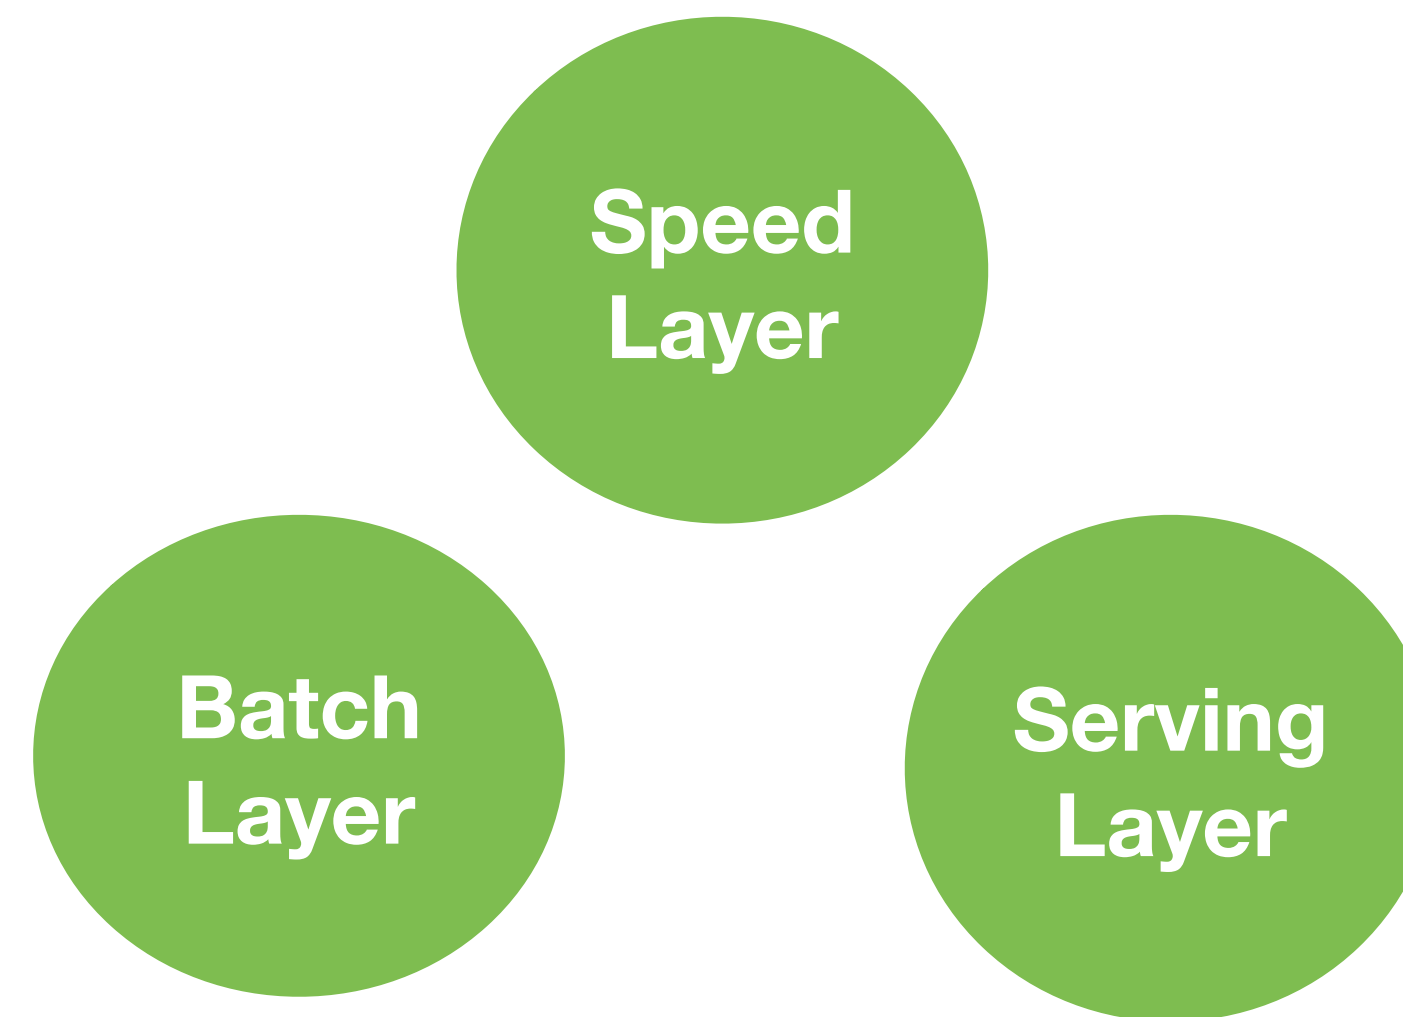
Be able to achieve low latency updates when necessary.

Be able to “ask” arbitrary questions to the system



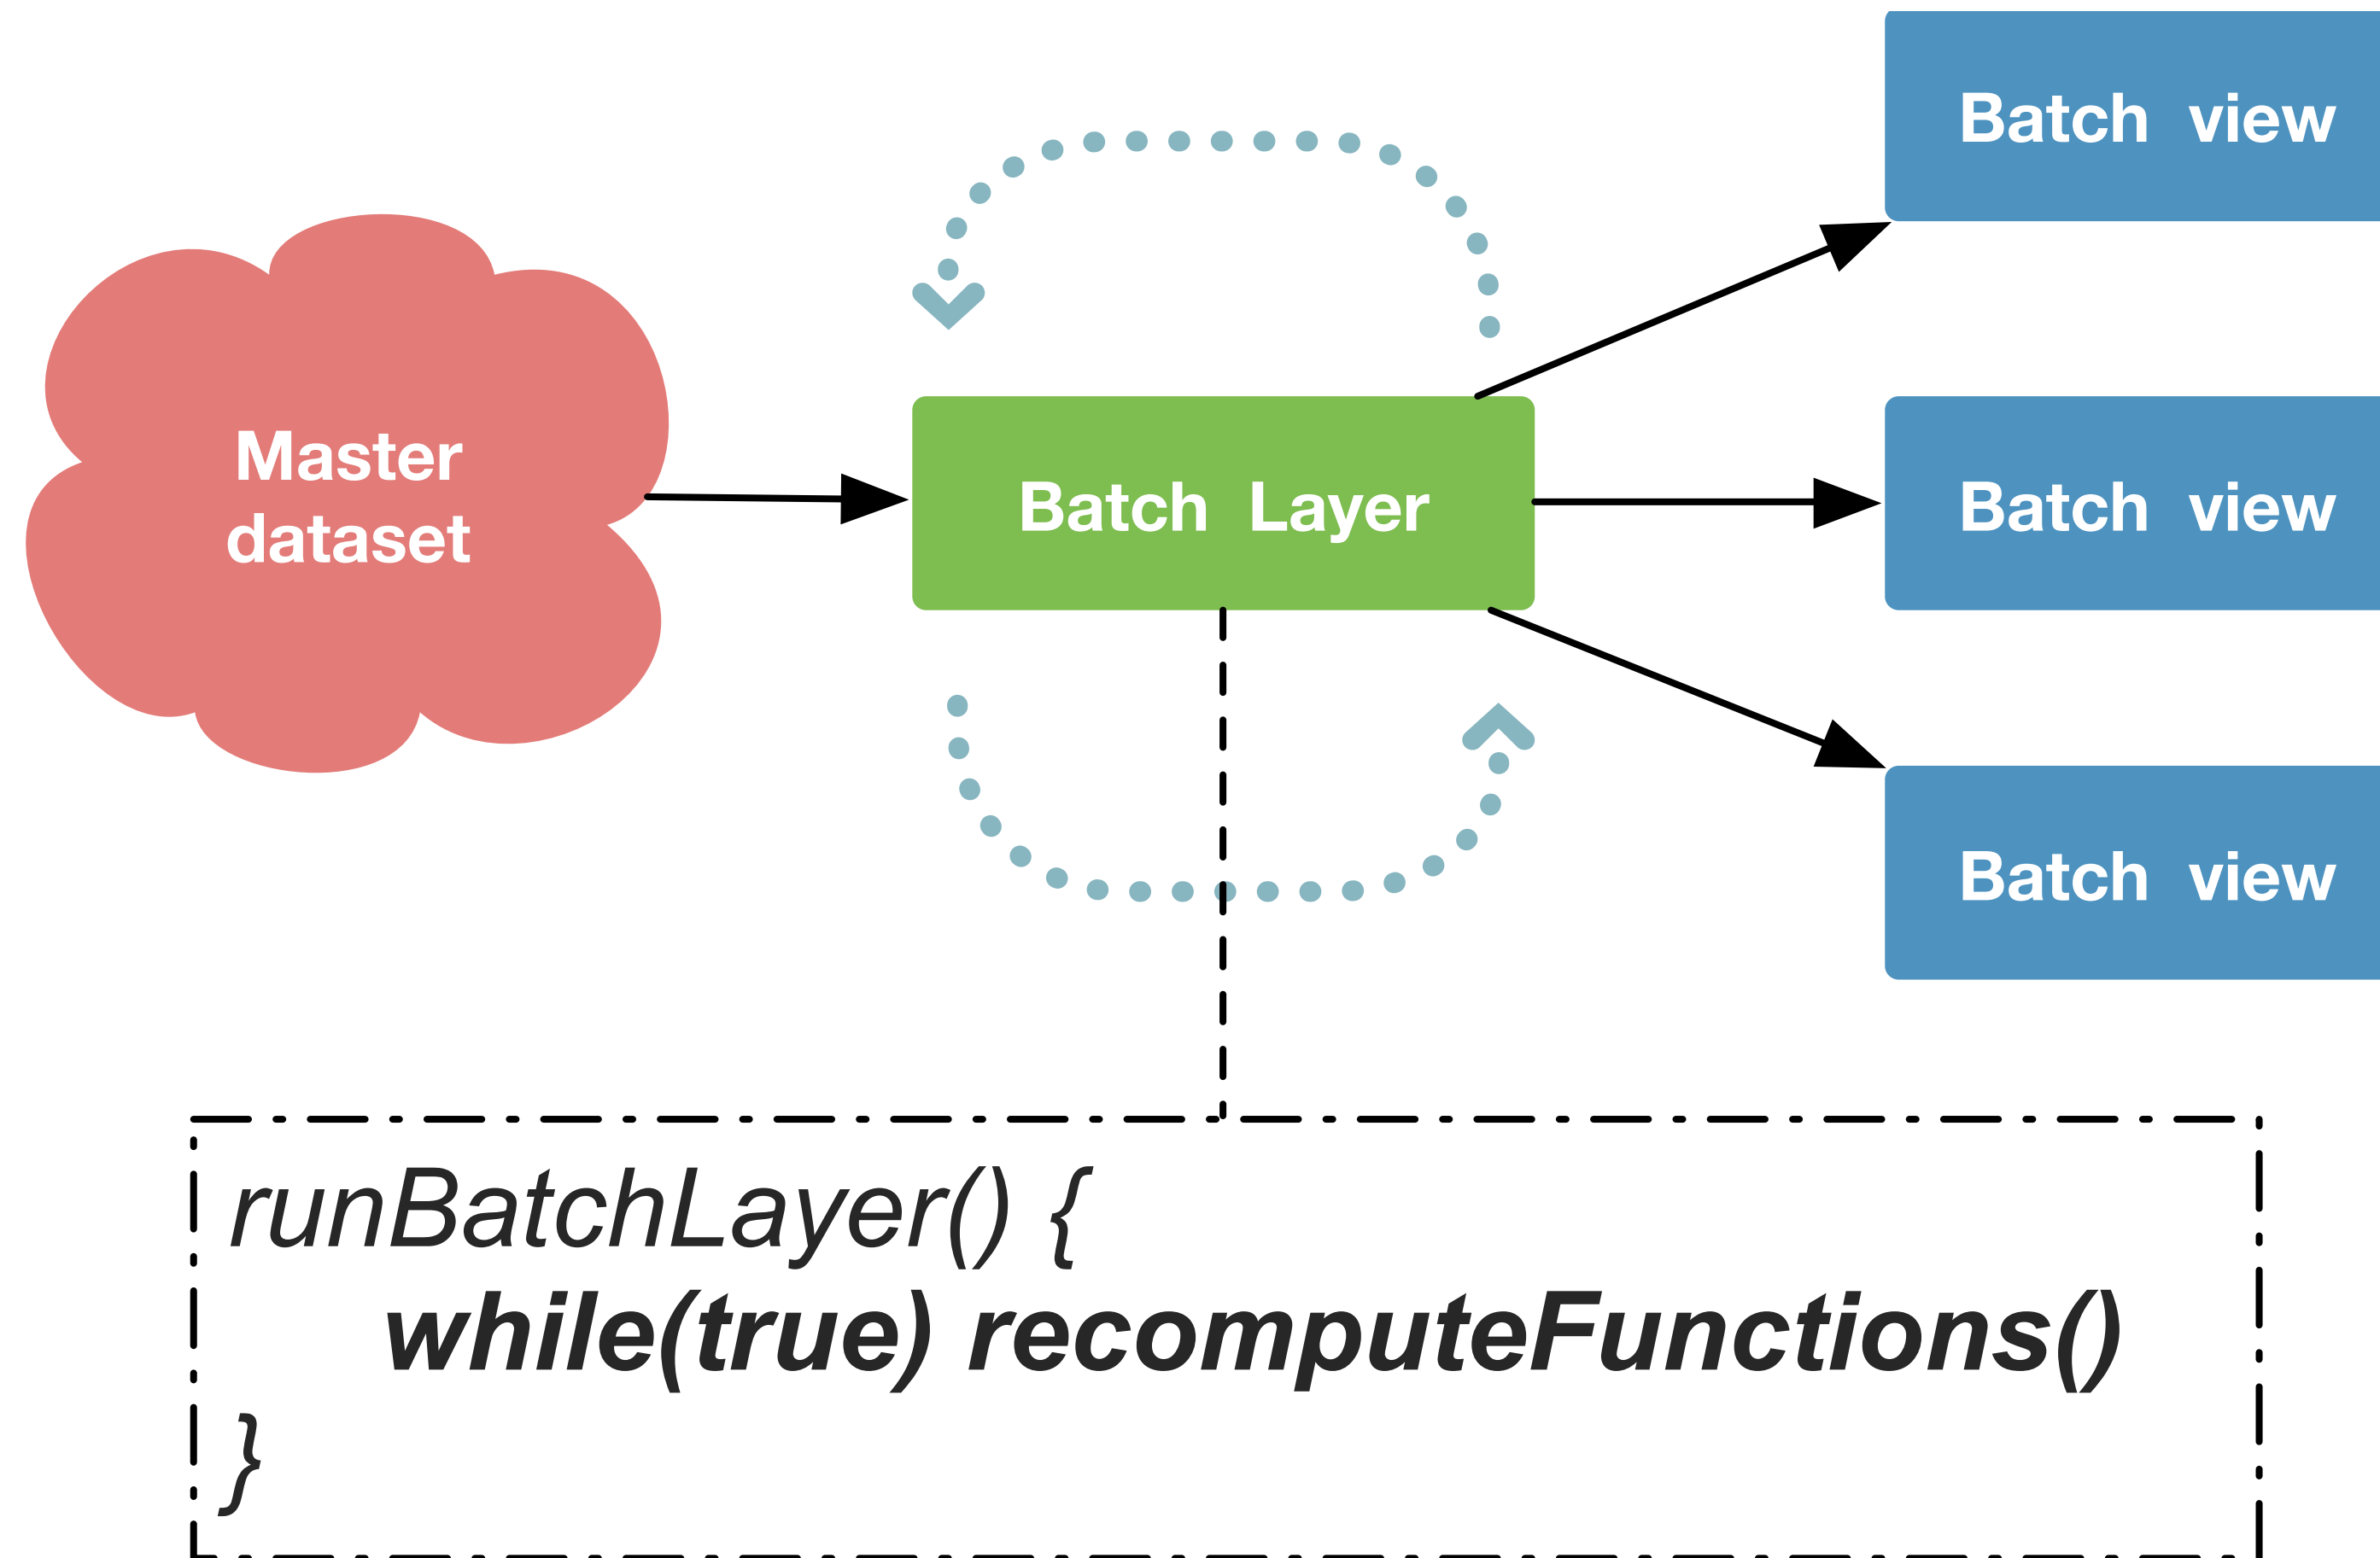
# How to support those properties?

Decompose the problem into pieces

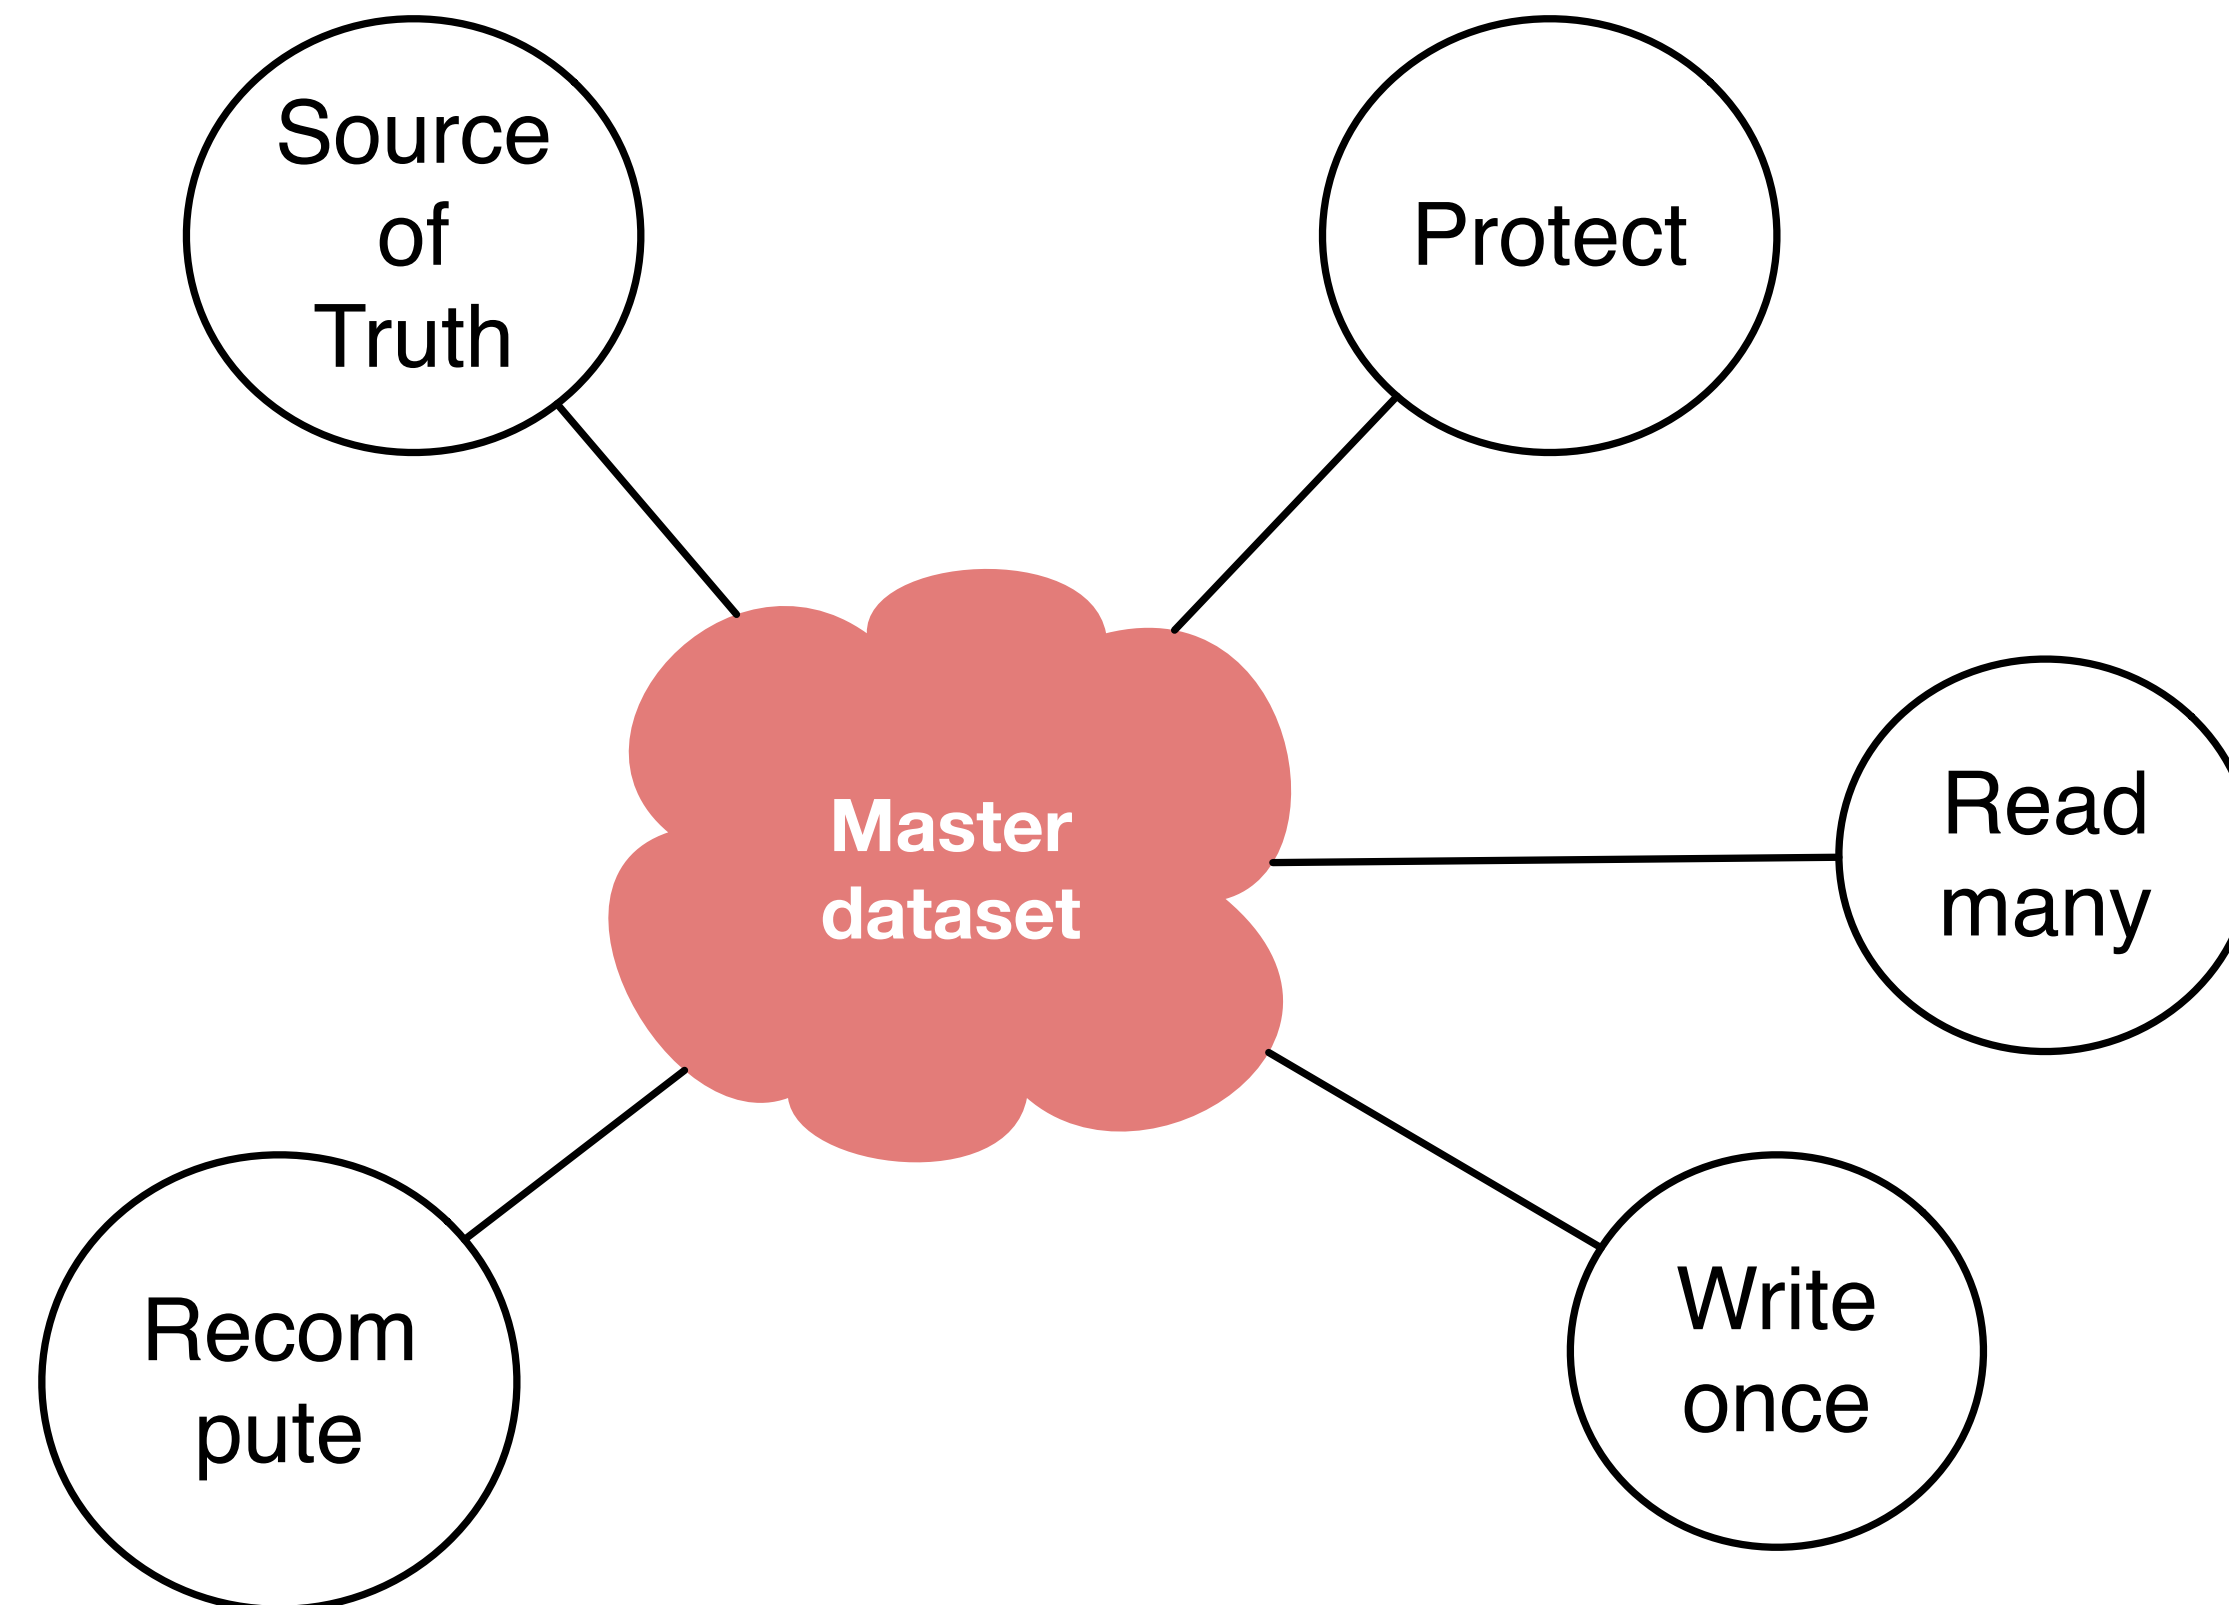




# Batch Layer



# The Master dataset



# Master dataset requirements

Operation	Requisite	Comments
<i>Writes</i>	<i>Efficient appends of new data</i>	<i>Basically add new pieces of data. Easy to append new data</i>
Writes	<b>Scalable storage</b>	Need to handle possibly, petabytes of data
<i>Reads</i>	<i>Support for parallel processing</i>	<i>Functions usually work on the entire dataset. Need to support handling large amounts of data</i>
Reads	<b>Vertically partition data</b>	Not necessary to look all the data all the time. Some functions may need to look at only relevant data (e.g. 1 week of calls)
<i>Writes/Reads</i>	<i>Costs for processing. Flexible storage</i>	<i>Storage costs money (a lot). Need flexibility on how to store and compress data.</i>

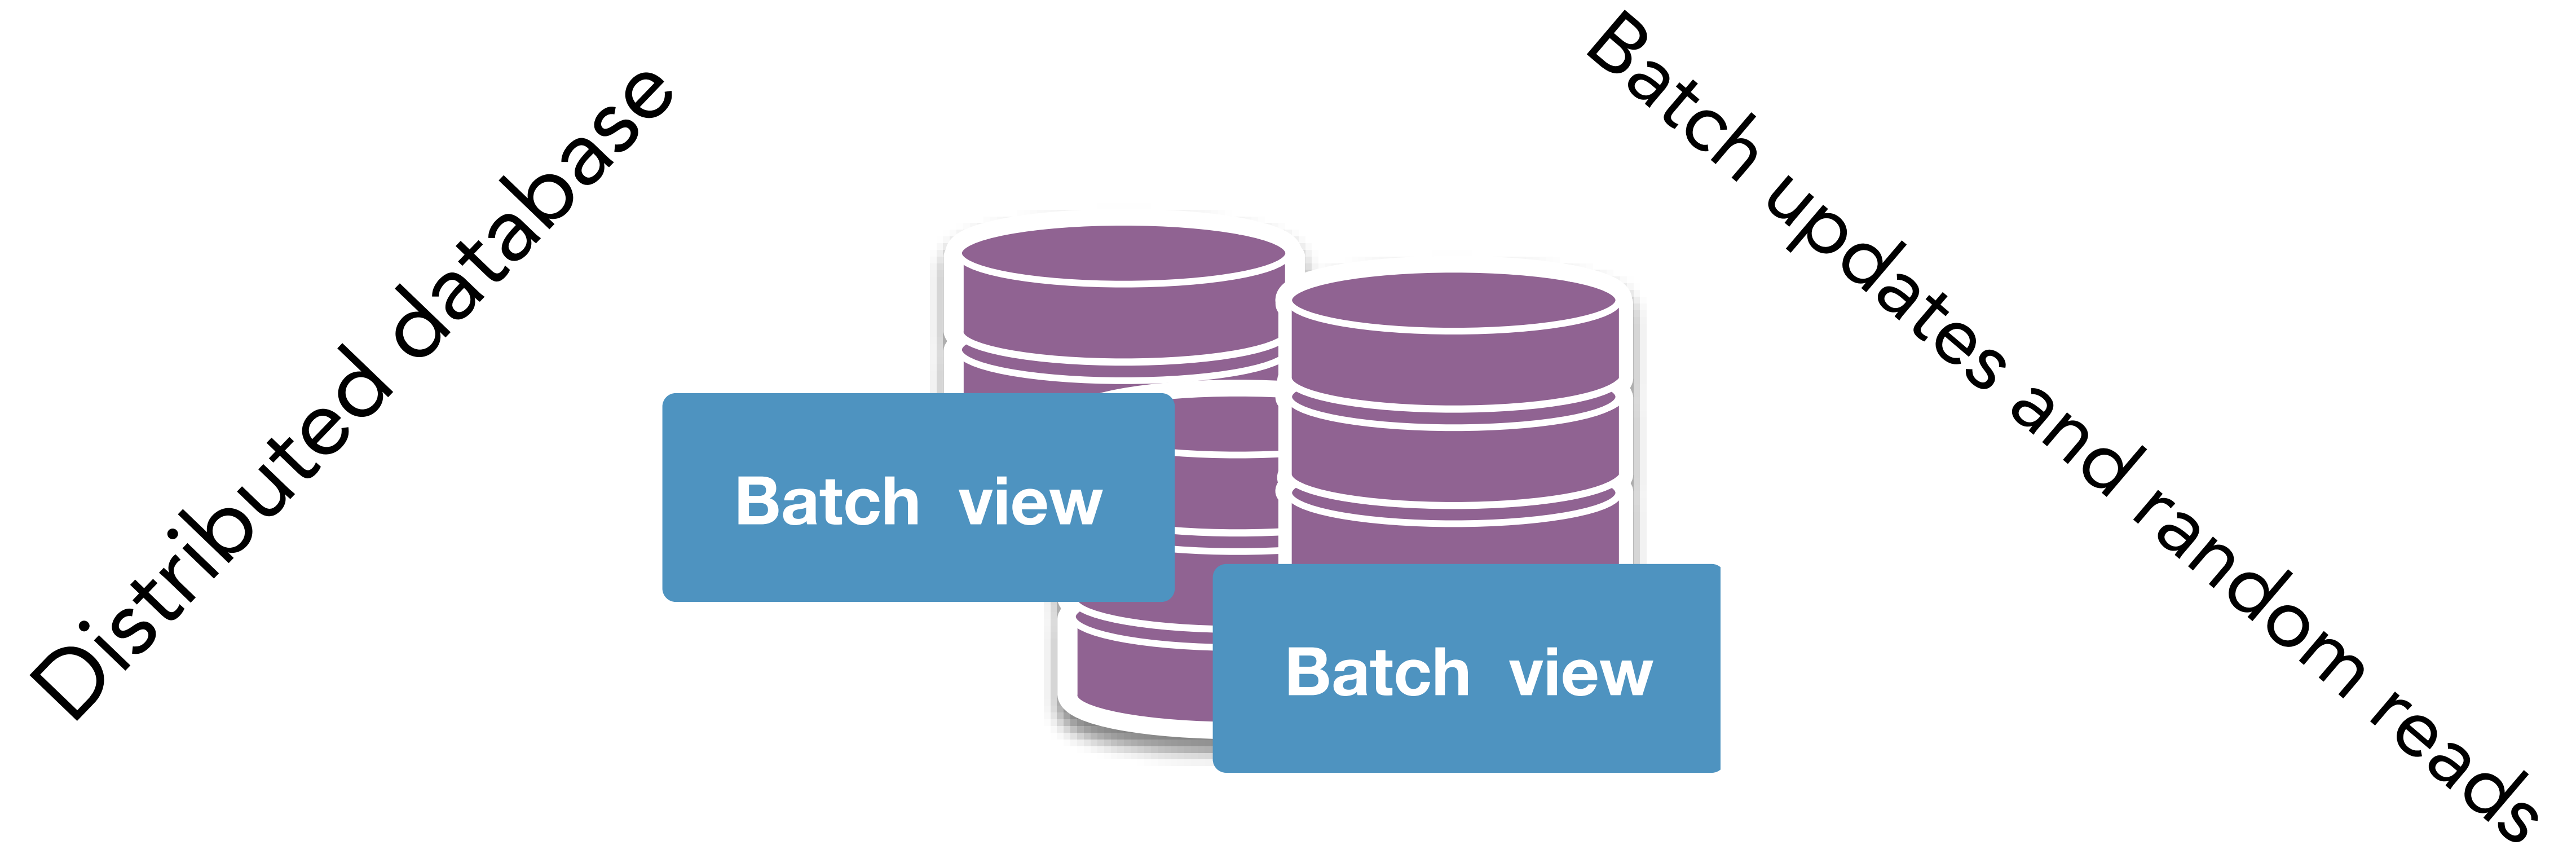
# The Master dataset



## HDFS - Hadoop Distributed File System

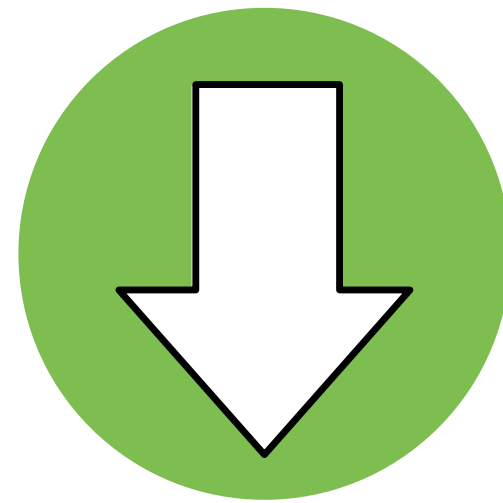
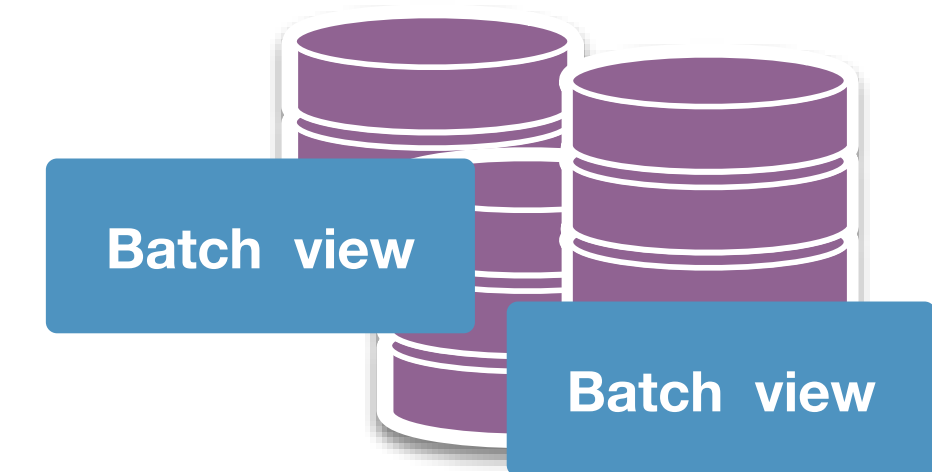


# Serving Layer

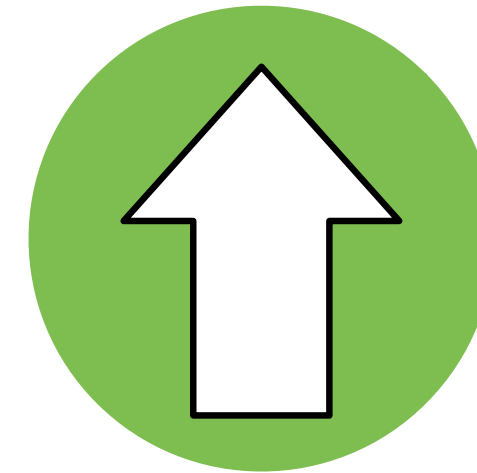


Makes the batch views “*queryable*”

# Desired properties on the Serving layer



Low Latency



High Throughput

# Serving layer database requirements

Fault-tolerant

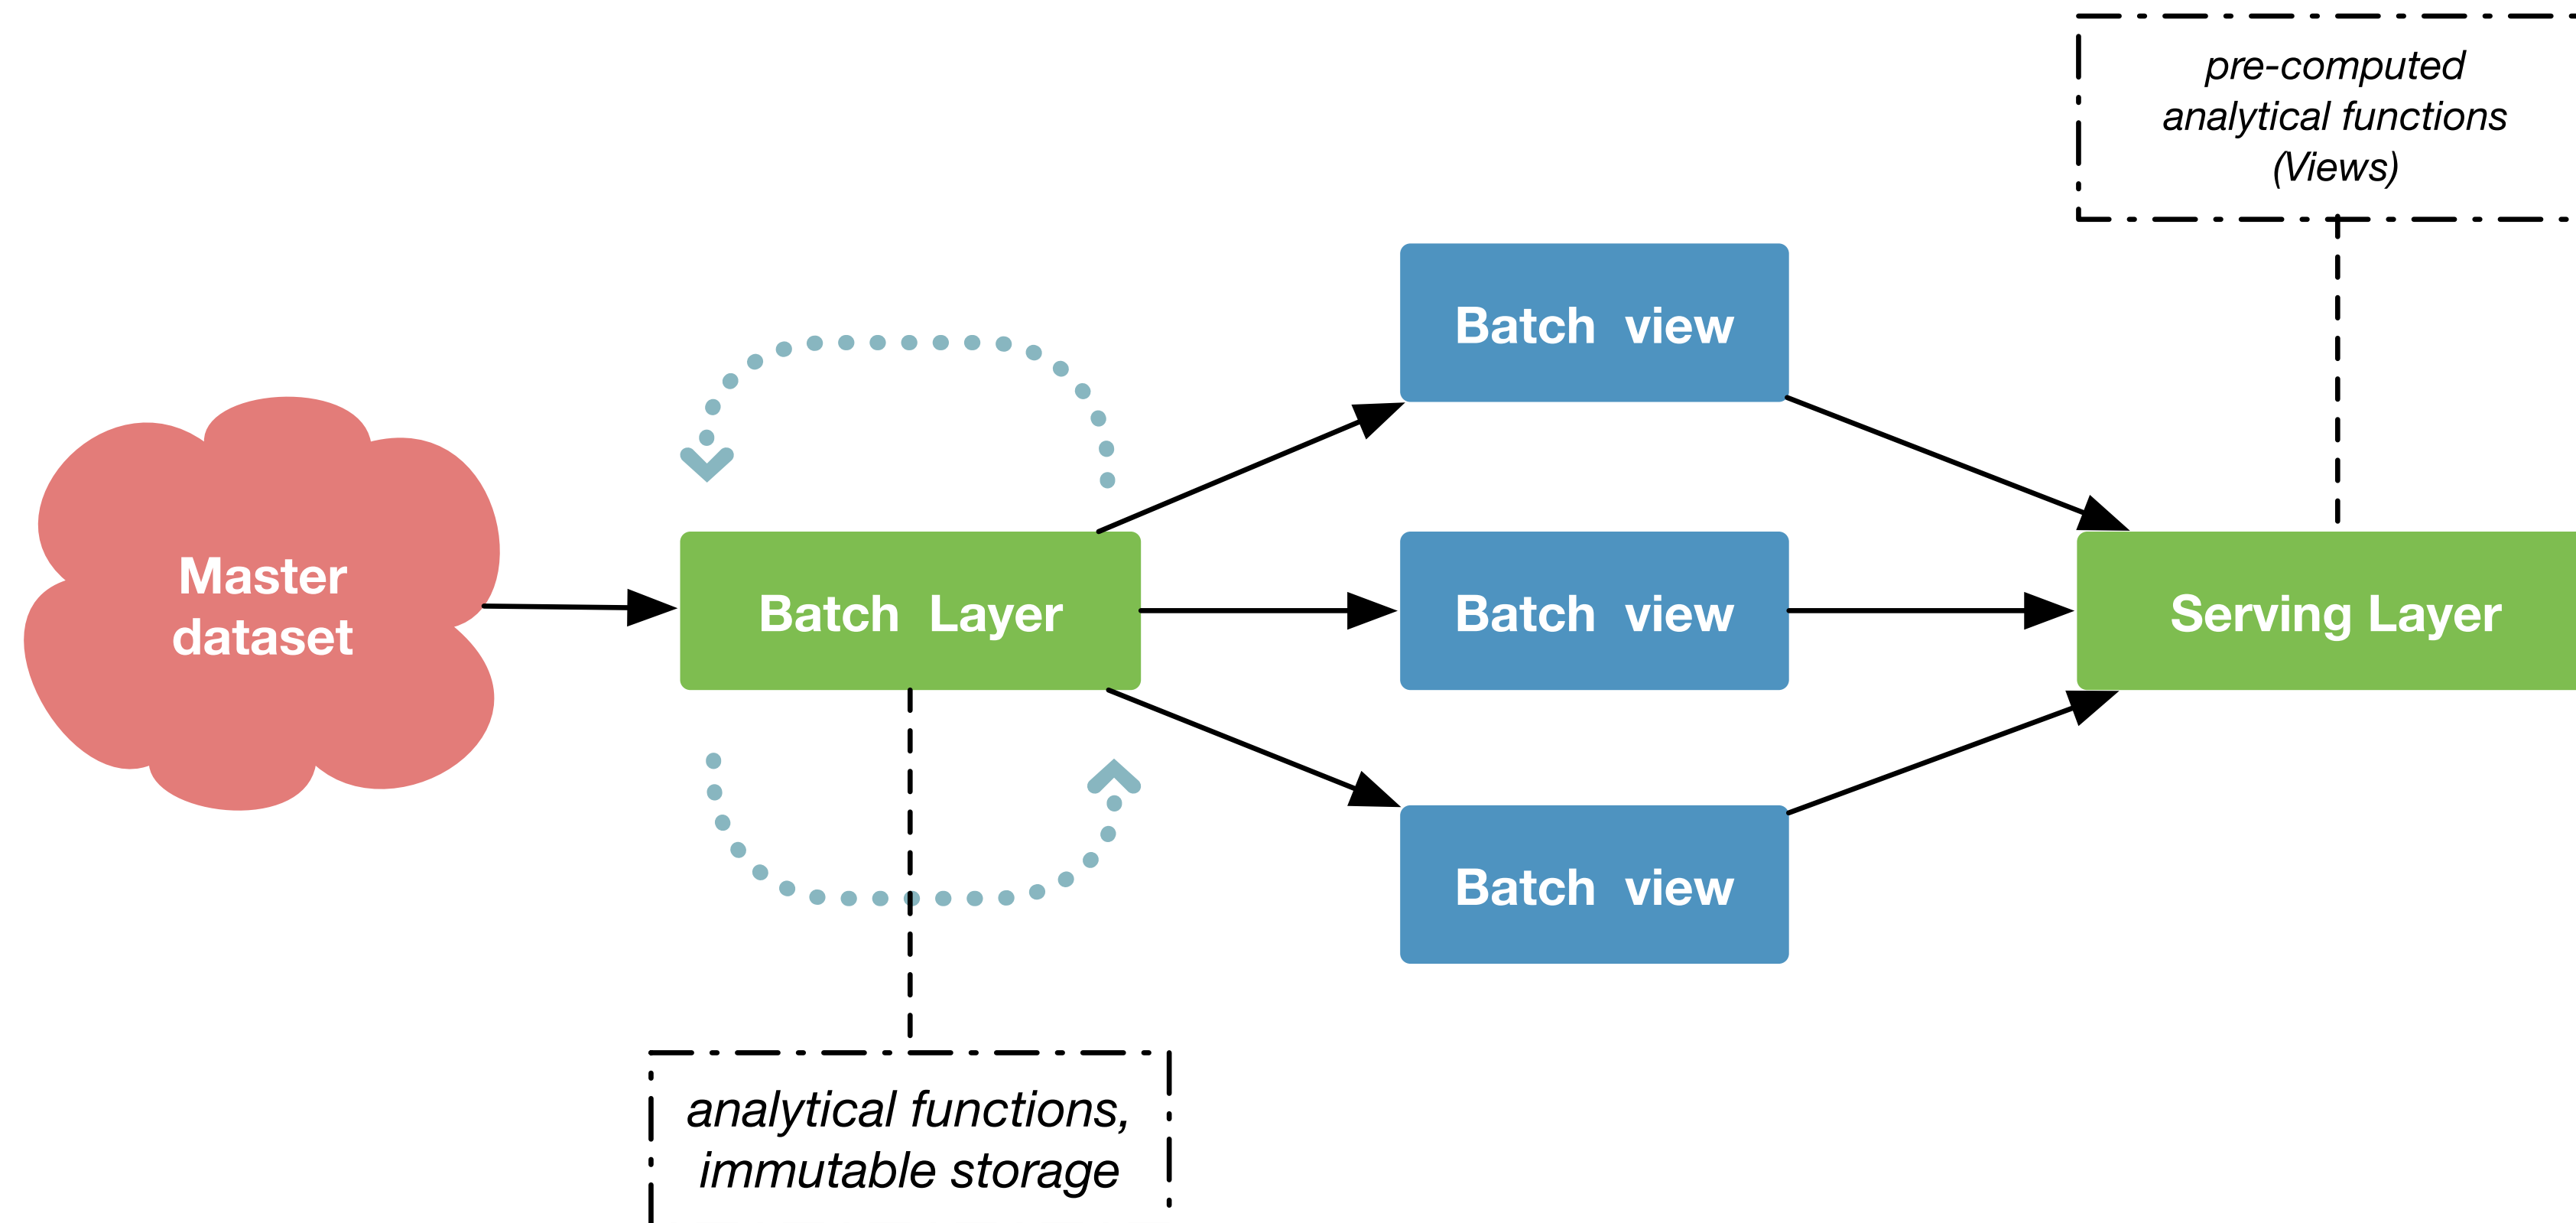
Batch writable



Random reads

Scalable

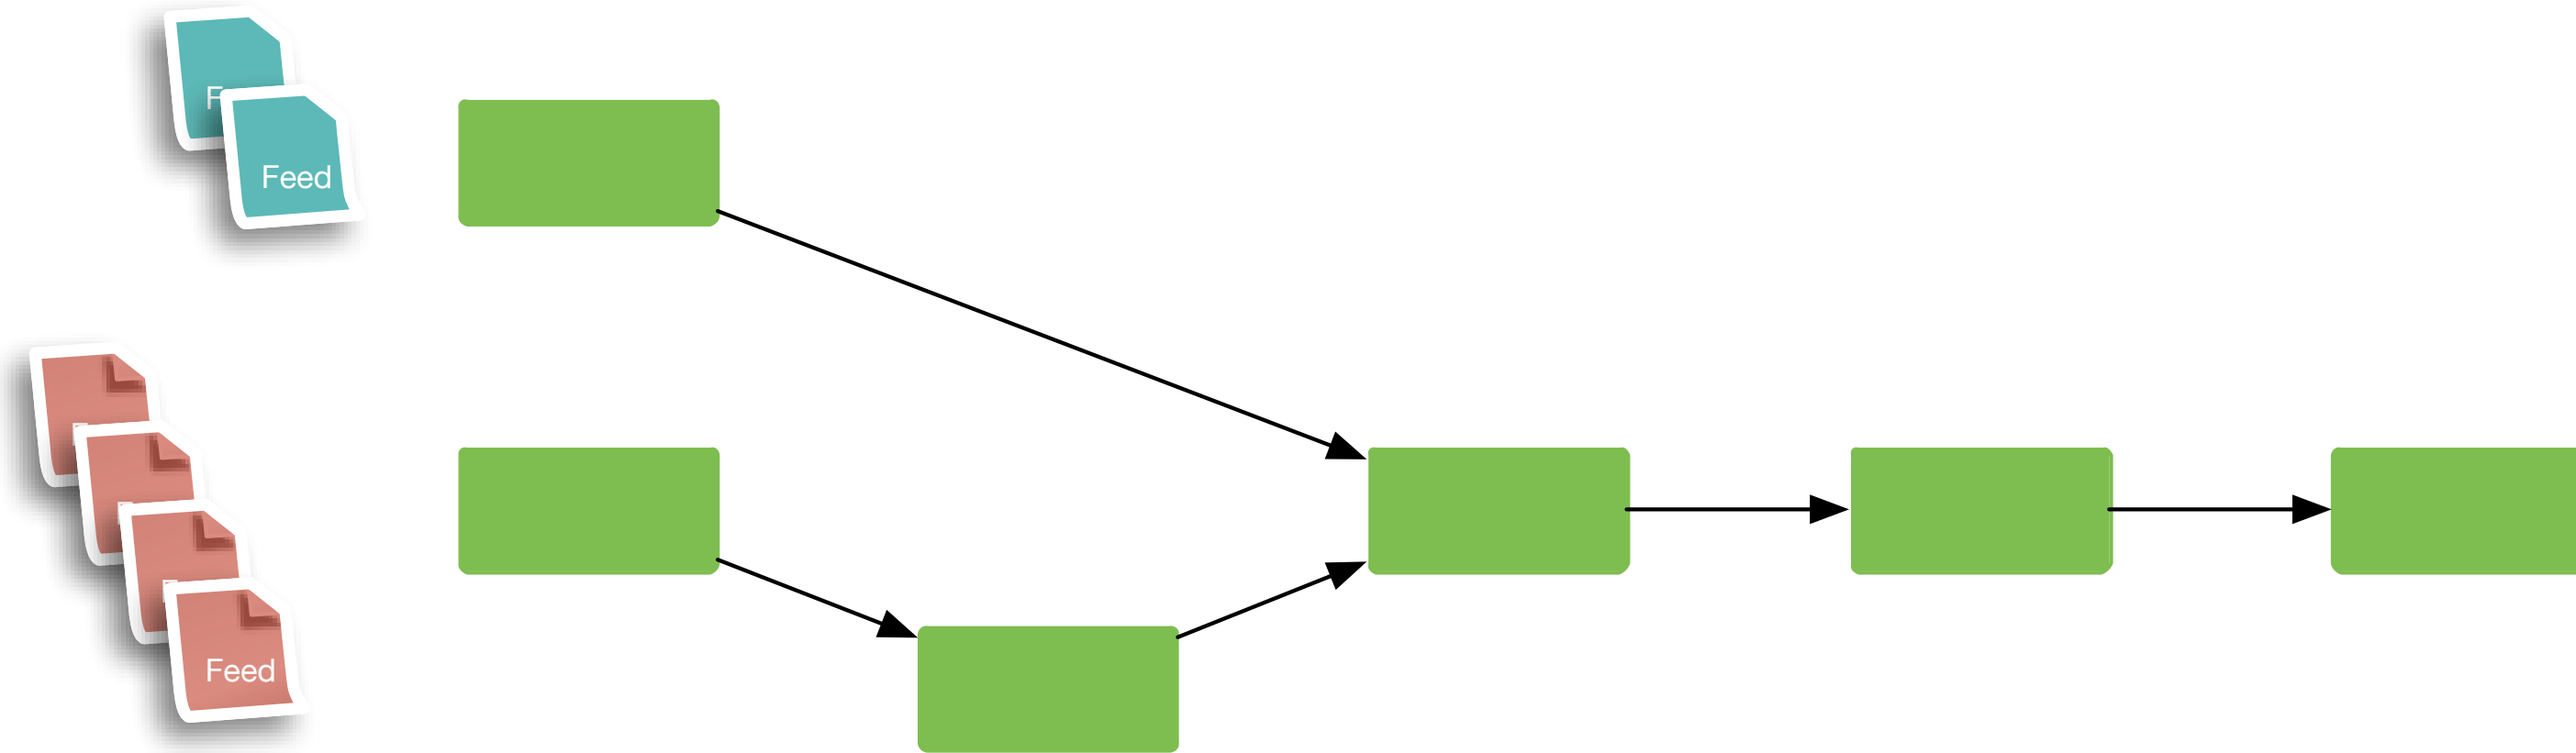
# Batch and serving layers



Only property missing - **low latency updates**



# Speed layer



**Allows arbitrary functions computed on arbitrary data  
on (near) real-time.**

# Speed layer



***Narrow but more up to date view***



***Depending of functions complexity incremental  
computation approach is recommended***

# Incremental computation



***realtime view = function(new data, realtime view)***

# Eventual Accuracy

Some computations are harder to compute

For such cases approximations are used. Results are approximate to the correct answer

Sophisticated queries such as realtime machine learning are usually done with eventual accuracy. Too complex to be computed exactly.



# Approximation & Randomisation

## Approximation

find an answer correct within some factor  
*(answer that is within 10% of correct result)*

## Randomisation

allow a small probability of failure  
*(1 in 10,000 give the wrong answer)*

# Synopses structures

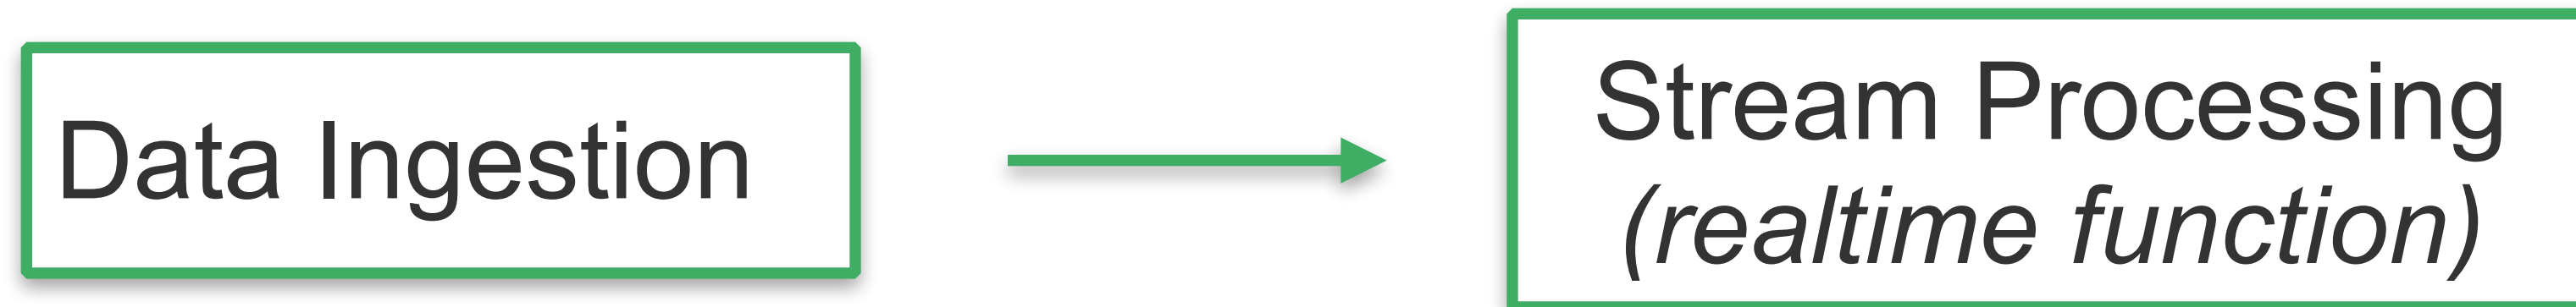
- Sampling
- Sketches
- Histograms
- Wavelets

*<http://charuaggarwal.net/synopsis.pdf>*

## Library implementations

- algebird (<https://github.com/twitter/algebird>)
- samoa (<https://github.com/yahoo/samoa/wiki>)

# Stream processing (real-time function)

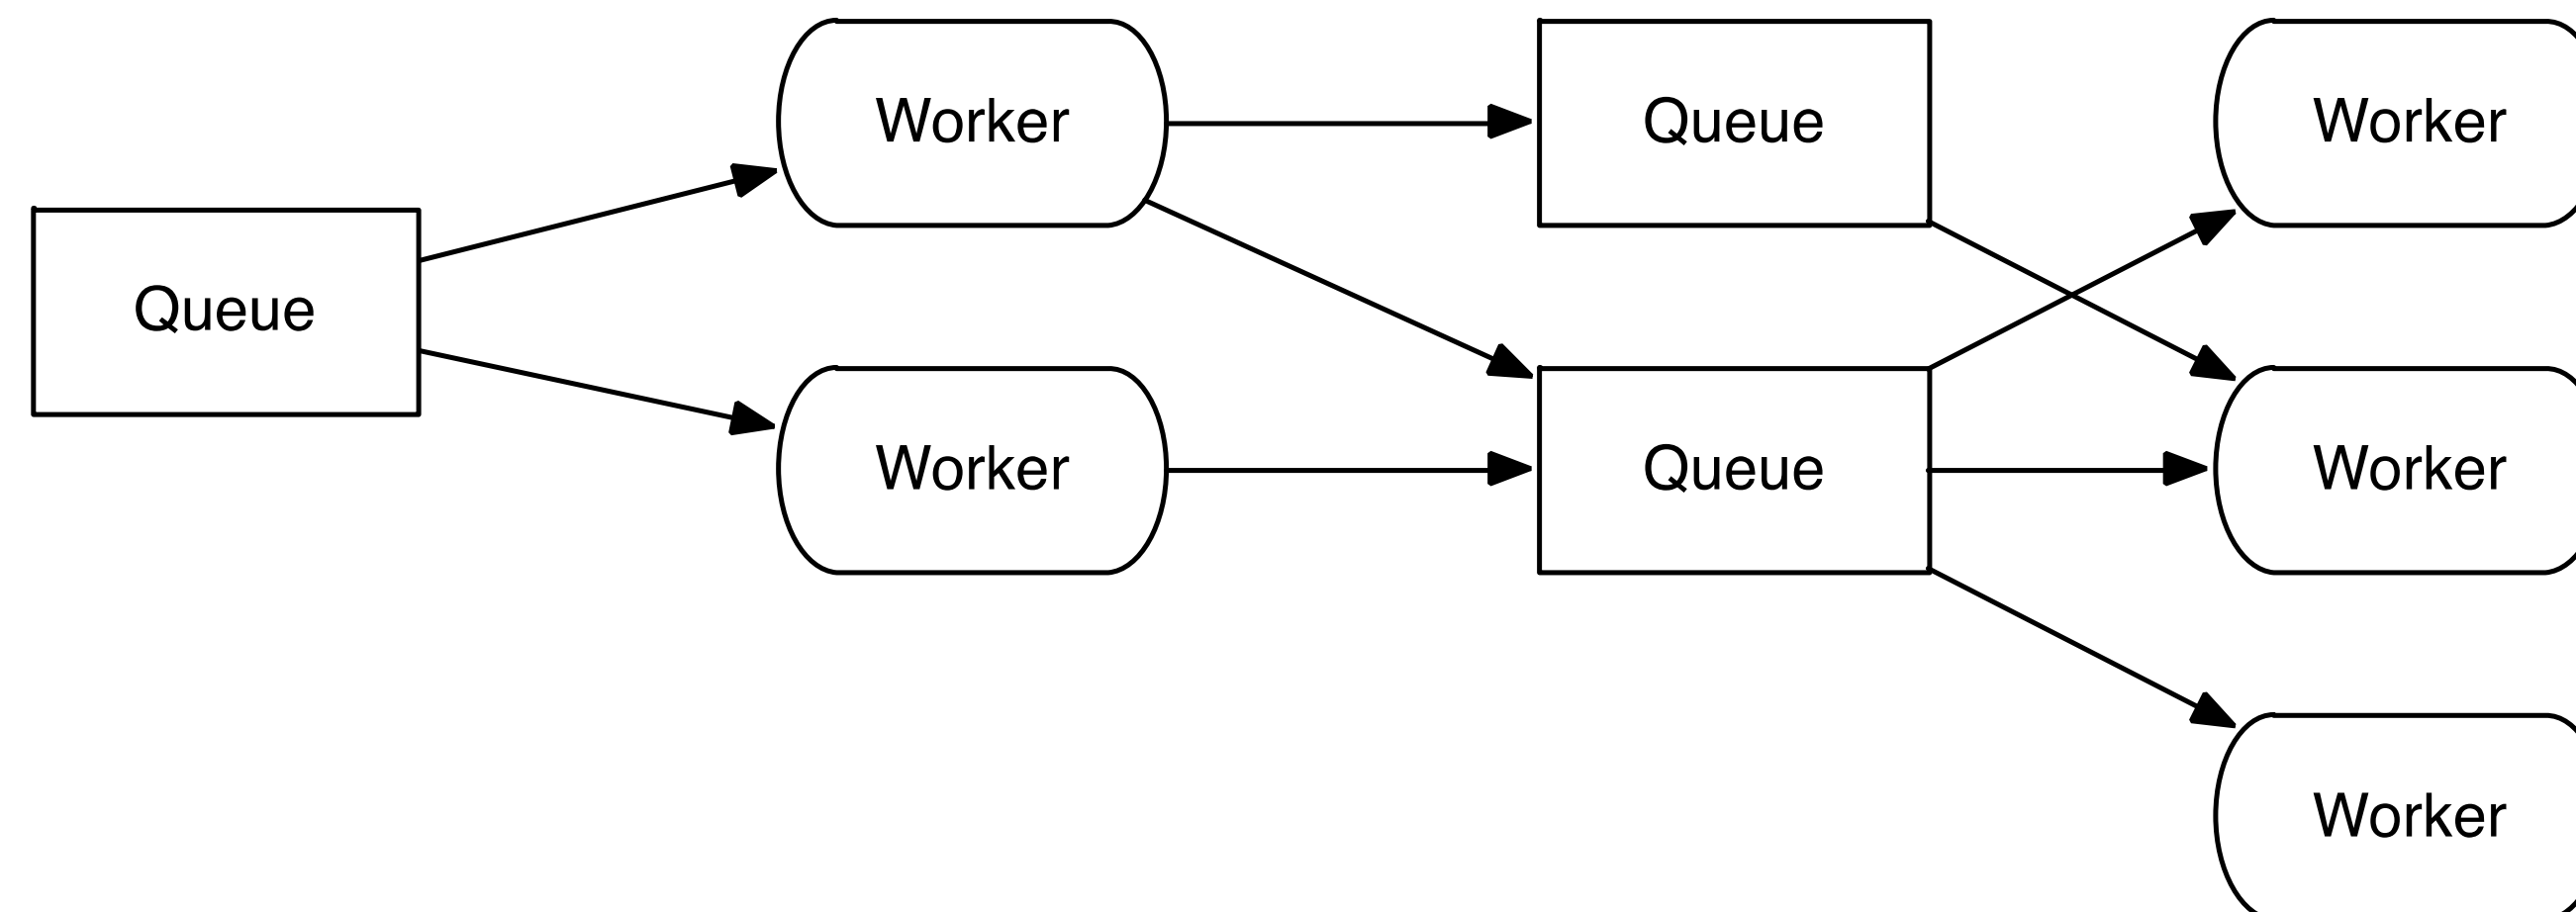


*Run the realtime functions to update the realtime views*

	one-at-a-time	micro-batched
lower latency	✓	✗
High throughput	✗	✓
At-least-once semantics	✓	✓
exactly-once semantics	some cases	✓
Simpler programming model	✓	✗

# One-at-a-time

Divide your processing into worker processes, and put queues between the worker processes

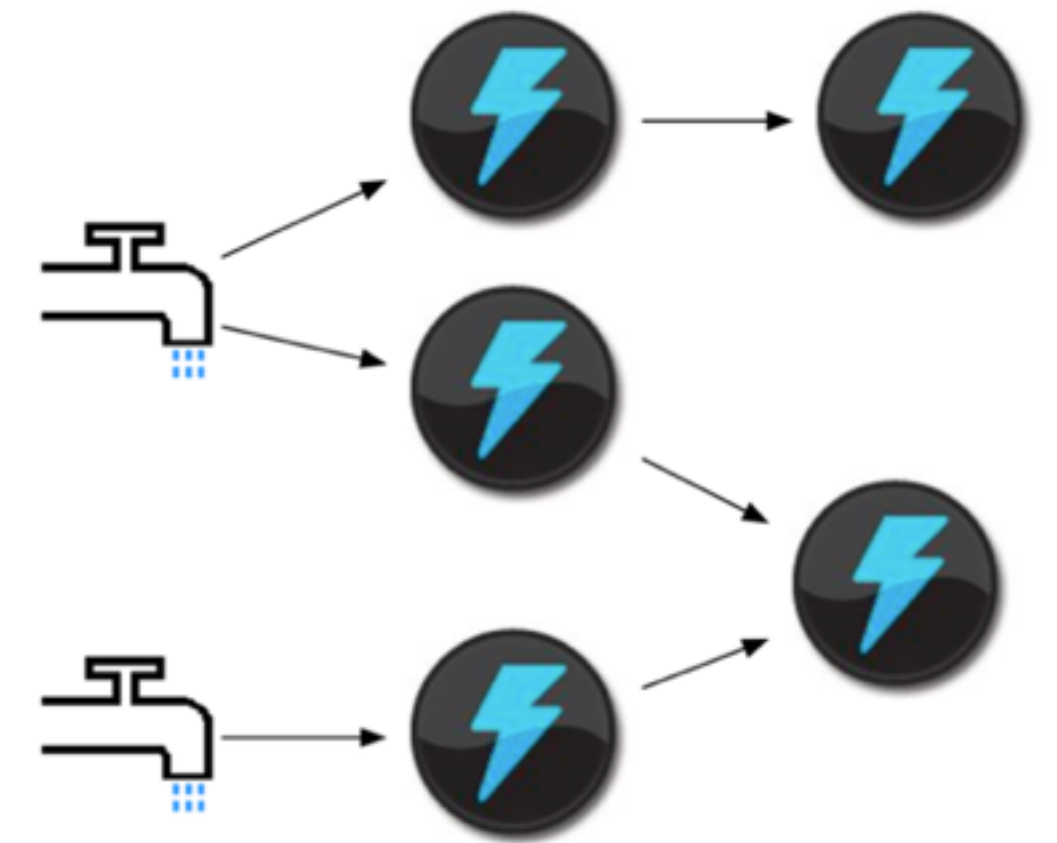


***Queues and workers paradigm***



# Generalised one-at-a-time approach

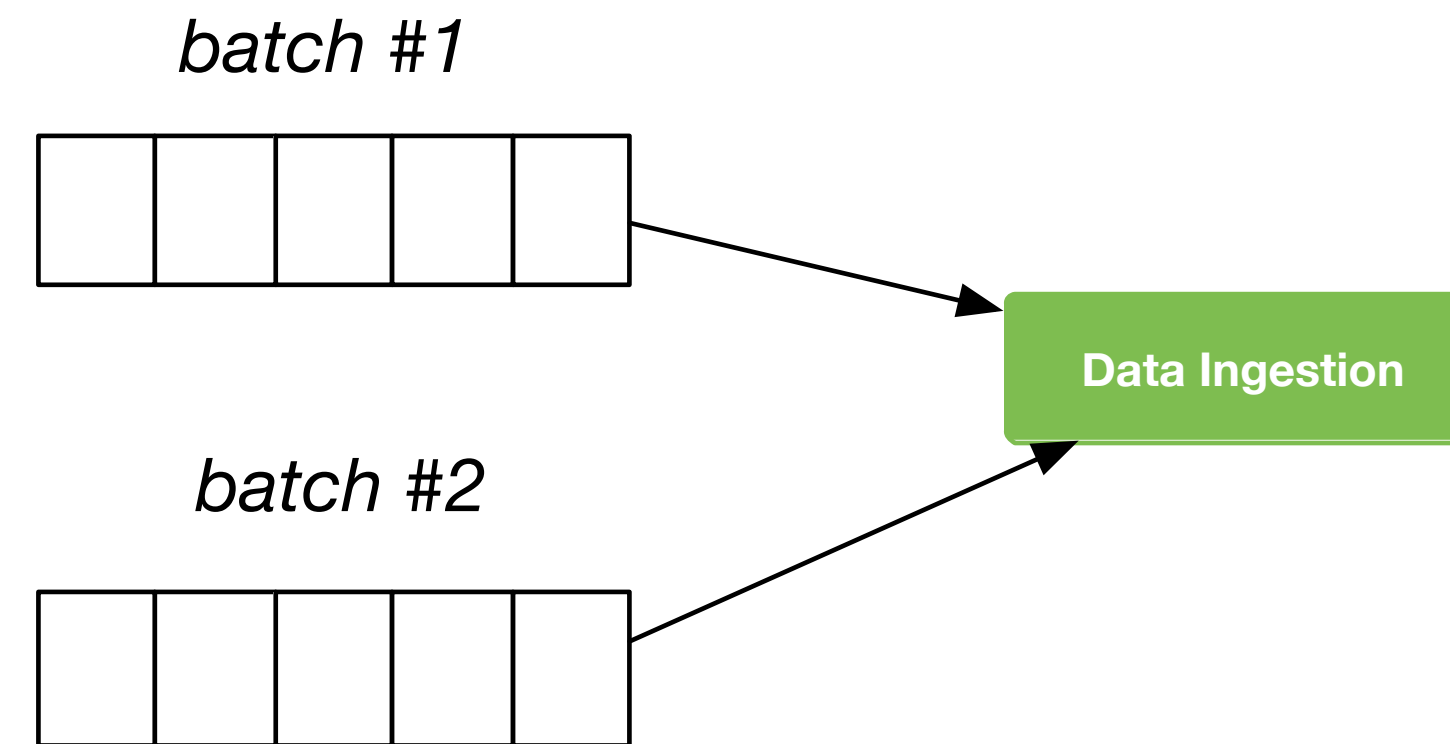
- Works on a higher level
- Stream computation defined as a graph (usually).
  - Storm, InfoSphere Streams models
- Filters and pipes
- Spring XD
- *At least once* in case of failures.



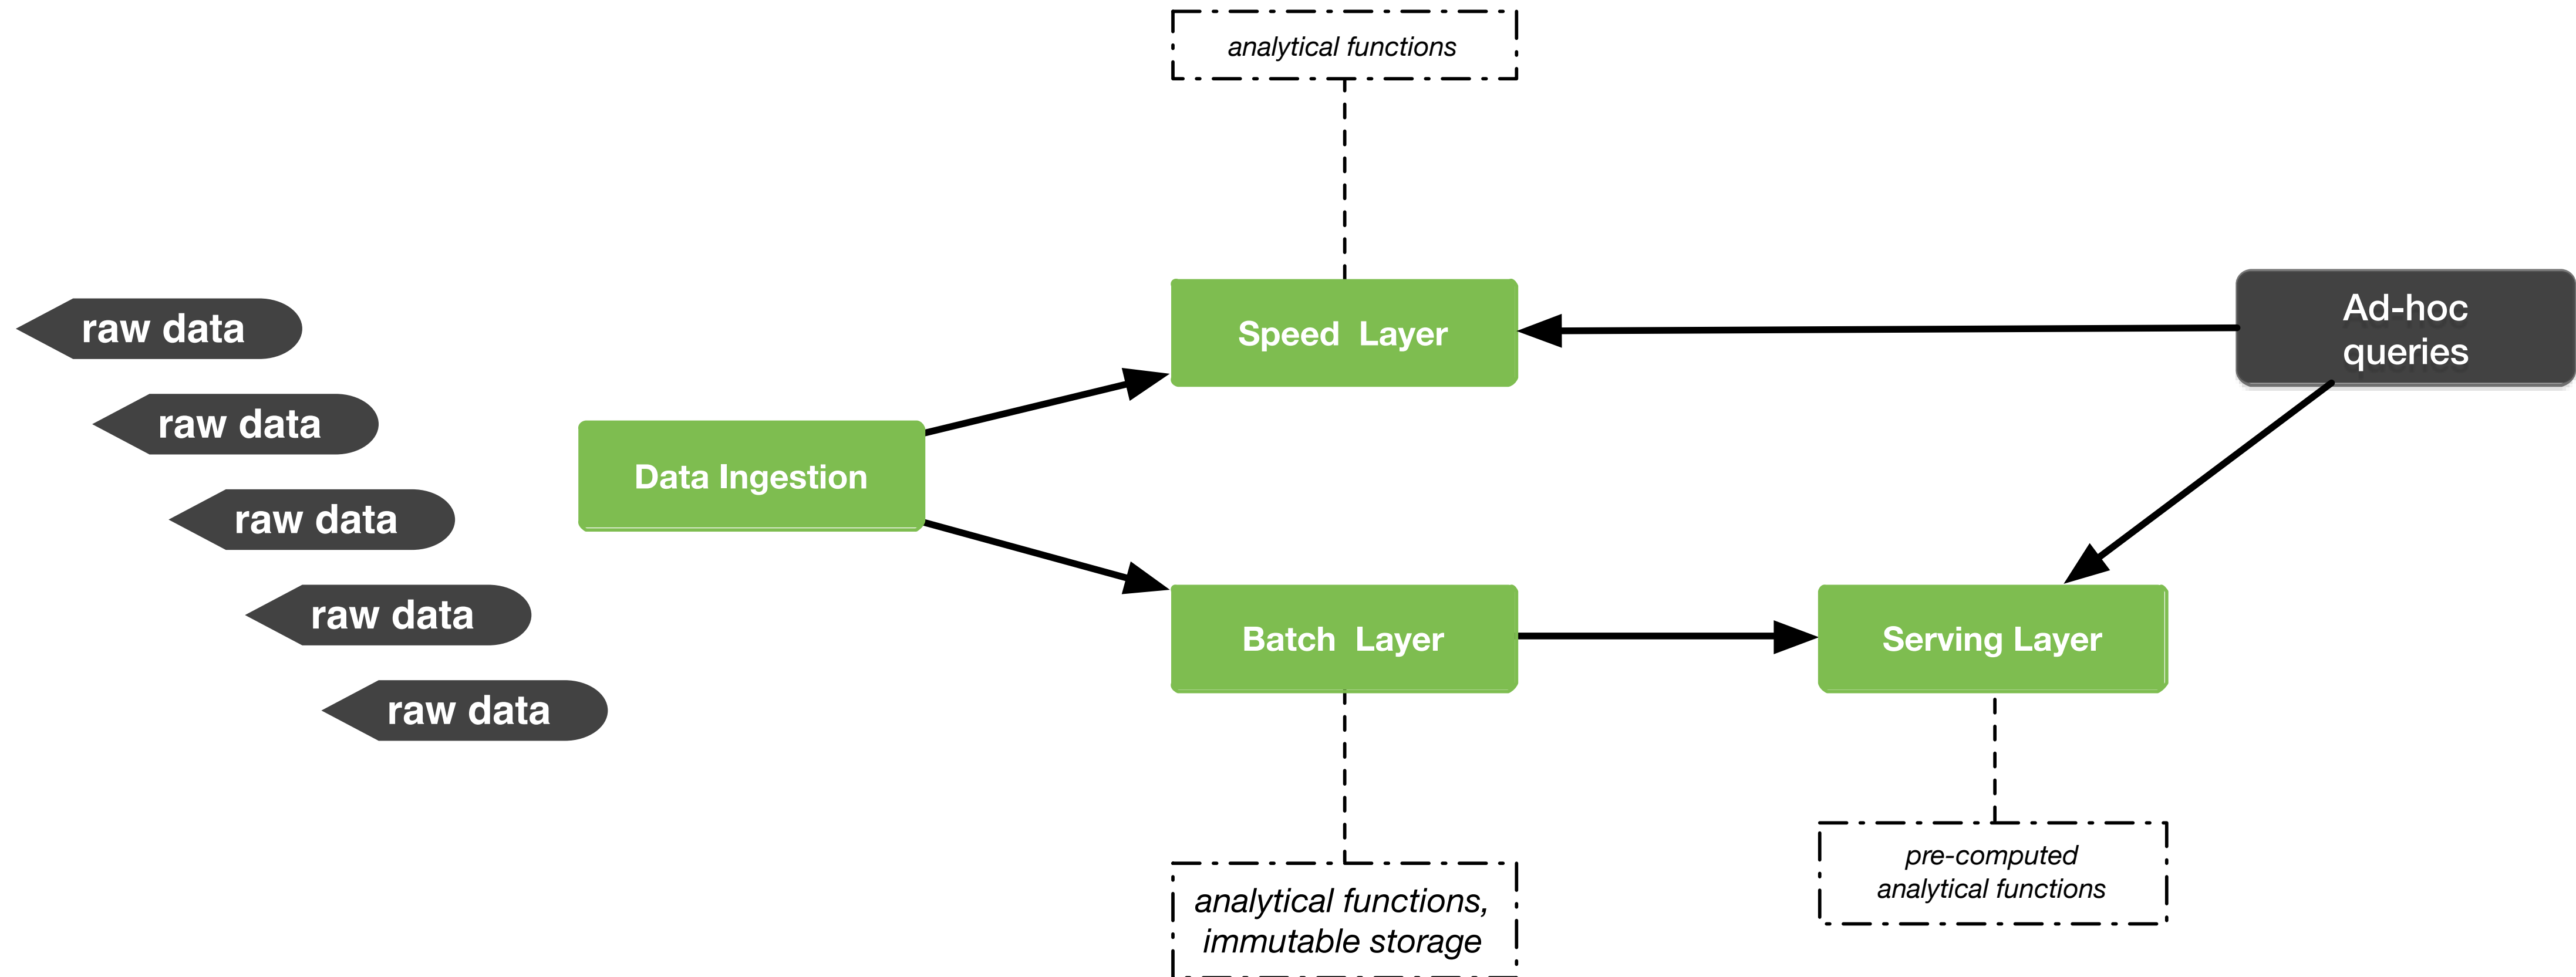
```
cut -d" " -f1 < access.log | sort | uniq -c | sort -rn | less
```

# Micro-batching

- Small batches of events processed one at a time
- Exactly one processing
- Implementations:
  - Spark Streaming



# Lambda Architecture





# Principles of Lambda Architecture

Store data in it's rawest form

Immutability and perpetuity

Re-computation

*Query = function(all data)*

# Implementing the Lambda Architecture

## The ACM DEBS 2014 Grand Challenge<sup>1</sup>

*To demonstrate the applicability of event-based systems to provide scalable, real-time analytics over high volume sensor data*

<sup>1</sup> <http://www.cse.iitb.ac.in/debs2014/>

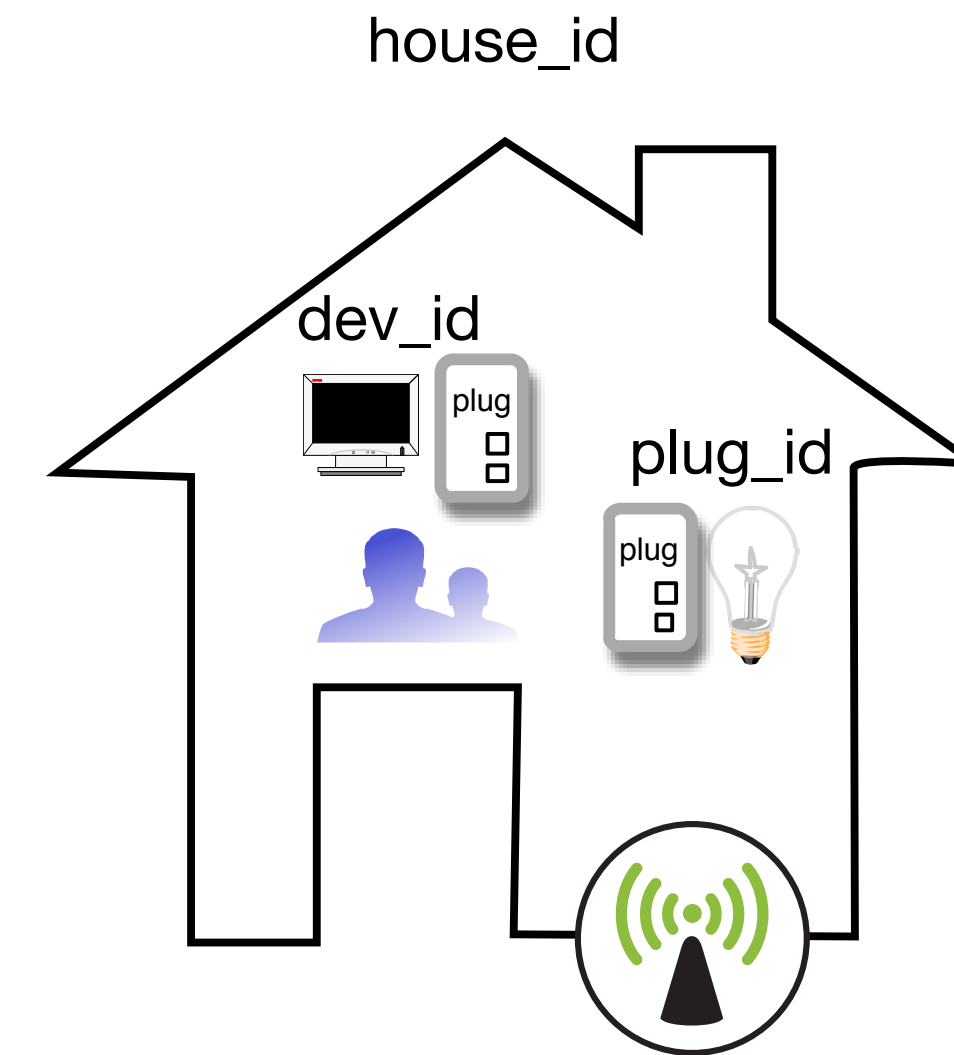


# ACM DEBS 2014 Challenge Application



## Scenario

Analysis of energy  
consumption measurement



## Short-term load forecasting

makes load forecasts based on current load and what was learned over historical data

## Load statistics for real-time demand management

finds outliers based on the energy consumption

# Data model



## Data source

Field	Comments
ID	UNIQUE IDENTIFIER
TIMESTAMP	TIMESTAMP OF MEASUREMENT
VALUE	MEASUREMENT
PROPERTY	TYPE: 0 - WORK, 1 - LOAD
PLUG_ID	UNIQUE IDENTIFIER OF A PLUG
HOUSEHOLD_ID	UNIQUE IDENTIFIER OF A HOUSEHOLD
HOUSE_ID	UNIQUE IDENTIFIER OF A HOUSE

## PL - House

Field	Comments
TS	TIMESTAMP OF STARTING TIME PREDICTION
HOUSE_ID	HOUSE ID
PREDICTED_LOAD	PREDICTED LOAD

## PL - Plug

Field	Comments
TS	TIMESTAMP OF STARTING TIME PREDICTION
HOUSE_ID	HOUSE ID
HOUSEHOLD_ID	
PLUG_ID	
PREDICTED_LOAD	PREDICTED LOAD

## Outliers

Field	Comments
TS_START	TIMESTAMP OF START TIME WINDOW
TS_STOP	TIMESTAMP OF STOP TIME WINDOW
HOUSE_ID	HOUSE ID
PERCENTAGE	% PLUGS LOAD HIGHER THAN NORMAL

# Analytical models

## Load prediction per house, per plug

It is based on the average load of the current window and the median of the average loads of windows covering the same time of all past days

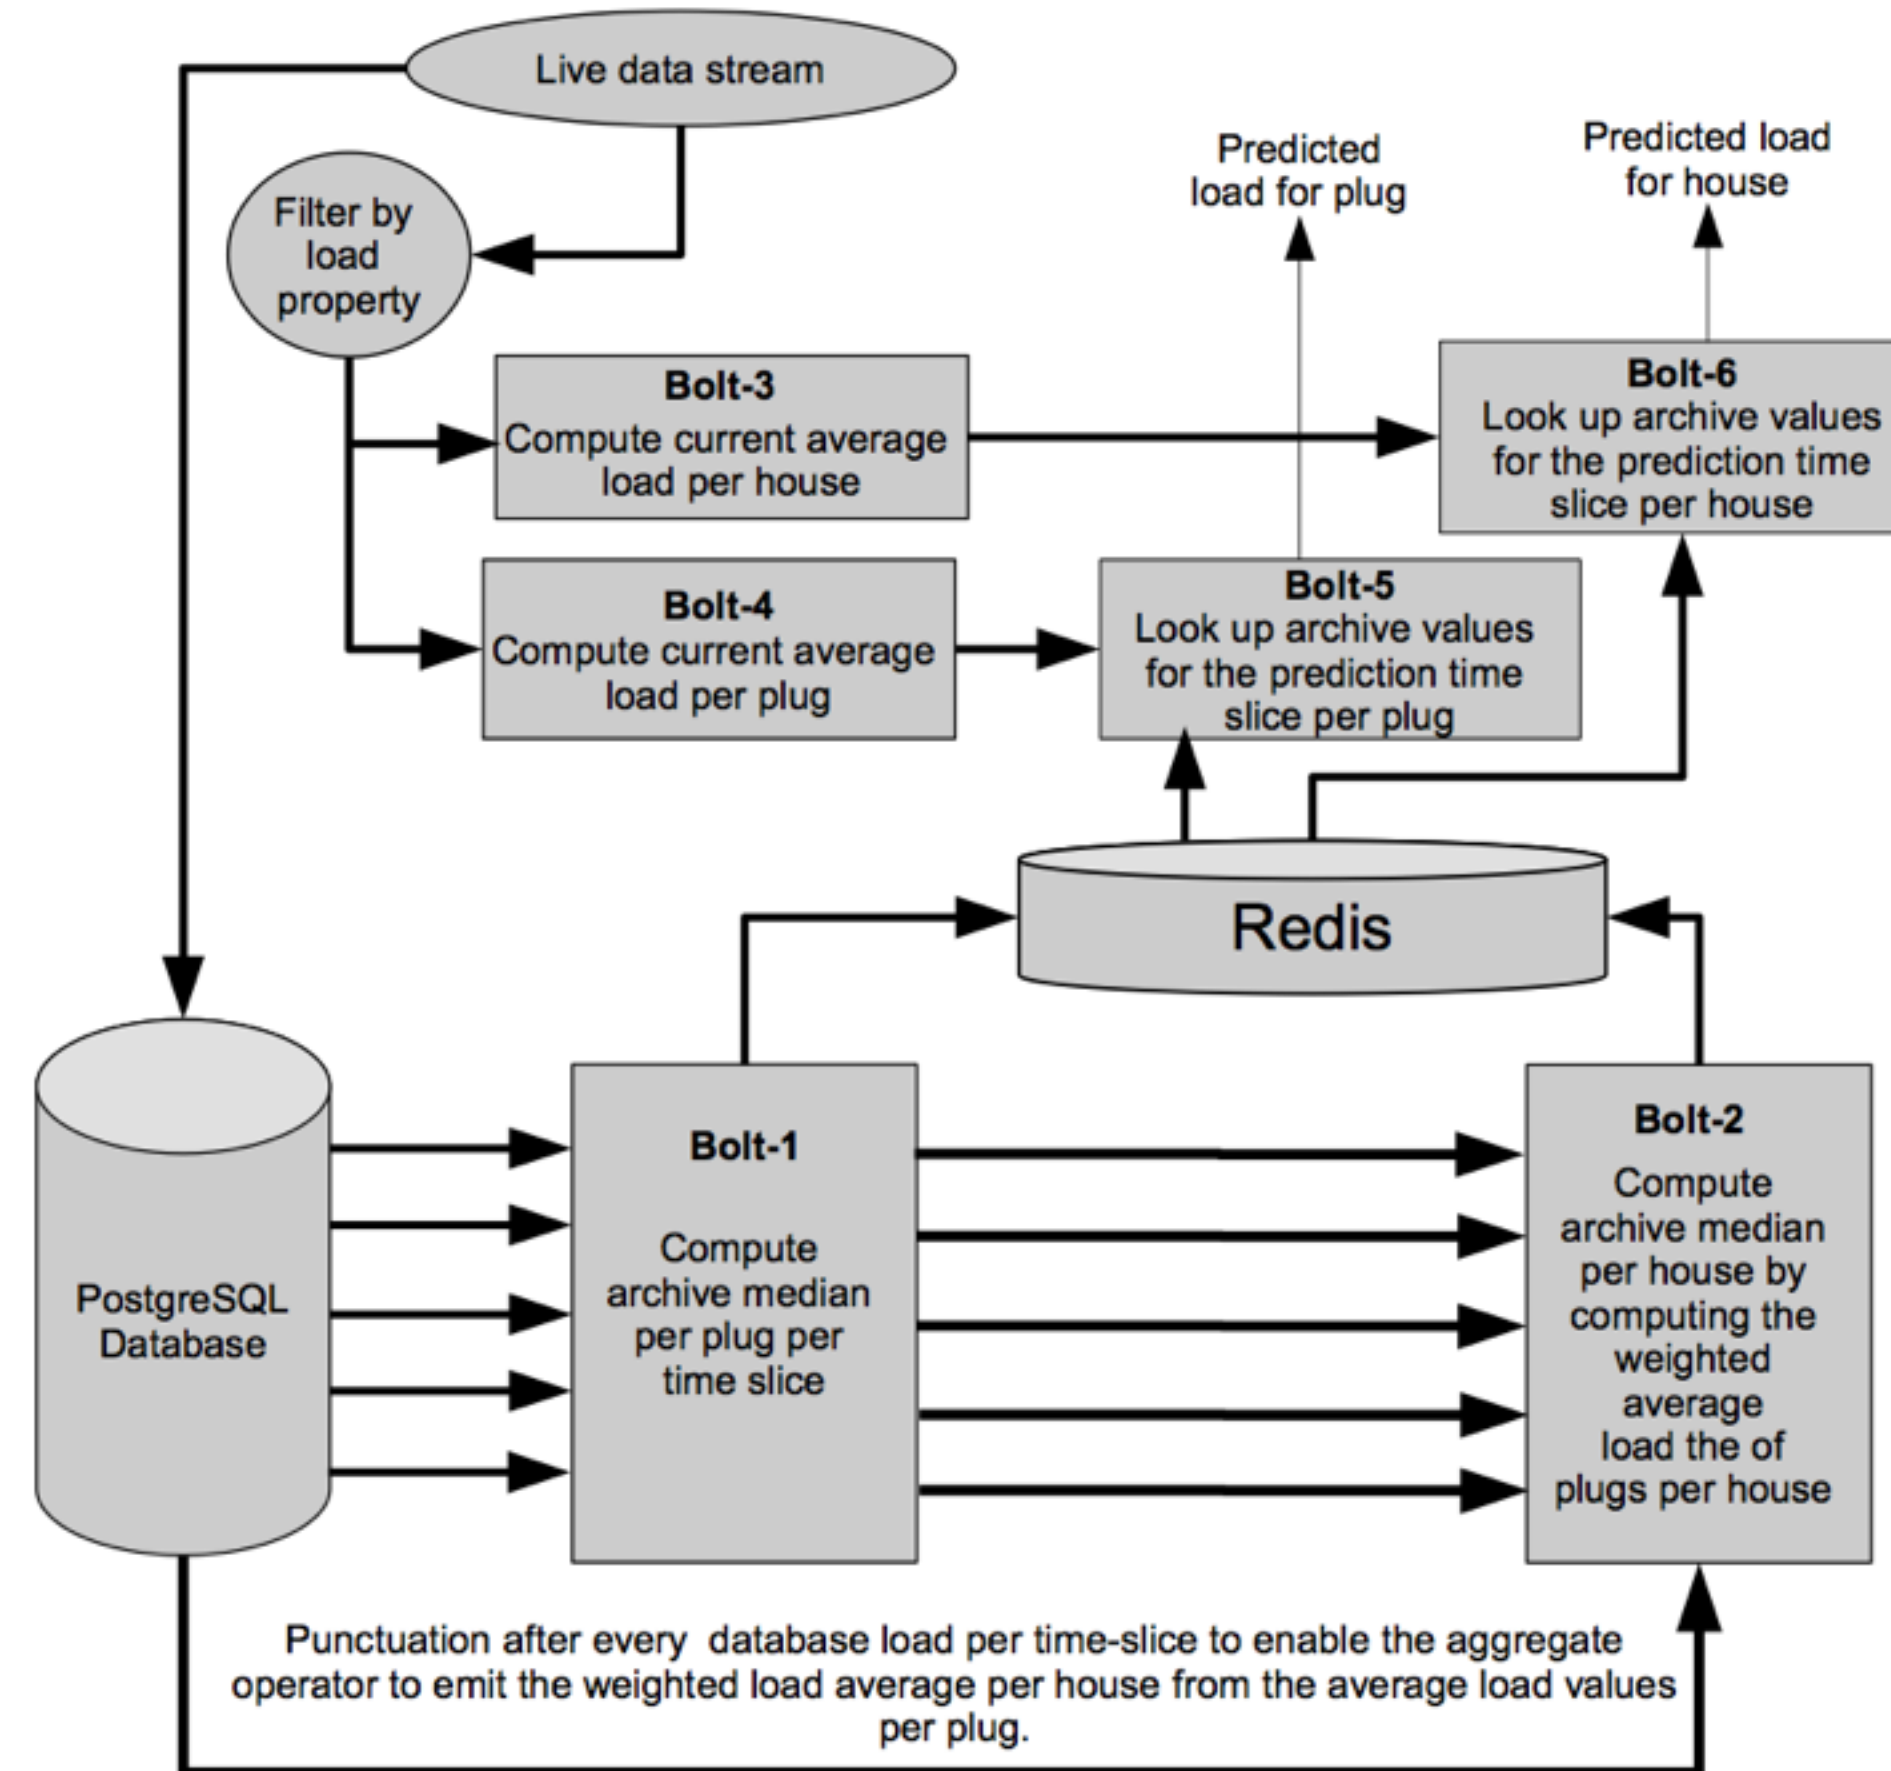
$$L(s_{i+2}) = (avgLoad(s_i) + median(avgLoad(s_j))) / 2$$

$$s_j = s_{(i+2-n*k)}$$

## Outliers

*For each house calculate the percentage of plugs which have a median load during the last hour greater than the median load of all plugs (in all households of all houses) during the last hour*

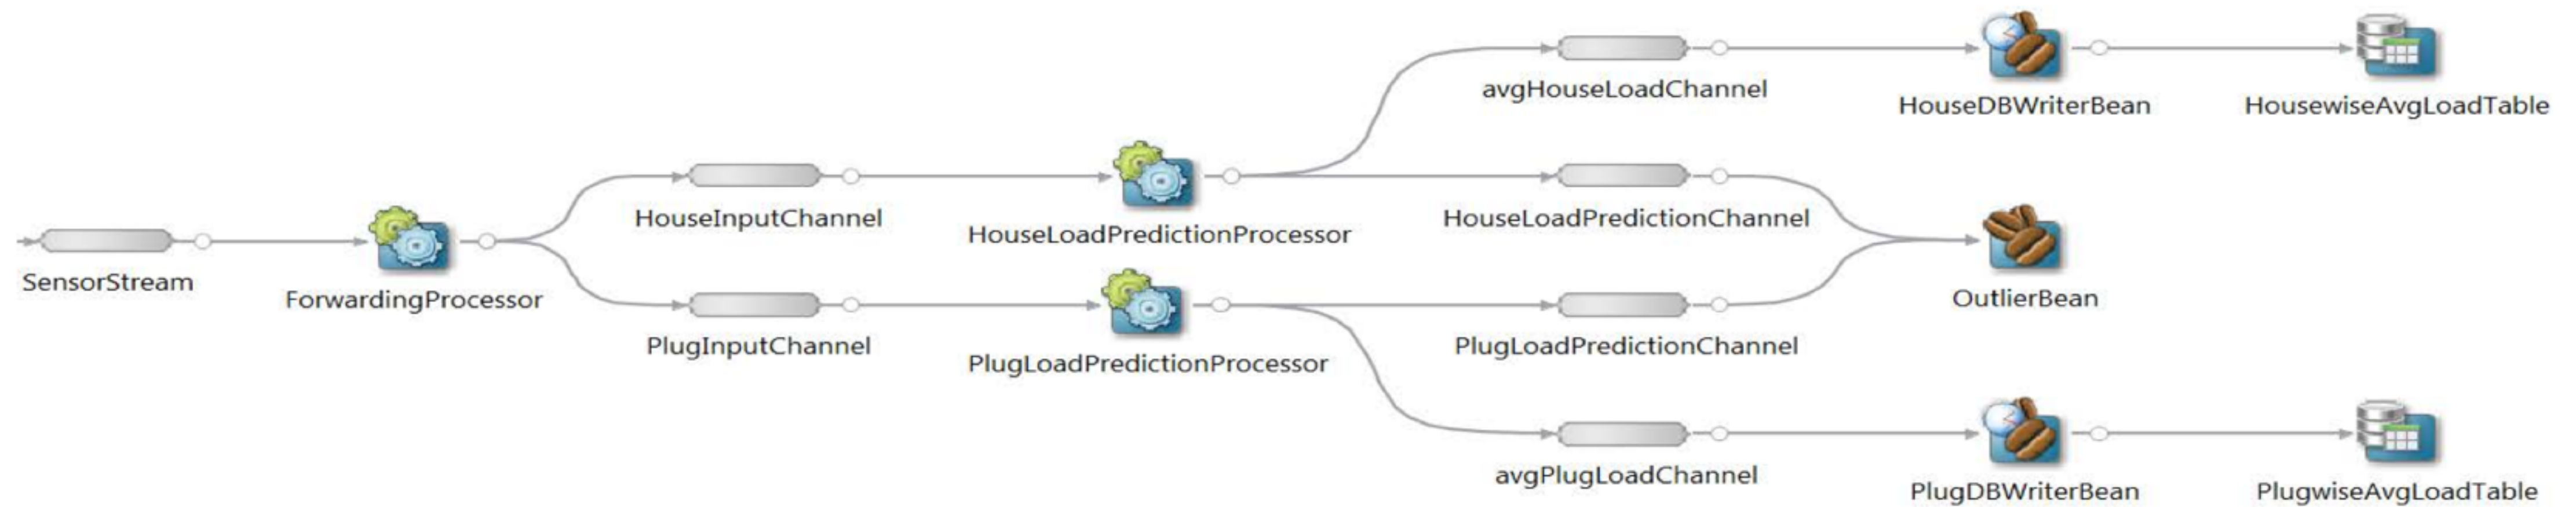
# How others did



*Storm-based  
implementation*



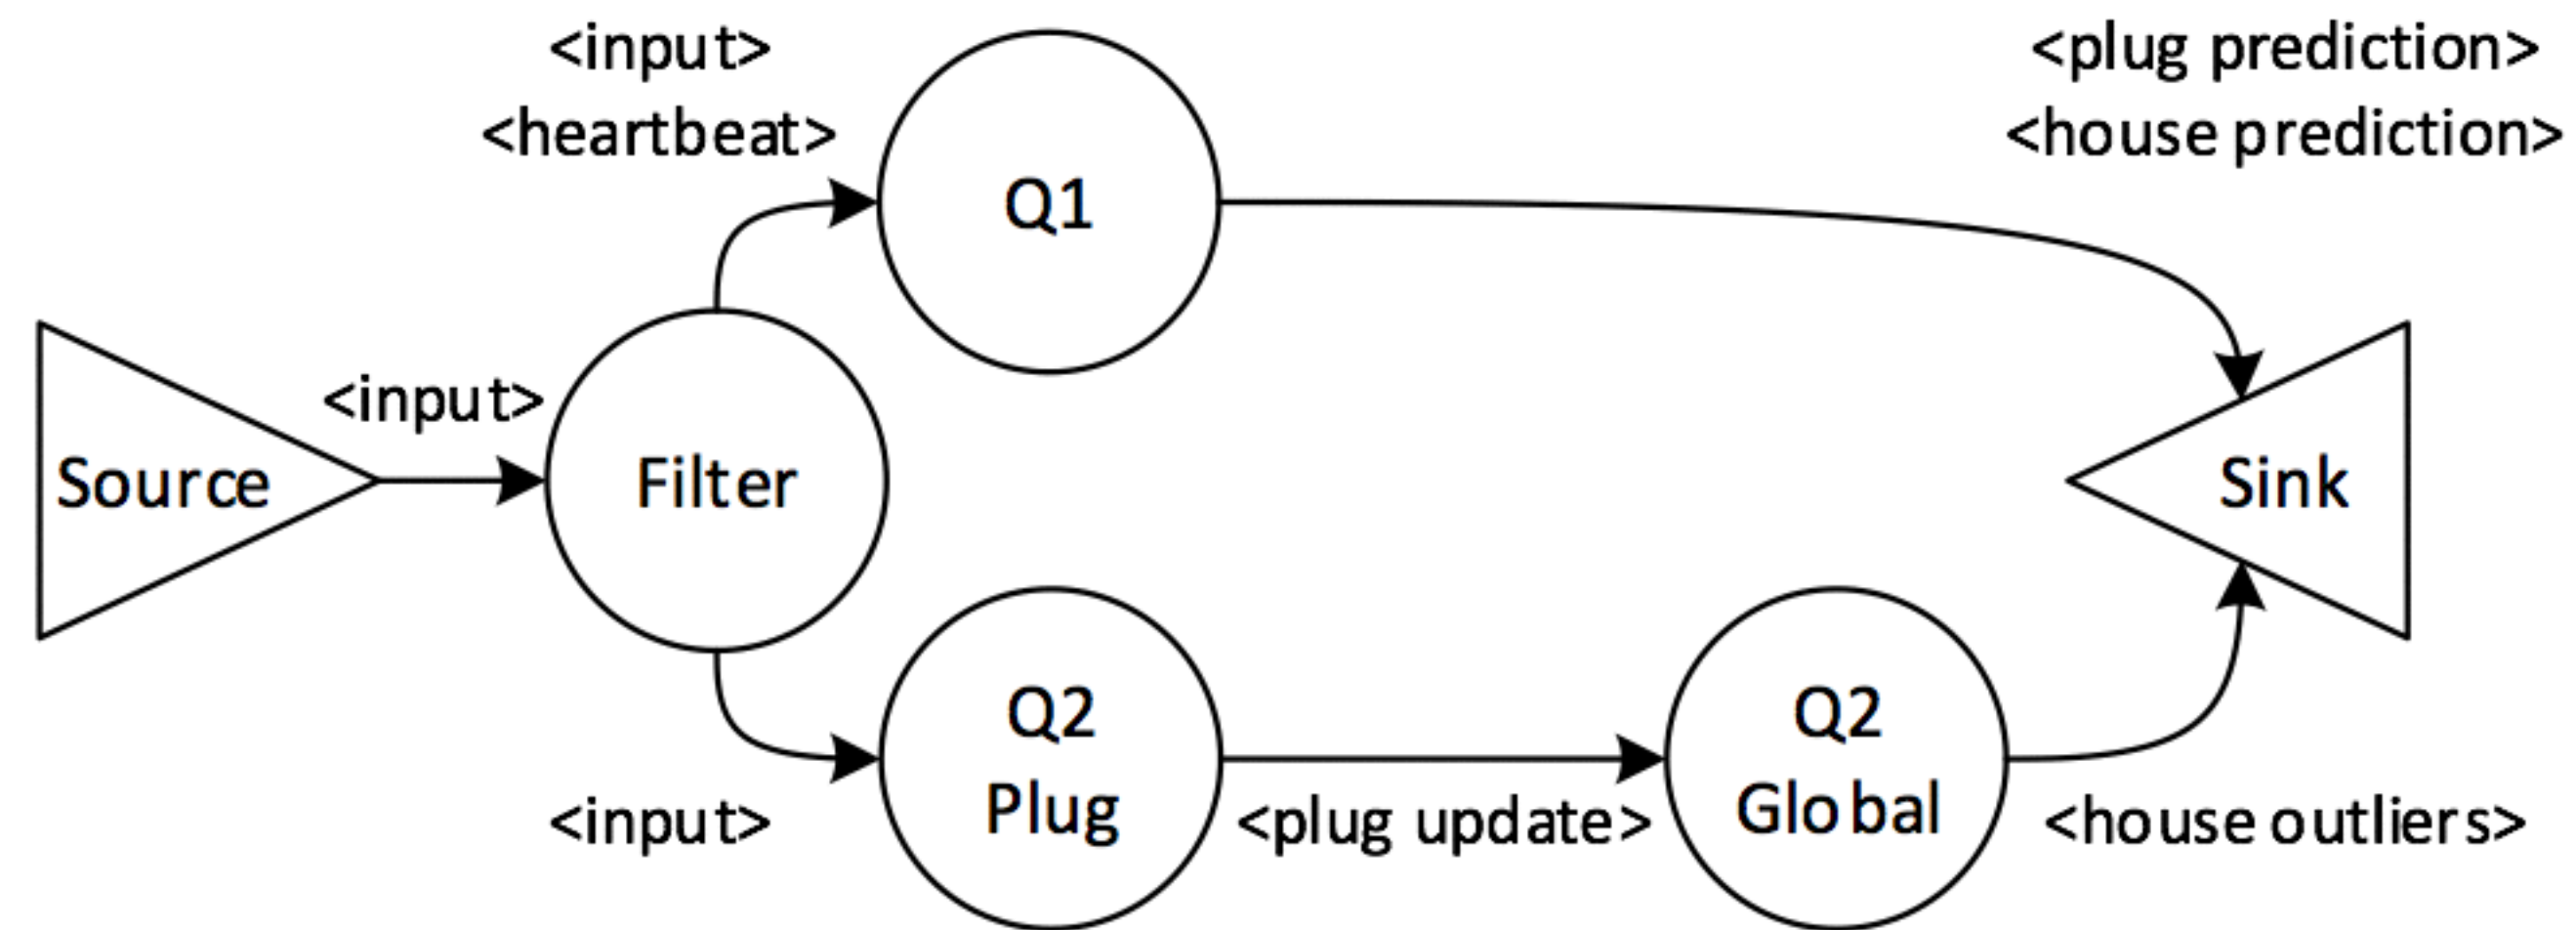
# How others did



Oracle DB + Oracle Event Processing (OEP)



# How others did

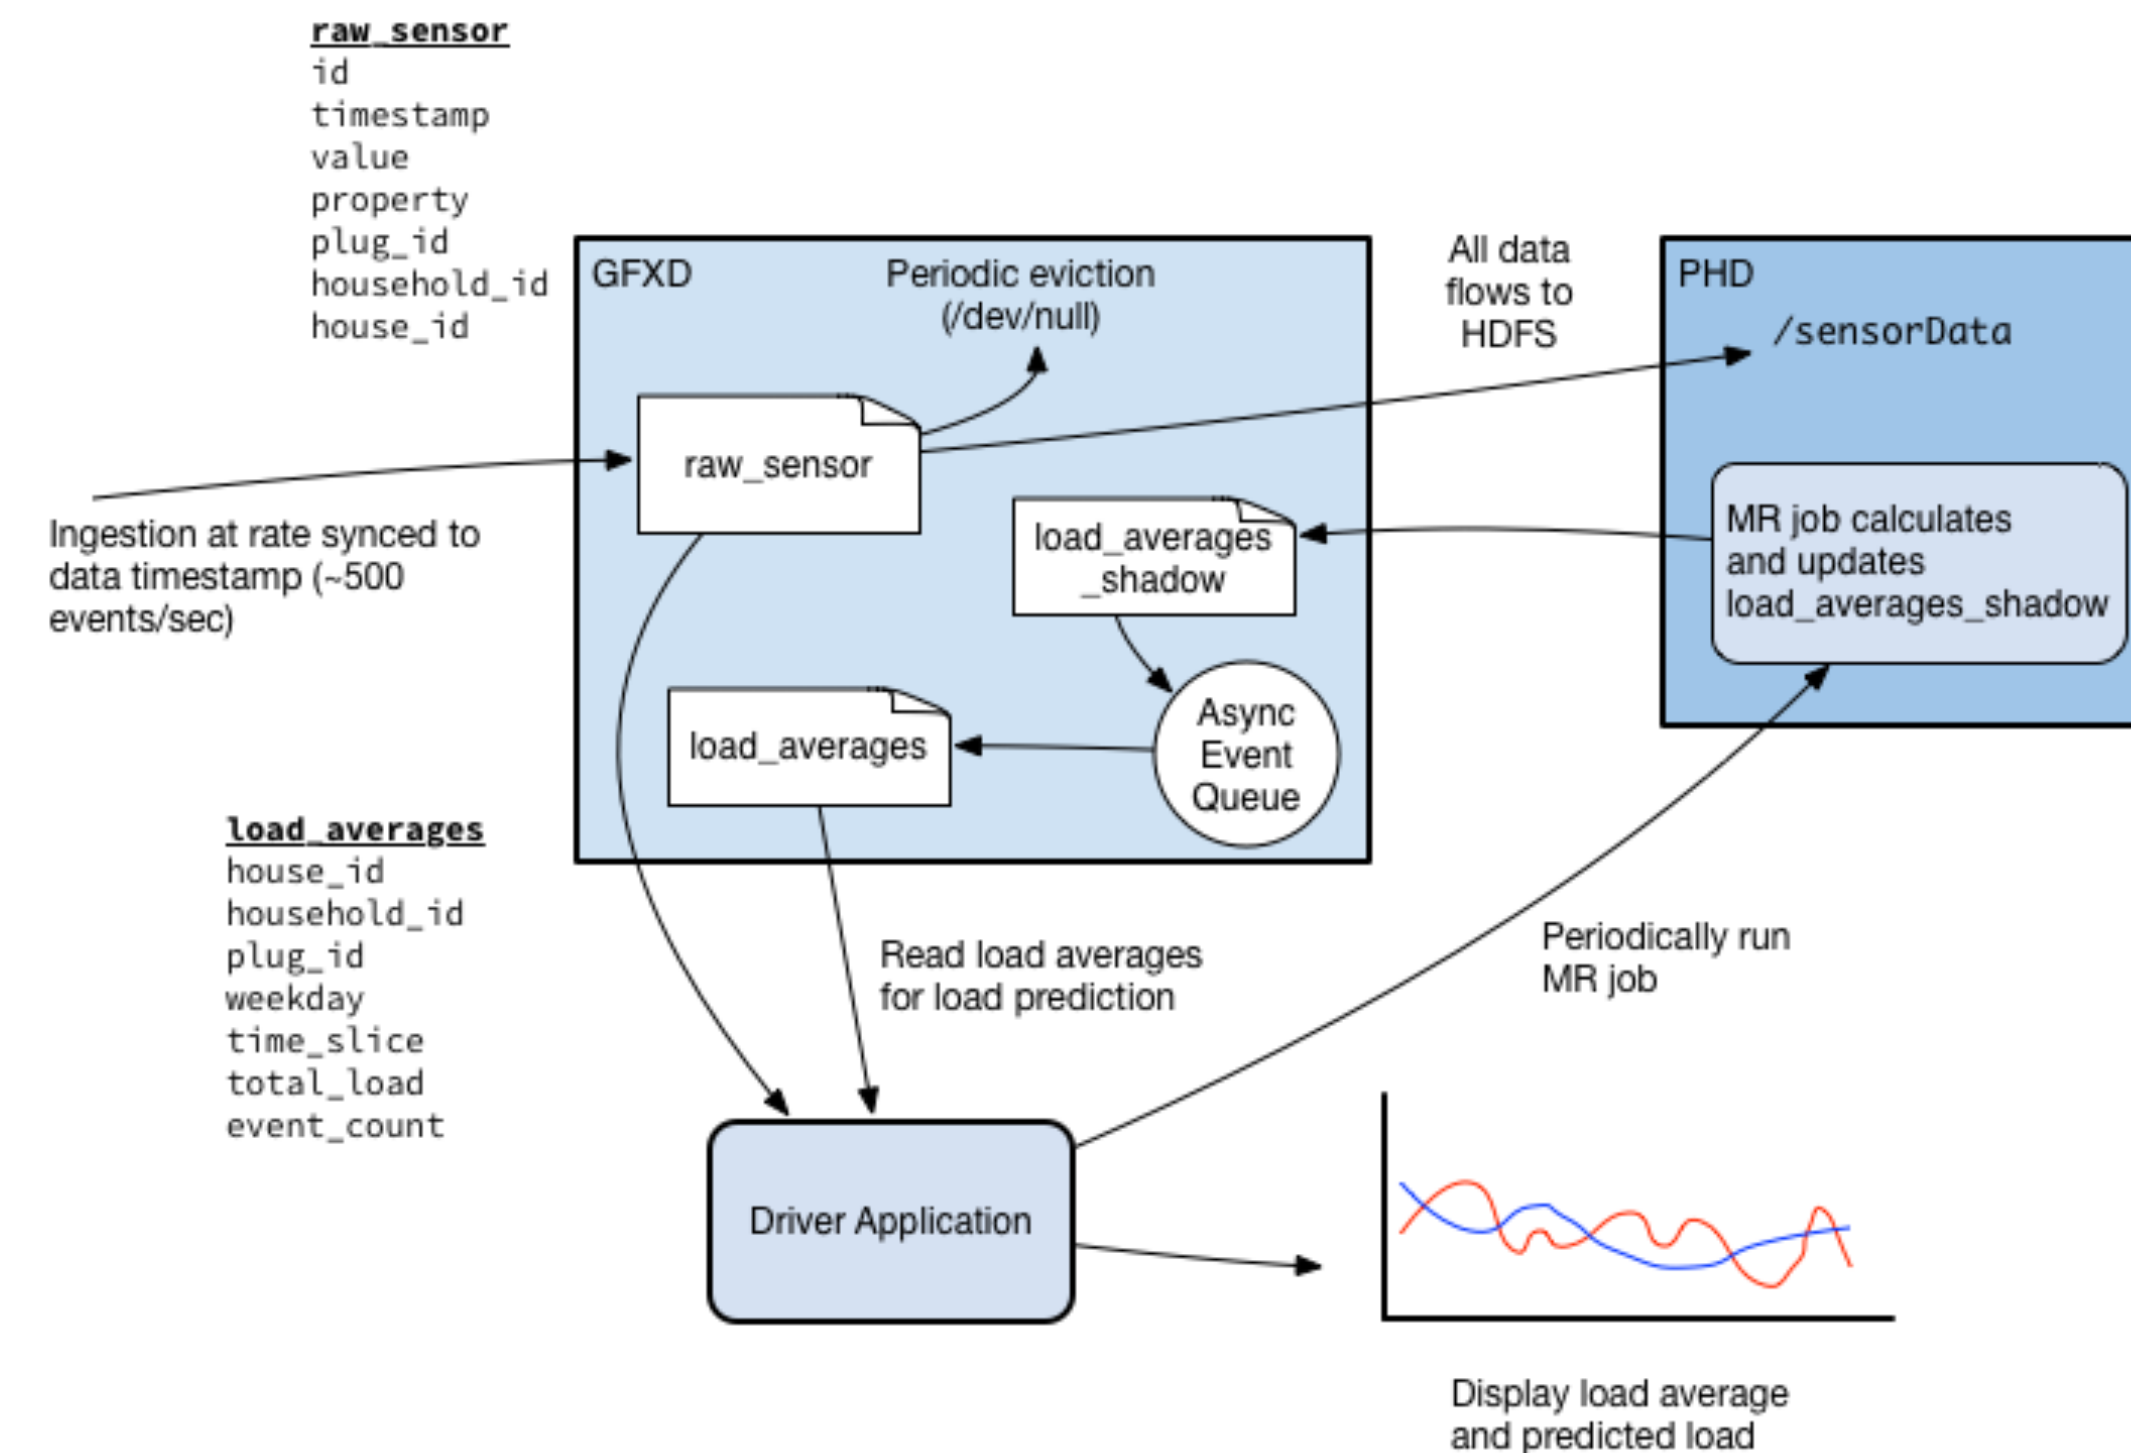


SEEP: Real-Time Stateful Big Data Processing

<http://lsds.doc.ic.ac.uk/projects/seep/debs14-gc>

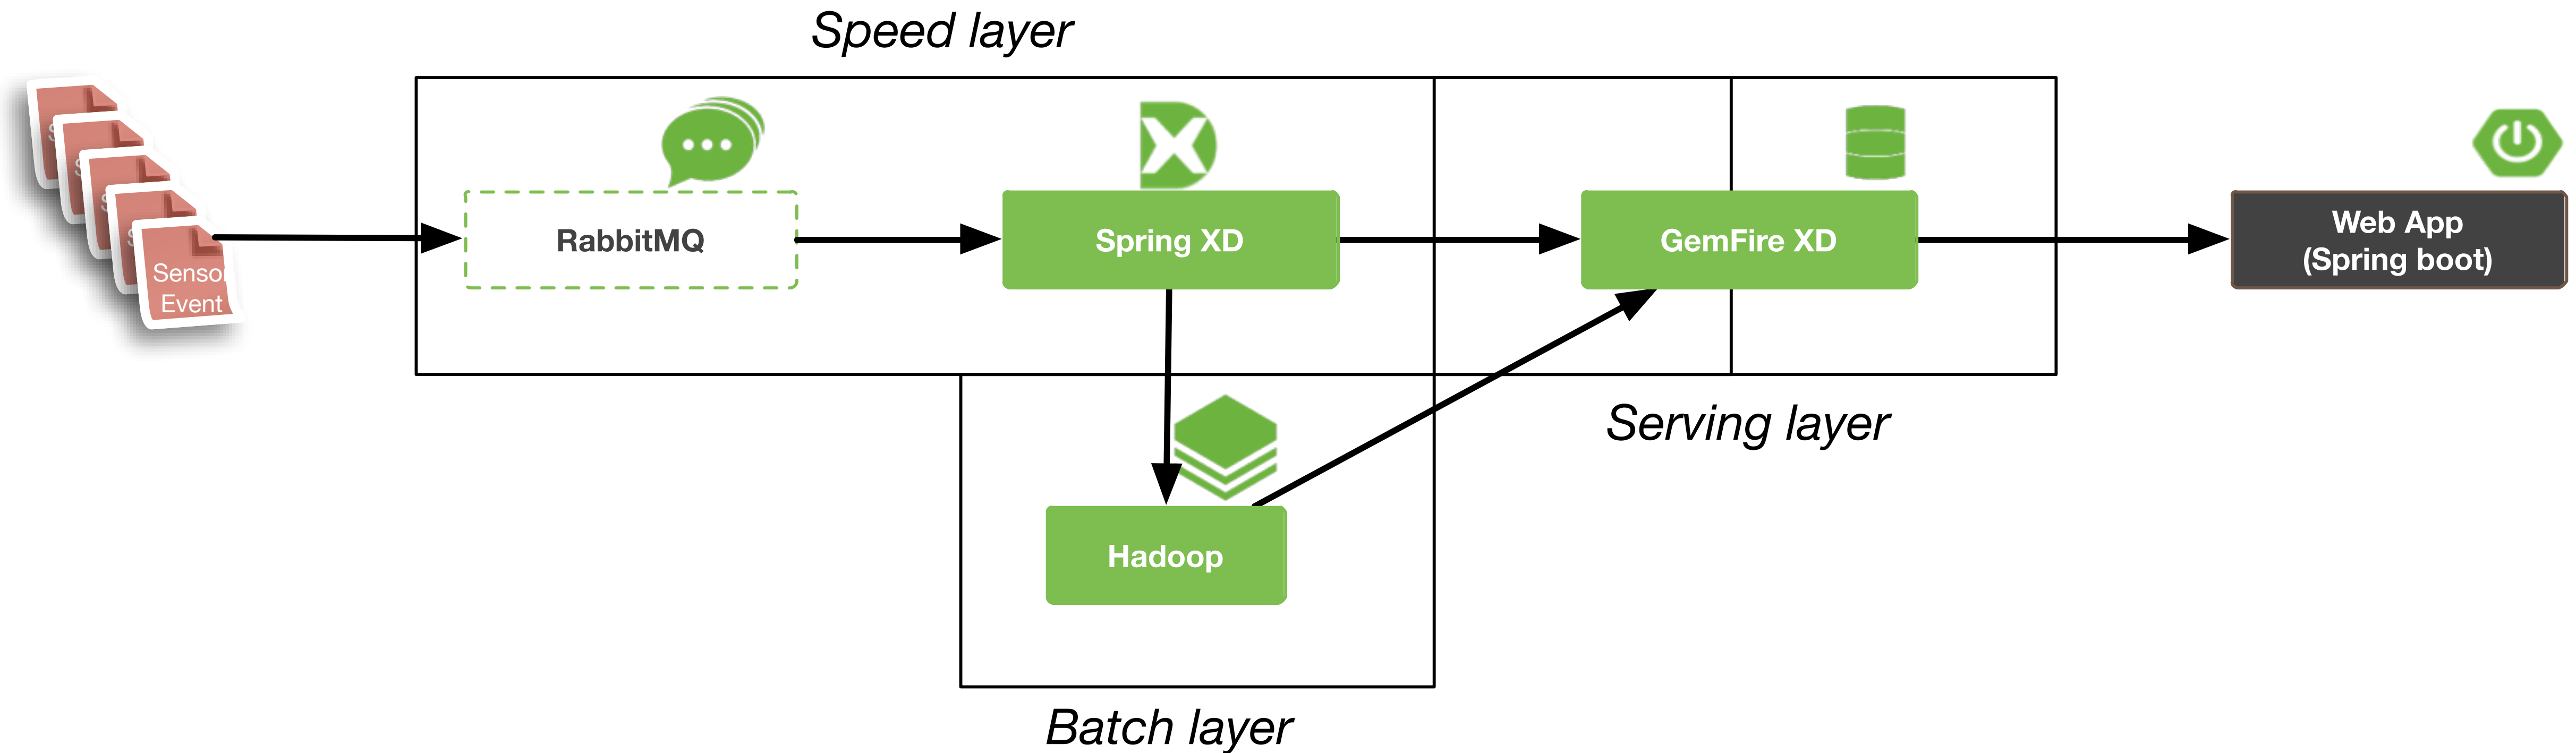
# Implementation details

Fork from <https://github.com/pivotalsoftware/gfxd-demo>

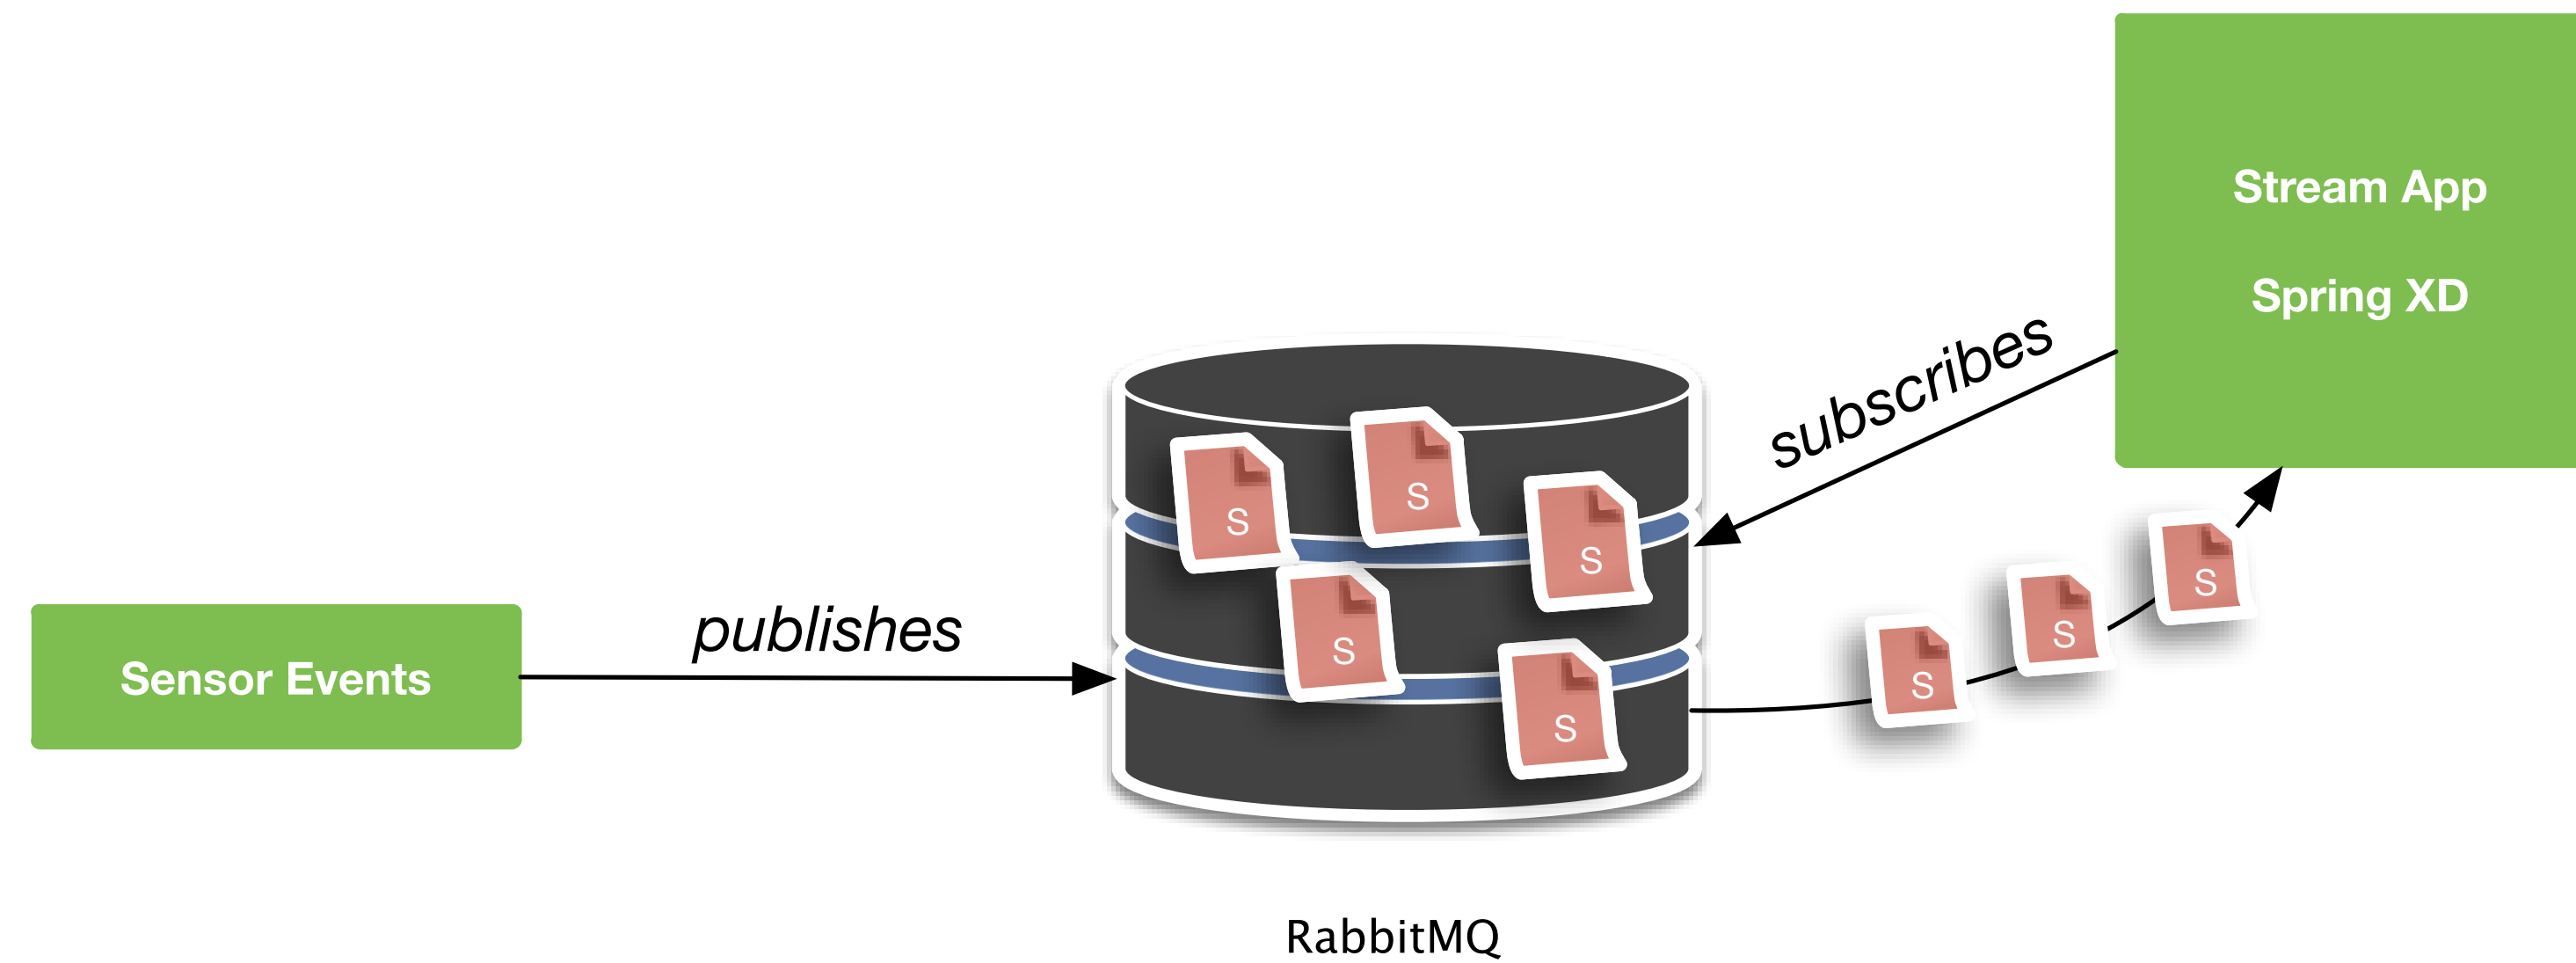


*Thanks to Jens Deppe and William Markito*

# Our approach - ACM DEBS 2014 system architecture



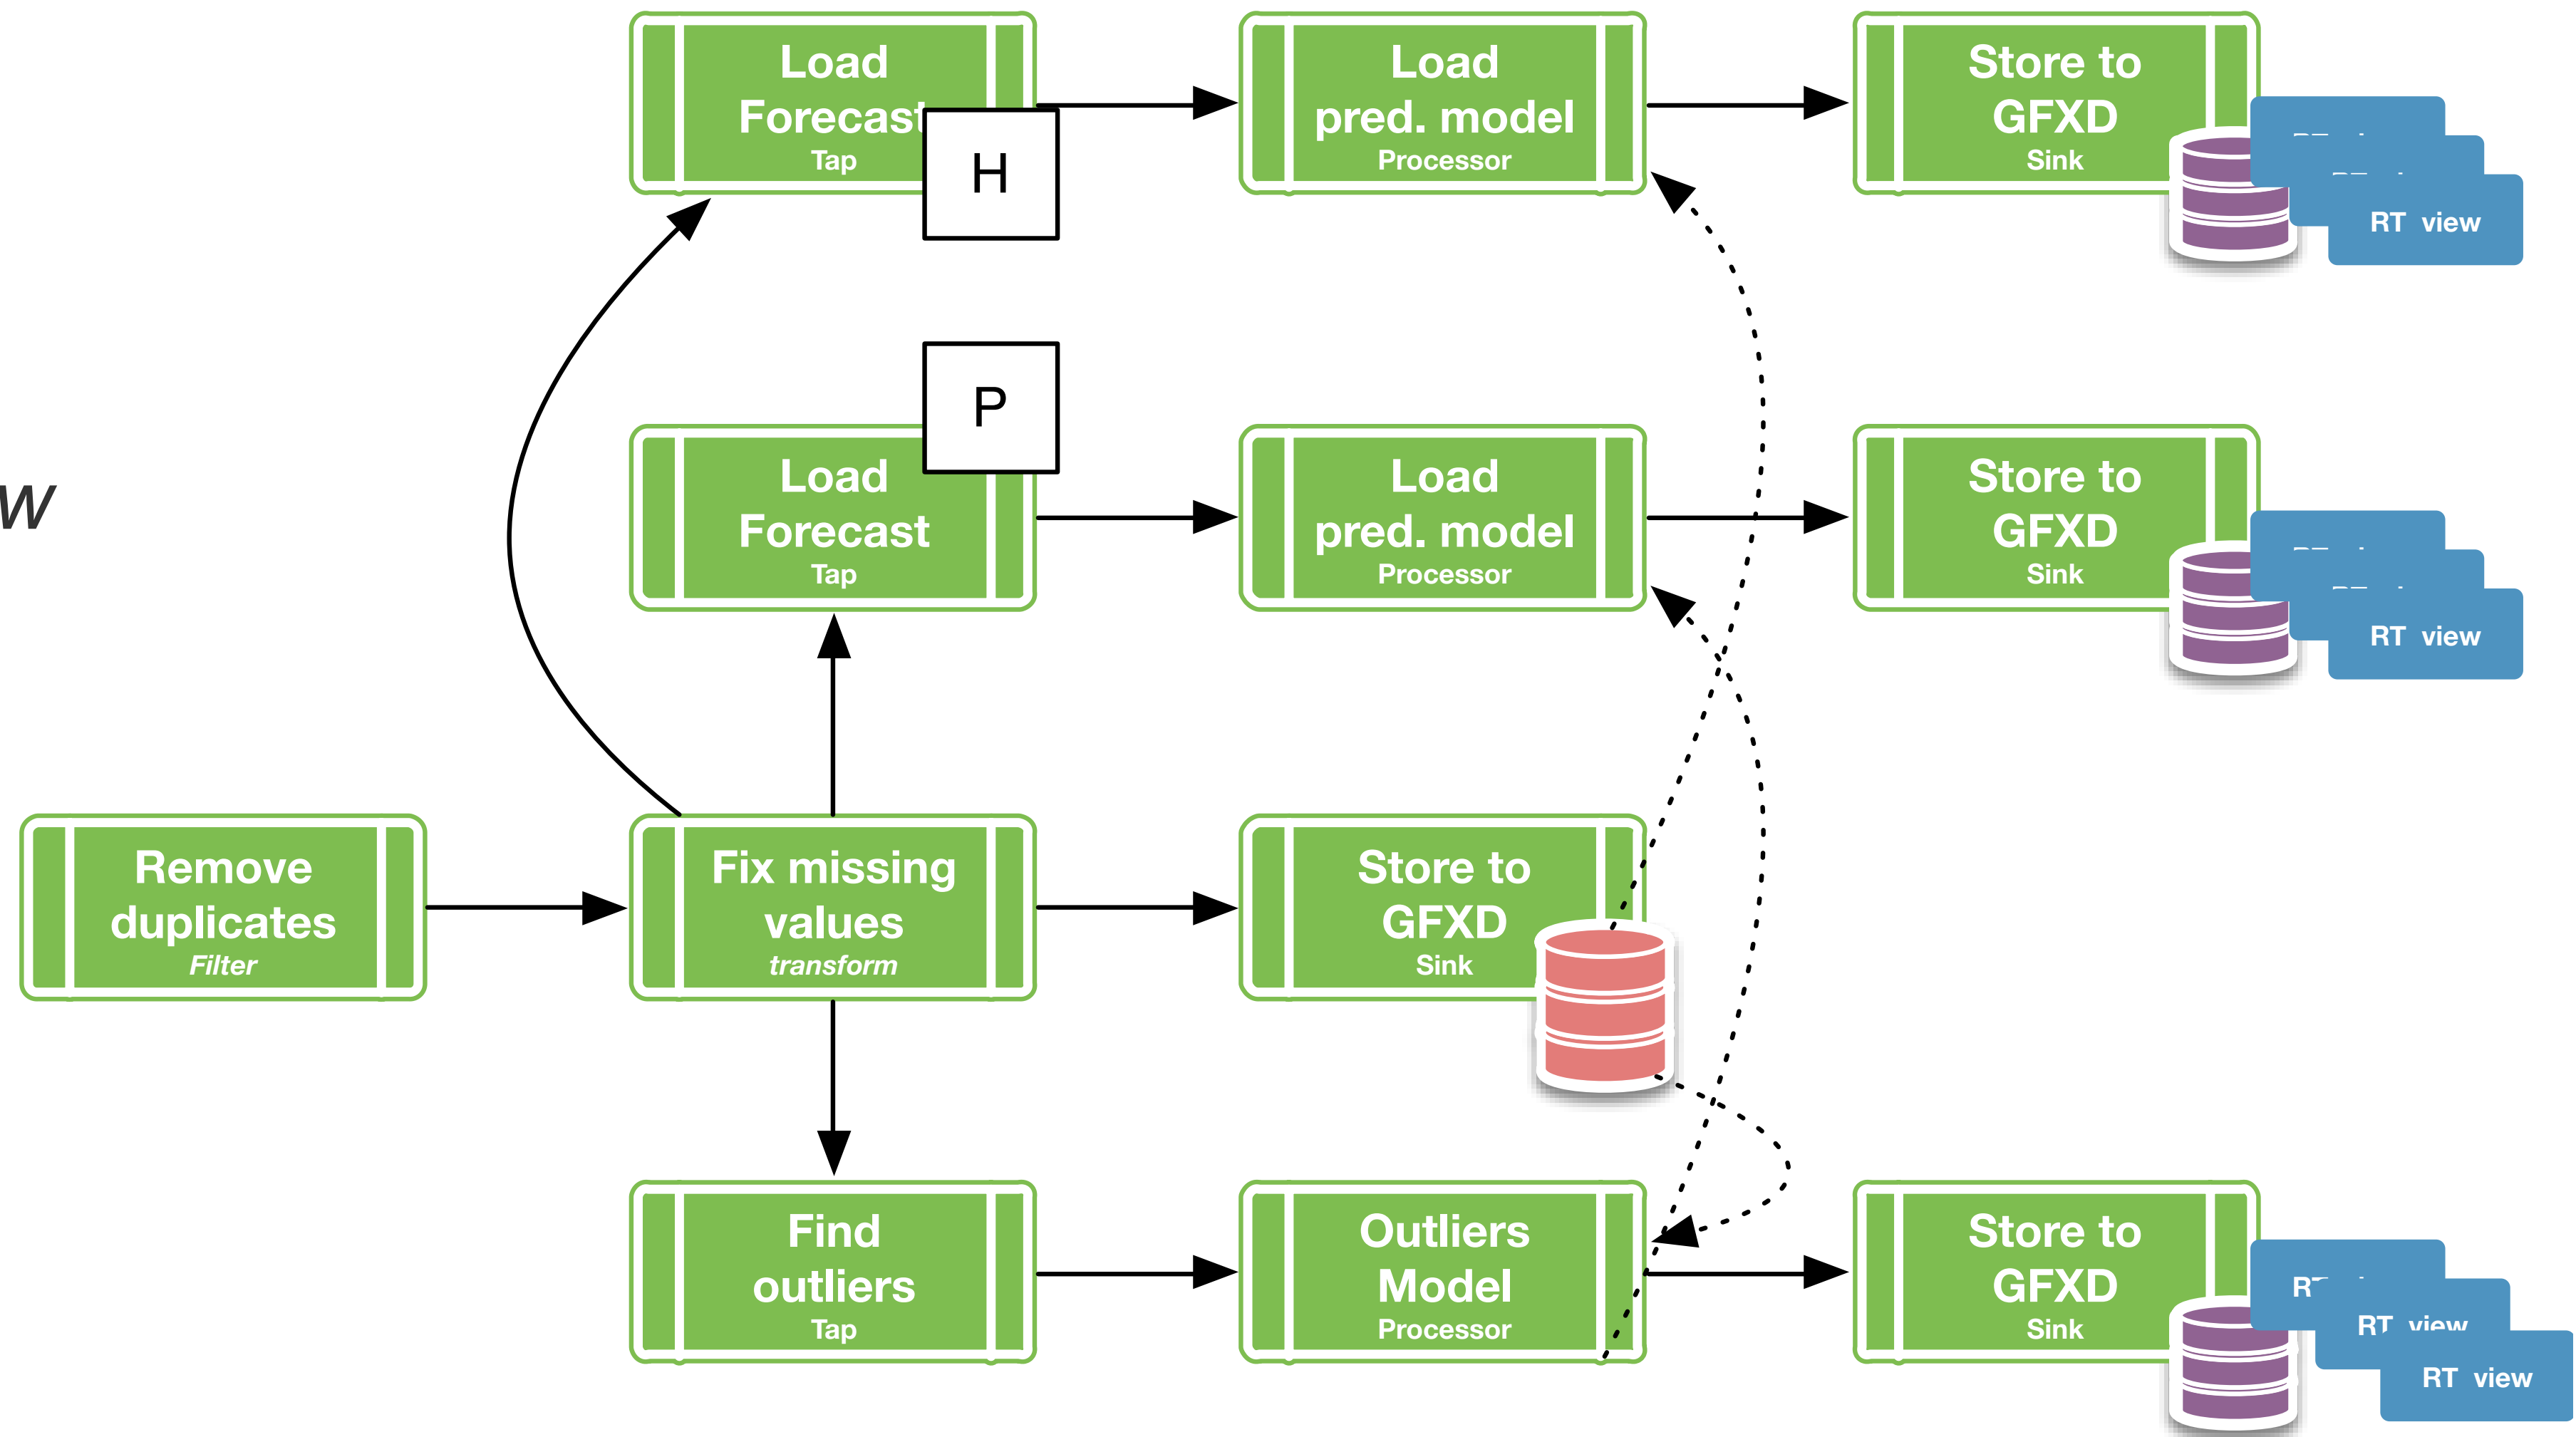
# Data ingestion



# Speed layer



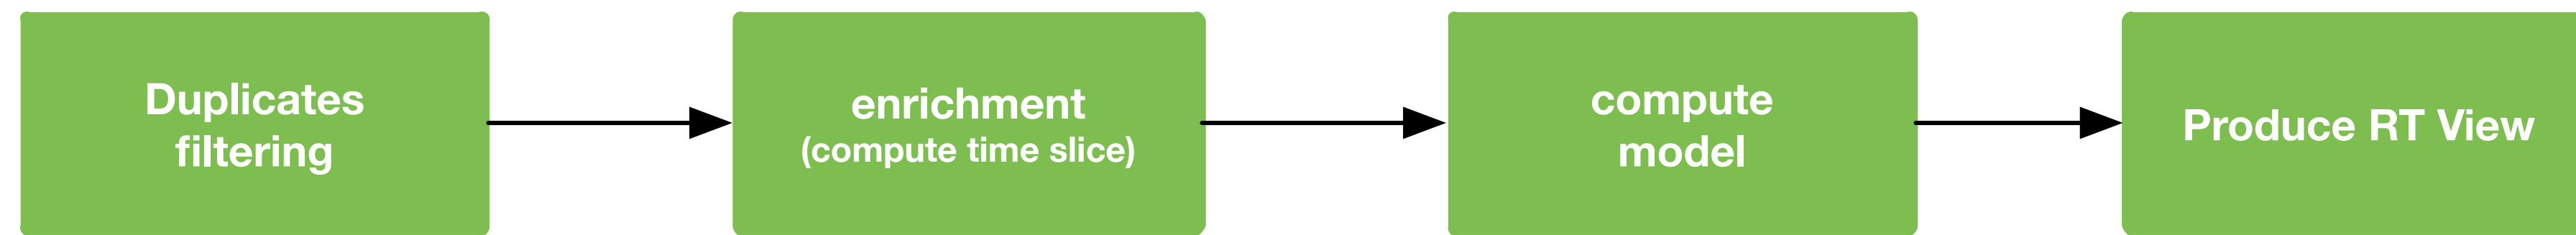
*Application flow*



*not a flow* .....



# Stream actions



# Spring XD streams



*pumpin* - sends all data to a master queue

*sensoreventenricher* - Consumes the data from the queue, filter and transform the data before store on HDFS

*findoutliers* - Taps from master\_ds stream to compute outlier model

*loadpred{h,p}* - Taps from master\_ds stream to compute load prediction for house and plug (2 streams)

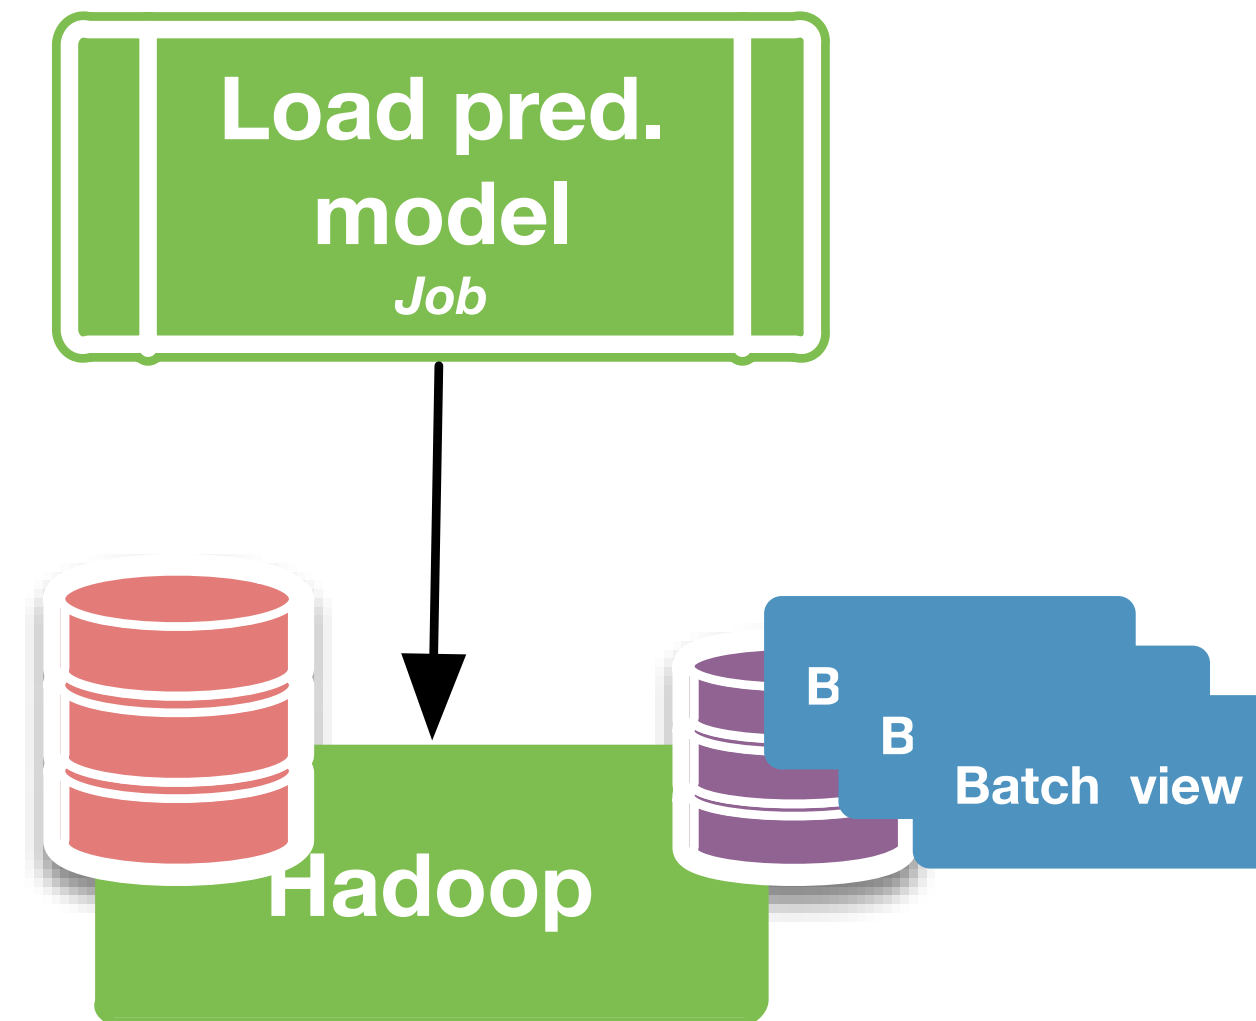
5 streams

# Batch Layer



*Hadoop based system*

Batch job that starts from Spring XD  
Uses cron to run every X hour/min/day?



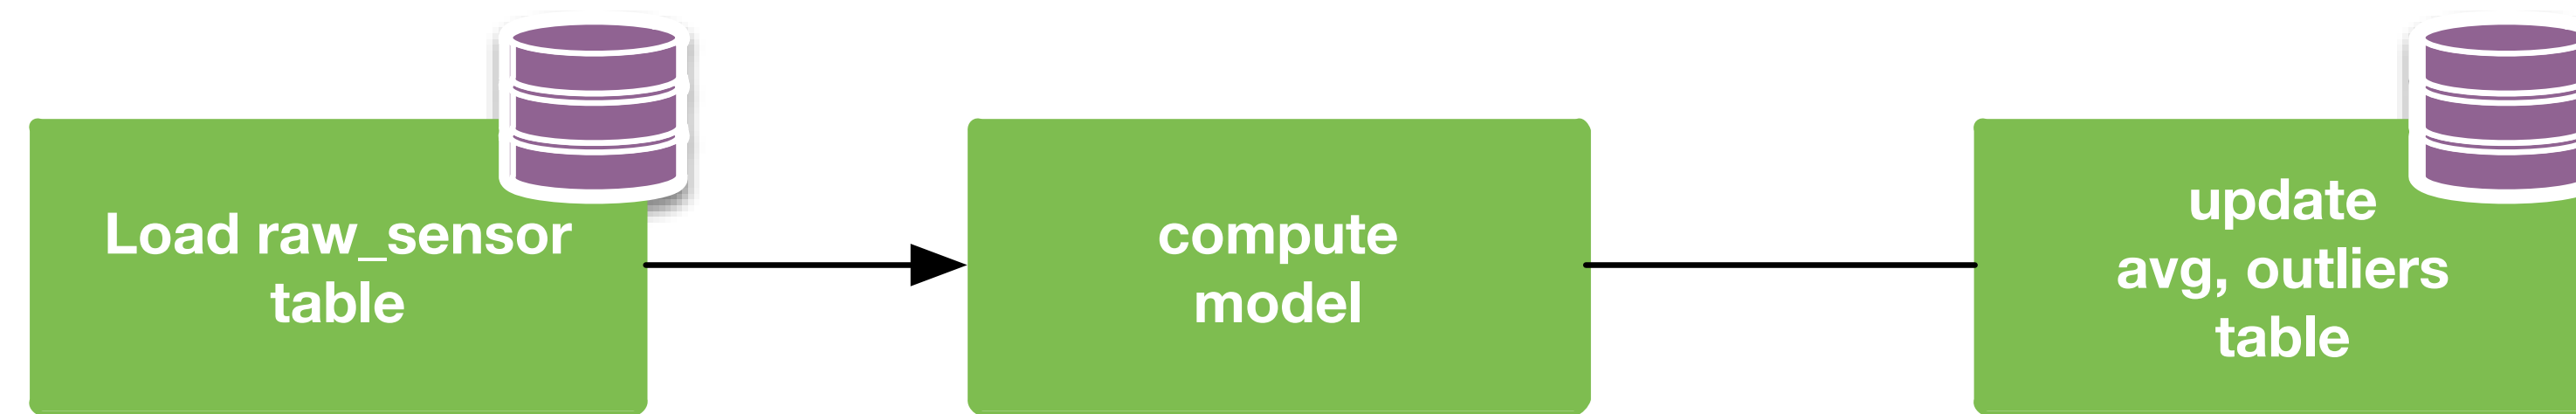
*Store an immutable and constantly growing master dataset (of all datasets)*

*Compute arbitrary functions (models) on the existing datasets.*

*Essentially, runs the MR models.*

# Spring XD jobs

- A unique job to run the MR model to compute the historical aspect of the models.
- MR job runs every “day/hour/min”
- Job is completely independent from the other streams



1 job

# Serving layer (RT and Batch views)



Gemfire XD

Field	Comments
TS	<i>TIMESTAMP OF STARTING TIME PREDICTION</i>
HOUSE_ID	<i>HOUSE ID</i>
PREDICTED_LOAD	<i>PREDICTED LOAD</i>



Field	Comments
TS_START	<i>TIMESTAMP OF START TIME WINDOW</i>
TS_STOP	<i>TIMESTAMP OF STOP TIME WINDOW</i>
HOUSE_ID	<i>HOUSE ID</i>
PERCENTAGE	<i>% PLUGS LOAD HIGHER THAN NORMAL</i>

Field	Comments
TS	<i>TIMESTAMP OF STARTING TIME PREDICTION</i>
HOUSE_ID	<i>HOUSE ID</i>
HOUSEHOLD_ID	
PLUG_ID	
PREDICTED_LOAD	<i>PREDICTED LOAD</i>

*Get updates from MR and SP models*

*Holds both batch and real-time views*

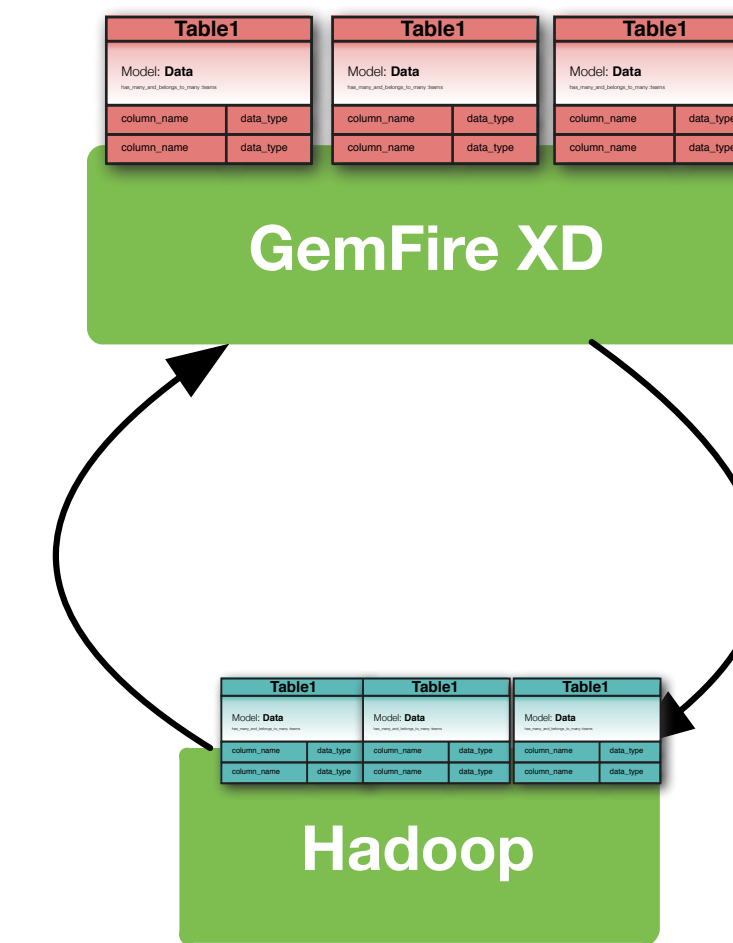


# Web UI



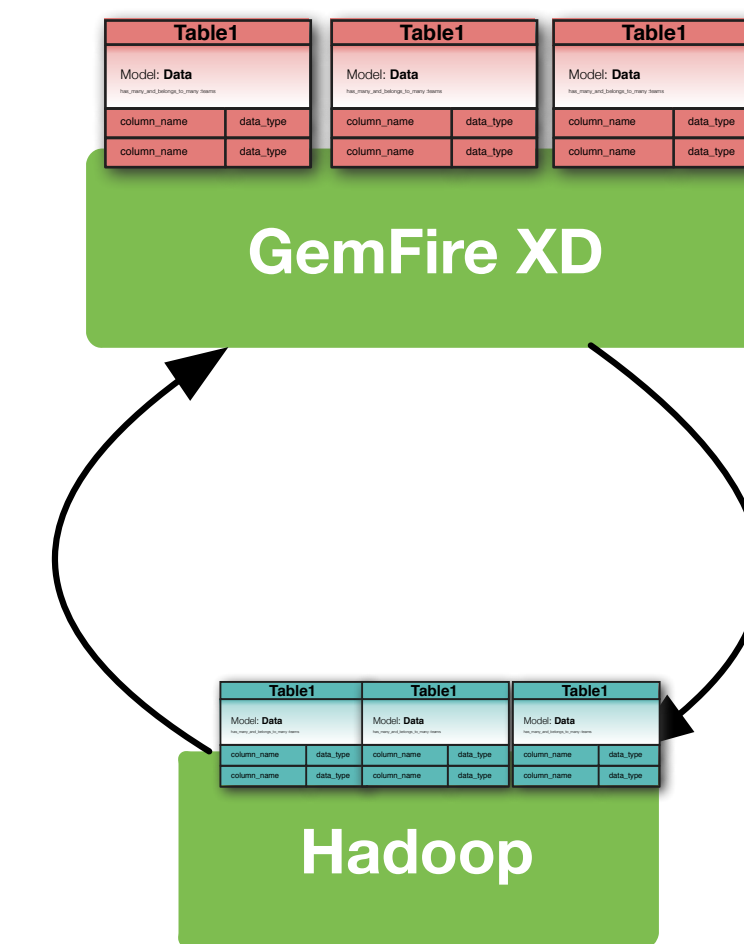
# Why *Gemfire XD*?

- ▶ Distributed database
- ▶ Tightly integrated with Pivotal Hadoop
- ▶ SQL support
- ▶ Fault-tolerant
- ▶ In-memory (fast access)



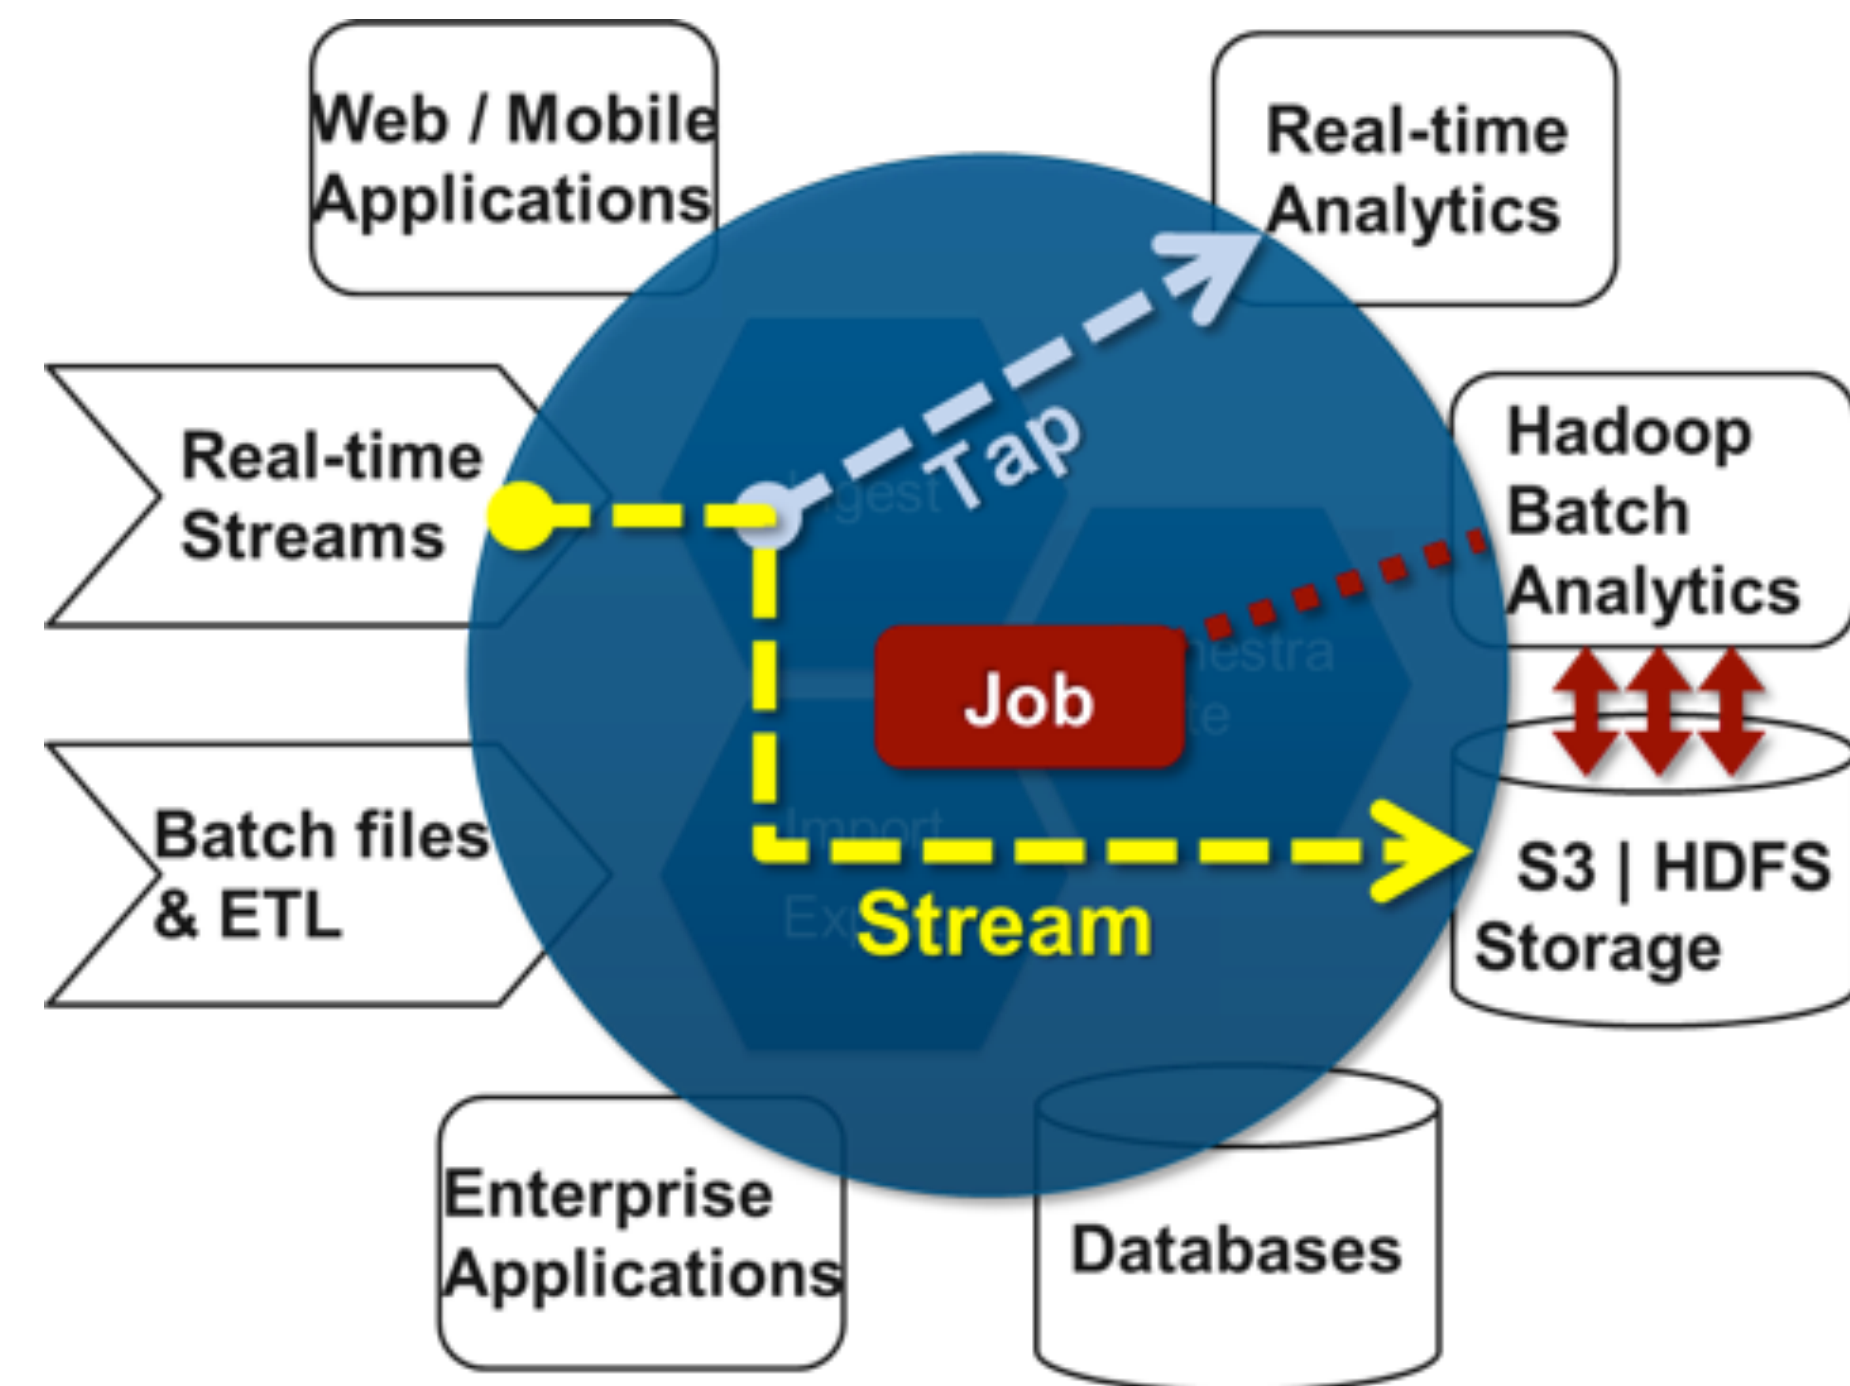
# Why Pivotal HD (PHD)

- ▶ Hadoop based
- ▶ Tightly integrated with *Gemfire XD*



# Why Spring XD?

- ▶ Data Ingestion
- ▶ Real-time Analytics
- ▶ Workflow Orchestration
- ▶ Integration



# Full Open source implementation



HBase??

Hadoop

SpringXD



# Is Lambda Architecture perfect?

- Suitable for specific use cases
- Doesn't contemplate reference data access.
- Not always possible to use same model for Real-time and Batch
- Other ideas to improve the lambda architecture
  - Multiple batch layers
  - incremental batch layers

# Source code



*<https://bitbucket.org/caxqueiroz/debs2014-challenge>*

Thank you!!